

# Jolt

SNARKs for virtual machines via lookups

**Arasu Arun<sup>\*</sup>, Srinath Setty<sup>†</sup>, Justin Thaler<sup>§‡</sup>**  
**Michael Zhu<sup>§</sup>, Sam Ragsdale<sup>§</sup>, Noah Citron<sup>§</sup>**

<sup>\*</sup>New York University    <sup>†</sup>Microsoft Research

<sup>§</sup>a16z crypto    <sup>‡</sup>Georgetown University

# Important Disclosures

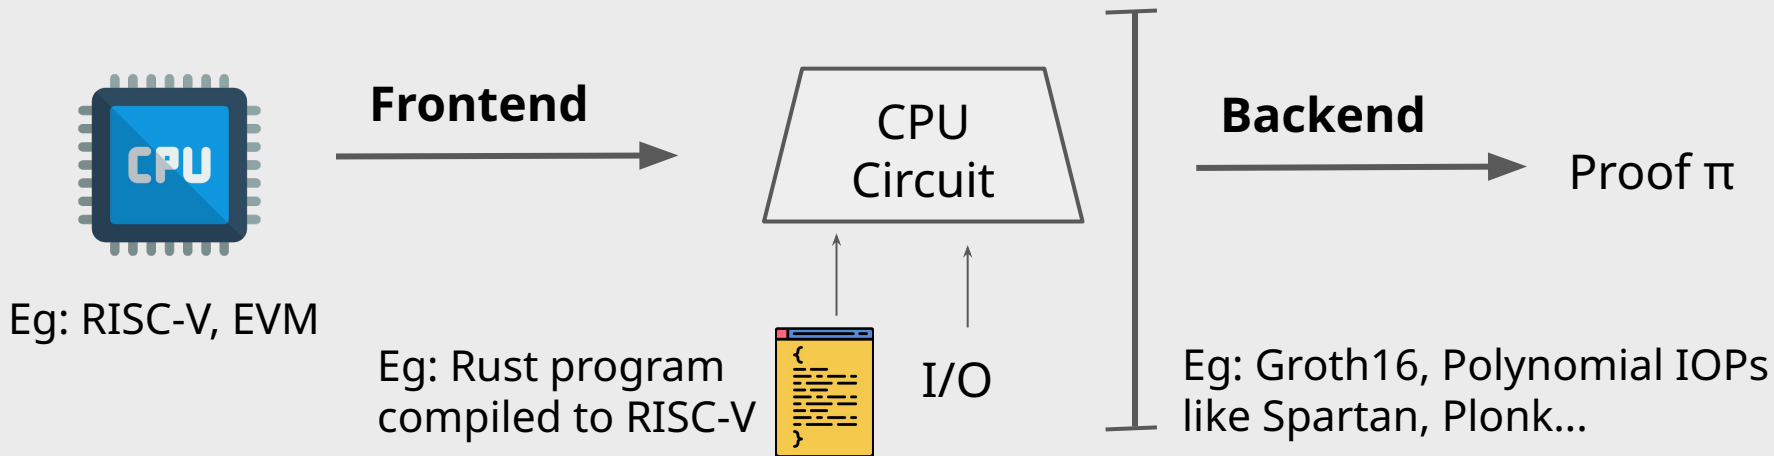
The views expressed here are those of the academic researcher and not necessarily of AH Capital Management, L.L.C. (“a16z”).

This content is provided for academic and informational purposes only, and should not be relied upon as business, investment, tax or legal advice and does not constitute an offer of advisory services.

Charts, graphs or other diagrams provided within are for academic and informational purposes only and should not be relied upon when making any investment decision. Past performance is not indicative of future results. The content speaks only as of the date indicated.

References to any securities, digital assets, and/or tokens (“assets”) are for illustrative purposes only and a16z may maintain holdings in the these assets. A list of investments made by funds managed by Andreessen Horowitz (excluding investments for which the issuer has not provided permission for a16z to disclose publicly as well as unannounced investments in publicly traded digital assets) is available at <https://a16z.com/investments/>.

# zkVM: frontends and backends



## Pros:

- One circuit for all programs.
- Re-use infrastructure: existing languages, compilers and tooling.
- Focus audit and optimization efforts on one circuit.

# Overheads involved in CPU circuits

```
switch (instr) {  
  case ADD: {...}  
  case SUB: {...}  
  ...  
  case XOR: {...}  
}
```

1. To handle arbitrary programs, each step must handle **all possible operations**.


RISC-V: ~ 50 operations

EVM: ~ 140 operations

These incur extra field operation and/or commitment costs.

2. Also, **bitwise operations** involved in primitive ISAs aren't efficiently performed with field elements (such as in a prime-order field).

# Jolt: just one lookup table



A zkVM that mostly performs **lookups** to tables outside of the circuit.

- Minimal: Just **60 constraints** and **80 field elements** per step of RISC-V!

How?

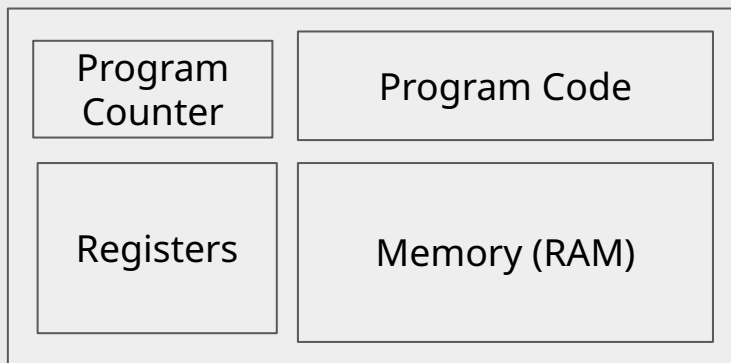
- Primitive assembly instructions have interesting **mathematical structure** (namely, efficient polynomial representations)
- We leverage this to perform structured lookups using **Lasso**.

This leads to a modular and extensible framework.

**Lasso** ([ia.cr/2023/1216](https://ia.cr/2023/1216)) was developed alongside Jolt.

# CPU state, transition arithmetization


**Machine state:**





**Transition function:**

1. **Fetch** instr.
2. **Decode** opcode, operands.
3. **Execute** instruction.
4. **Update** registers.

**Frontend**  
 $\Rightarrow$

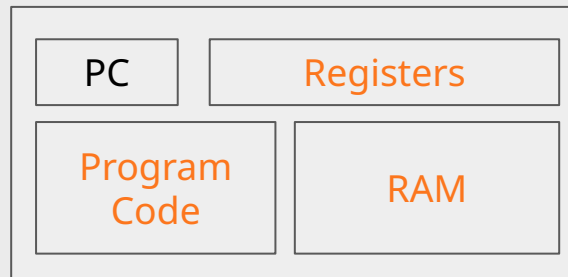


...




n-step  
program

# Memory-checking



We use “offline memory checking” (BEGKN91). Adapted to SNARKs in Spice (SAGL18).

Reduces consistency of memory operations to a **multiset equality check**. Jolt performs this using GKR-style grand product arguments.



Prover is **linear time** in the number of memory operations and memory size.

[BEGKN91] - Checking the correctness of memories - Blum et al., 1991

[SAGL18] - Proving the correct execution of concurrent services in zero-knowledge - Setty et al., 2018

# Offload instruction execution.

Suppose, for each operation, we build a table containing all possible executions:  $T_{op}(x, y) = OP(x, y)$  for all  $x, y$ .

1. Fetch instr.
2. Decode **opcode**, **operands**.
3. **Lookup** the instruction.
4. Update registers/RAM.

**opcode, operands**

z

$T_{RISC-V}$

$T_{ADD}$

$T_{XOR}$

...

$T_{MUL}$

This table would be way too big. Most instructions take two operands, which leads to  $2^{128}$  entries for 64-bit operands.



# But these tables are highly structured

We never have to materialize the tables because they each have a **succinct representation**.



Each operation's output is an efficient-to-evaluate **multilinear polynomial** over its input bits.

Why is this exciting?

Because polynomials are the language of SNARK backends.

Some examples:

$$T_{\text{XOR}}(x, y) = \sum_{i=0}^{63} 2^i (x_i \cdot y_i + (1 - x_i) \cdot (1 - y_i))$$


$$T_{\text{LT}}(x, y) = \sum_{i=0}^{63} (1 - x_i) \cdot y_i \cdot \widetilde{\text{EQ}}(x_{>i}, y_{>i})$$

$$T_{\text{SLL}}(x, y) = \sum_{k=0}^{63} \widetilde{\text{EQ}}(y, k) \cdot \sum_{j=k}^{63} 2^j x_{j-k}$$

$$\widetilde{\text{EQ}}(x, y) = \prod_{i \in n} (x_i \cdot y_i + (1 - x_i) \cdot (1 - y_i))$$

# The tables can be decomposed even further.

Each table is a simple collation of smaller **subtables**, each represented by an MLE over a chunk of the original inputs.

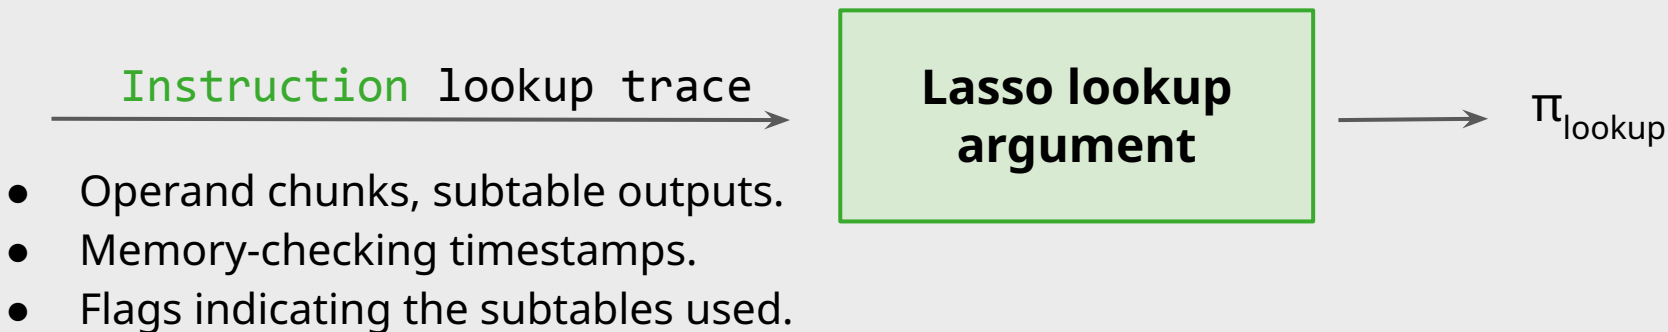


We only need **23 unique subtables** MLEs to represent **all RISC-V base** instructions.

# Lasso efficiently looks up decomposed tables

**Core tools:** sum-checks and memory-checking. Built on Spark from Spartan.

(Setty19: [ia.cr/2019/550](https://ia.cr/2019/550))



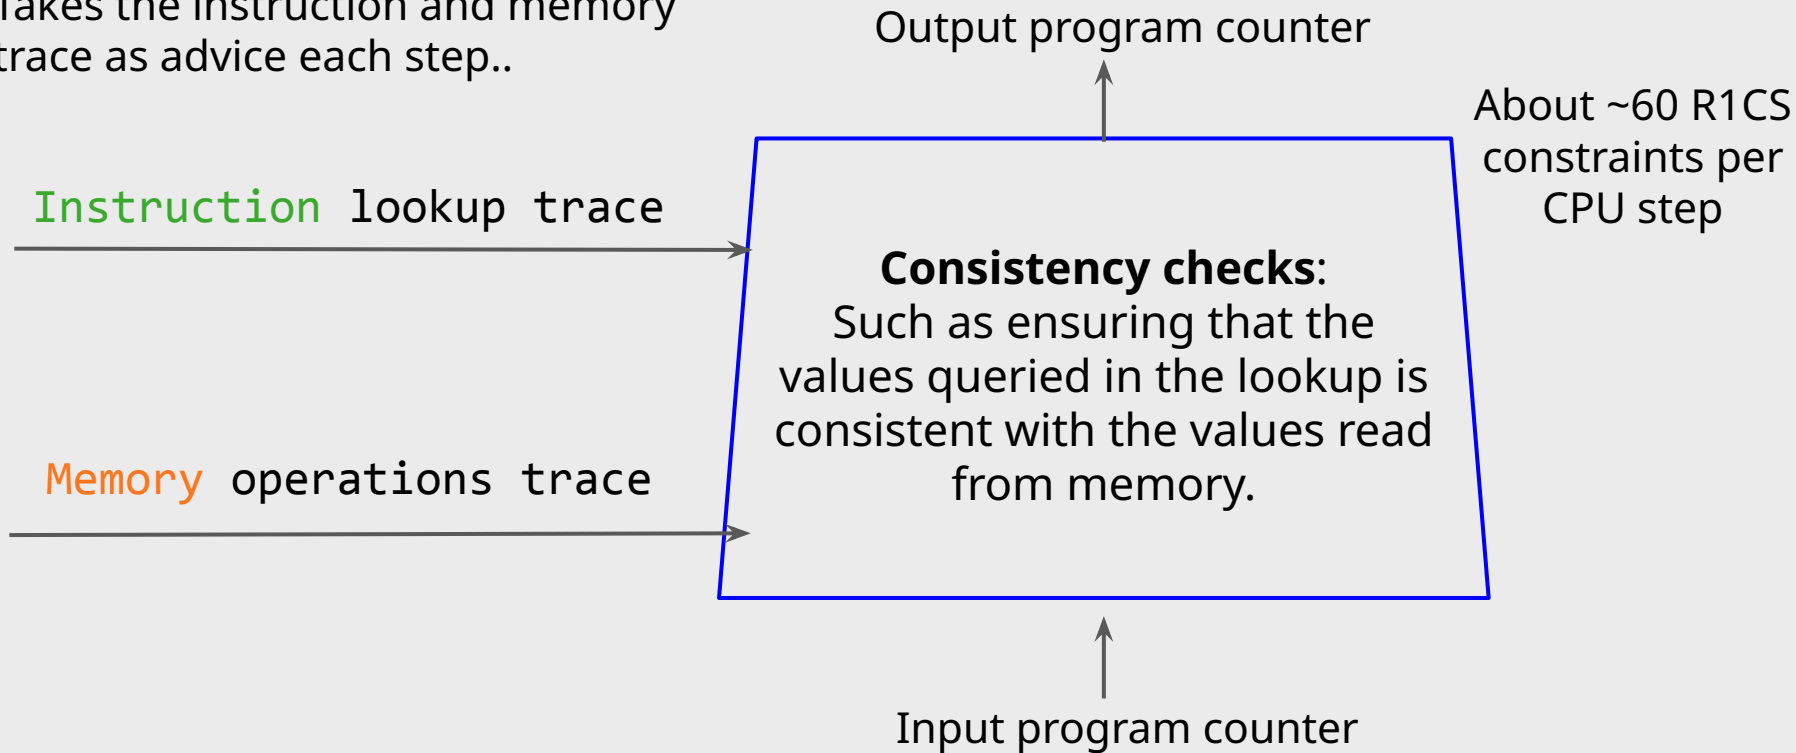
Performing  $m$  lookup s into an  $N$ -sized table with  $c$  chunks =  $O(cm + N^{1/c})$   
prover commitment costs.

$N = 2^{128}$ ,  $c = 8 \Rightarrow$  second term is  $2^{16}$

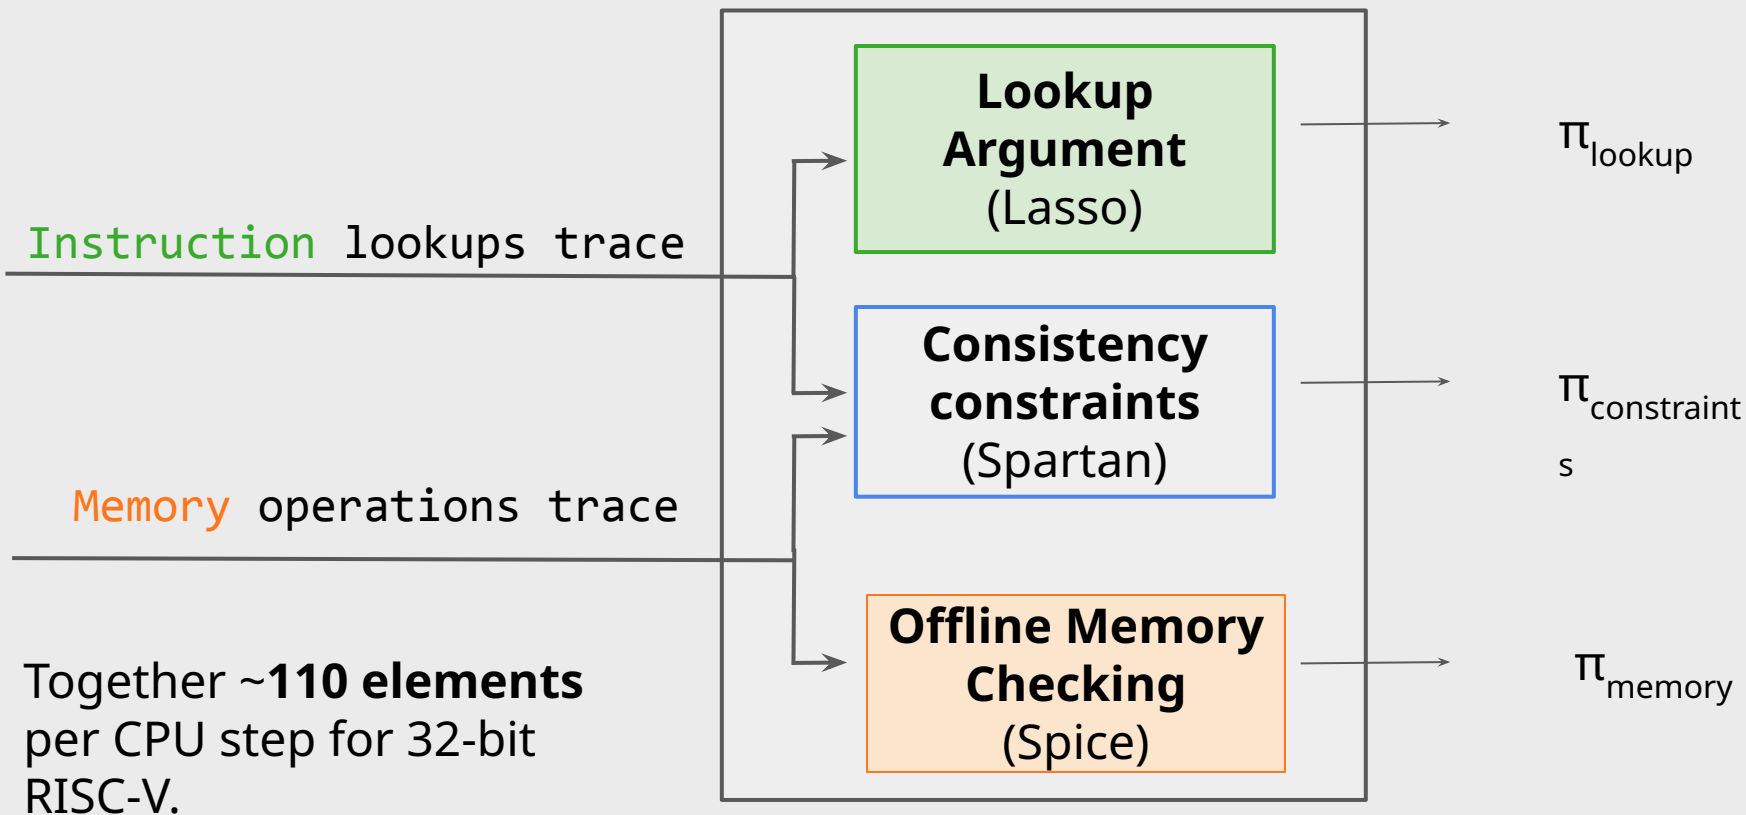
(Lasso - STW23: [ia.cr/2023/1216](https://ia.cr/2023/1216))

# So what does the CPU circuit do now?

Takes the instruction and memory trace as advice each step..



# Putting it all together: the Jolt prover modules



# The Jolt prover's costs

Commitment costs: Using Hyrax. As most elements are small, it's equivalent to **11 arbitrary** (256-bit long) field elements per step when using Pippenger's MSM algorithm for 32-bit RISC-V.

Jolt's backend: just **sum-checks** and multilinear polynomial evaluations

Module	Main operation	P cost
Lookups (Lasso)	1 sum-check, 2 GKR's	$O(c^2n)$
Constraints (Spartan)	2 sum-checks	$O(n)$
Memory-checking (Spice)	2 GKR's	$O(n +  \text{memory} )$

# Initial Jolt implementation

- Open source: <https://github.com/a16z/jolt>
- Instruction set: RISC-V 32-bit Base Integer Instruction Set (RV32I)
- Polynomial commitment scheme: Hyrax
- Elliptic curve: BN254 (interchangeable)
- Fork of [Spartan2](#) (optimized to leverage uniformity) to prove R1CS


# VM Instructions in Jolt



# VM Instructions in Jolt

XOR(11001001, 01111000)

# VM Instructions in Jolt

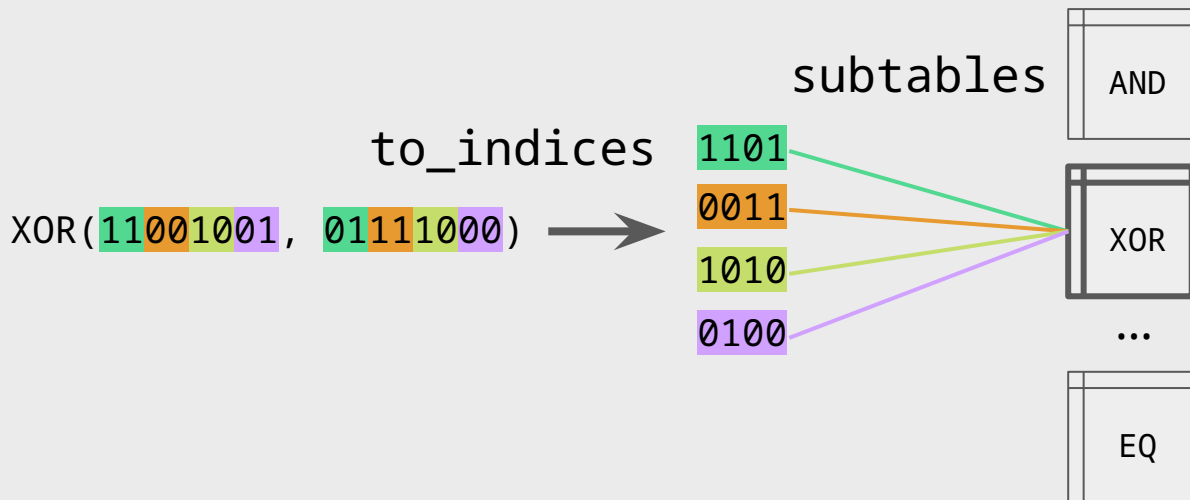


```
fn to_indices(&self, C: usize, log_M: usize) -> Vec<usize> {  
    chunk_and_concatenate_operands(self.0, self.1, C, log_M)  
}
```

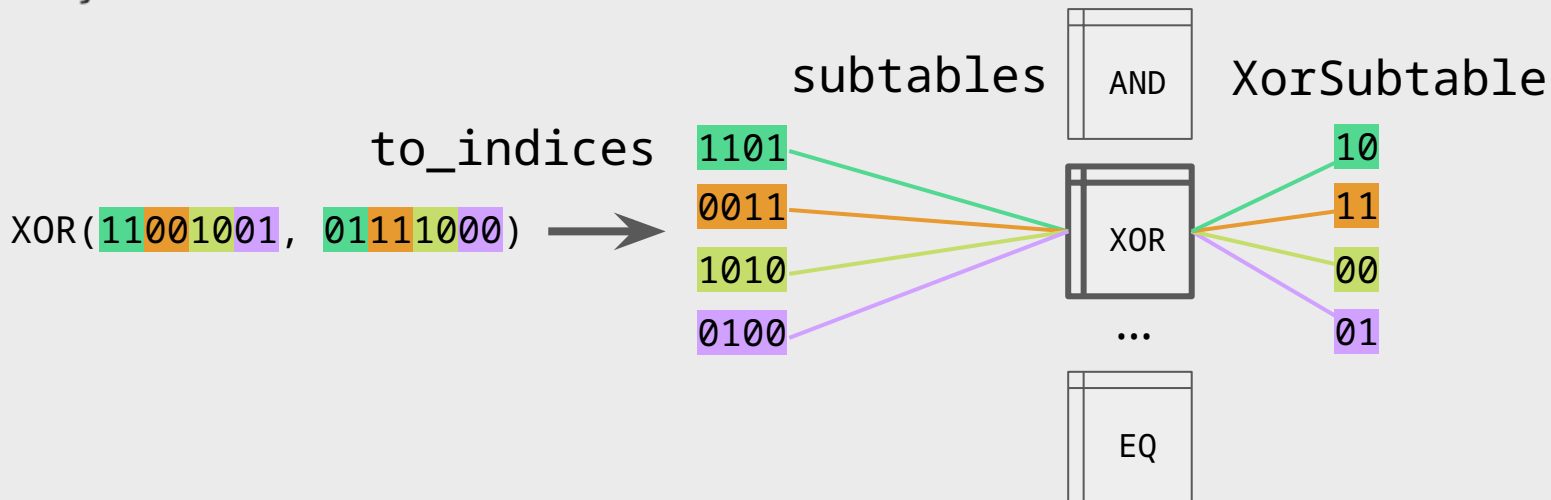
XOR(11001001, 01111000)  $\rightarrow$  to\_indices

1101
0011
1010
0100

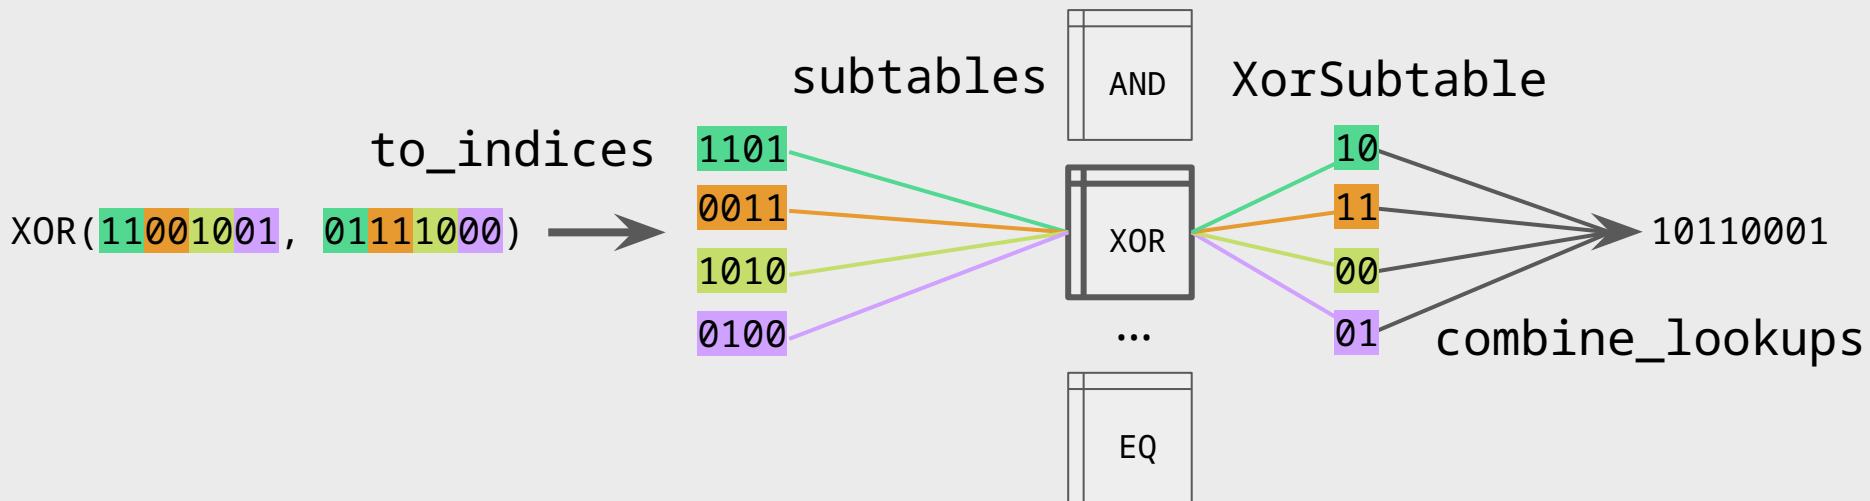
```
fn subtables<F: JoltField>(
    &self,
    C: usize,
    _: usize,
) -> Vec<(Box<dyn LassoSubtable<F>>, SubtableIndices)> {
    vec![(Box::new(XorSubtable::new()), SubtableIndices::from(0..C))]
}
```



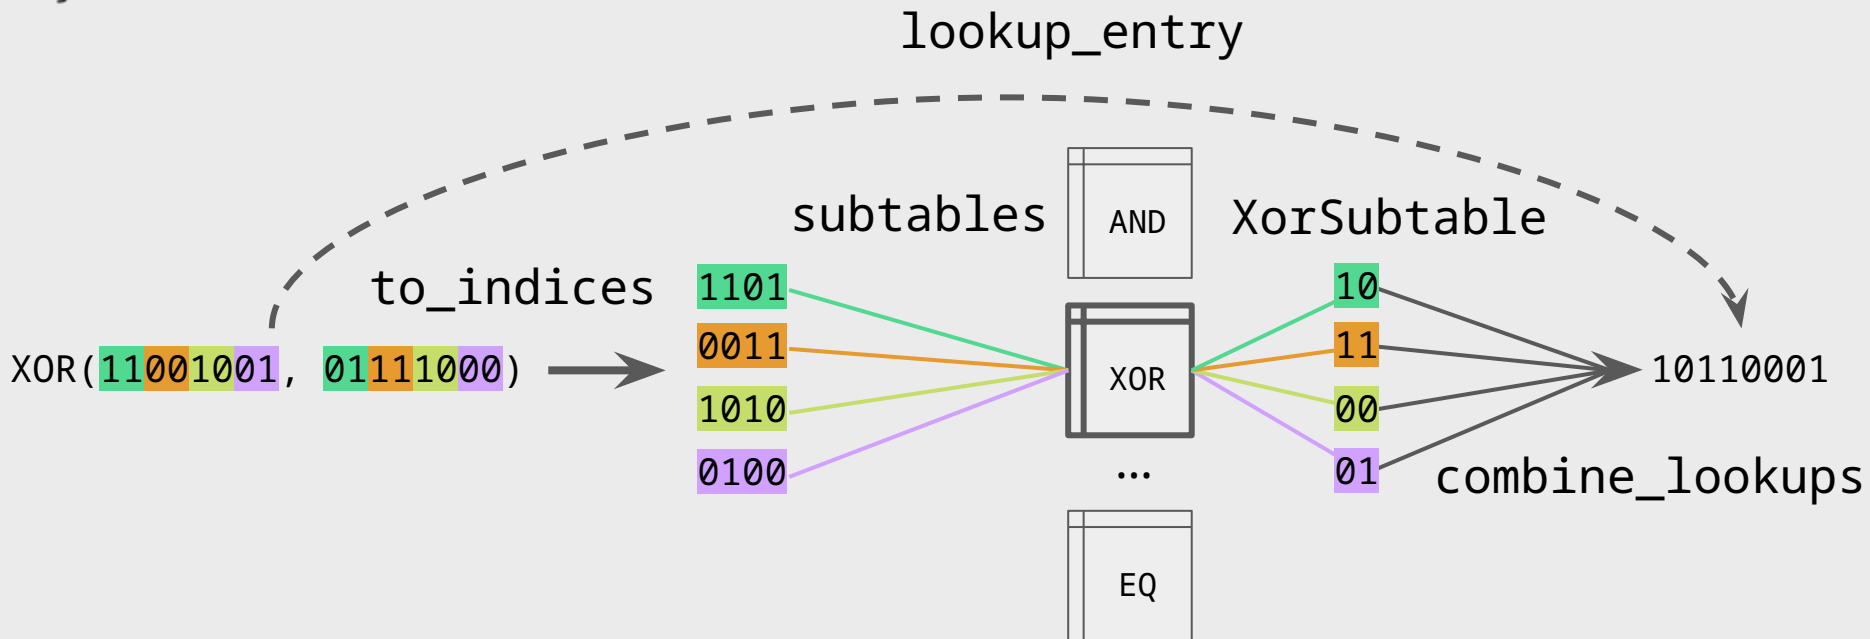
```
pub trait LassoSubtable<F: JoltField>: 'static + Sync {  
    fn subtable_id(&self) -> SubtableId {  
        TypeId::of::<Self>()  
    }  
    fn materialize(&self, M: usize) -> Vec<F>;  
    fn evaluate_mle(&self, point: &[F]) -> F;  
}
```

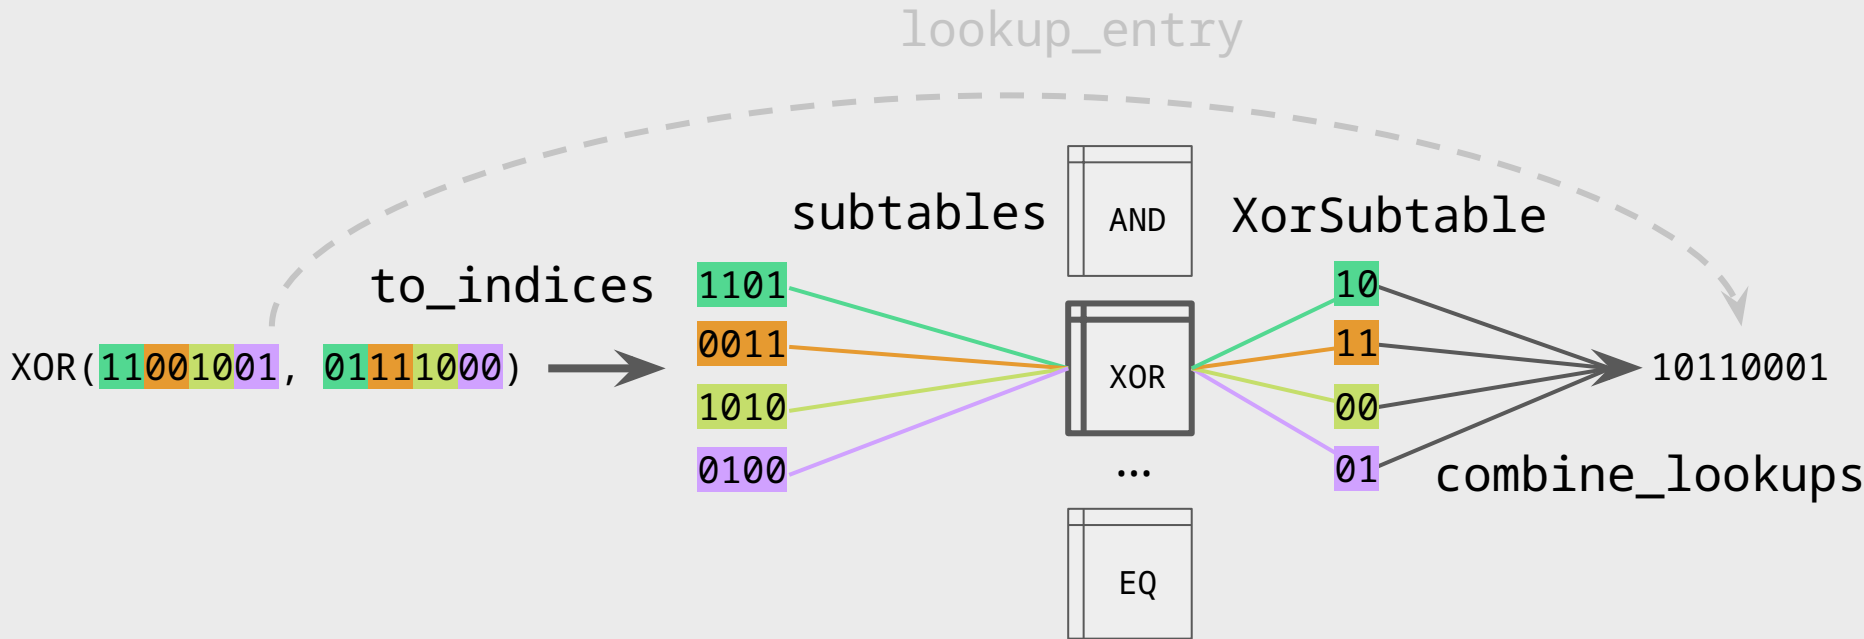


```
fn combine_lookups<F: JoltField>(&self, vals: &[F], C: usize, M: usize) -> F {  
    concatenate_lookups(vals, C, log2(M) as usize / 2)  
}
```

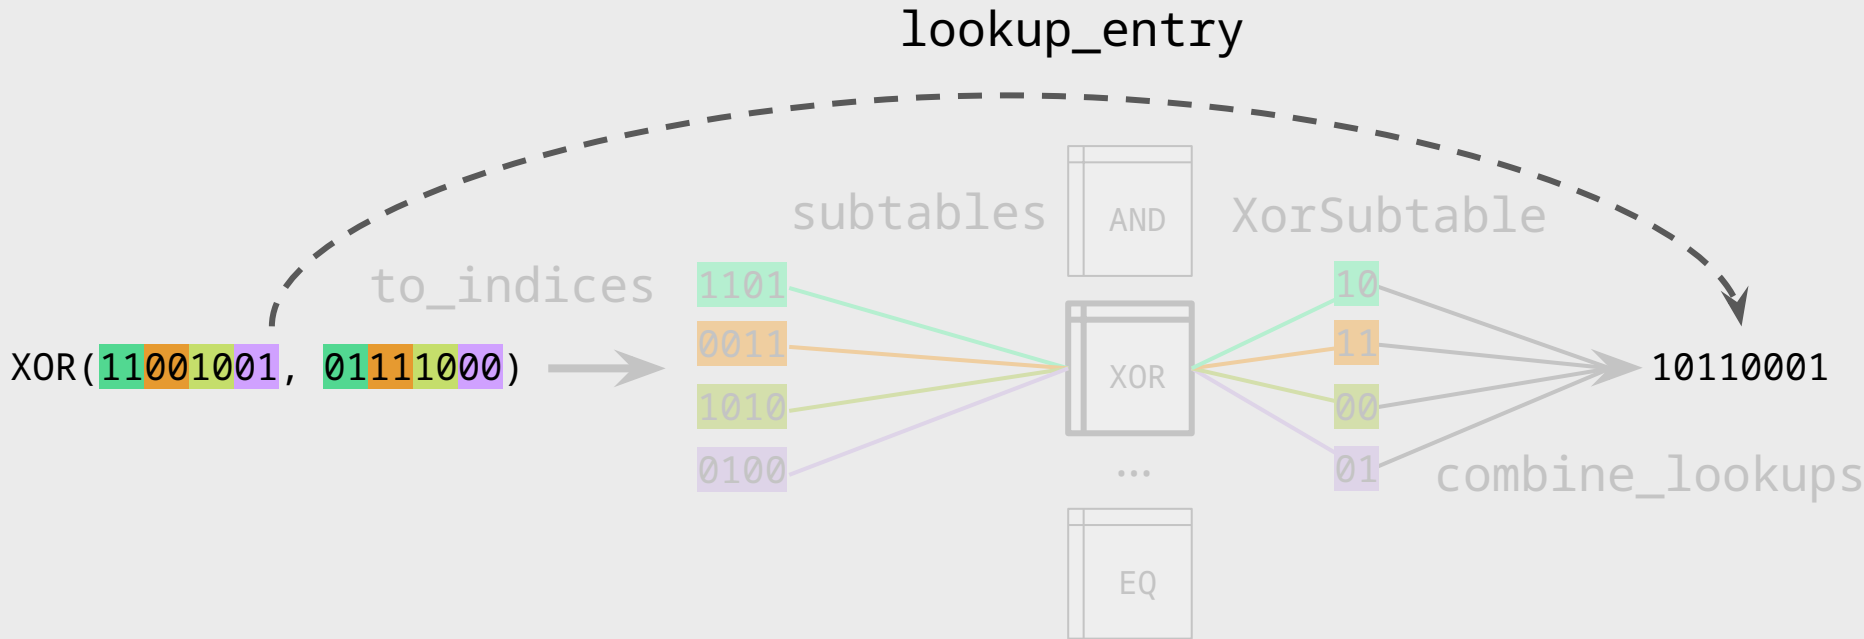


```
fn lookup_entry(&self) -> u64 {  
    self.0 ^ self.1  
}
```









# Performance

Jolt proves ~100,000 RV32I instructions per second (**100 kHz**) on M3 Max Macbook Pro (16-core CPU, 128GB RAM)

# Performance

Jolt proves ~100,000 RV32I instructions per second (**100 kHz**) on M3 Max Macbook Pro (16-core CPU, 128GB RAM)

$$\text{Proving overhead} = \frac{\text{CPU clock speed}}{\text{Proving speed}}$$

# Performance

Jolt proves ~100,000 RV32I instructions per second (**100 kHz**) on M3 Max Macbook Pro (16-core CPU, 128GB RAM)

$$\text{Proving overhead} = \frac{\text{CPU clock speed}}{\text{Proving speed}} = \frac{12 \cdot 4.05 \text{ GHz} + 4 \cdot 2.75 \text{ GHz}}{100 \text{ kHz}}$$

# Performance

Jolt proves ~100,000 RV32I instructions per second (**100 kHz**) on M3 Max Macbook Pro (16-core CPU, 128GB RAM)


$$\text{Proving overhead} = \frac{\text{CPU clock speed}}{\text{Proving speed}} = \frac{\begin{array}{c} \text{M3 Max "performance" cores} \\ 12 \cdot 4.05 \text{ GHz} \end{array} + \begin{array}{c} \text{M3 Max "efficiency" cores} \\ 4 \cdot 2.75 \text{ GHz} \end{array}}{100 \text{ kHz}}$$

# Performance


Jolt proves ~100,000 RV32I instructions per second (**100 kHz**) on M3 Max Macbook Pro (16-core CPU, 128GB RAM)

$$\begin{aligned} \text{Proving overhead} &= \frac{\text{CPU clock speed}}{\text{Proving speed}} = \frac{\begin{array}{c} \text{M3 Max "performance" cores} \\ 12 \cdot 4.05 \text{ GHz} \end{array} + \begin{array}{c} \text{M3 Max "efficiency" cores} \\ 4 \cdot 2.75 \text{ GHz} \end{array}}{100 \text{ kHz}} \\ &= \mathbf{596,000x} \text{ overhead} \\ &\quad \text{vs native CPU execution} \end{aligned}$$

# Performance



# Performance

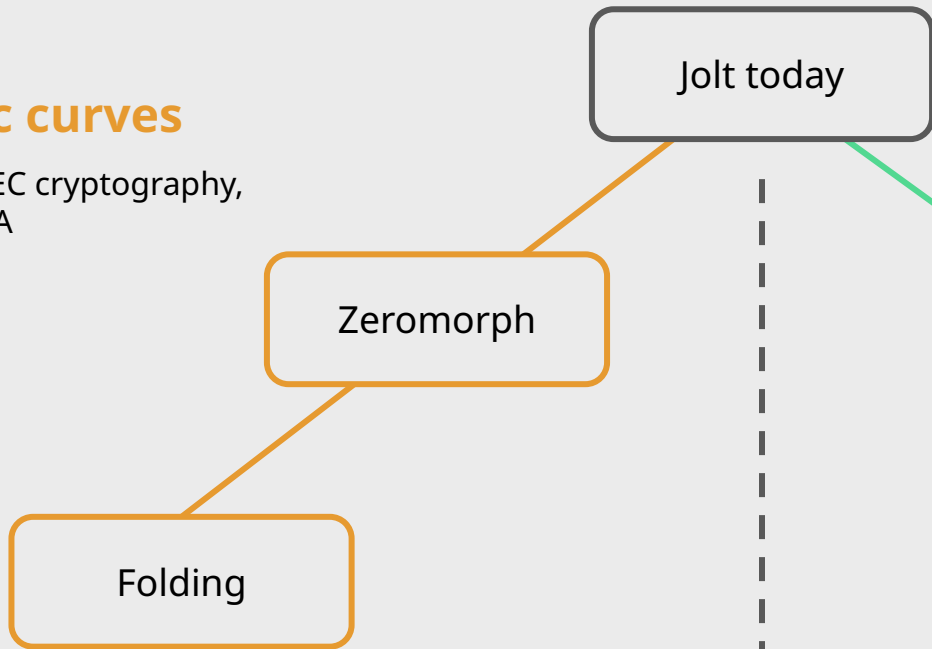




# What's next?


## elliptic curves

good for EC cryptography,  
e.g. ECDSA



## binary extension fields

good for general purpose



# Thank you!