


Memory checking in IVC-based zkVMs

ZKProof, Berlin 2024


**Jens Groth, Nexus - nexus.xyz
based on discussions with Yinuo Zhang**

Virtual Machine




- Program P specified according to an Instruction Set Architecture
- Input x to the program (public)
- Auxiliary input w to the program (maybe not public)
- Output y

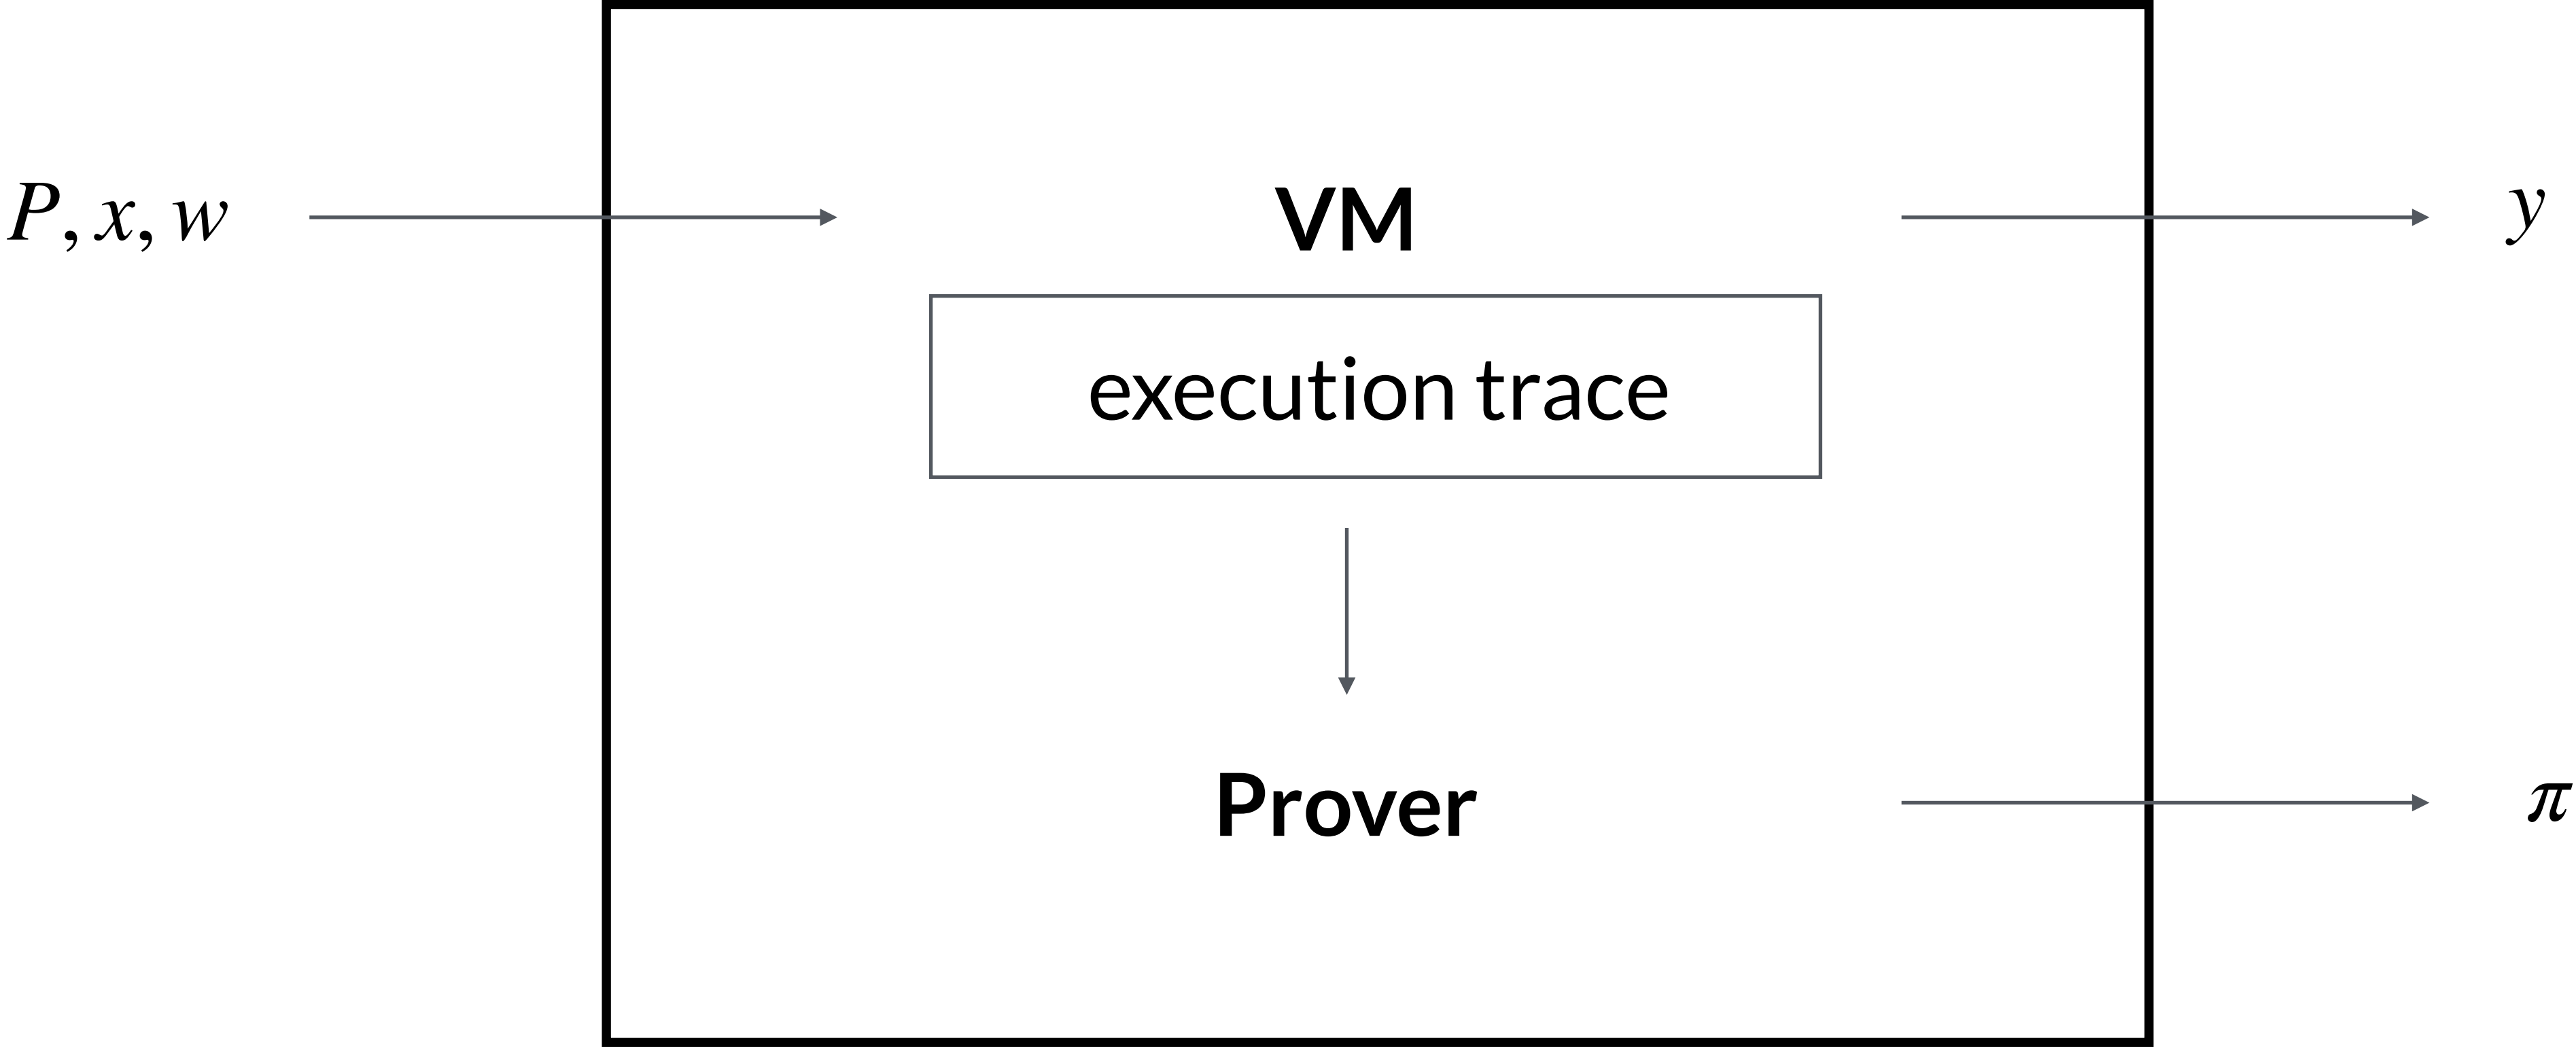
Zero Knowledge Virtual Machine



Zero Knowledge Virtual Machine & matching verification algorithm



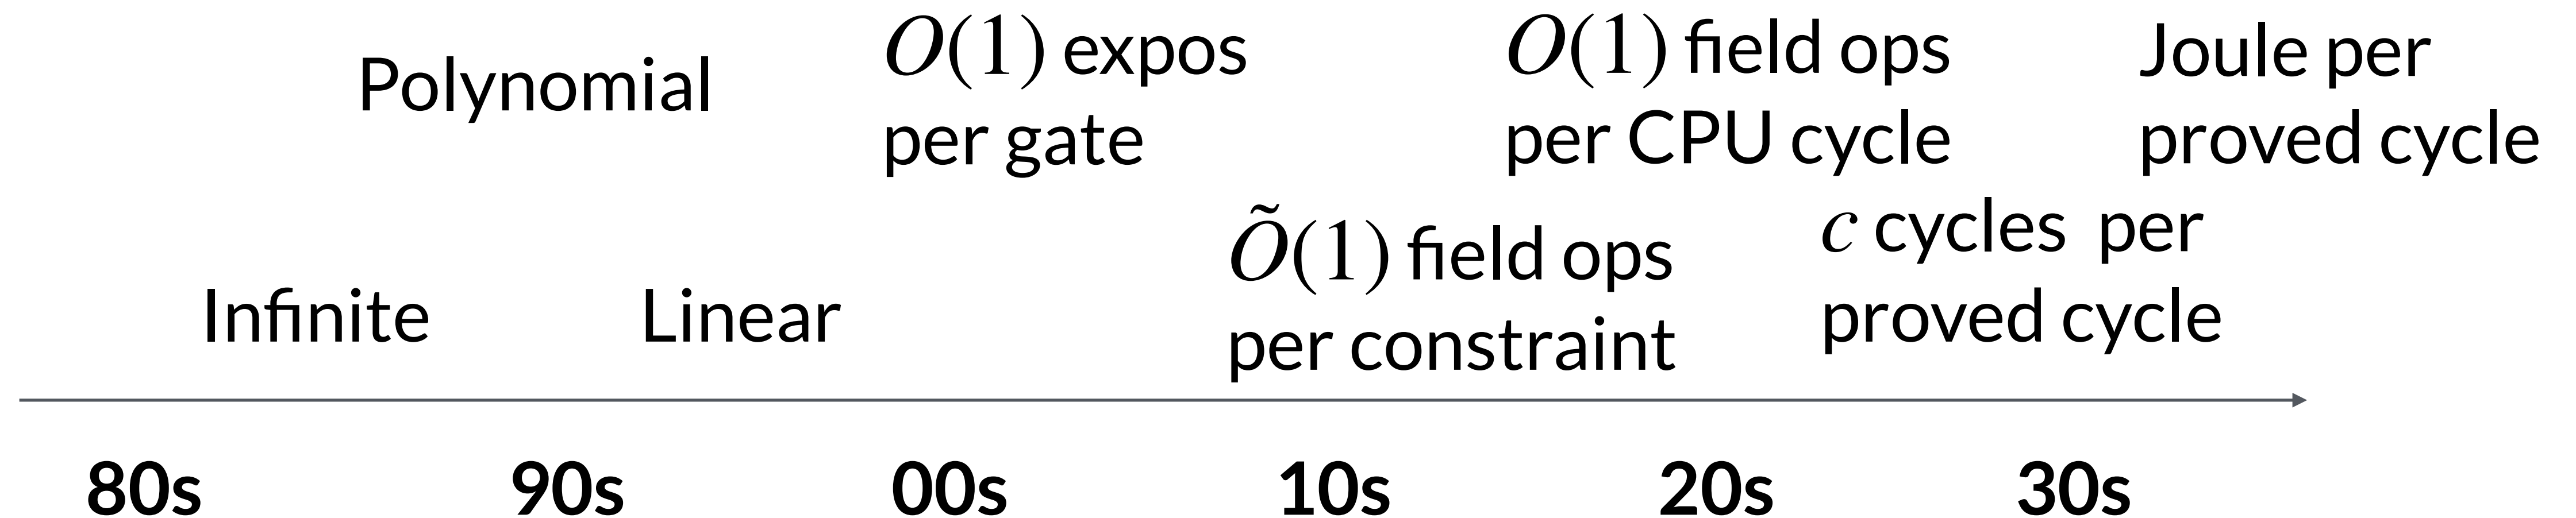
Inside the zkVM



Prover efficiency

- CPUs: #transistors per IC increases $\sim 50\%$ /year (Moore's law)
- zkVMs: #proved CPU cycles per second increases $> 1000\%$ /year (in 2024)


Polite language amongst cryptographers when describing prover time



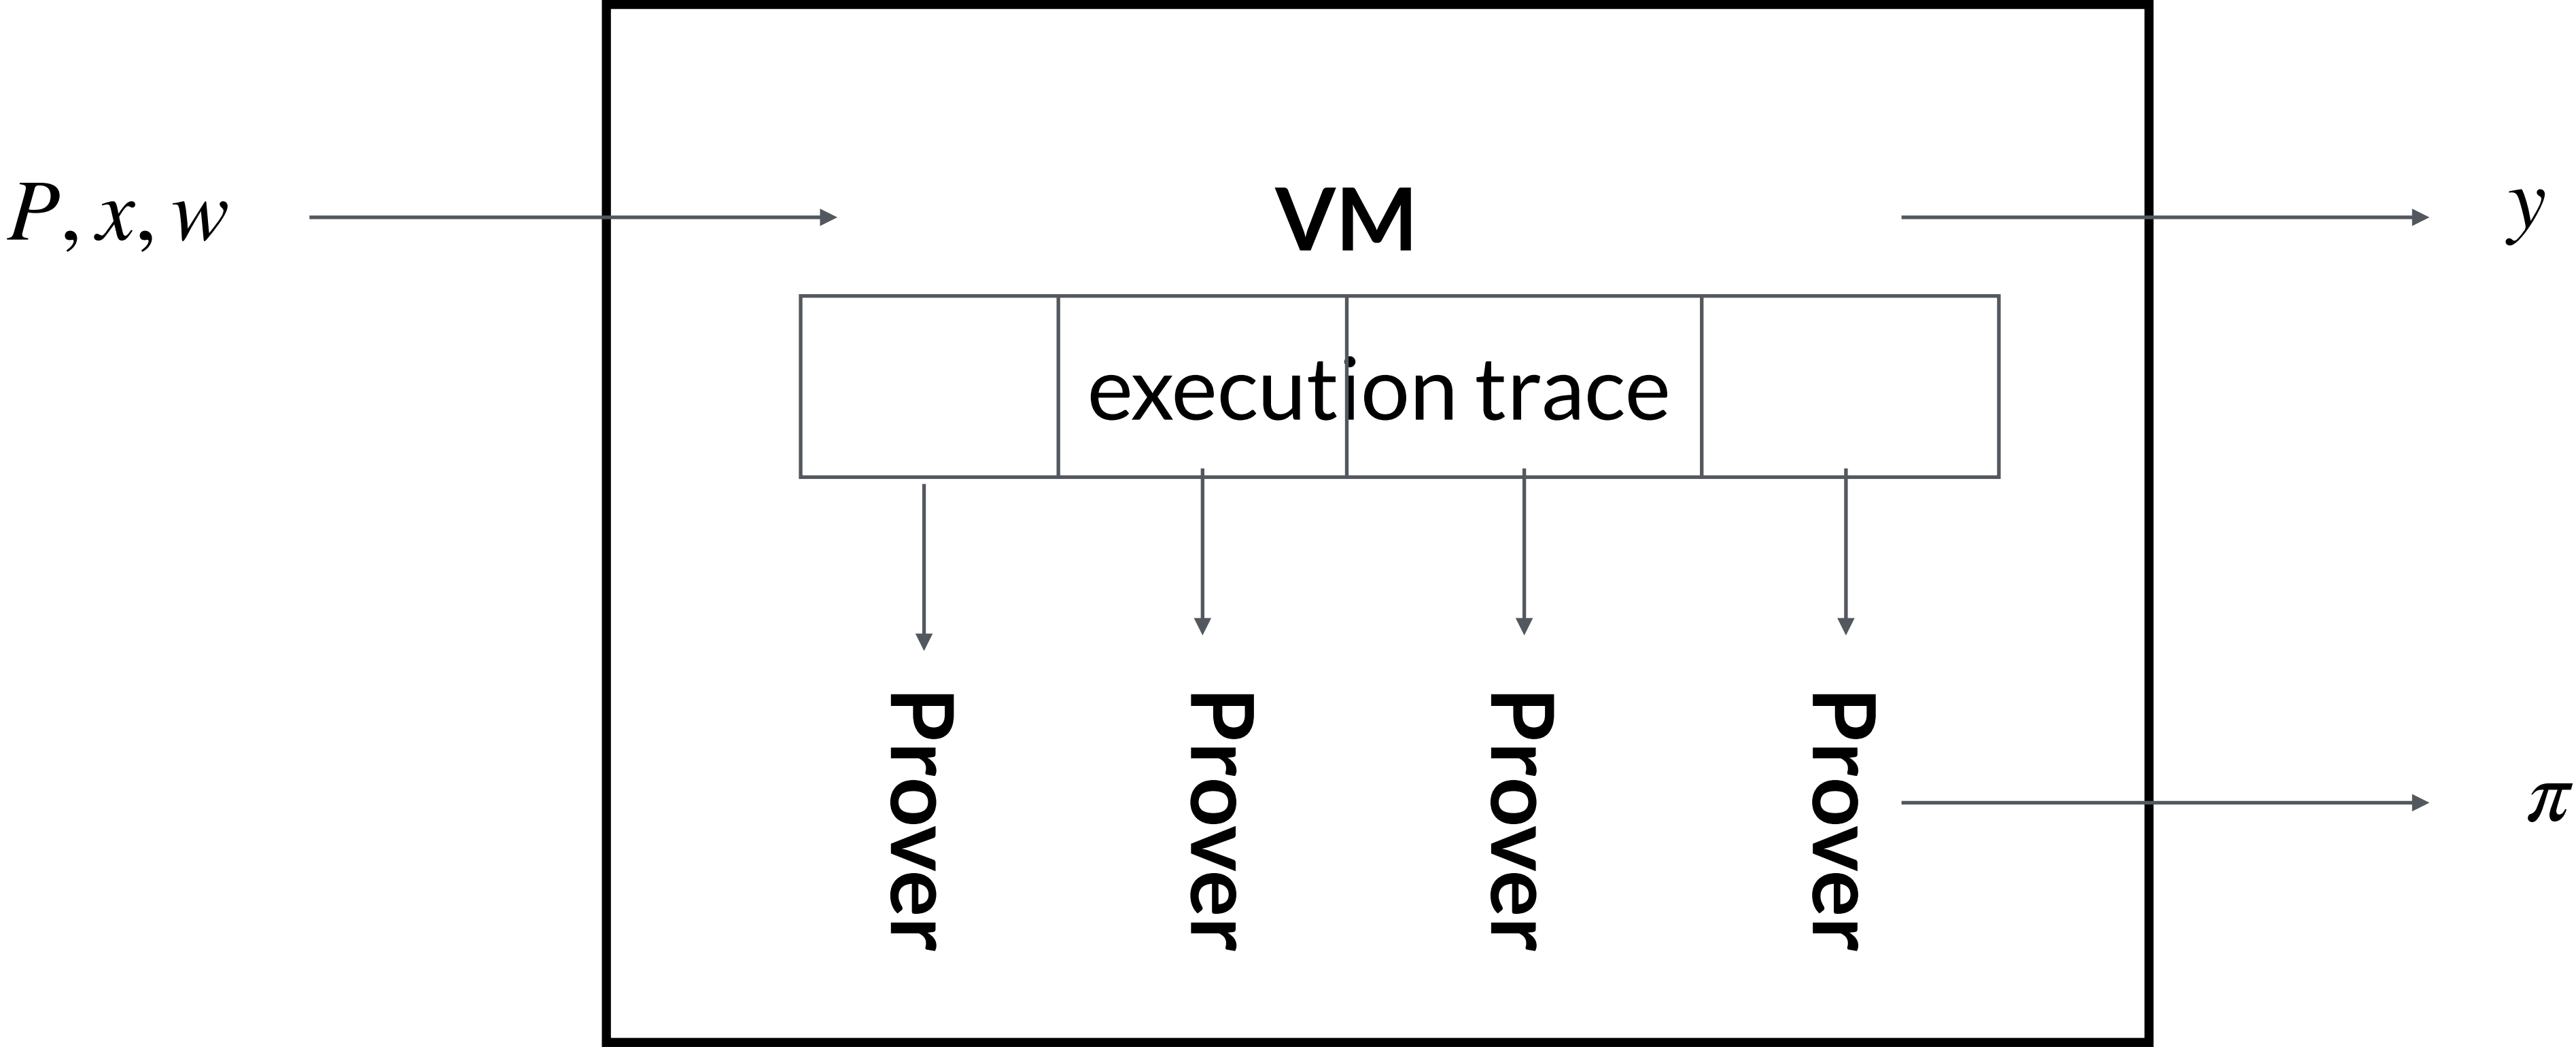
Prover efficiency is still the bottleneck

- Proof size ✓
- Verification time ✓
- Proving time ?


Timeline of complicated big things to prove




Slicing a long computation



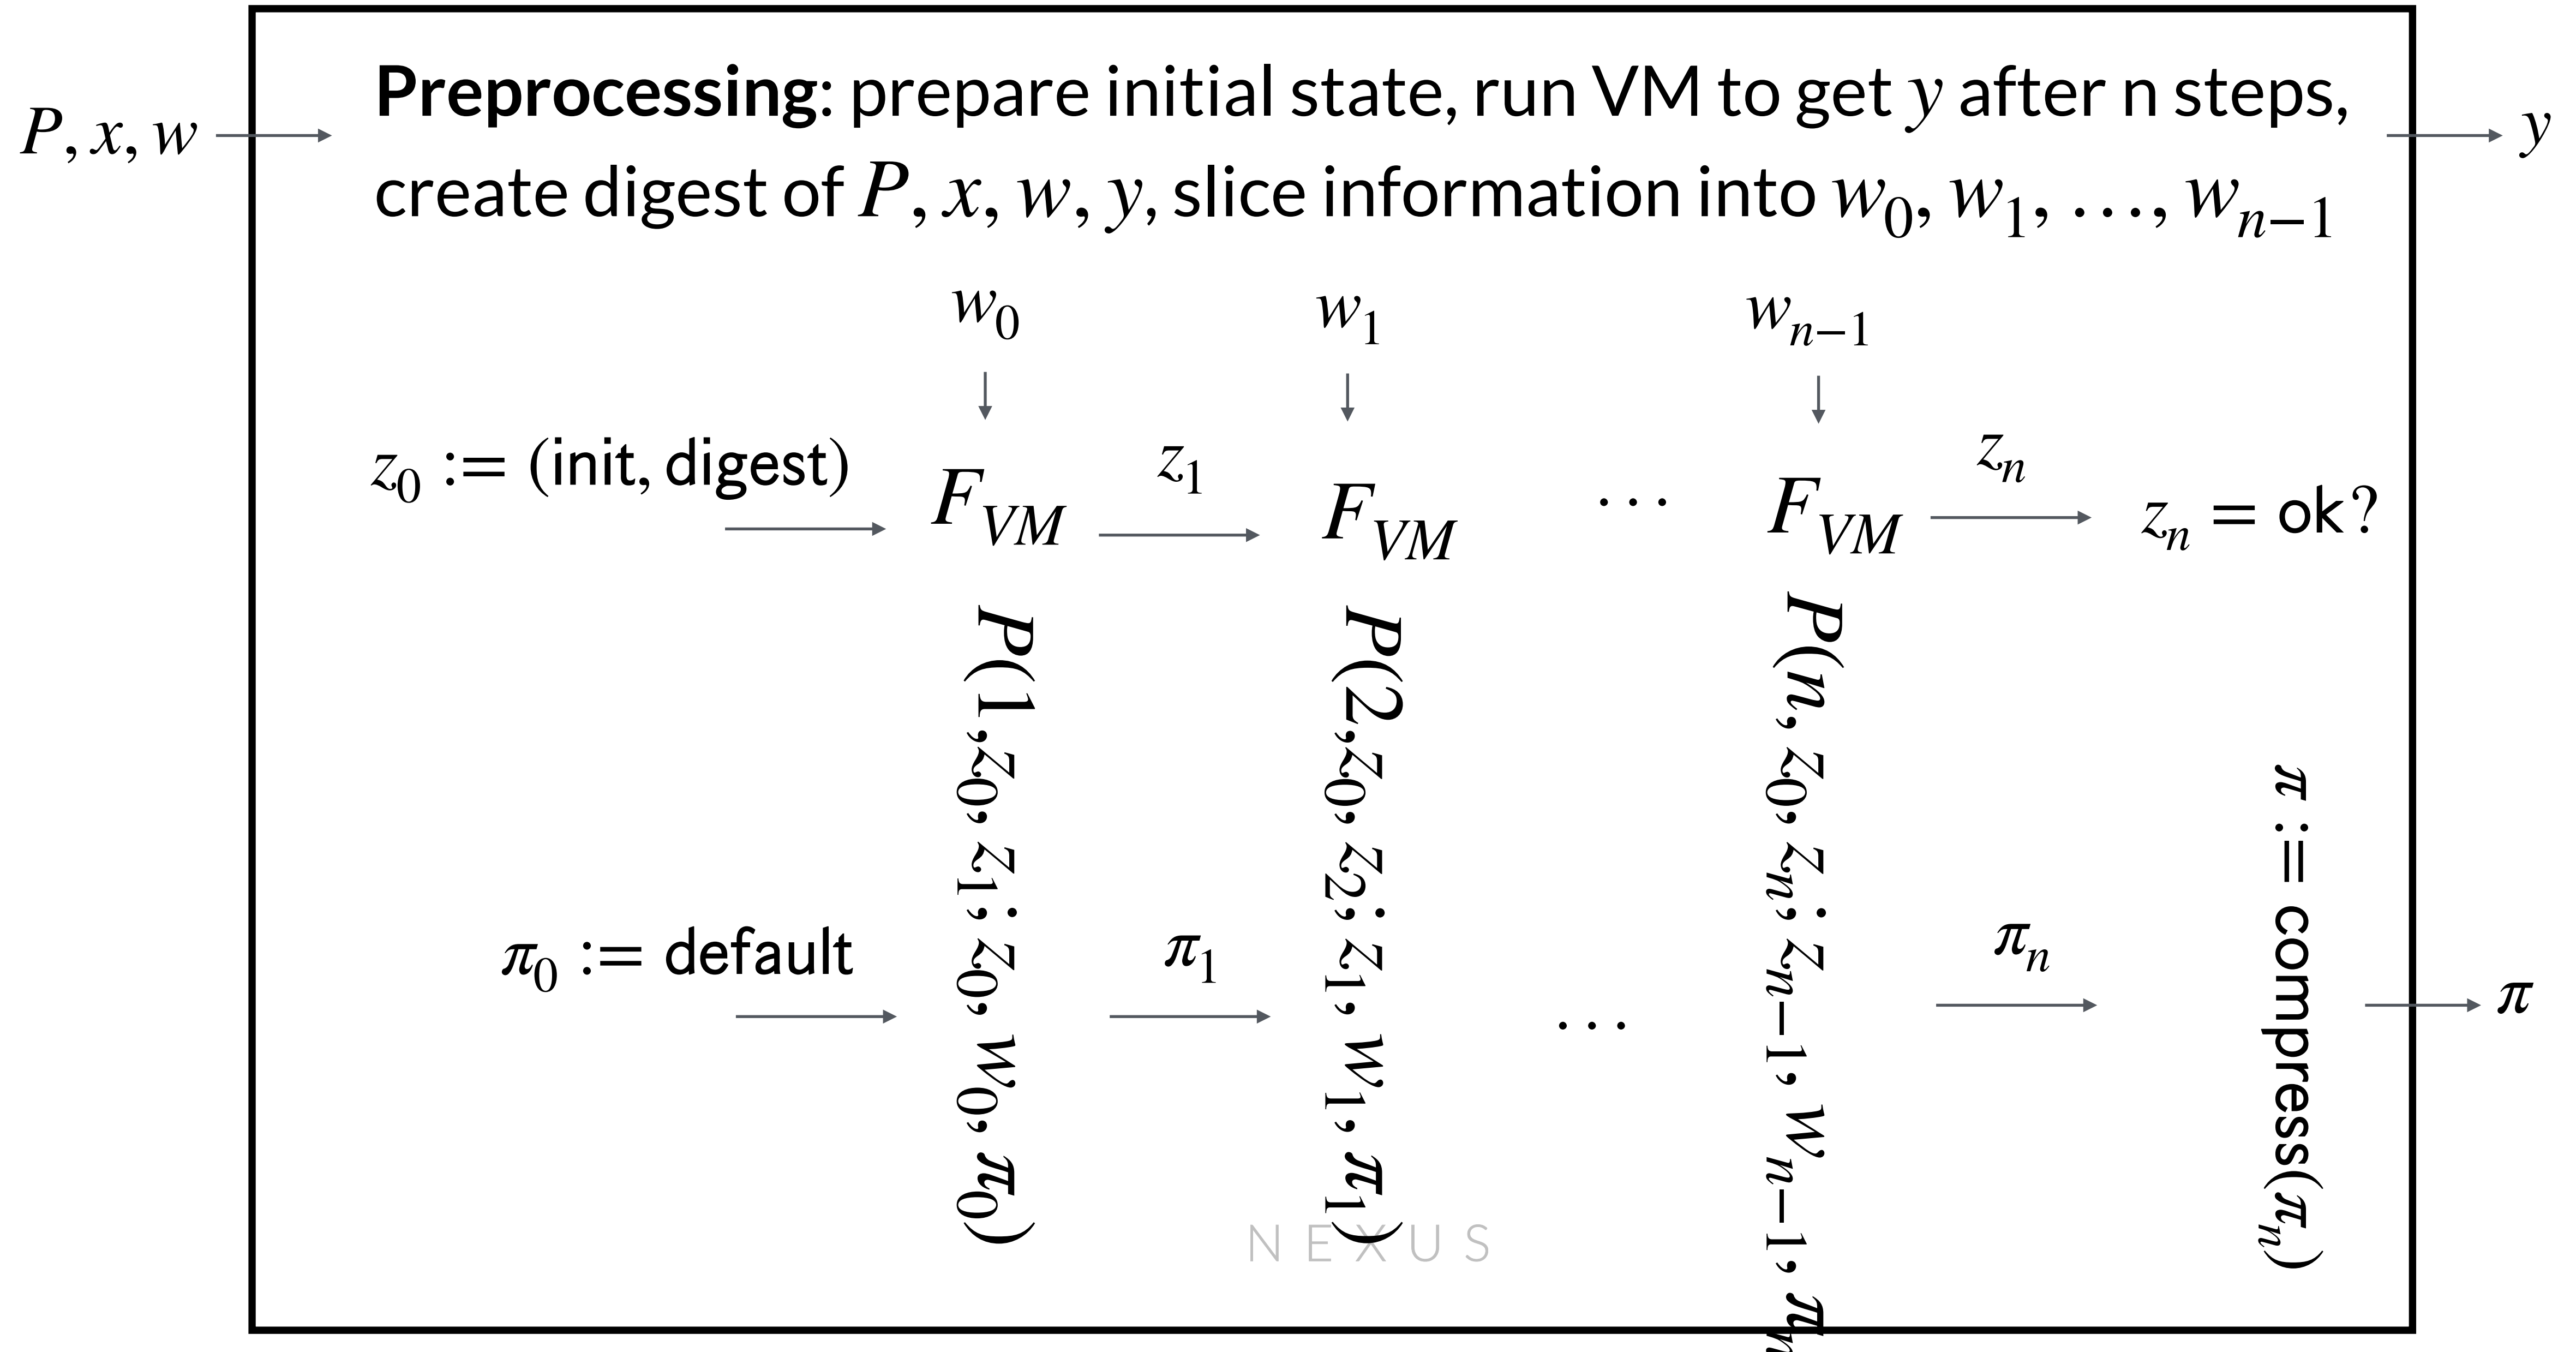
Incrementally verifiable computation




IVC-based zkVM



IVC-handling by the zkVM



The cost of maintaining memory

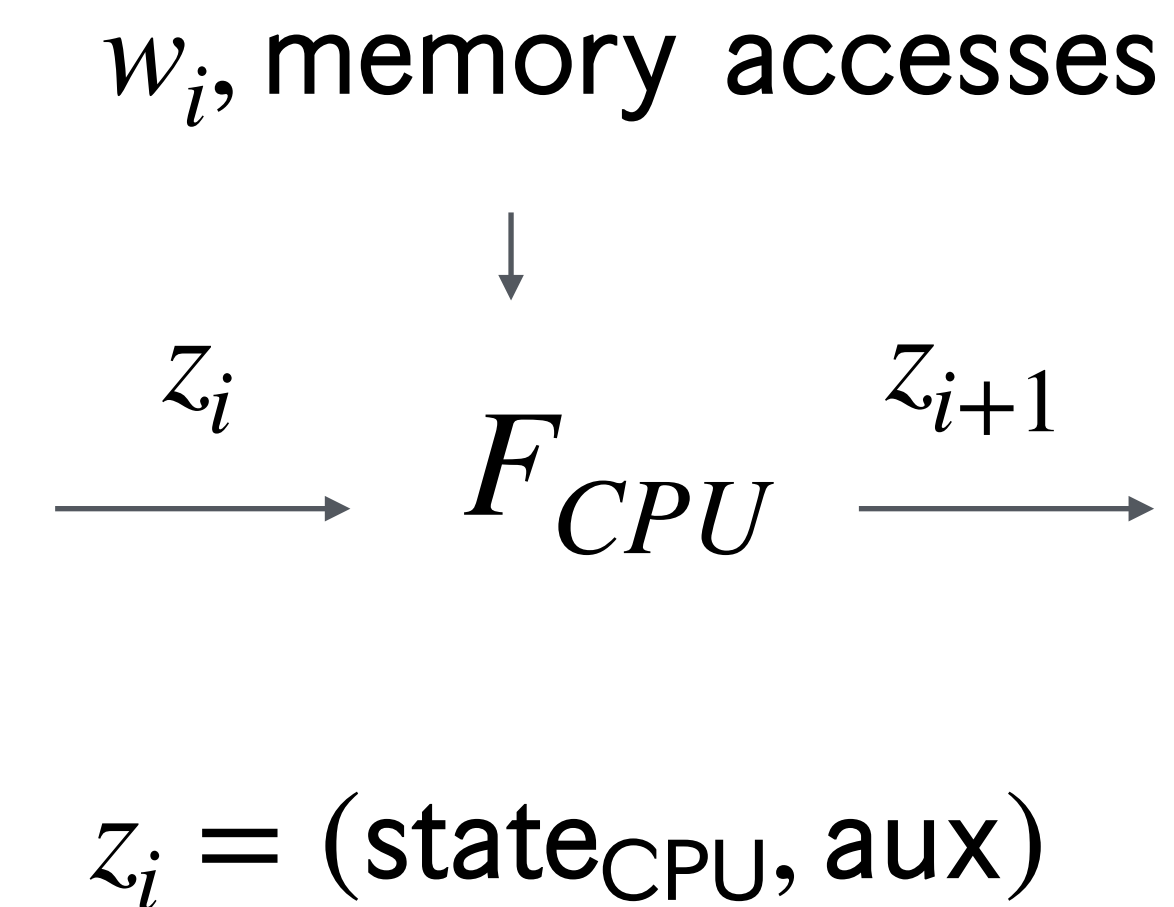
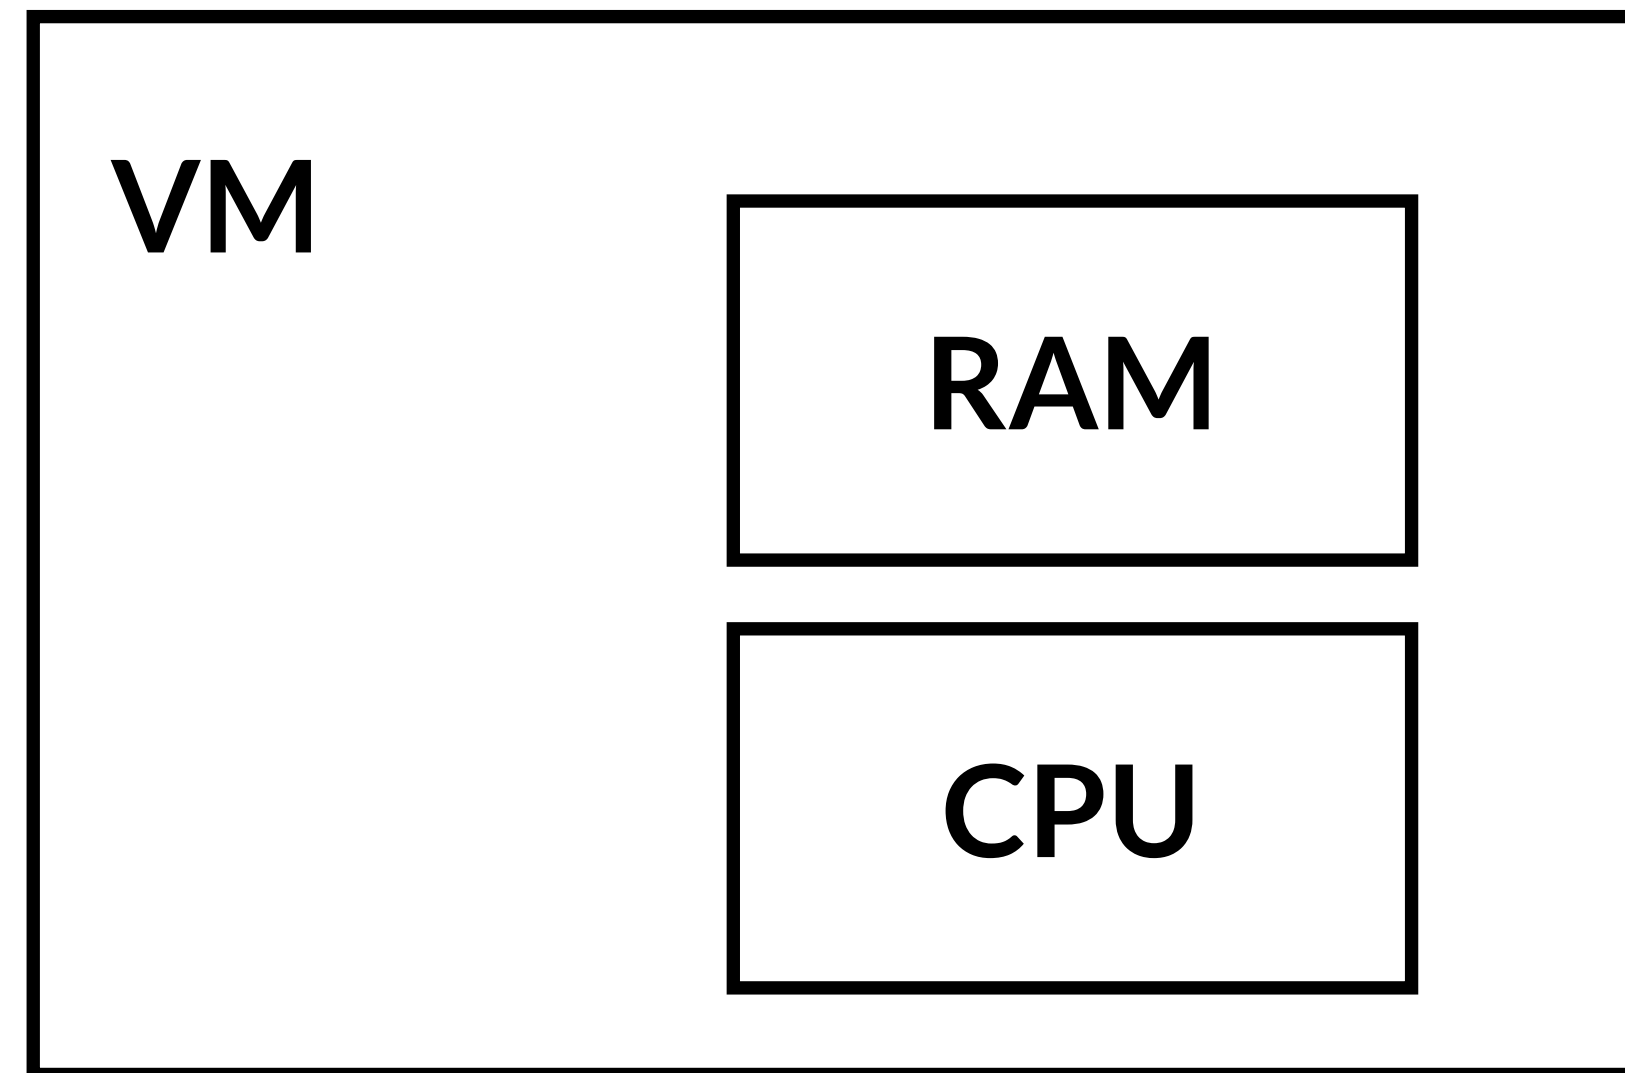


$$\begin{array}{ccccc} & & w_i & & \\ & & \downarrow & & \\ \xrightarrow{z_i} & & F_{VM} & & \xrightarrow{z_{i+1}} \end{array}$$

$z_i = (\text{memory}, \text{state}_{\text{CPU}}, \text{aux})$


- Prover cost: constraints for memory transition :- (

Externalizing the memory

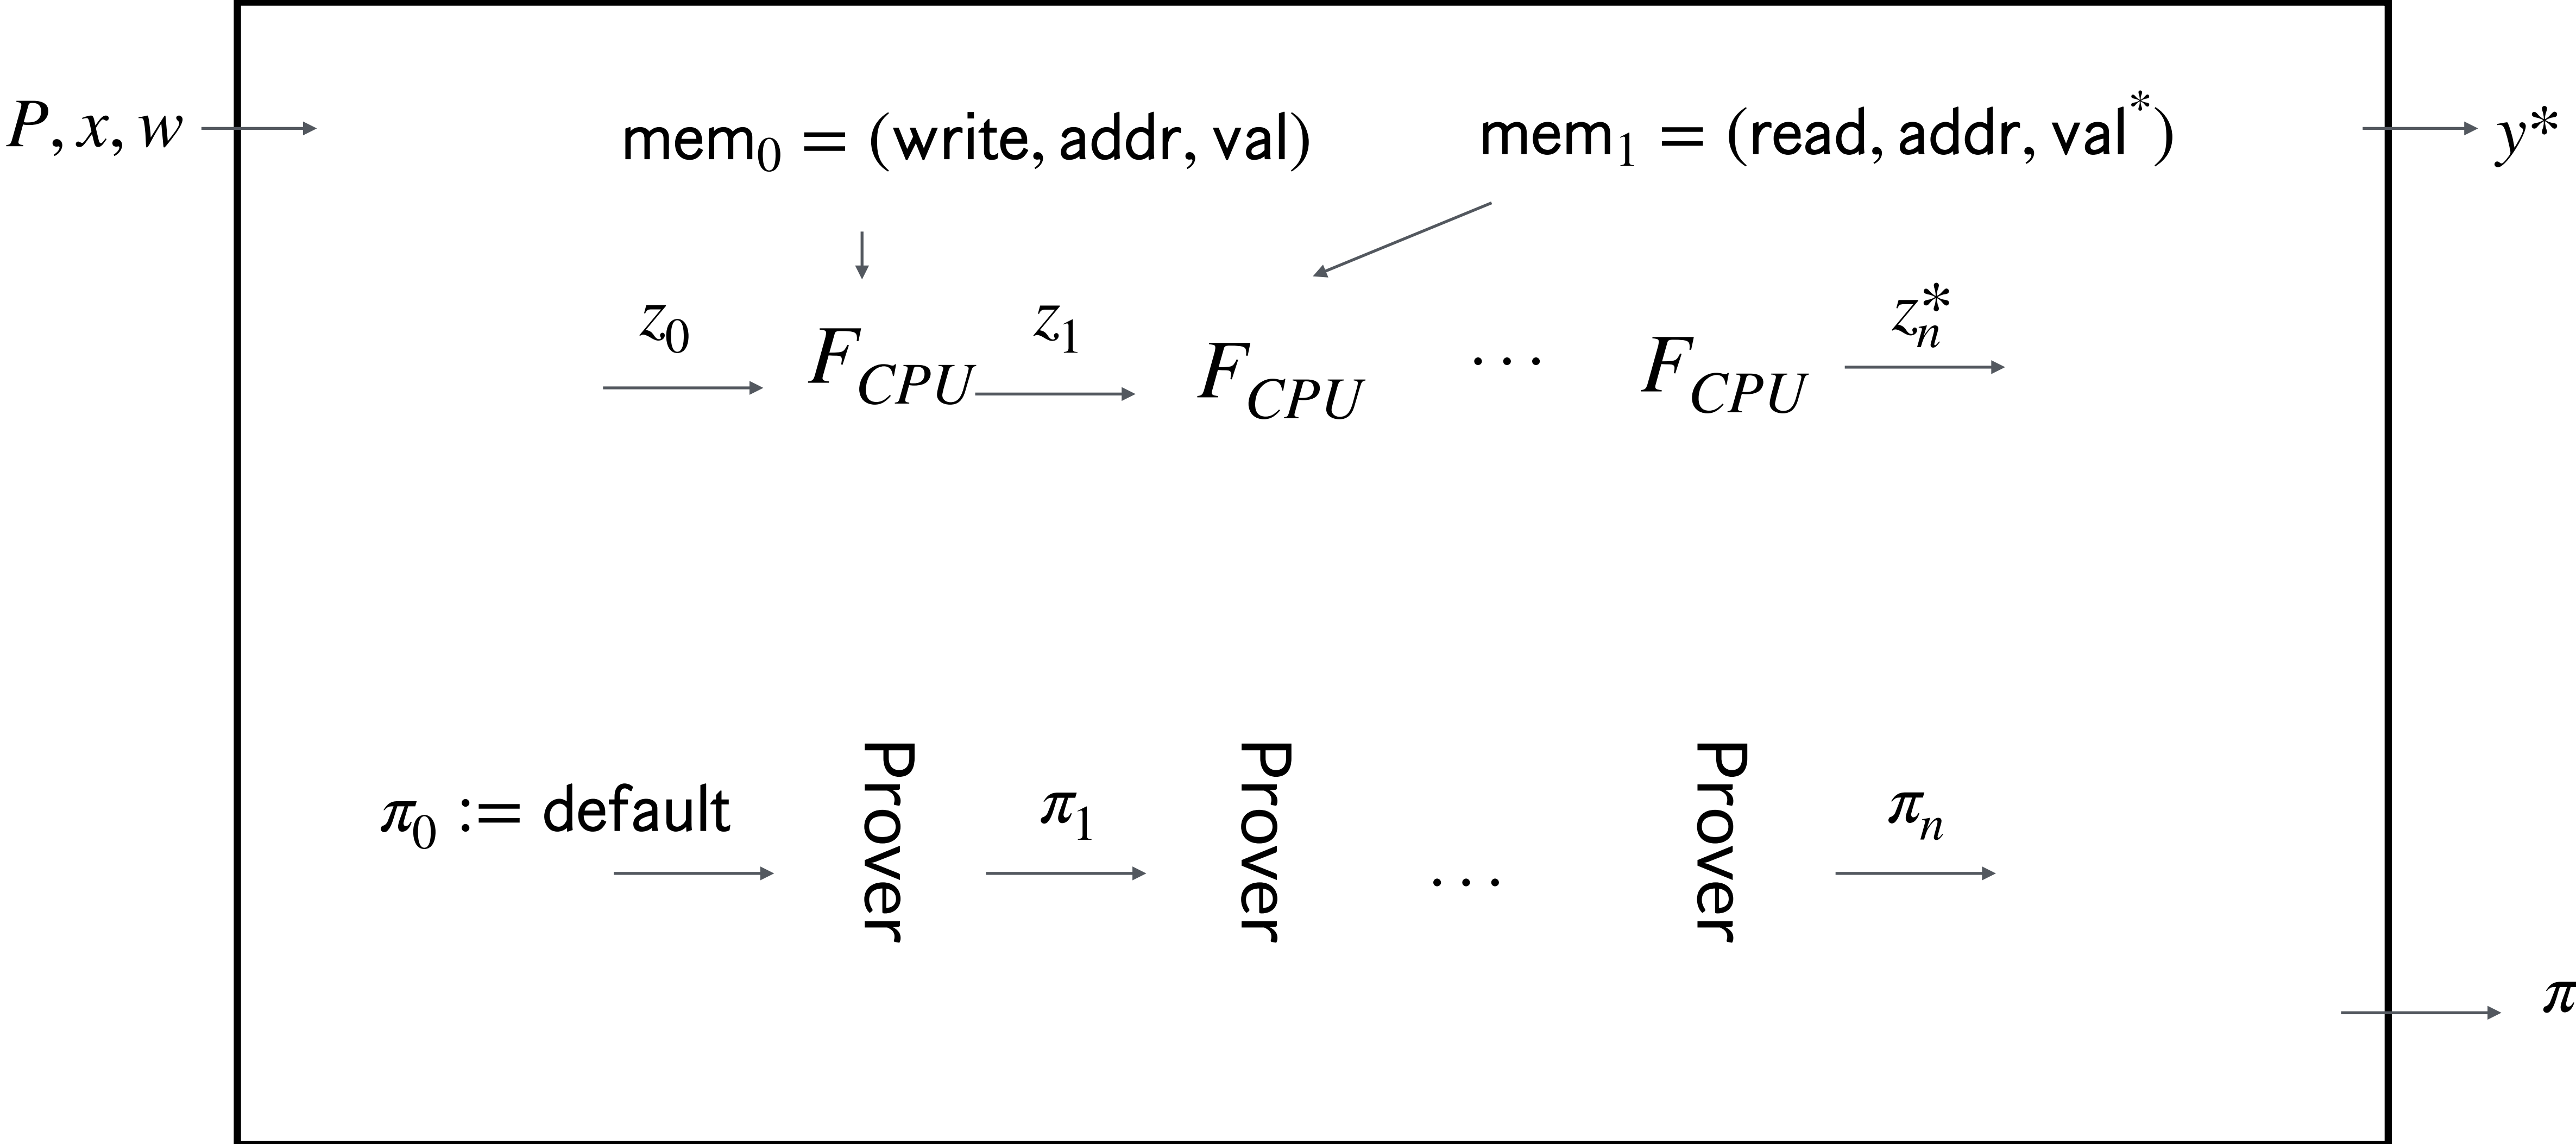


- Prover cost: only pay for constraints on CPU state changes such as program counter and registers, and the *actual* changes to memory :-)

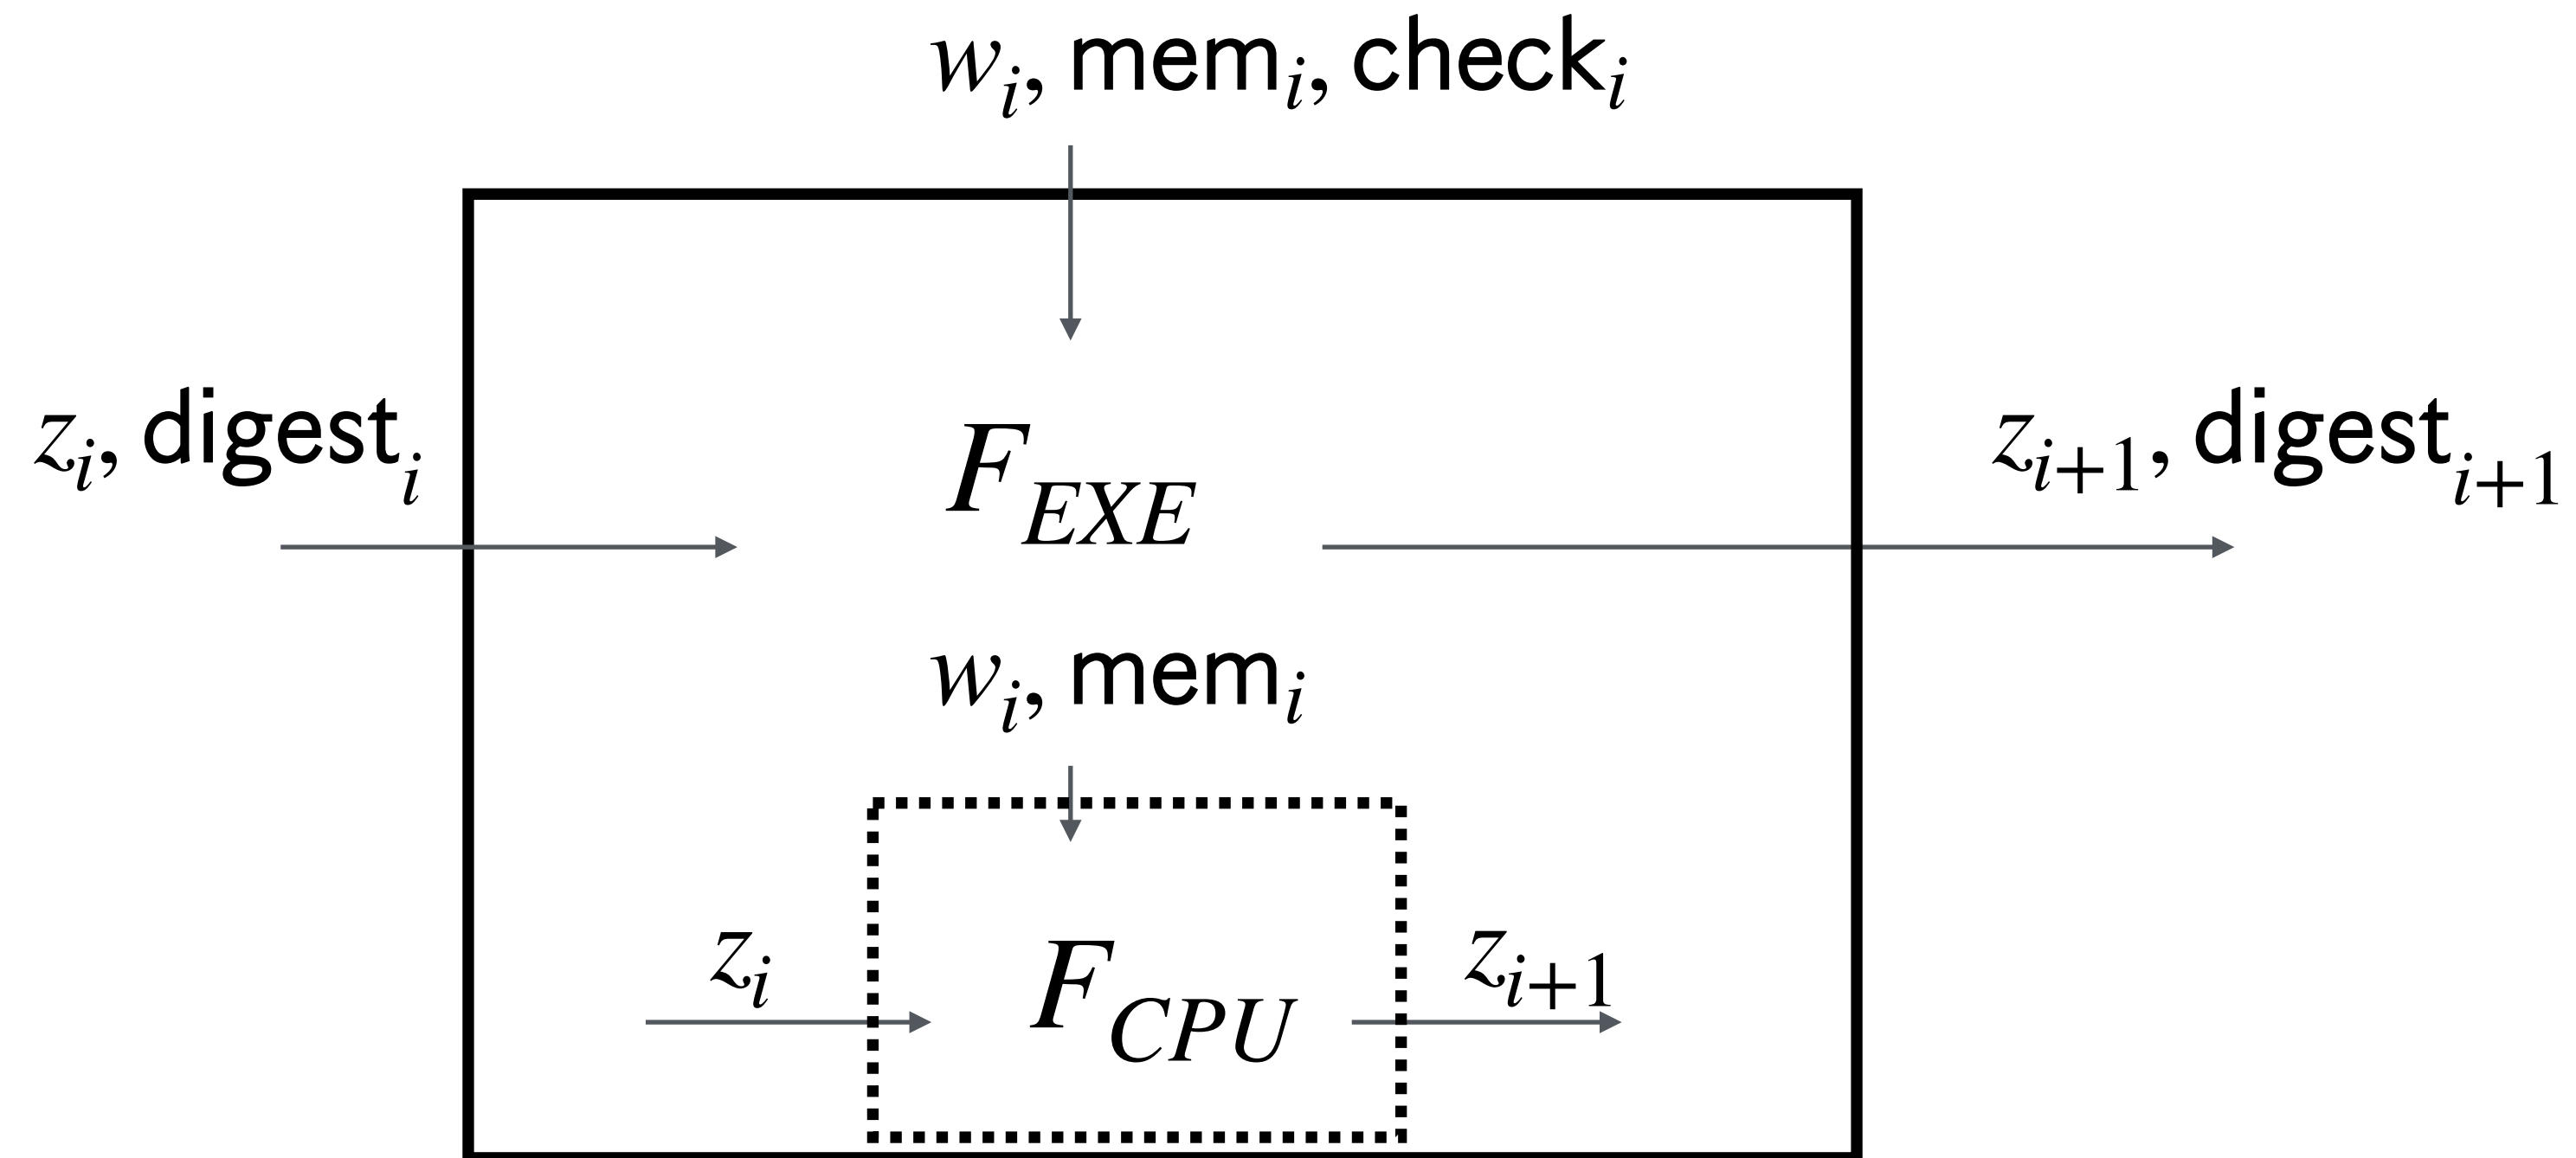
Strawman zkVM with externalized memory



Cheating zkVM prover exploiting inconsistent memory

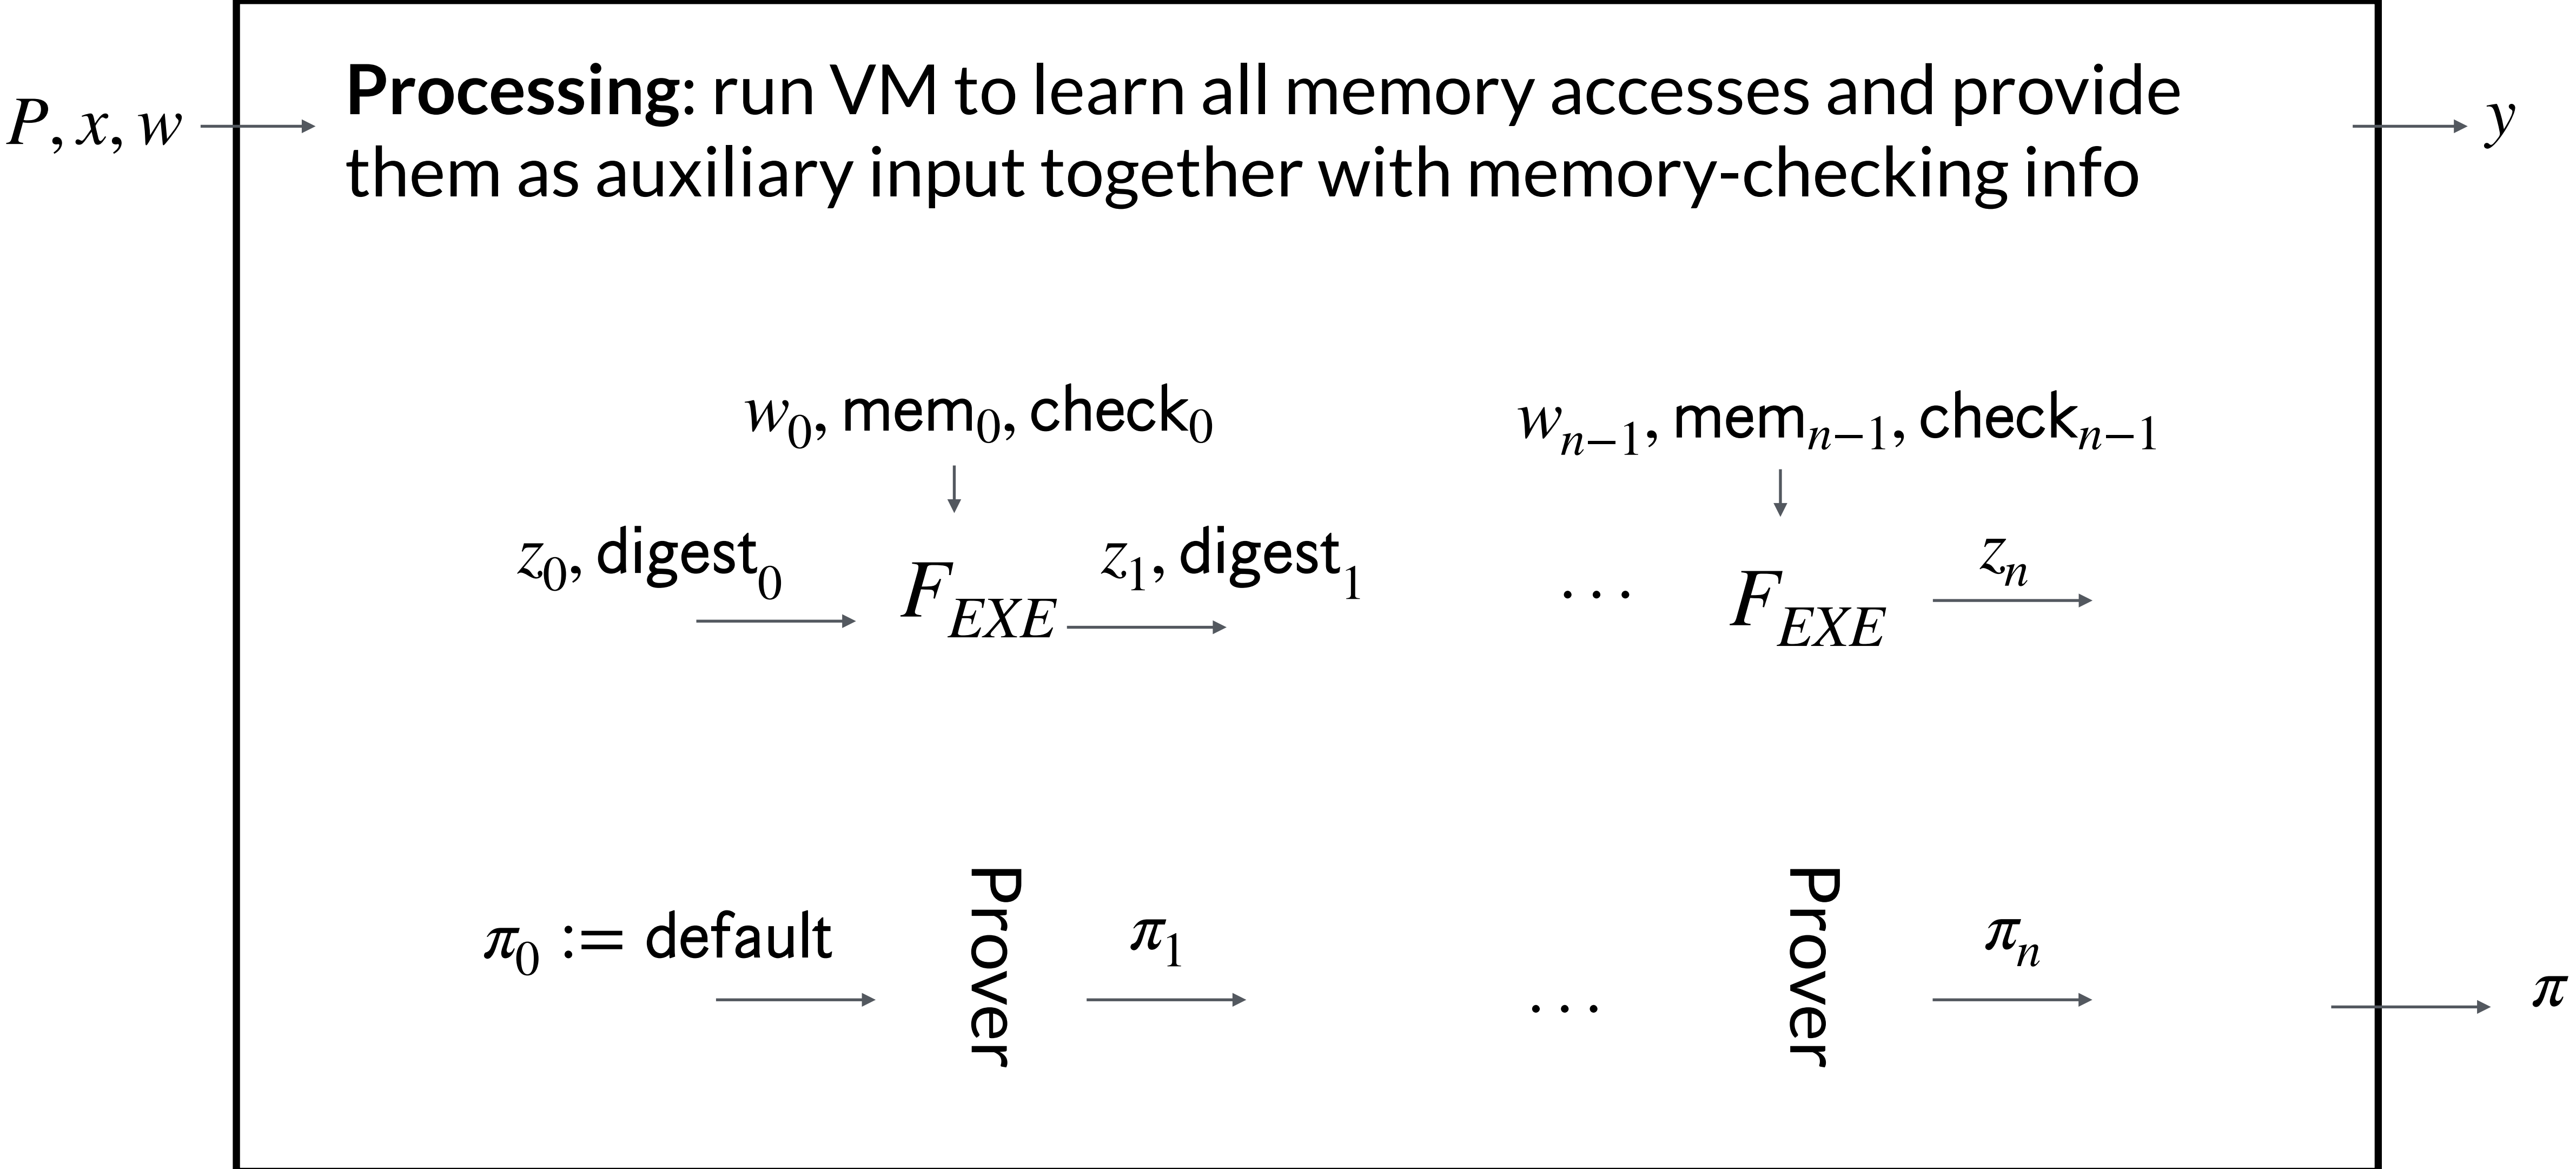


Memory protection for the CPU




F_{EXE} uses digest_i , check_i to check the memory accesses in mem_i are consistent with a global memory access pattern across the IVC computation. Then calls F_{CPU} as a subroutine to get z_{i+1}

zkVM with externalized memory



Memory checking via merkle trees of the memory



Proving a Merkle-path costs $O(\log M)$ hashes per access :-(

Keeping track of memory accesses

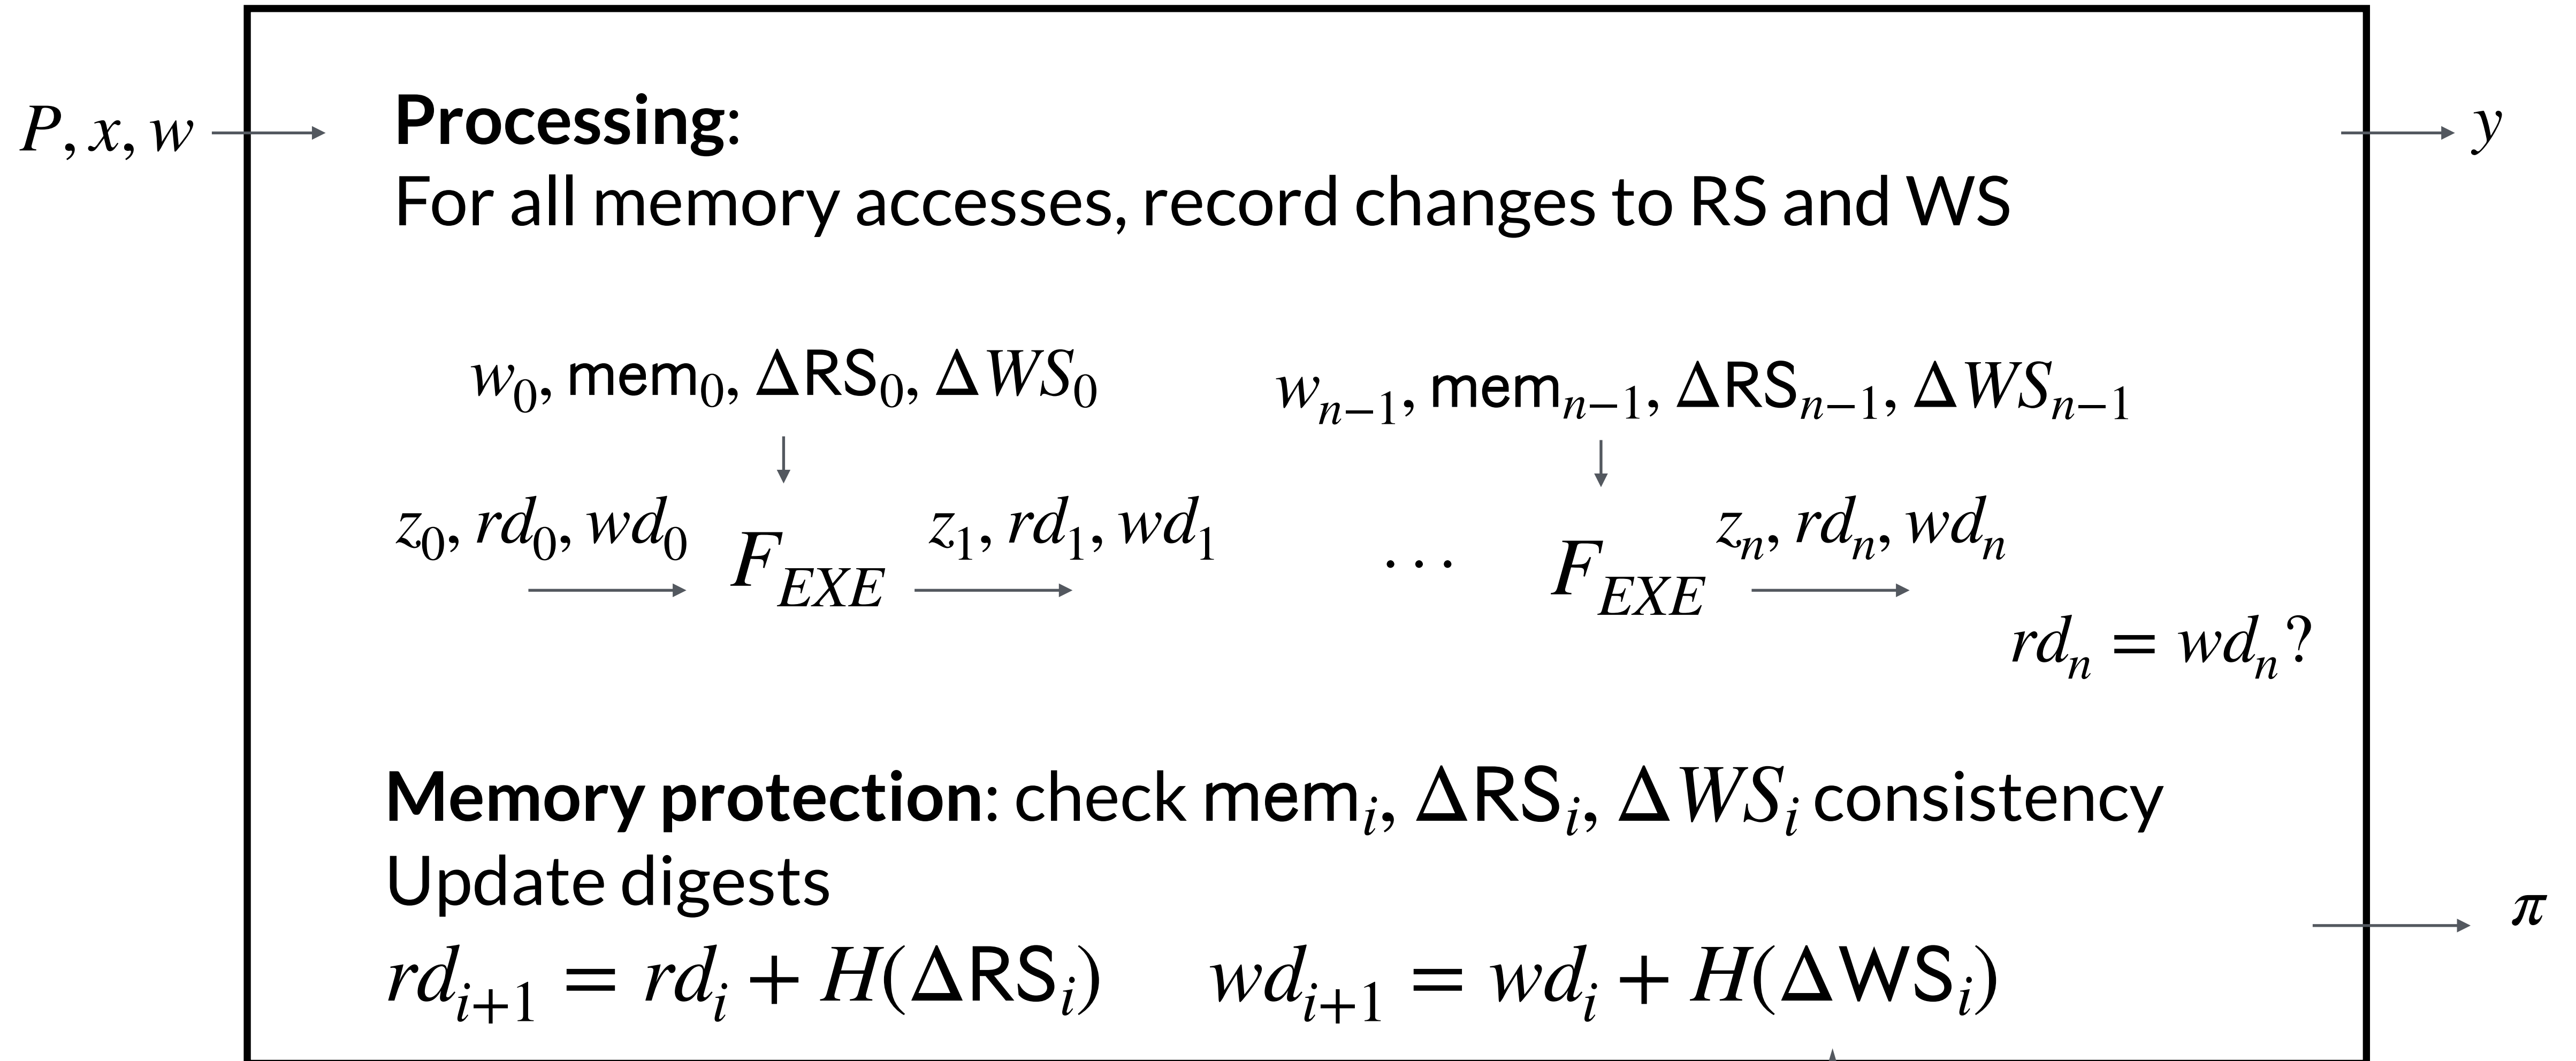
- Read and write set RS, WS used to track memory accesses
- **Invariant:** RS trails WS by the last values written to memory
- Initialize *WS* with tuples $(a, v, t = 0)$ // starting memory
- To read/write value v from/to address a at time t
 - find $(a, v_{\text{old}}, t_{\text{old}})$ in $WS \setminus RS$ // if a read, check $v = v_{\text{old}}$
 - add $(a, v_{\text{old}}, t_{\text{old}})$ to RS // now RS has caught up with WS in address a
 - add (a, v, t) to WS // now RS trails WS in address a again
- When the program is done, add for each address one tuple $(a, v_{\text{old}}, t_{\text{old}})$ to RS
- **Theorem:** memory is consistent if $RS \equiv WS$

Multiset hashing

- \mathbb{G} (additive) group where dlog is hard
- H maps a multiset S (of elements in a domain D) to an element in \mathbb{G}
- Homomorphic: $H(S_1) + H(S_2) = H(S_1 \cup S_2)$
- Collision resistant: infeasible to find $S_1 \neq S_2$ s.t. $H(S_1) = H(S_2)$
- Instantiation: Let $h : D \rightarrow \mathbb{G}$ be a hash function (modeled as a random oracle) and define

$$H(s_1, \dots, s_n) = h(s_1) + \dots + h(s_n)$$

Memory checking via multi-set hashing



Proving 2 hashes per access is less expensive than a Merkle path

Should memory checking be expensive in IVC?

- Monolithic SNARKs pay $O(1)$ field op constraints per memory access
- Recall we can ensure memory is consistent by tracking reads and writes
Write set WS tracks writes (a, v, t)
The read set RS tracks reads (a, v, t)
We access memory by adding $(a, v_{\text{old}}, t_{\text{old}})$ to RS and (a, v, t) to WS
We always need $t_{\text{old}} < t$ and when reading $v_{\text{old}} = v$
- After a final read of the memory, we're happy if $RS = WS$
- So all we need is an efficient check that $RS = WS$

Memory checking via LogUp arguments [EagenKRN22, Haböck22]

- RS and WS consist of tuples (a, v, t) appearing in different orders in the execution
- Given a map from tuples (a, v, t) to field elements in \mathbb{F} , the RS tuples are represented as $r_1, \dots, r_n \in \mathbb{F}$ and the WS tuples as $w_1, \dots, w_n \in \mathbb{F}$
- **Theorem:** r_1, \dots, r_n is a permutation of w_1, \dots, w_n if
$$\sum \frac{1}{r_i + H} = \sum \frac{1}{w_j + H}$$
- Monolithic SNARKs use a Fiat-Shamir challenge $h \leftarrow \mathbb{F}$ to test this equality


A suitable challenge

- We would like to check a LogUp test

$$\sum_i \frac{1}{r_i + H} = \sum_j \frac{1}{w_j + H} \text{ by plugging in a pseudorandom challenge } h$$

- Problem 1
 - Generate h early: the prover can cheat by choosing r_i, w_j that depend on h
 - Generate h late: now the IVC is concluded and we missed our chance to use it
- Problem 2
 - The IVC is incremental, no step has the full memory view to compute the sums


Recall we have split the step function into memory checking and CPU checking



The execution trace is independent of the challenge h . It contains CPU state and memory accesses (a, v, t) that can be determined by what F_{CPU} sees.

Incrementally building the challenge [Soukhanov23]

- Preprocessing by the zkVM prover
 - run the VM to learn the CPU's view
 - memory accesses are part of the CPU's view, it knows the tuples (a, v, t)
 - commit to the CPU's view in each step c_0, \dots, c_{n-1}
 - compute the challenge h via a hash chain (starting at some default h_0)



Incrementally verifying the challenge [Soukhanov23]

P, x, w → **Preprocessing:** create commitments c_0, \dots, c_{n-1} to the values seen by each step, based on those commitments compute a hash chain to get h

$$\begin{array}{ccc}
 w_0, \text{mem}_0, \Delta \text{RS}_0, \Delta \text{WS}_0, c_0 & & w_{n-1}, \text{mem}_{n-1}, \Delta \text{RS}_{n-1}, \Delta \text{WS}_{n-1}, c_{n-1} \\
 \downarrow & & \downarrow \\
 z_0, h_0, h \xrightarrow{\quad} F_{EXE} \xrightarrow{\quad} z_1, h_1, h & \cdots & F_{EXE} \xrightarrow{\quad} z_n, h_n, h \quad h_n = h?
 \end{array}$$

F_{EXE} on h_i, c_i computes $h_{i+1} = H(h_i, c_i)$
 During execution it also checks c_i is correct, e.g.,
 all memory accesses appear in the commitment

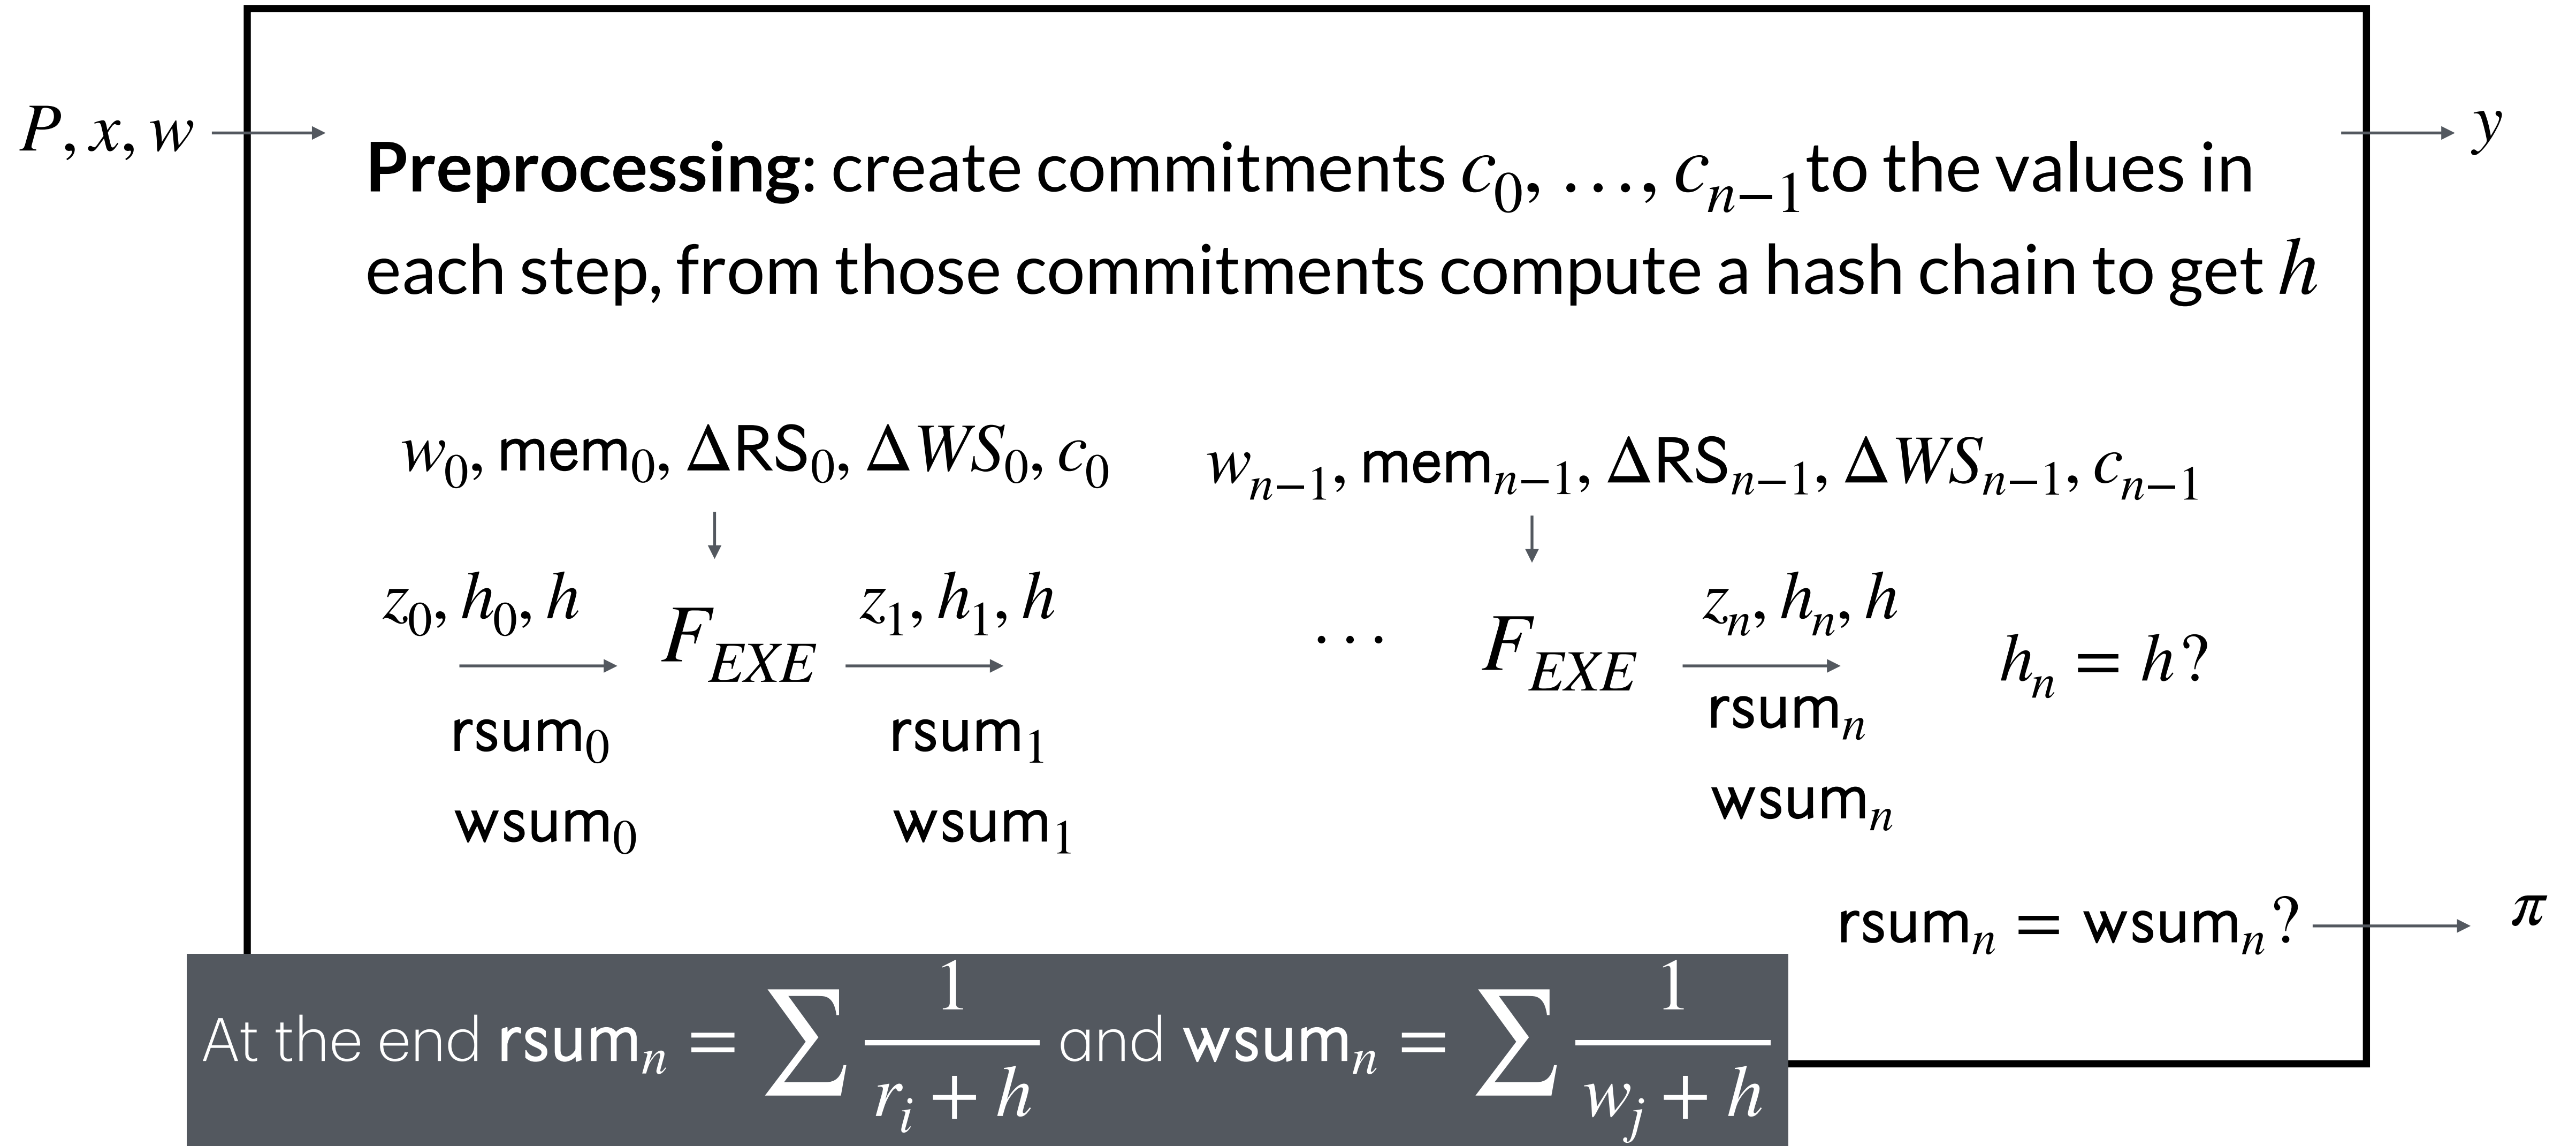
A suitable challenge

- We would like to conduct a LogUp test

$$\sum_i \frac{1}{r_i + H} = \sum_j \frac{1}{w_j + H} \text{ by plugging in a pseudorandom challenge } h$$

- The zkVM uses the precomputed hash-chain challenge h , which it gives as input to the IVC from the start
- The IVC verifies h is correctly computed
- The zkVM prover also keeps track of partial LogUp sums for the read and write sets. At each step it updates the partial LogUp sums according to the memory accesses in this step

Checking the LogUp equality



Comparing the memory-checking techniques

Technique	Efficiency (computation to verify per memory access)	Incrementality (how IVC-like)
Merkle	$O(\log M)$ hashes	High. Can get proof at any step and keep going
Multiset hash	$O(C)$ hashes	Medium. Can keep going and get proof at any step at the cost of finalizing the memory
2-pass offline memory checking	$O(1)$ field ops	Low. The first pass dictates the bound on the computation, you cannot get intermediate proofs.

Nexus - nexus.xyz

- Nexus 1.0
 - VM inspired by RISC-V
 - Can compile Rust to VM
 - Open source on GitHub, MIT and Apache licensed
 - Not so fast yet: 100 proved CPU cycles/second
- Future
 - 2024 goal: 1T proved CPU cycles/second
 - Path to get there
 - a) fast core prover (single threaded)
 - b) prover network (massive parallelization)

Folding

Instance aggregation (prover and verifier)



Witness aggregation (prover)



Relaxed R1CS

Notation: $\overline{m} \in \mathbb{G}$ is a Pedersen commitment (non-randomized) to $m \in \mathbb{F}^n$

Setup: $A, B, C \in \mathbb{F}^{m \times n}$


Instance: $(\overline{w}, x, r, \overline{e}) \in \mathbb{G} \times \mathbb{F}^\ell \times \mathbb{F} \times \mathbb{G}$

Witness: $(w, e) \in \mathbb{F}^{n-1-\ell} \times \mathbb{F}^m$ openings of commitments such that
 $z = (x, r, w)$ satisfies $Az \circ Bz = rCz + e$

Note: R1CS (non-relaxed) has $r = 1$ and $\overline{e} = \overline{0}$

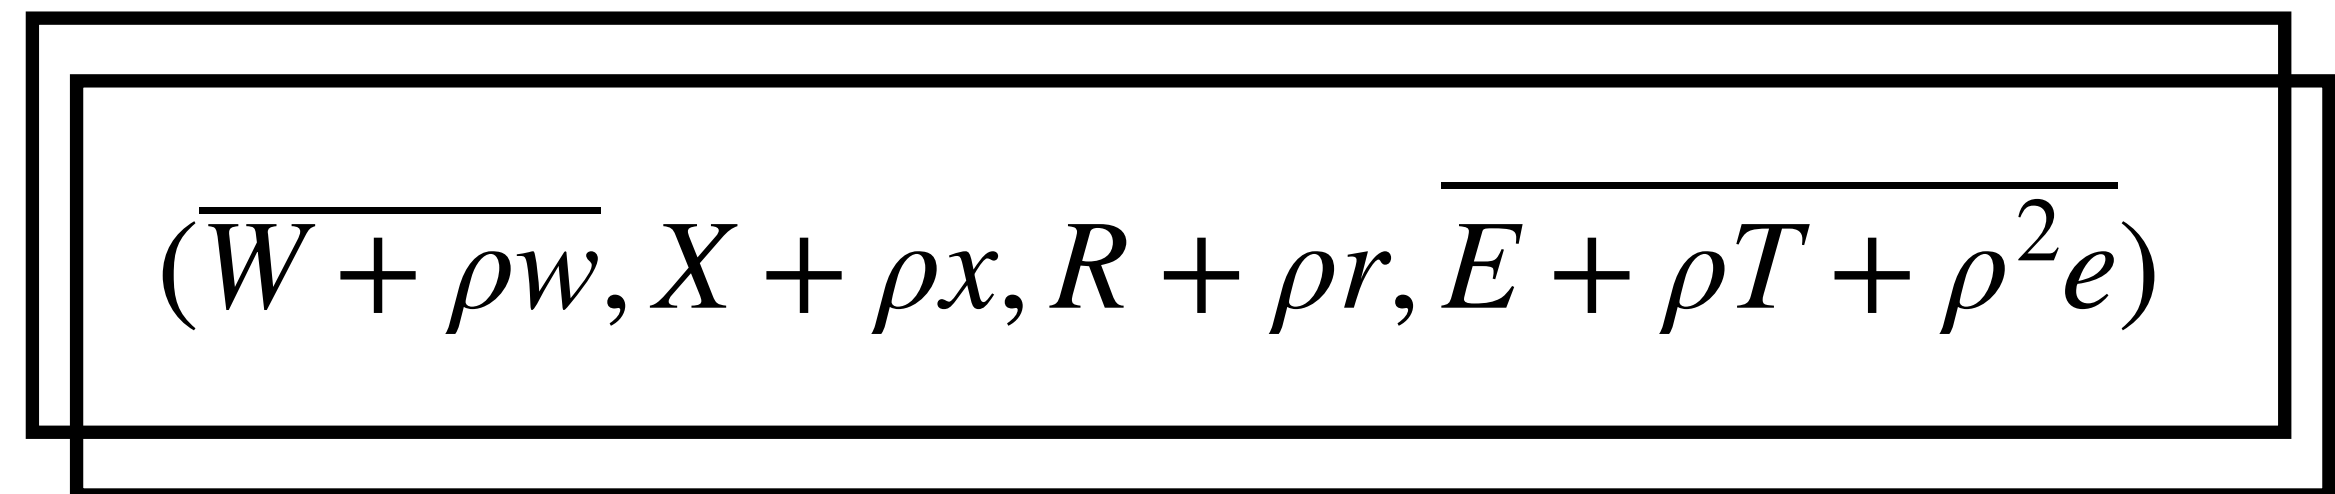
Example: Nova folding scheme

Instance aggregation




Prover and verifier

- P: Let $Z = (X, R, W)$, $z = (x, r, w)$
- P: Let $\bar{T} = \overline{AZ \circ Bz + Az \circ BZ - RCz - rCZ}$
- P&V: Fiat-Shamir challenge ρ



Witness aggregation



NEXUS

Nova modified to split pre-challenge values and post-challenge-values


Commitments to values verified in Nova are split into two Pedersen commitments, one to values known during preprocessing before h is known and before IVC has started.

The corresponding RR1CS variant is:

Instance: $(\bar{w}_{\text{pre}}, \bar{w}_{\text{post}}, x, r, \bar{e}) \in \mathbb{G}^2 \times \mathbb{F}^\ell \times \mathbb{F} \times \mathbb{G}$


Witness: $(w_{\text{pre}}, w_{\text{post}}, e) \in \mathbb{F}^{n-1-\ell} \times \mathbb{F}^m$ openings of commitments such that $z = (x, r, w_{\text{pre}}, w_{\text{post}})$ satisfies $Az \circ Bz = rCz + e$

Incrementally verifiable computation based on folding - intuition




- u_i R1CS instance claiming the current step i gives the correct z_i
- U_i RR1CS instance containing aggregated folding of u_1, u_2, \dots, u_{i-1}
- w_i witness for correct computation in step i
- W_i aggregated RR1CS witness for U_i

Incrementally verifiable computation based on folding - intuition



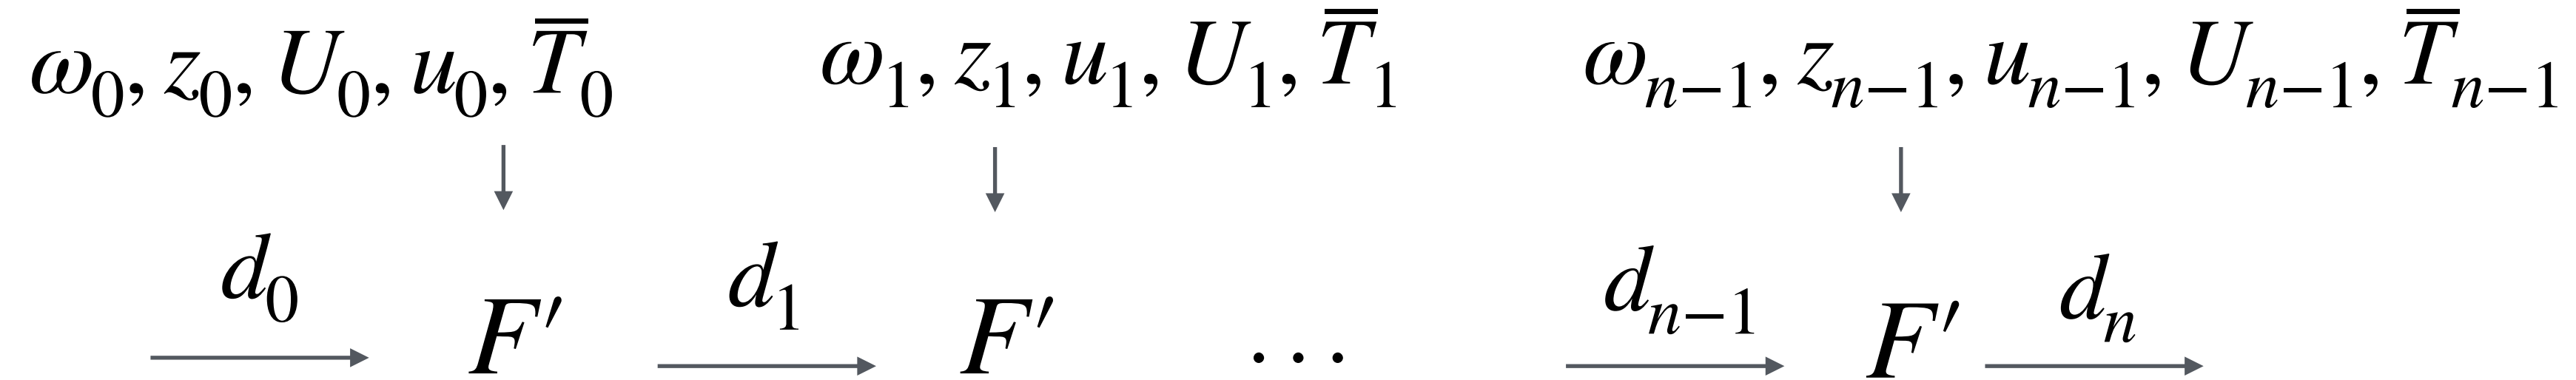
Does not work, the prover could cheat with U_n

Incrementally verifiable computation based on folding - augmented step function



Better, but still sweeping things under the rug

Incrementally verifiable computation based on Nova folding



- u_i R1CS instance claiming the current step i is correctly computed and that U_i was correctly computed as a fold of U_{i-1}, u_{i-1}
- U_i RR1CS instance containing aggregated folding of u_1, u_2, \dots, u_{i-1}
- d_i A digest (hash) of z_i, U_i // compact state, larger auxiliary input
- F' Augmented step function that unpacks digest, folds u_{i-1}, U_{i-1} to get U_i , computes $z_i = F(z_{i-1}, \omega_{i-1})$, and the digest d_i of z_i, U_i