



Ligero

Ligetron and the Llama Inference

Muthu Venkitasubramaniam

CEO, Ligero Inc.

Georgetown University

Joint work with Ruihan Wang (Ligero Inc.) and Carmit Hazay (Ligero Inc., Bar-Ilan University)

This research was developed with funding from the Defense Advanced Research Projects Agency (DARPA). The views, opinions and/or findings expressed are those of the author and should not be interpreted as representing the official views or policies of the Department of Defense or the U.S. Government.
Distribution Statement A. Approved for public release. Distribution Unlimited.

Thanks
Stealth Software Tech.

ZK evolution

- **At the Beginning:** Reduce to an NP-complete problem (eg, Graph Hamiltonicity, Graph 3-Coloring)
- **Then:** Boolean and Arithmetic circuits
- **Since Blockchains:** R1CS, Circom, Cairo, Gnark, etc...
- **Today:** zkVMs

Instrumenting ZK today - zkVMs


1. Code application in a high-level language (C,C++,Rust)
2. Compile down to popular VMs. Eg, RISC-V, WASM
3. Prove correct execution of VM

Challenges in Scaling ZK

First, it was (circuit) representation

Then, optimizing prover running time

Today, optimize prover memory



Main Question

Do there exist **time and space preserving** ZK-SNARKs from minimal assumptions?

YES* - Based on hash functions [BBHV 2022]

Prover time $\tilde{O}(T(n))$ and space $\tilde{O}(S(n))$. Proof Length $\tilde{O}(T(n)/S(n))$


Today: Concretely Efficient!

Introducing Ligetron

A Time and Space Efficient ZK-SNARK

Key Ingredients

1. WASM as an intermediate representation – **Ligetron is a zkWASM**
2. A **space-efficient** variant of the Ligerio ZKP [Ames et al, 2017]



Ligetron

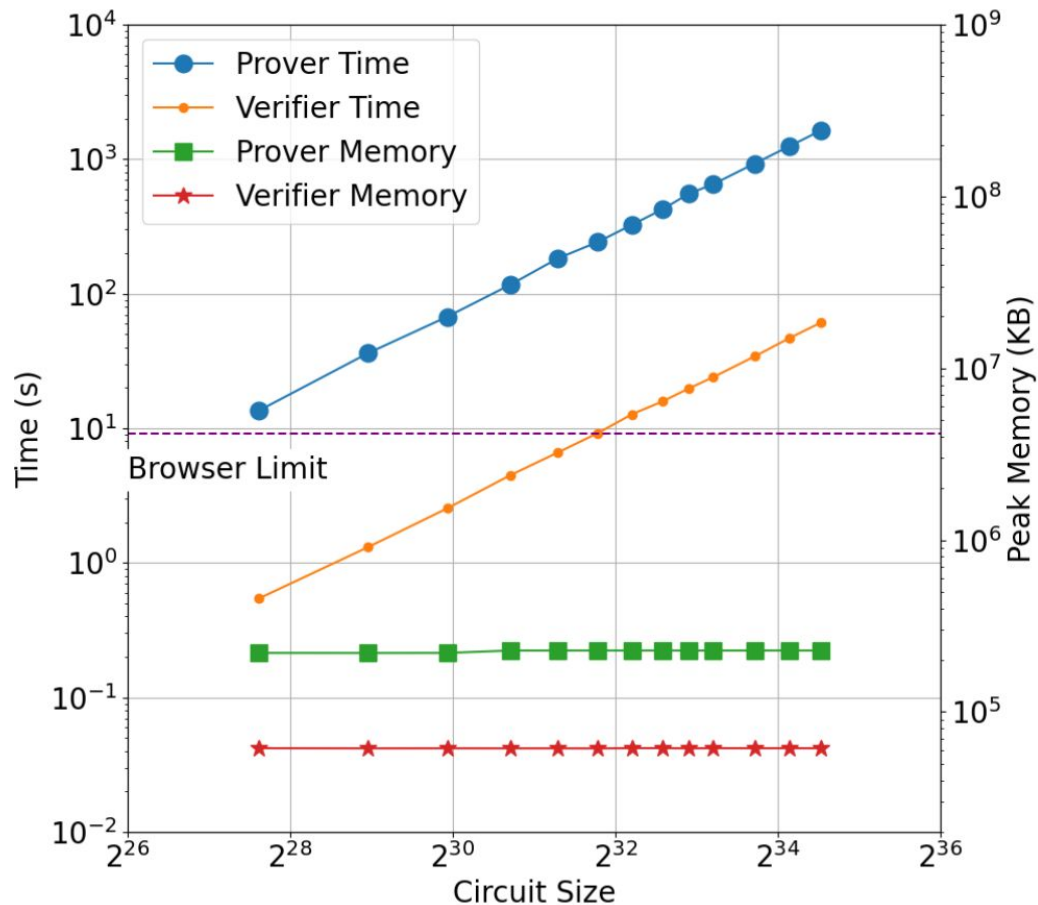
Non-interactive	Yes
Succinct proof	Sublinear
Succinct verification	No*
Time efficient	Yes (quasilinear)
Space efficient	Yes (linear)
Implementable	Yes (runs from a browser)

Why WASM?

- Space/time efficiency from operational semantics of WASM (stacks, blocks)
- Few set of instructions to compile
- Numerous compilers from high-level languages to WASM (C,C++,Rust,etc)
- Clean sandbox (no I/Os, no system calls, no malloc*)
- Extern function calls for gadgets / pre-compiles

Performance

- ❖ Prover time: 500 ns/g
- ❖ Verifier time: 250 ns/g
- ❖ Prover memory: < 100 MB
- ❖ Verifier memory: < 10 MB



Performance

	Prover Time Batched	Verifier Time Batched	Prover Time Random	Verifier Time Random
Browser	1.2 us/g	20 ns/g	5 us/g	2.5 us/g
Macbook	65 ns/g	3 ns/g	500 ns/g	250 ns/g

Where we are?

Proof Length

- **Unstructured circuit:** Square root circuit size. But, parameterizable
- **Structured circuit (M copies of size T circuit):** $O(M+k.T)$

Verification

- **Unstructured circuit:** Quasilinear in circuit size
- **Structured circuit:** $O(M+k.T)$

Extent of WASM integration

- All 32-bit and 64-bit integer operations
- Oblivious control flow
- More recently, RAM, 32-bit floating operations

Are ZKVMs Scalable?

Jolt

Risco

SP1

VM	<input type="radio"/> RISC-V	<input type="radio"/> RISC-V	<input type="radio"/> RISC-V
Prover Speed	<input type="radio"/> 150 kHz int32	<input type="radio"/> 25 kHz int32	<input type="radio"/> 40-150 kHz int32
Succinct Proof/Verification	<input type="radio"/> Yes/Yes	<input type="radio"/> Yes/Yes	<input type="radio"/> Yes/Yes
Hardware	<input type="radio"/> 64 Core 512 GB RAM	<input type="radio"/> 64 Core 512 GB RAM	<input type="radio"/> 64 Core 512 GB RAM
Memory Efficient	<input type="radio"/> No	<input type="radio"/> No	<input type="radio"/> No

Are ZKVMs Scalable?

Current Techniques for Memory Efficiency

- Recursive Composition
 - Incrementally Verifiable Computation (IVC)
1. Break computation into small units
 2. Prove units in distributed manner
 3. Compose

Bottlenecks

- Need a lot of hardware (each unit needs to be run in parallel)
- Gadgets/pre-compiles are harder in RISC-V
- zkVMs today are really succinctVMs

Benchmark 1 (Structured Circuits): Rollups

Simple payment circuit

- EdDSA signatures
- Poseidon hash function for Merkle Trees



Starknet: 200 nodes with 400MB RAM
Benchmark: 2000 tps*



Ligetron: g5.xlarge (single A10 GPU)
Benchmark: 500tps

WASM extern for 256-bit

```
160 struct fp256_class {
161     fp256_class() { _fp256_init(data_); }
162
163     fp256_class(int i) { _fp256_init(data_); _fp256_set_ui(data_, i); }
164     fp256_class(uint32_t i) { _fp256_init(data_); _fp256_set_ui(data_, i); }
165     fp256_class(const char *str, int base = 10)
166         { _fp256_init(data_); _fp256_set_str(data_, str, base); }
167
168     fp256_class(const fp256_t o)
169         { _fp256_init(data_); _fp256_set_fp256(data_, o); }
170     fp256_class(const fp256_class& o)
171         { _fp256_init(data_); _fp256_set_fp256(data_, o.data_); }
172
173     // fp256_class(fp256_class&& o)
174     //     { _fp256_init(data_); std::swap(*data_, *o.data_); }
175
176     fp256_class& operator=(const fp256_t o)
177         { _fp256_set_fp256(data_, o); return *this; }
178     fp256_class& operator=(const fp256_class& o)
179         { _fp256_set_fp256(data_, o.data_); return *this; }
180     // fp256_class& operator=(fp256_class&& o)
181     //     { std::swap(*data_, *o.data_); return *this; }
182
183     ~fp256_class() { _fp256_clear(data_); }
184
185     __fp256_backend* data() { return data_; }
186     const __fp256_backend* data() const { return data_; }
187
188 protected:
189     fp256_t data_;
190 };
191
```


WASM extern for 256-bit

Implementations:

- EdDSA signatures
- Poseidon hash function

```
196 // Vector Initialization
197 // -----
198 WASM_EXTERN(ligetronebatch, fp256vec_get_size)
199 uint64_t _fp256vec_get_size();
200
201 WASM_EXTERN(ligetronebatch, fp256vec_init)
202 void _fp256vec_init(fp256vec_t v);
203
204 WASM_EXTERN(ligetronebatch, fp256vec_clear)
205 void _fp256vec_clear(fp256vec_t v);
206
207 WASM_EXTERN(ligetronebatch, fp256vec_set_ui)
208 void _fp256vec_set_ui(fp256vec_t v, uint32_t* num, uint64_t len);
209
210 WASM_EXTERN(ligetronebatch, fp256vec_set_ui_scalar)
211 void _fp256vec_set_ui(fp256vec_t v, uint32_t num);
212
213 WASM_EXTERN(ligetronebatch, fp256vec_set_str)
214 int _fp256vec_set_str(fp256vec_t v, const char *str[], uint64_t len, int base = 0);
215
216 WASM_EXTERN(ligetronebatch, fp256vec_set_str_scalar)
217 int _fp256vec_set_str(fp256vec_t v, const char *str, int base = 0);
218
219 WASM_EXTERN(ligetronebatch, fp256vec_copy)
220 void _fp256vec_copy(fp256vec_t out, const fp256vec_t in);
221
222 WASM_EXTERN(ligetronebatch, fp256vec_print)
223 void _fp256vec_print(const fp256vec_t v);
```



Benchmark 2 (“Unstructured”): LLM Inference



Modulus Labs: 200+ hours for verifying GPT2-XL 1.5B
using 128-core 1TB CPU with 10TB disk space



Ligetron: 14 hours for verifying the Llama 7B
using 8-core CPU with peak memory 10GB*

*In collaboration with  **nim**

Our Approach for Llama inference



Pure C implementation of Llama 7B

<https://github.com/karpathy/llama2.c>

Main Challenge
floating point ops

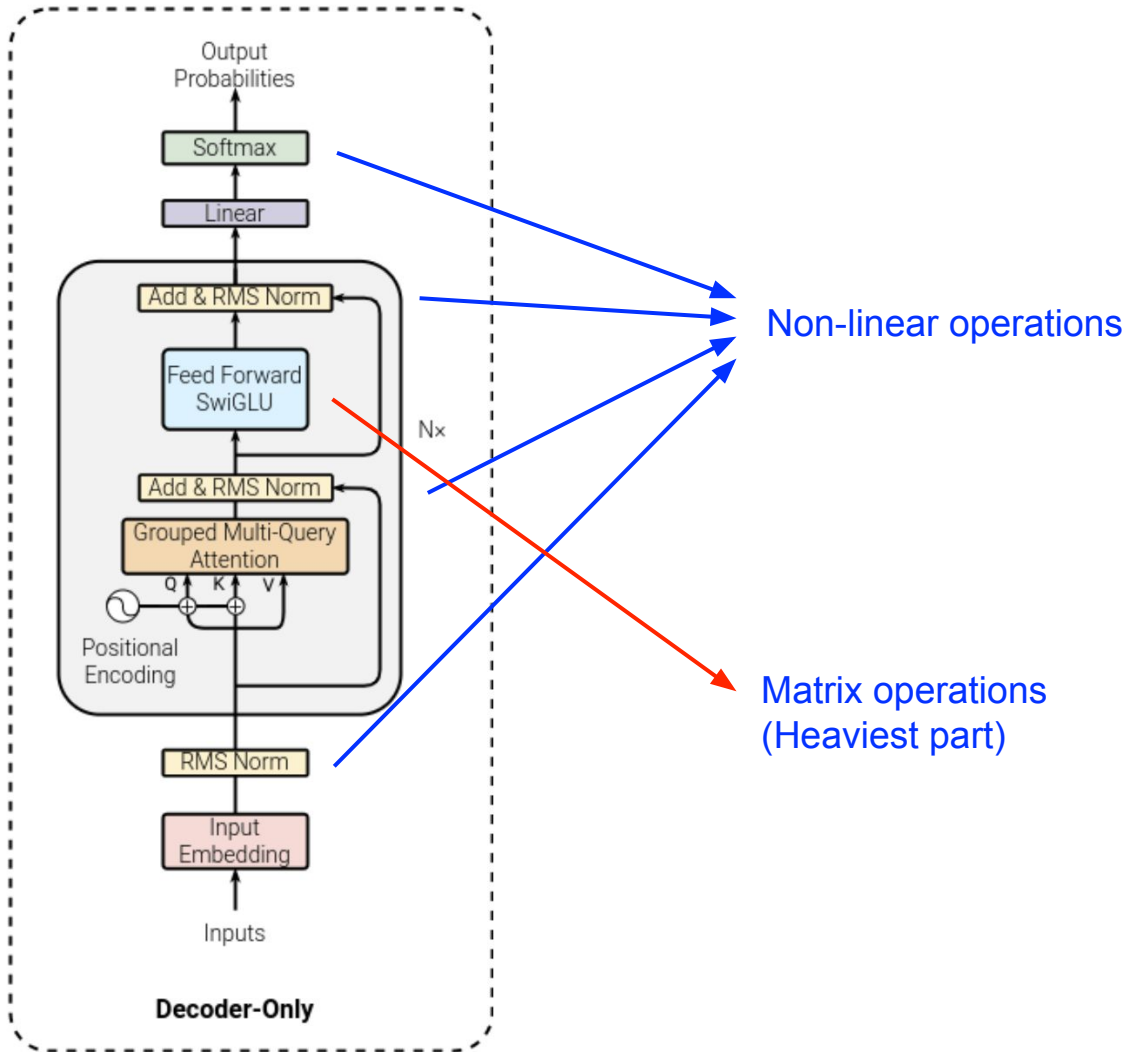
llama2.c



Have you ever wanted to inference a baby [Llama 2](#) model in pure C? No? Well, now you can!

Train the Llama 2 LLM architecture in PyTorch then inference it with one simple 700-line C file ([run.c](#)). You might think that you need many billion parameter LLMs to do anything useful, but in fact very small LLMs can have surprisingly strong performance if you make the domain narrow enough (ref: [TinyStories](#) paper). This repo is a "fullstack" train + inference solution for Llama 2 LLM, with focus on minimalism and simplicity.

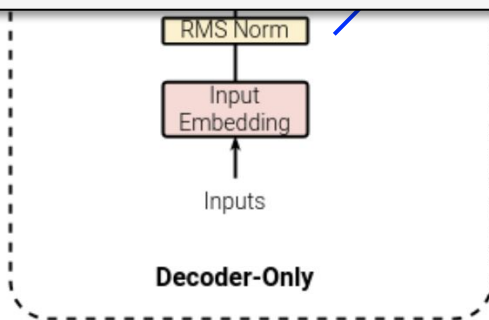

Llama 7B



Llama 7B

```
/* minimax approximation to sin on [-pi/4, pi/4] with rel. err. ~= 5.5e-12 */
double sin_core (double x)
{
    double x4, x2, t;
    x2 = x * x;
    x4 = x2 * x2;
    /* evaluate polynomial using a mix of Estrin's and Horner's scheme */
    return ((2.7181216275479732e-6 * x2 - 1.9839312269456257e-4) * x4 +
            (8.3333293048425631e-3 * x2 - 1.6666666640797048e-1)) * x2 * x + x;
}
```

(eg, sin, exp, etc)



Llama 7B

Operation distribution

fp32 add (100 million)

~1000 constraints

fp32 mul (200 million)

~200 constraints


Total constraints

~130 billion add

~78 billion mul

```
Num Linear constraints:      125591511068
Num quadratic constraints:   78174961265
Num quadratic constraints (padded): 78174968624
-----
f32.fnn_le: 32001
f32.convert_i32_u: 1
f32.fnn_ne: 1
f32.convert_i32_s: 103235588
f32.const : 6875505
i32.reinterpret_f32: 749634
i32.trunc_f32_u: 385280
f32.fnn_abs: 1499136
i32.trunc_f32_s: 749568
f32.reinterpret_i32: 385346
f32.fnn_div: 1178626
f32.fnn_mul: 209896969
f32.fnn_add: 104927810
f32.fnn_gt: 401281
f32.fnn_lt: 2285694
f32.fnn_sub: 164097
f32.fnn_ge: 385281
```

Benchmark 2 (“Unstructured”): LLM Inf.




Modulus Labs: 200+ hours for verifying GPT2-XL 1.5B
using 128-core 1TB CPU with 10TB disk space



Ligetron: 14 hours for verifying the Llama 7B
using 8-core CPU with peak mem. 10GB

Time to benchmark: 15 days

*In collaboration with  **nim**

Ligetrn 1.0 Compared to Other zkVMs



	<input checked="" type="radio"/> Ligetrn 1.0	<input type="radio"/> Jolt	<input type="radio"/> Risc0	<input type="radio"/> SP1
VM	<input checked="" type="radio"/> WASM	<input type="radio"/> RISC-V	<input type="radio"/> RISC-V	<input type="radio"/> RISC-V
Prover Speed	<input checked="" type="radio"/> 50 kHz int32, int64, fp32	<input type="radio"/> 150 kHz int32	<input type="radio"/> 25 kHz int32	<input type="radio"/> 40-150 kHz int32
Succinct Proof/Verification	<input checked="" type="radio"/> Yes/No	<input type="radio"/> Yes/Yes	<input type="radio"/> Yes/Yes	<input type="radio"/> Yes/Yes
Hardware	<input checked="" type="radio"/> 8 Core 16GM RAM	<input type="radio"/> 64 Core 512 GB RAM	<input type="radio"/> 64 Core 512 GB RAM	<input type="radio"/> 64 Core 512 GB RAM
Memory Efficient	<input checked="" type="radio"/> Yes	<input type="radio"/> No	<input type="radio"/> No	<input type="radio"/> No

Demo

Build a ZK Application



Sample Applications ▾

Compile

Add Input

Prove

Verify

Send Proof

C/C++ Code

WASM

Inputs:

Argument 1:

×

Private input ☐

Build a ZK Application

Edit Distance ▾

Compile

Add Input

Prove

Verify

Send Proof

C/C++ Code

```
/* APPLICATION DESCRIPTION:
```

```
Proves the knowledge that the private string is within the edit distance  
of 3 characters from the public string.
```

```
Arguments:
```

```
Argument 1 - private string
```

```
Argument 2 - public string
```

```
Argument 3 - SHA256 hash of the private string
```

```
*/
```

```
#include "edit_distance.hpp"
```

```
#include "sha256.hpp"
```

```
#include "convert.hpp"
```

WASM

Inputs:

Argument 1: ×Private input ☐

Build a ZK Application

Edit Distance ▾

Compile

Add Input

Prove

Verify

Send Proof

C/C++ Code

/* APPLICATION DESCRIPTION:

Proves the knowledge that the private string is within the edit distance of 3 characters from the public string.

Arguments:

Argument 1 - private string

Argument 2 - public string

Argument 3 - SHA256 hash of the private string

*/

#include "edit_distance.hpp"

#include "sha256.hpp"

#include "convert.hpp"

WASM

Inputs:

Argument 1: ✕

abcd

Private input ☒

Argument 2: ✕

abcdef

Build a ZK Application

Edit Distance ▾

Compile

Add Input

Prove

Verify

Send Proof



C/C++ Code

/* APPLICATION DESCRIPTION:

Proves the knowledge that the private string is within the edit distance of 3 characters from the public string.

Arguments:

Argument 1 - private string

Argument 2 - public string

Argument 3 - SHA256 hash of the private string

*/

#include "edit_distance.hpp"

#include "sha256.hpp"

#include "convert.hpp"

WASM

```
(module
  (type (;0;) (func (param i32)))
  (type (;1;) (func (param i32 i32) (result i32)))
  (type (;2;) (func))
  (type (;3;) (func (result i32)))
  (type (;4;) (func (param i32 i32 i32) (result i32)))
  (type (;5;) (func (param i32) (result i32)))
  (import "wasi_snapshot_preview1" "args_sizes_get" (func (;0;) (type 1)))
  (import "wasi_snapshot_preview1" "args_get" (func (;1;) (type 1)))
  (import "env" "assert_constant" (func (;2;) (type 0)))
  (import "env" "assert_one" (func (;3;) (type 0)))
  (import "wasi_snapshot_preview1" "proc_exit" (func (;4;) (type 0)))
  (func (;5;) (type 2))
  (func (;6;) (type 4) (param i32 i32 i32) (result i32)
    (local i32 i32 i32)
```

Inputs:

Argument 1: ✕

abcd

Private input ☒

Argument 2: ✕

abcdef

Build a ZK Application

Edit Distance ▾

Compile

Add Input

Prove

Verify

Send Proof

C/C++ Code

/* APPLICATION DESCRIPTION:

Proves the knowledge that the private string is within the edit distance of 3 characters from the public string.

Arguments:

- Argument 1 - private string
- Argument 2 - public string
- Argument 3 - SHA256 hash of the private string

*/

```
#include "edit_distance.hpp"
#include "sha256.hpp"
#include "convert.hpp"
```

WASM

```
(module
  (type (;0;) (func (param i32)))
  (type (;1;) (func (param i32 i32) (result i32)))
  (type (;2;) (func))
  (type (;3;) (func (result i32)))
  (type (;4;) (func (param i32 i32 i32) (result i32)))
  (type (;5;) (func (param i32) (result i32)))
  (import "wasi_snapshot_preview1" "args_sizes_get" (func (;0;) (type 1)))
  (import "wasi_snapshot_preview1" "args_get" (func (;1;) (type 1)))
  (import "env" "assert_constant" (func (;2;) (type 0)))
  (import "env" "assert_one" (func (;3;) (type 0)))
  (import "wasi_snapshot_preview1" "proc_exit" (func (;4;) (type 0)))
  (func (;5;) (type 2))
  (func (;6;) (type 4) (param i32 i32 i32) (result i32)
    (local i32 i32 i32)
```

Inputs:

Argument 1:

abcd

Private input ☒

Argument 2:

abcdef

Results:

Prove result: true

Prove execution time: 5.61 s

Circuit size: 425,015

Verify result: true

Verify execution time: 2.71 s

Our Roadmap

1. You don't have succinct verification, duh!

Pre-processing Ligetron for succinct verification

2. But your proofs are not short, duh!

Compose with SNARKs to verify on-chain (Groth16/Halo2)

3. But you can't handle non-oblivious code, duh!

- <https://eprint.iacr.org/2023/1257> (CCS 2023 – Distinguished Paper Award)
- <https://eprint.iacr.org/2024/456> (Precompiles)

4. What about lookup arguments?

We can do that as well



Ligero

Develop on Ligetron:

Build your ZK app at ligetron.com

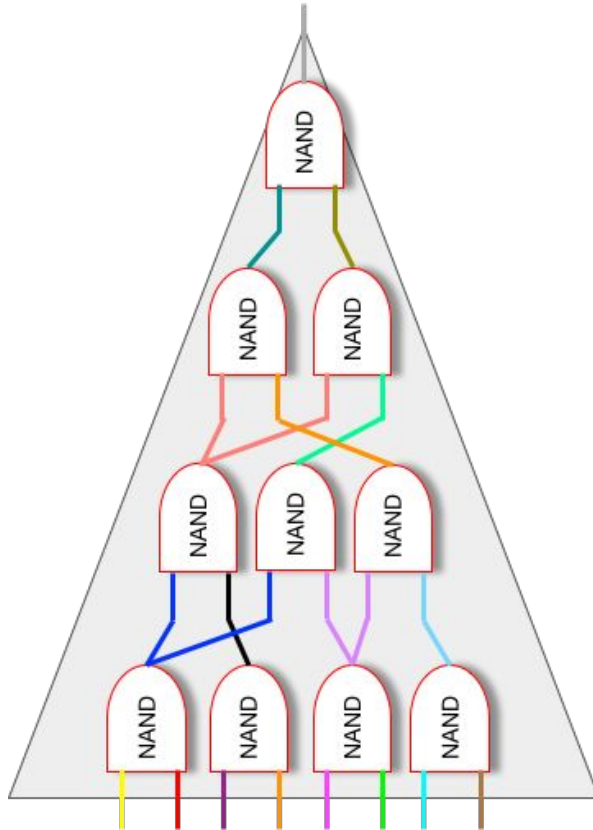
Develop with Ligetron:

We are looking for developers


Please email muthu@ligero-inc.com


Our Techniques

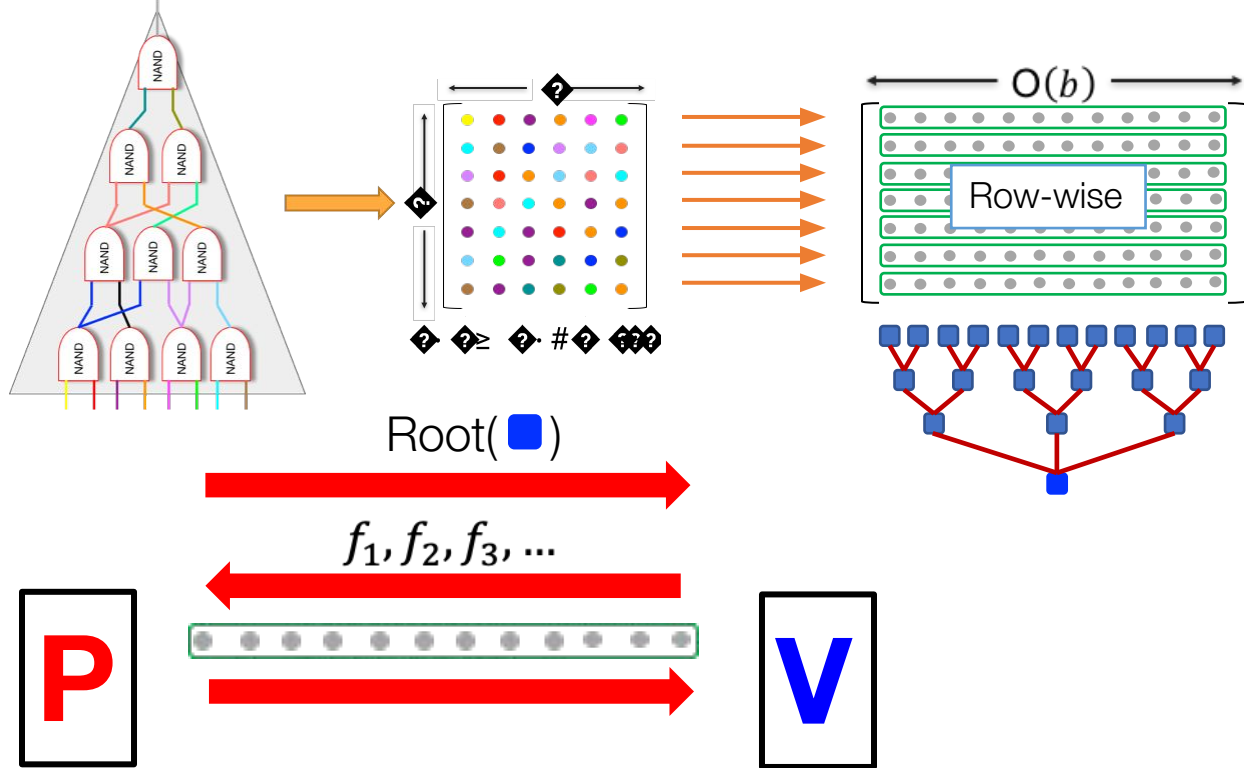
Extended witness \bar{w}




Extended witness \bar{w}













Root(■)





Root(■)





Root(■)

P

V

f_1, f_2, f_3, \dots

i_1, i_2, i_3, \dots

Proof Length:


$O(b + \kappa \cdot a)$

Prover Computation:


$O(a)$ FFTs of $O(b)$

What about row aggregates?

1. Code test – the prover encoded each row correctly
2. Quadratic test – the prover computed multiplication gates correctly
3. Linear test – the prover computed “linear” gates correctly




1. Code test – prover encoded correctly



2. Multiplication gates were correct


Arrange values so that all constraints for multiplication gates are so that $\bar{w}_i \odot \bar{w}_{i+1} = \bar{w}_{i+2}$ for some set of values i



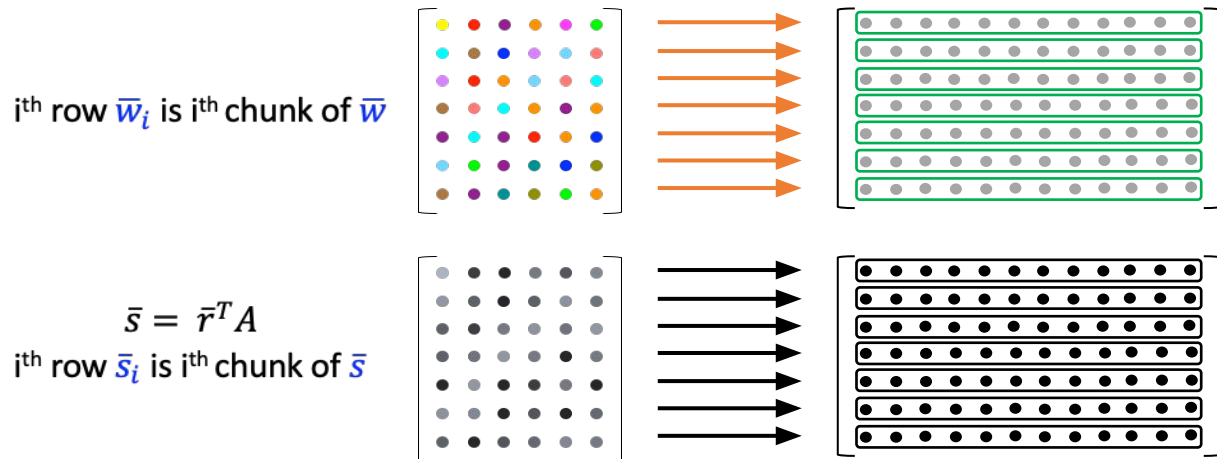
Check if \bar{q} is a valid codeword **encoding** 0s and agrees on revealed columns

2. Linear gates were correct

Linear constraints can be expressed as $A\bar{w} = \bar{b}$




Check if response
encodes $\bar{r}^T \bar{b}$




3. Linear gates were correct

Linear constraints can be expressed as $A\bar{w} = \bar{b}$




Check if response
encodes $\bar{r}^T \bar{b}$






3. Linear gates were correct

Linear constraints can be expressed as $A\bar{w} = \bar{b}$



Check if \bar{q} is a valid
codeword **encoding \bar{z} s.t.**
 $\mathbf{1}^T \bar{z} = \bar{r}^T \bar{b}$
 and agrees on revealed
columns





Root(■)

P

V

f_1, f_2, f_3, \dots

i_1, i_2, i_3, \dots


Proof Length:

$O(b + \kappa \cdot a)$

Prover Computation:

$O(a)$ FFTs of $O(b)$

Set $a = T/S$ and $b = S$



Root(■)

P

V

f_1, f_2, f_3, \dots

i_1, i_2, i_3, \dots


Proof Length:

$\tilde{O}(T/S)$

Prover Computation:

$\tilde{O}(T)$

Set $a = T/S$ and $b = S$



Root(■)

P

V

f_1, f_2, f_3, \dots

i_1, i_2, i_3, \dots


Proof Length:

$O(T/S + \kappa \cdot S)$

Prover Computation:

$O(T/S)$ FFTs of $O(S)$

Set $a = T/S$ and $b = S$



Root(■)

P

V

f_1, f_2, f_3, \dots

i_1, i_2, i_3, \dots

Proof Length:

$\tilde{O}(T/S)$

Prover Computation:

$\tilde{O}(T)$

Set $a = T/S$ and $b = S$

**Build your ZK app
at ligetron.com**



Scan this to verify proof on your device