

Security Bounds for Proof-Carrying Data from Straightline Extractors

Alessandro Chiesa, Ziyi Guan, Shahar Samocha, Eylon Yogev



<https://eprint.iacr.org/2023/1646>



TL;DR

TL;DR

What is *proof-carrying data* (PCD)?

TL;DR

What is *proof-carrying data* (PCD)?

- Recursive compositions of SNARKs.

TL;DR

What is *proof-carrying data* (PCD)?

- Recursive compositions of SNARKs.
- It's useful for efficiently verifying distributed computations.

TL;DR

What is *proof-carrying data* (PCD)?

- Recursive compositions of SNARKs.
- It's useful for efficiently verifying distributed computations.

Problem:

TL;DR

What is *proof-carrying data* (PCD)?

- Recursive compositions of SNARKs.
- It's useful for efficiently verifying distributed computations.

Problem:

- PCD is deployed under the assumption "security of PCD" = "security of underlying SNARK".

TL;DR

What is *proof-carrying data* (PCD)?

- Recursive compositions of SNARKs.
- It's useful for efficiently verifying distributed computations.

Problem:

- PCD is deployed under the assumption "security of PCD" = "security of underlying SNARK".
- BUT existing security analyses show a huge gap in security ("PCD is far less secure than underlying SNARK").

TL;DR

What is *proof-carrying data* (PCD)?

- Recursive compositions of SNARKs.
- It's useful for efficiently verifying distributed computations.

Problem:

- PCD is deployed under the assumption "security of PCD" = "security of underlying SNARK".
- BUT existing security analyses show a huge gap in security ("PCD is far less secure than underlying SNARK").

This work:

TL;DR

What is *proof-carrying data* (PCD)?

- Recursive compositions of SNARKs.
- It's useful for efficiently verifying distributed computations.

Problem:

- PCD is deployed under the assumption "security of PCD" = "security of underlying SNARK".
- BUT existing security analyses show a huge gap in security ("PCD is far less secure than underlying SNARK").

This work:

- We propose an **idealized PCD** that models hash-based PCD in practice.

TL;DR

What is *proof-carrying data* (PCD)?

- Recursive compositions of SNARKs.
- It's useful for efficiently verifying distributed computations.

Problem:

- PCD is deployed under the assumption "security of PCD" = "security of underlying SNARK".
- BUT existing security analyses show a huge gap in security ("PCD is far less secure than underlying SNARK").

This work:

- We propose an **idealized PCD** that models hash-based PCD in practice.
- We prove that this idealized PCD is **as secure as its underlying SNARK**.

What is proof-carrying data (PCD)? [1/2]

What is proof-carrying data (PCD)? [1/2]

Proof-carrying data (PCD)

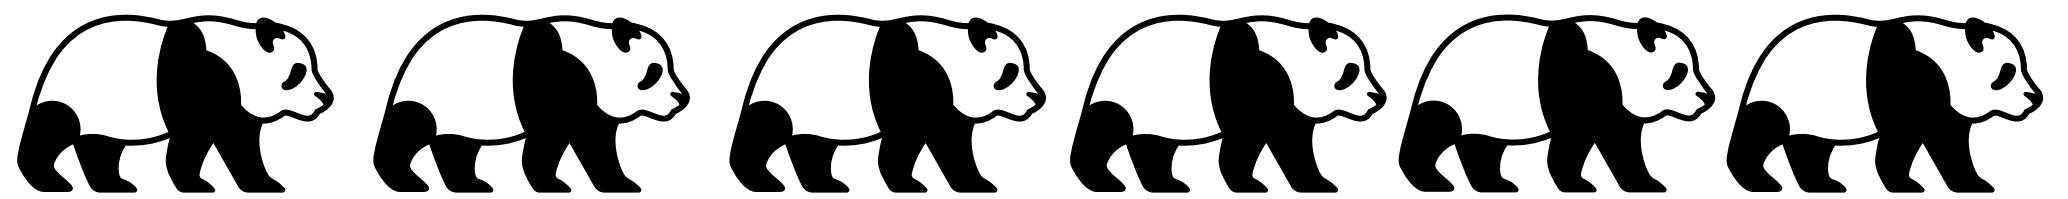
- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**

What is proof-carrying data (PCD)? [1/2]

Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**

E.g. A simple distributed computation: summing six numbers

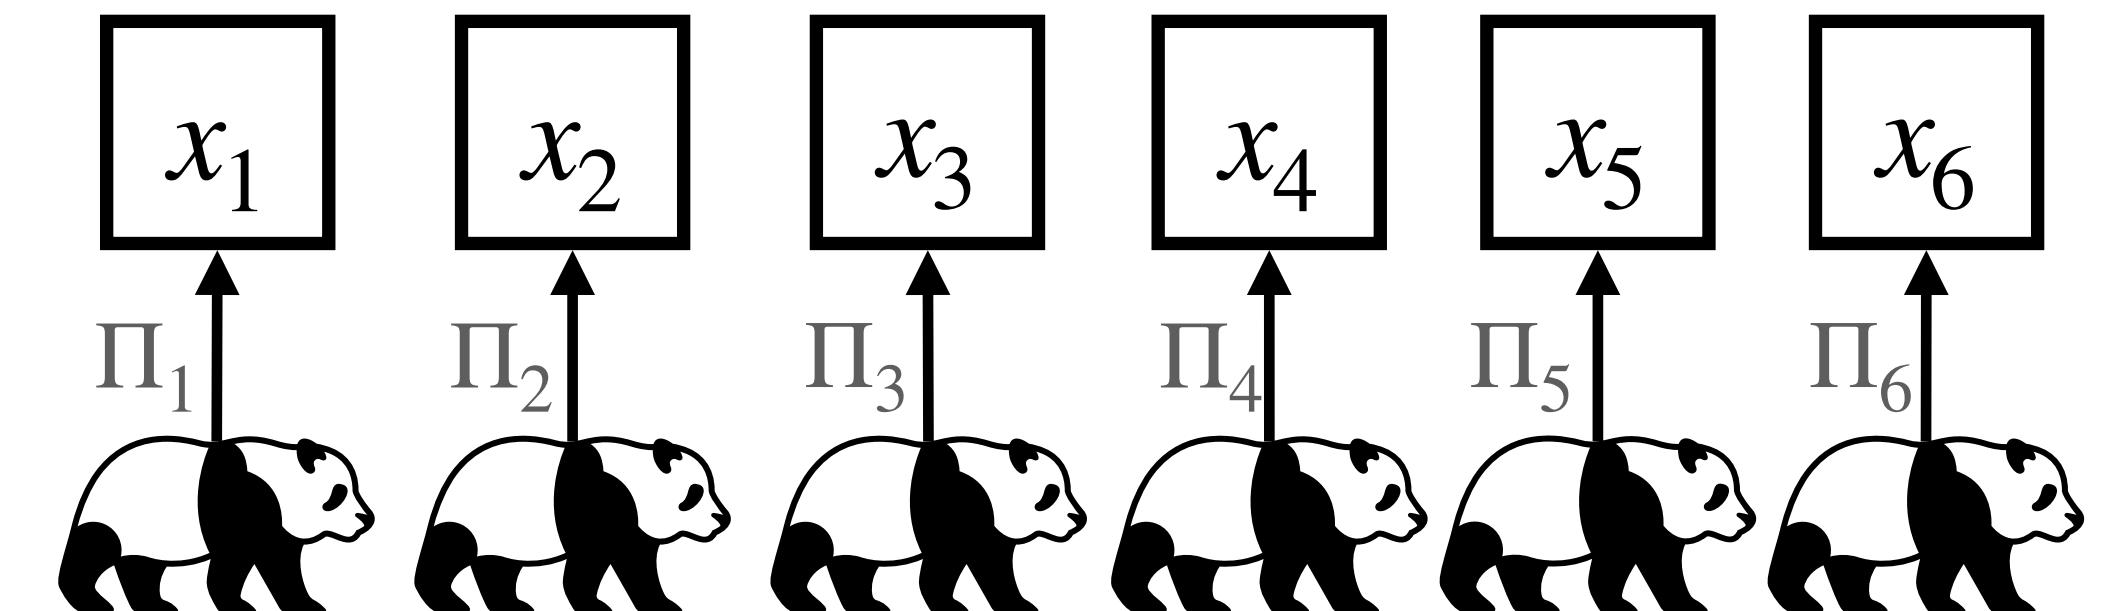


What is proof-carrying data (PCD)? [1/2]

Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**

E.g. A simple distributed computation: summing six numbers

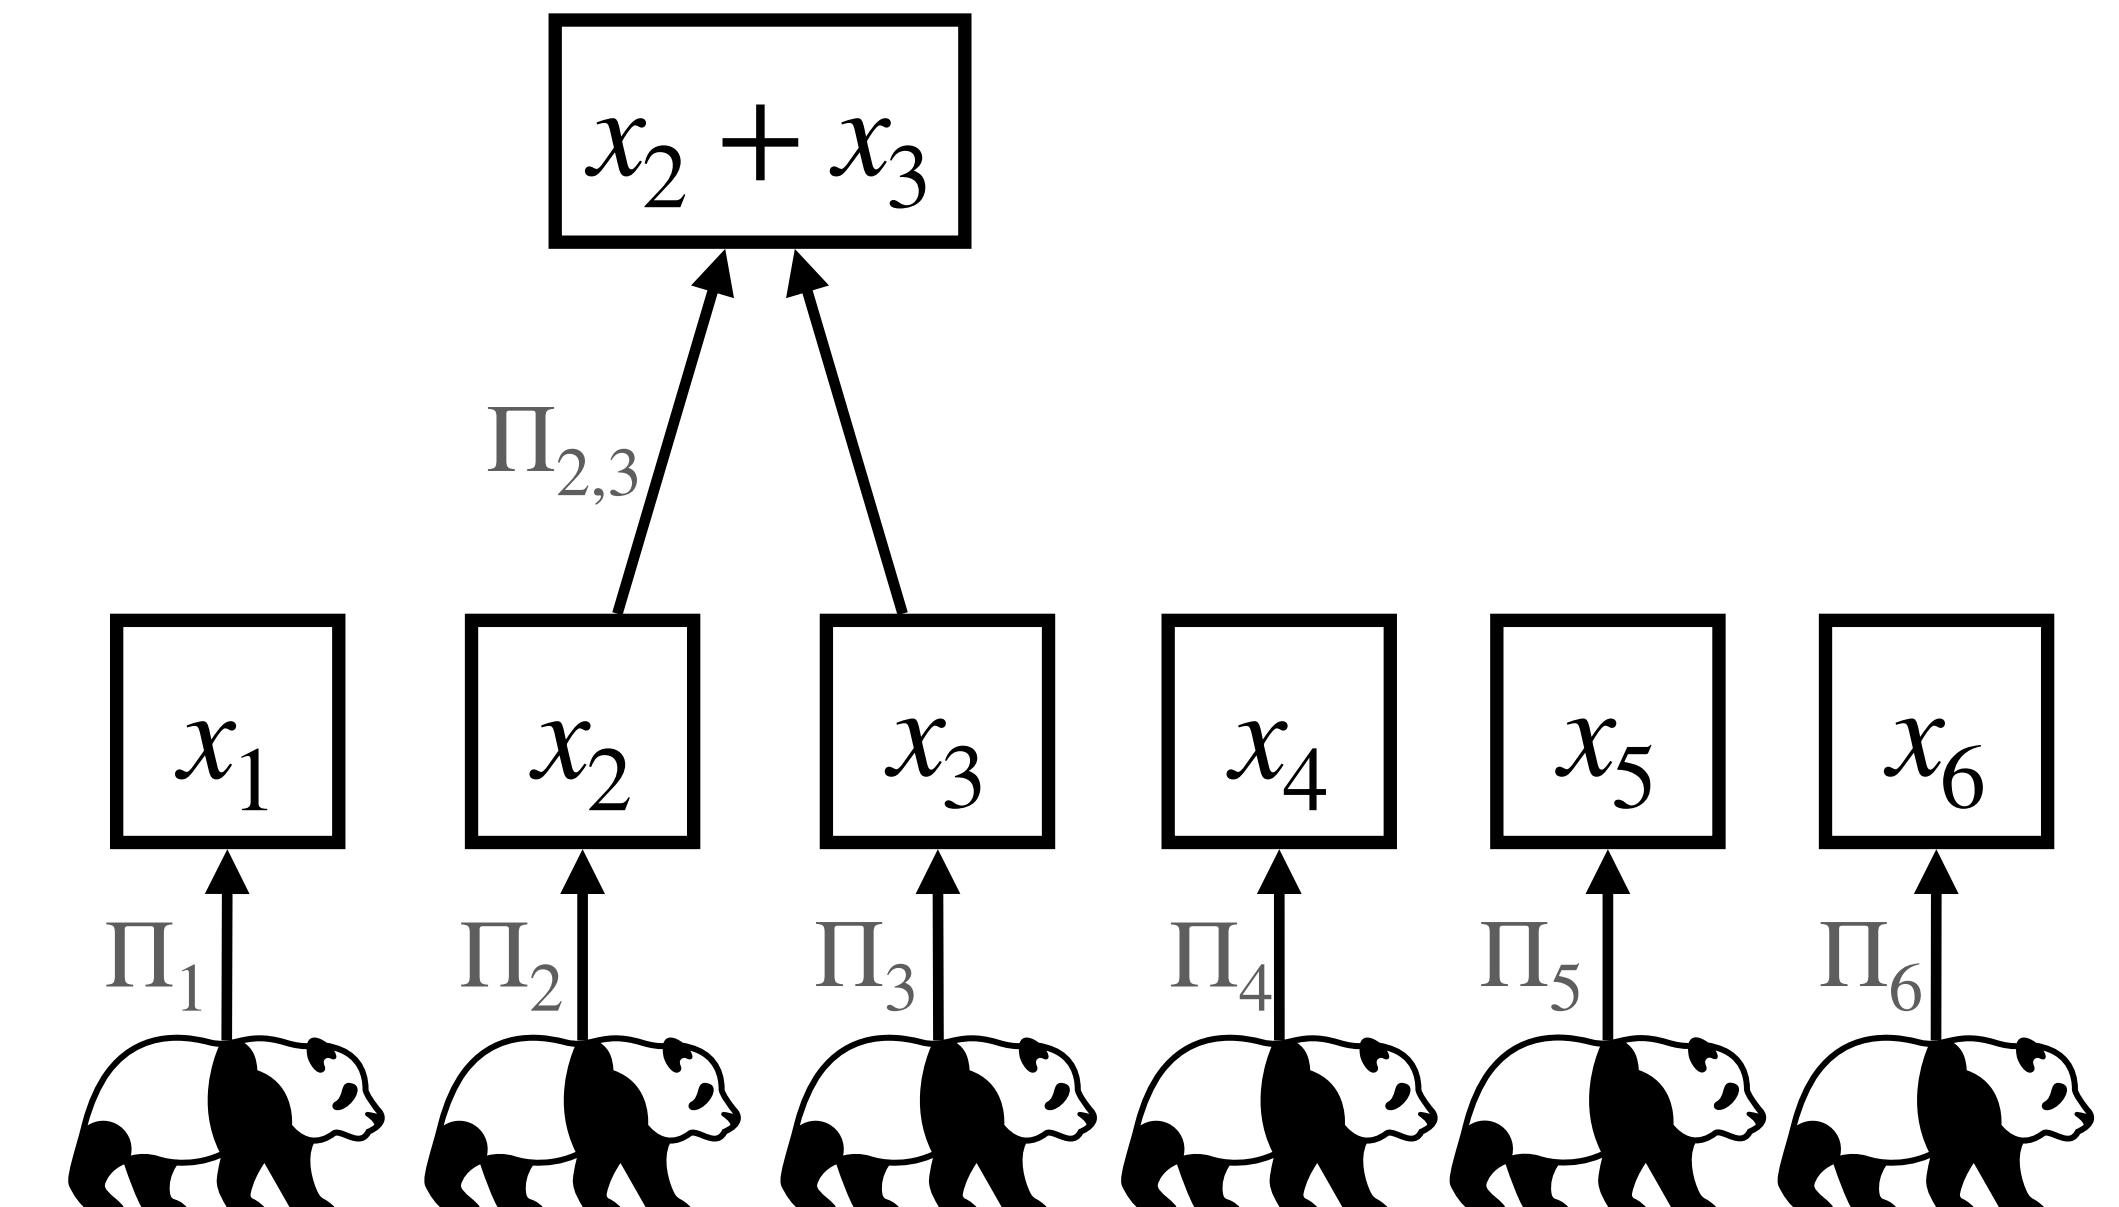


What is proof-carrying data (PCD)? [1/2]

Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**

E.g. A simple distributed computation: summing six numbers

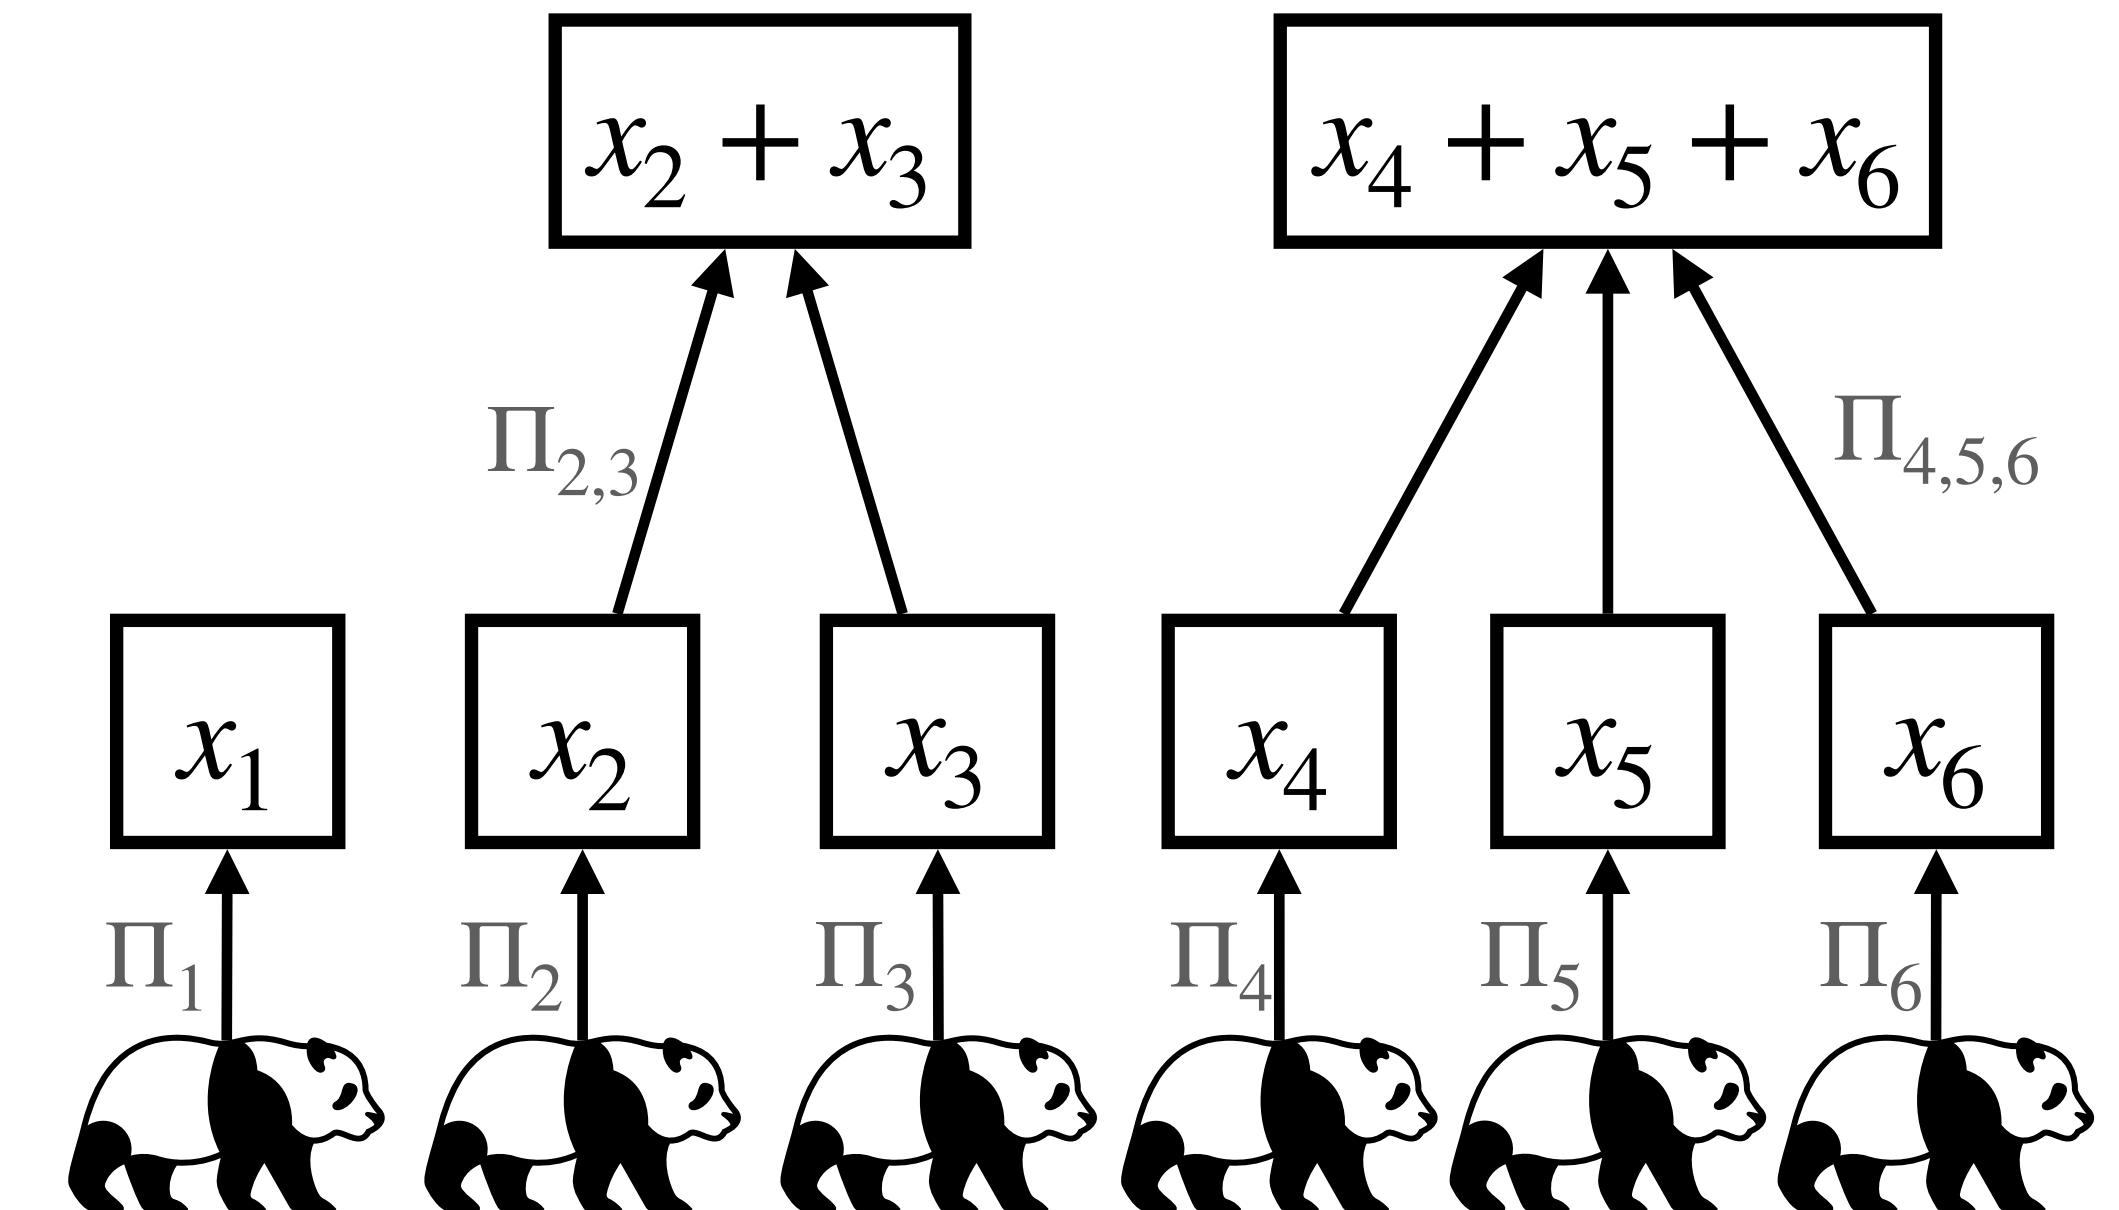


What is proof-carrying data (PCD)? [1/2]

Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**

E.g. A simple distributed computation: summing six numbers

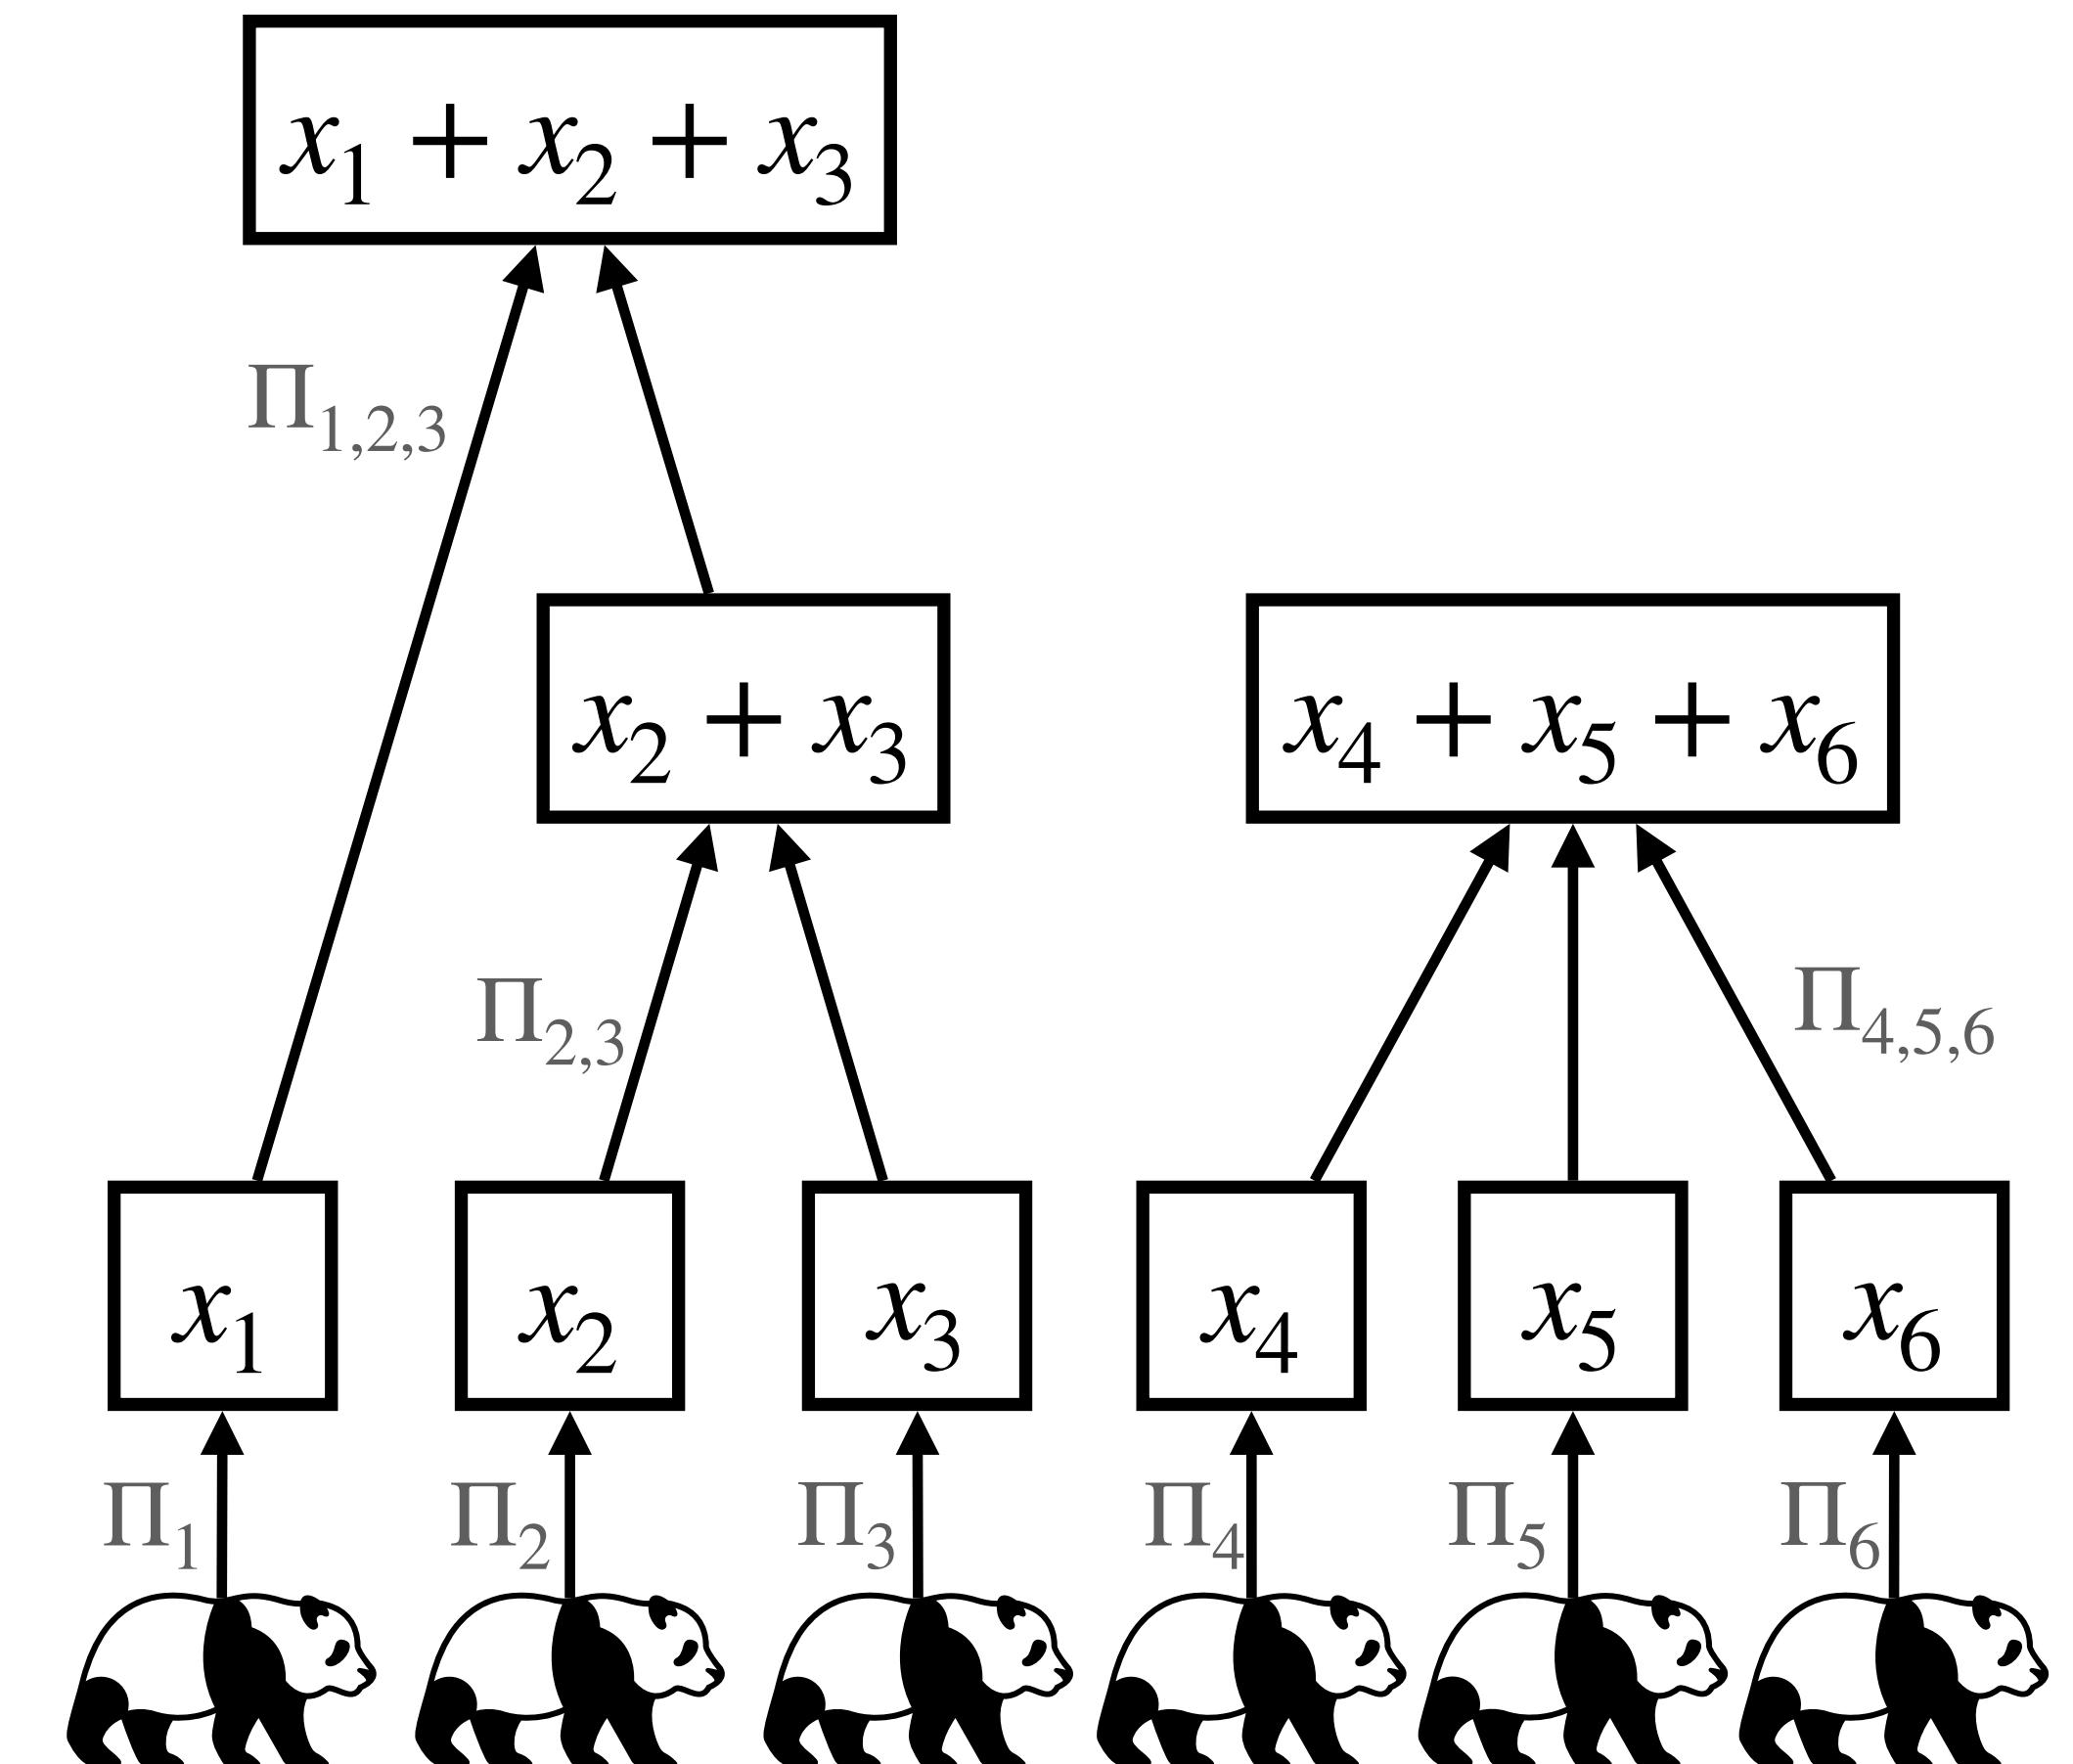


What is proof-carrying data (PCD)? [1/2]

Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**

E.g. A simple distributed computation: summing six numbers

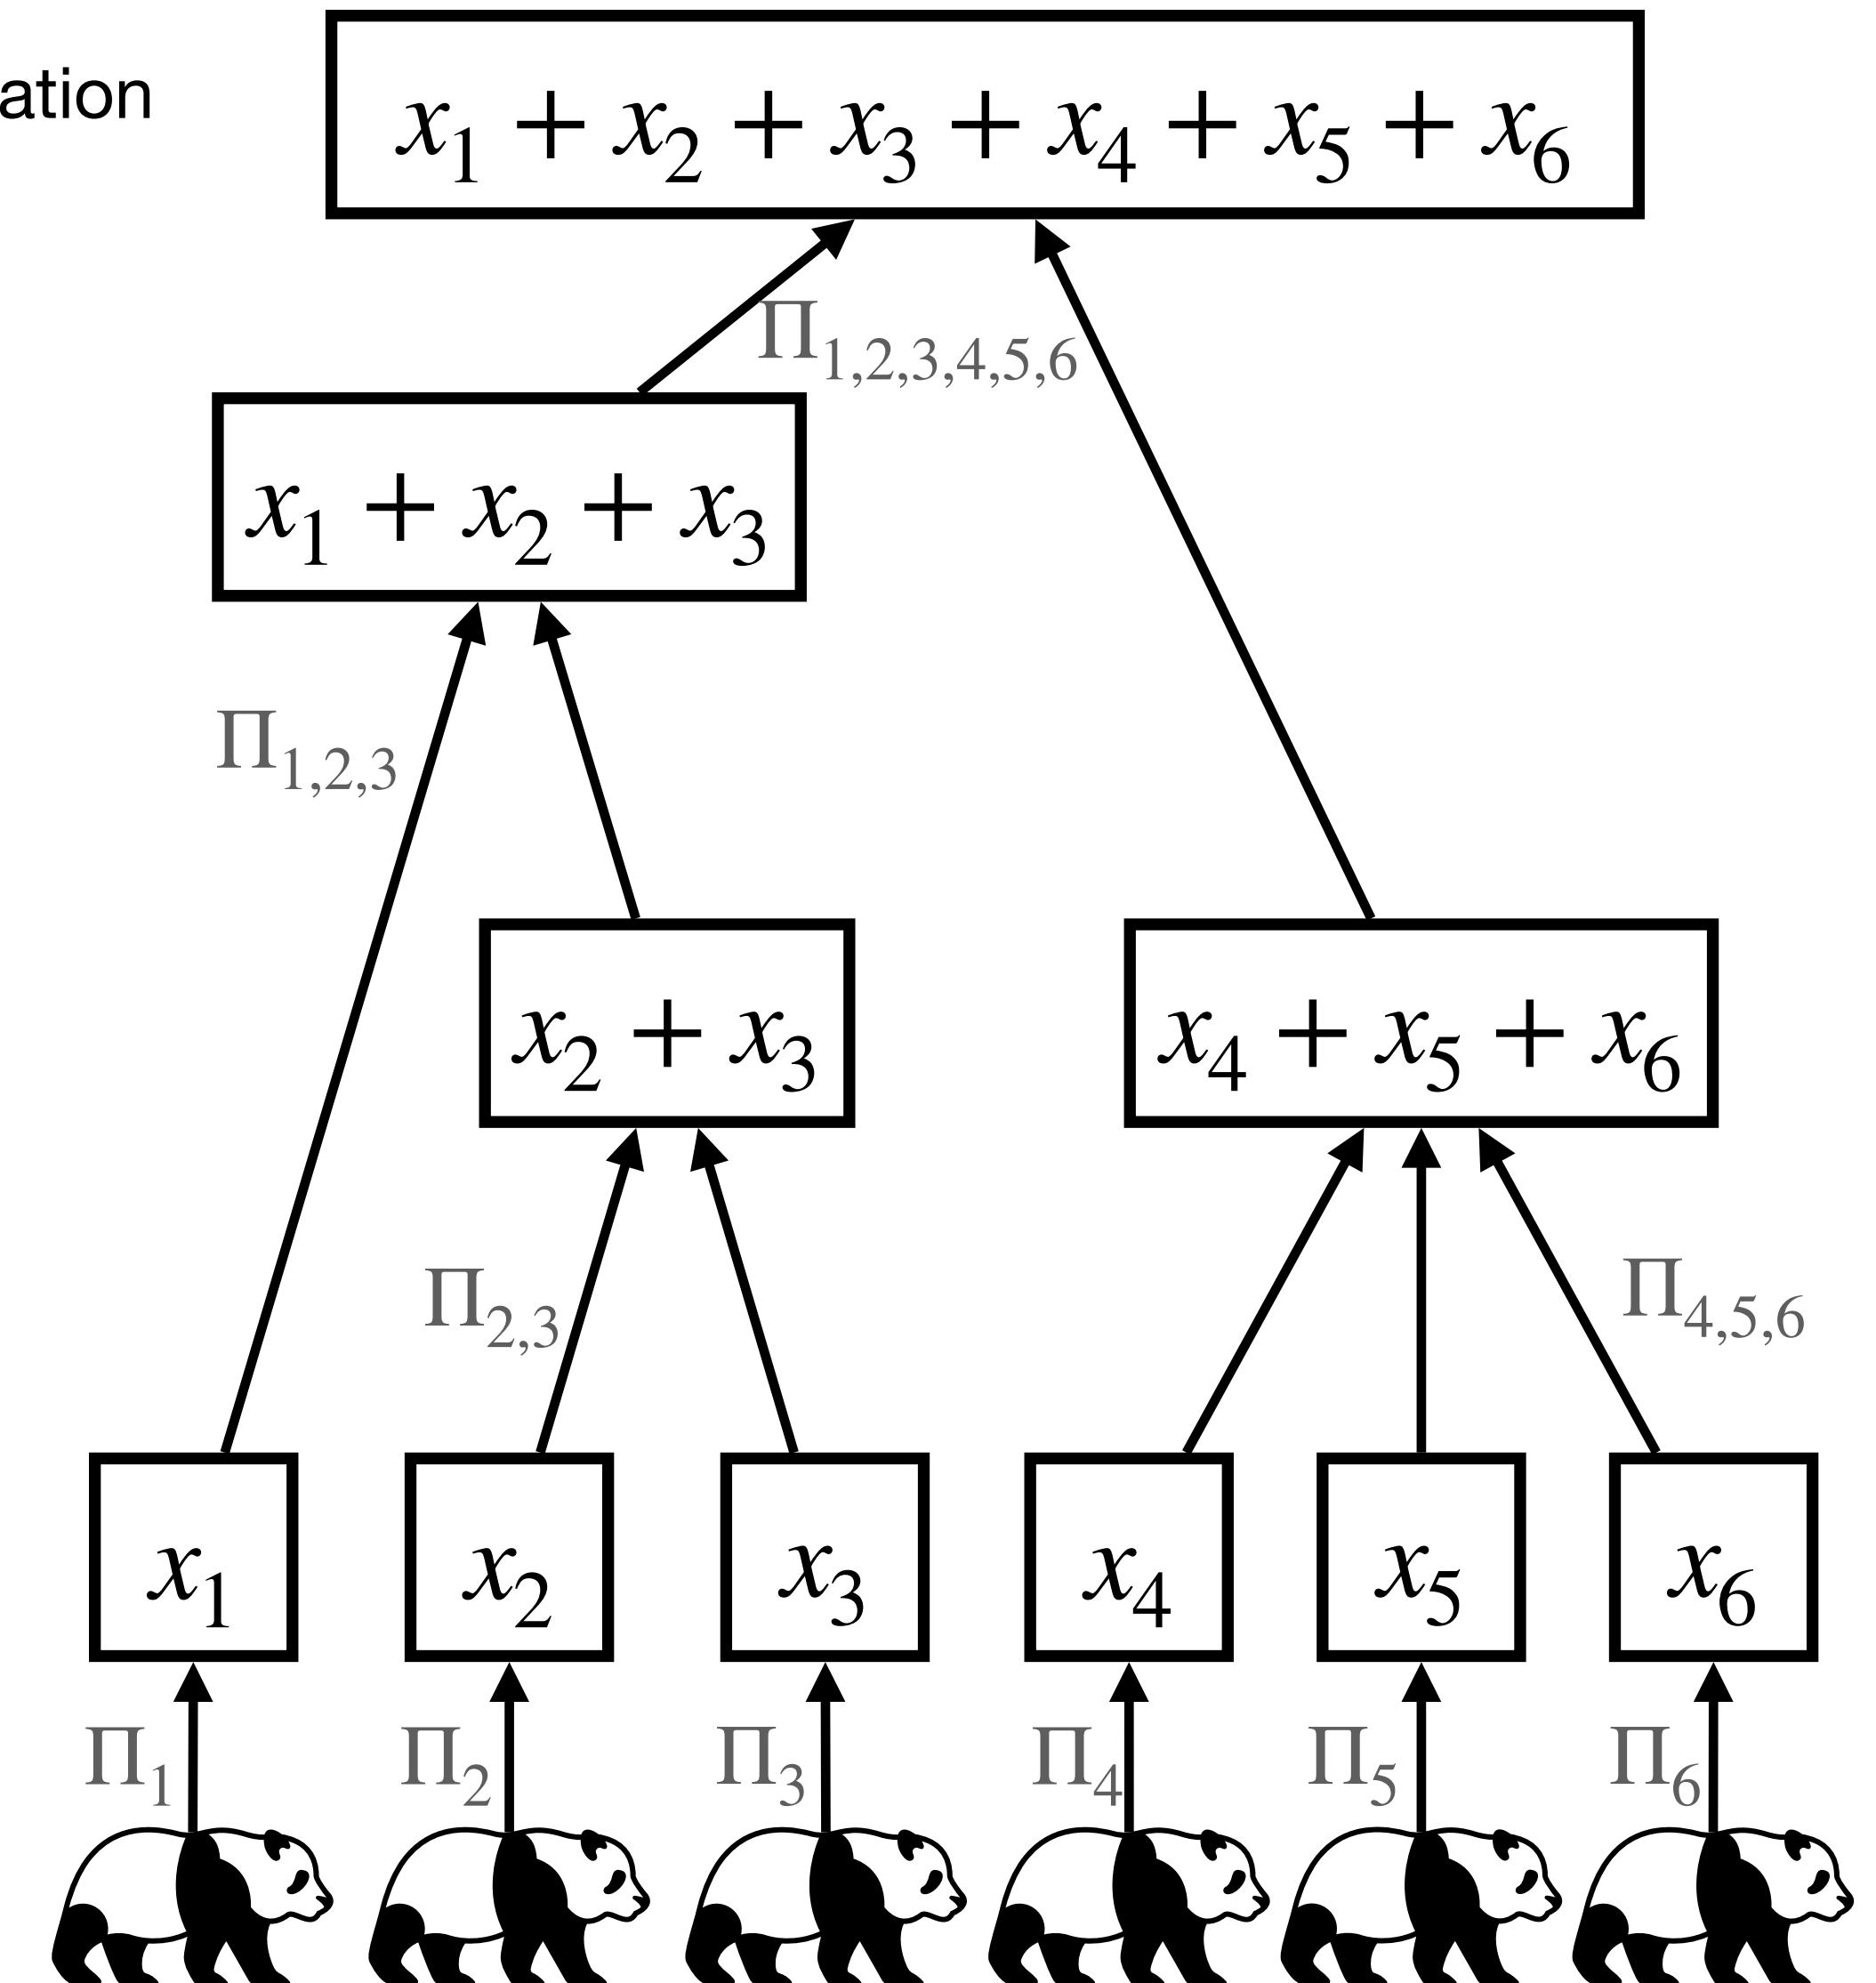


What is proof-carrying data (PCD)? [1/2]

Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**

E.g. A simple distributed computation: summing six numbers



What is proof-carrying data (PCD)? [2/2]

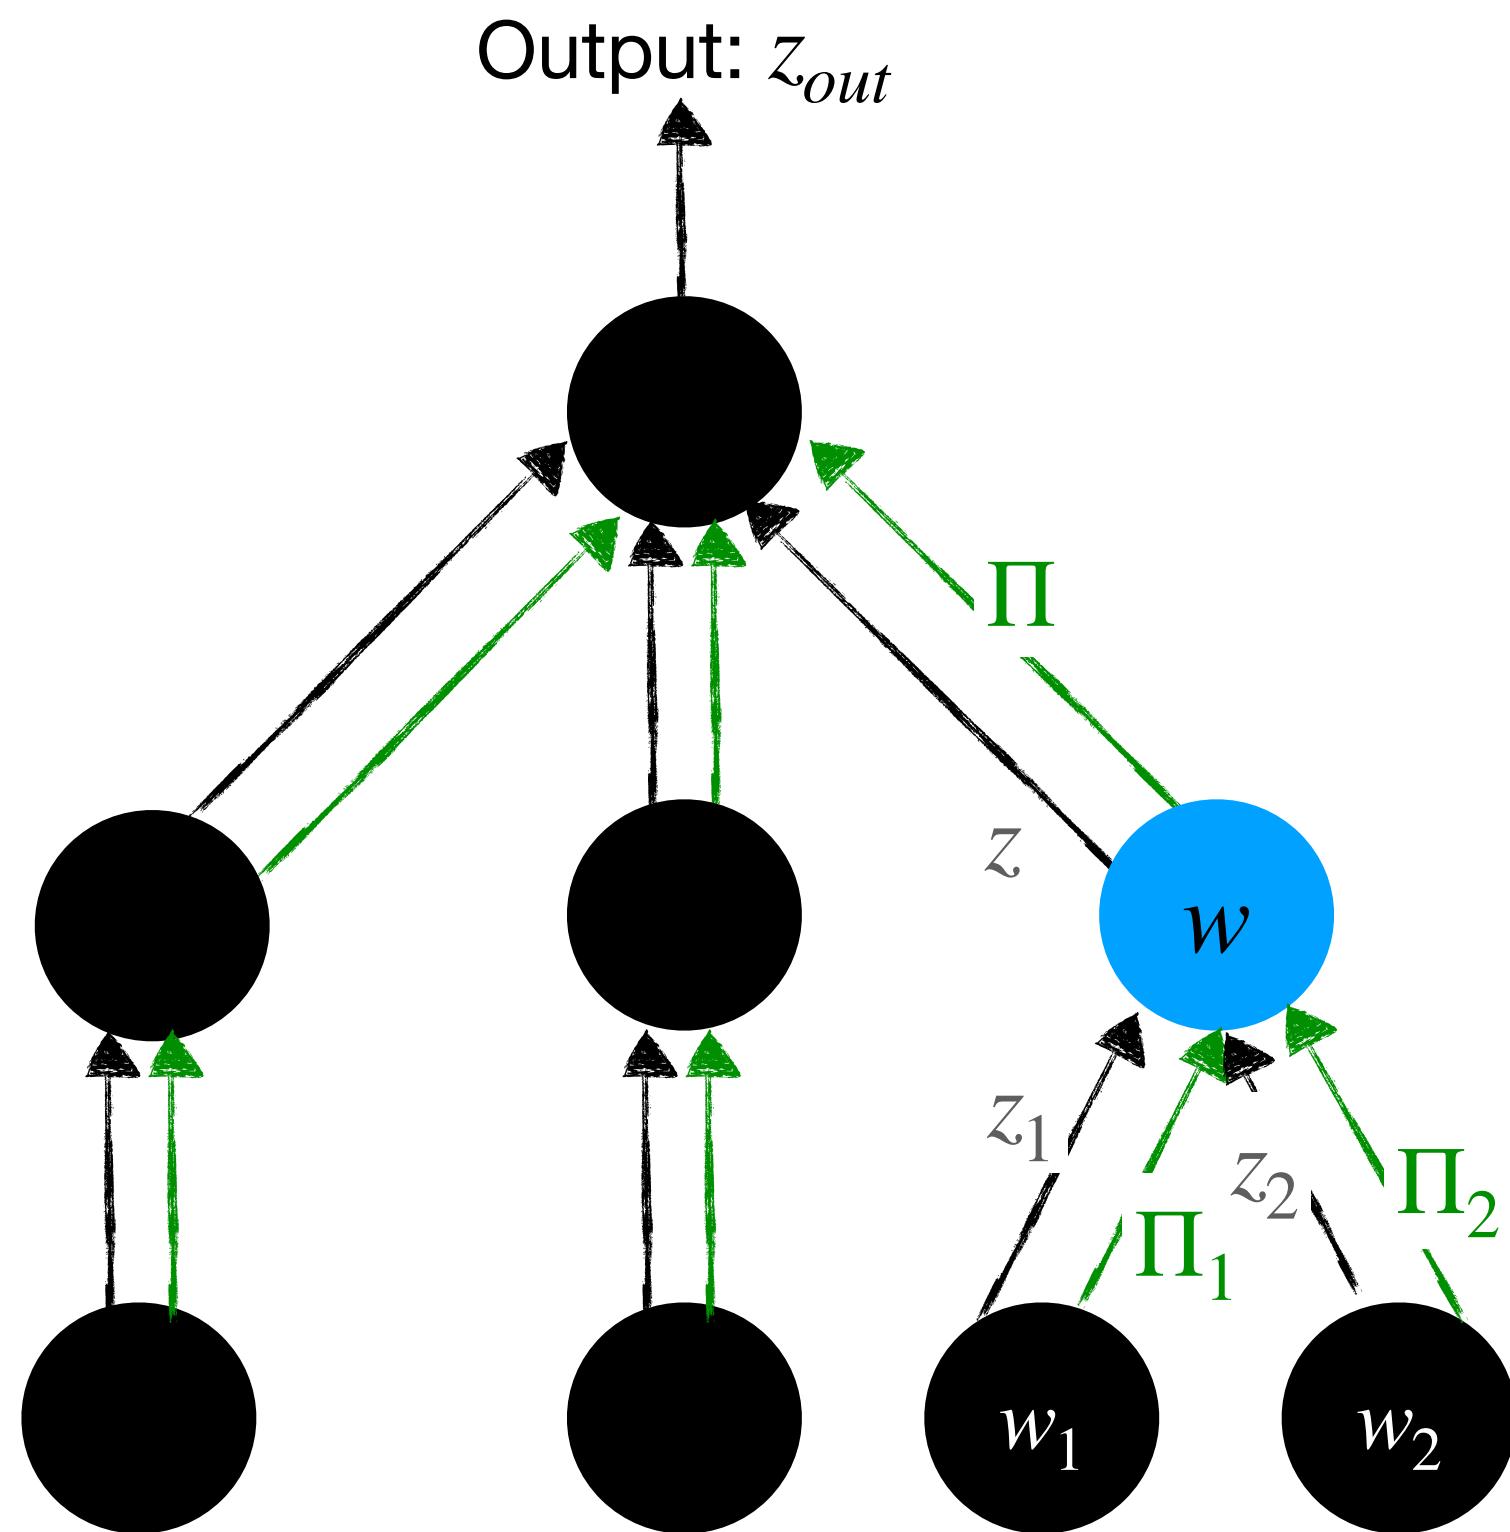
Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**

What is proof-carrying data (PCD)? [2/2]

Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**



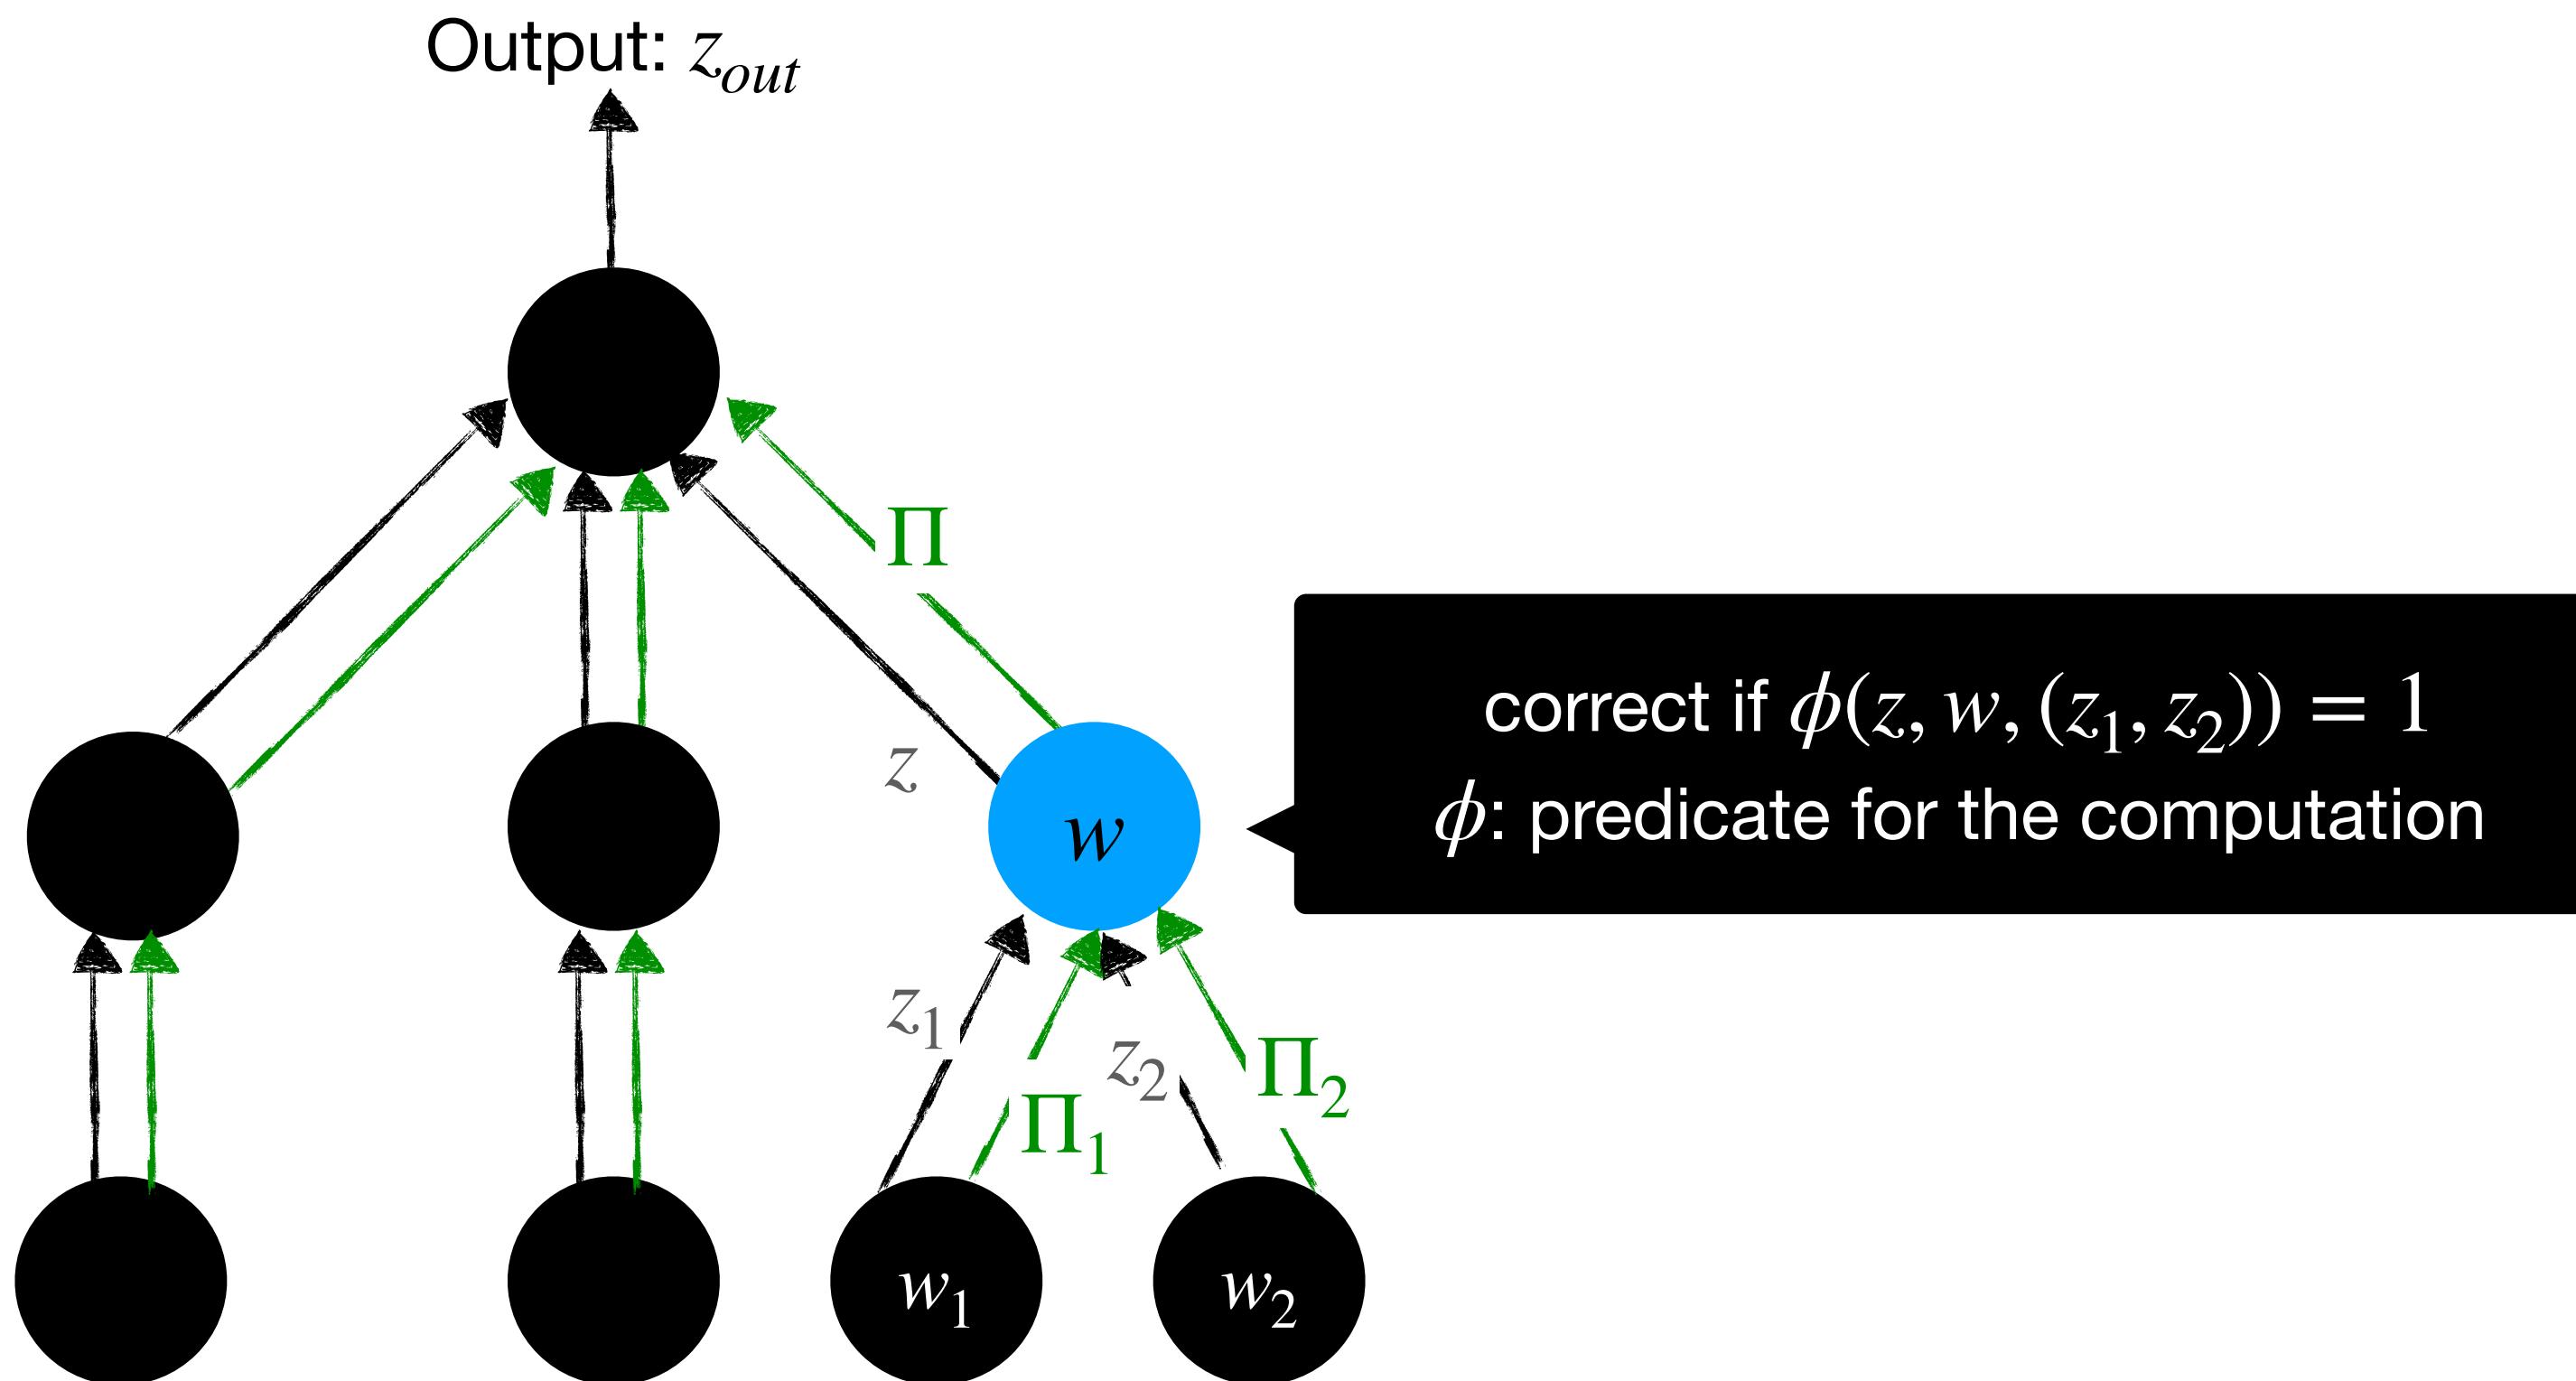
PCD transcript T

size $N = 8$, depth $D = 3$

What is proof-carrying data (PCD)? [2/2]

Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**



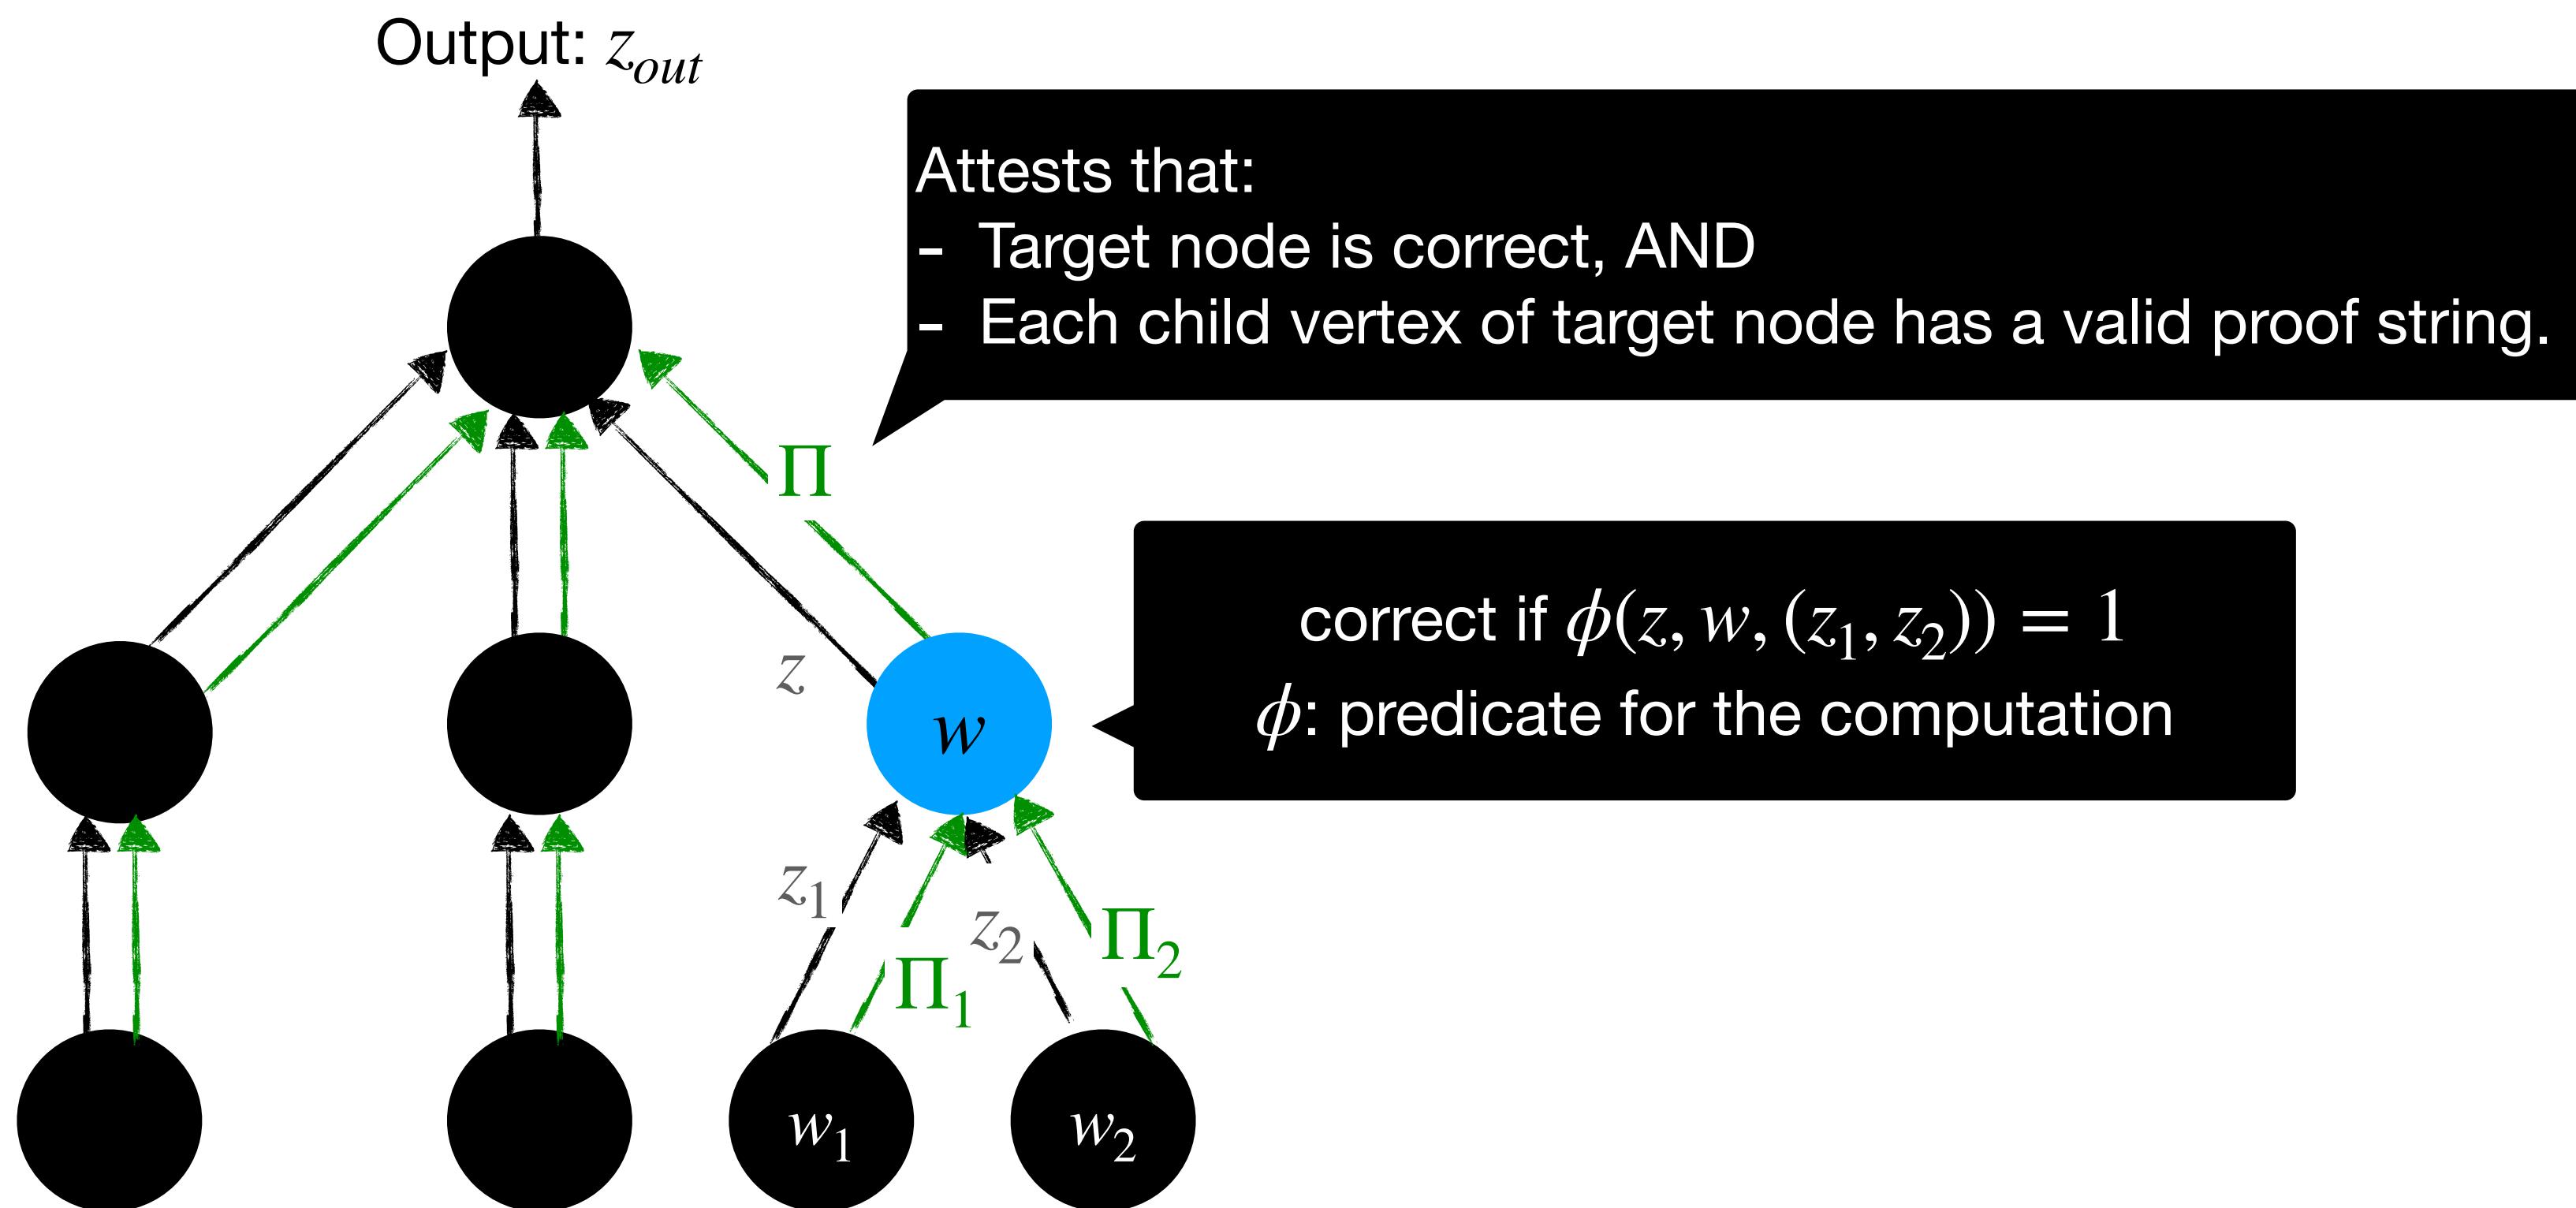
PCD transcript T

size $N = 8$, depth $D = 3$

What is proof-carrying data (PCD)? [2/2]

Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**



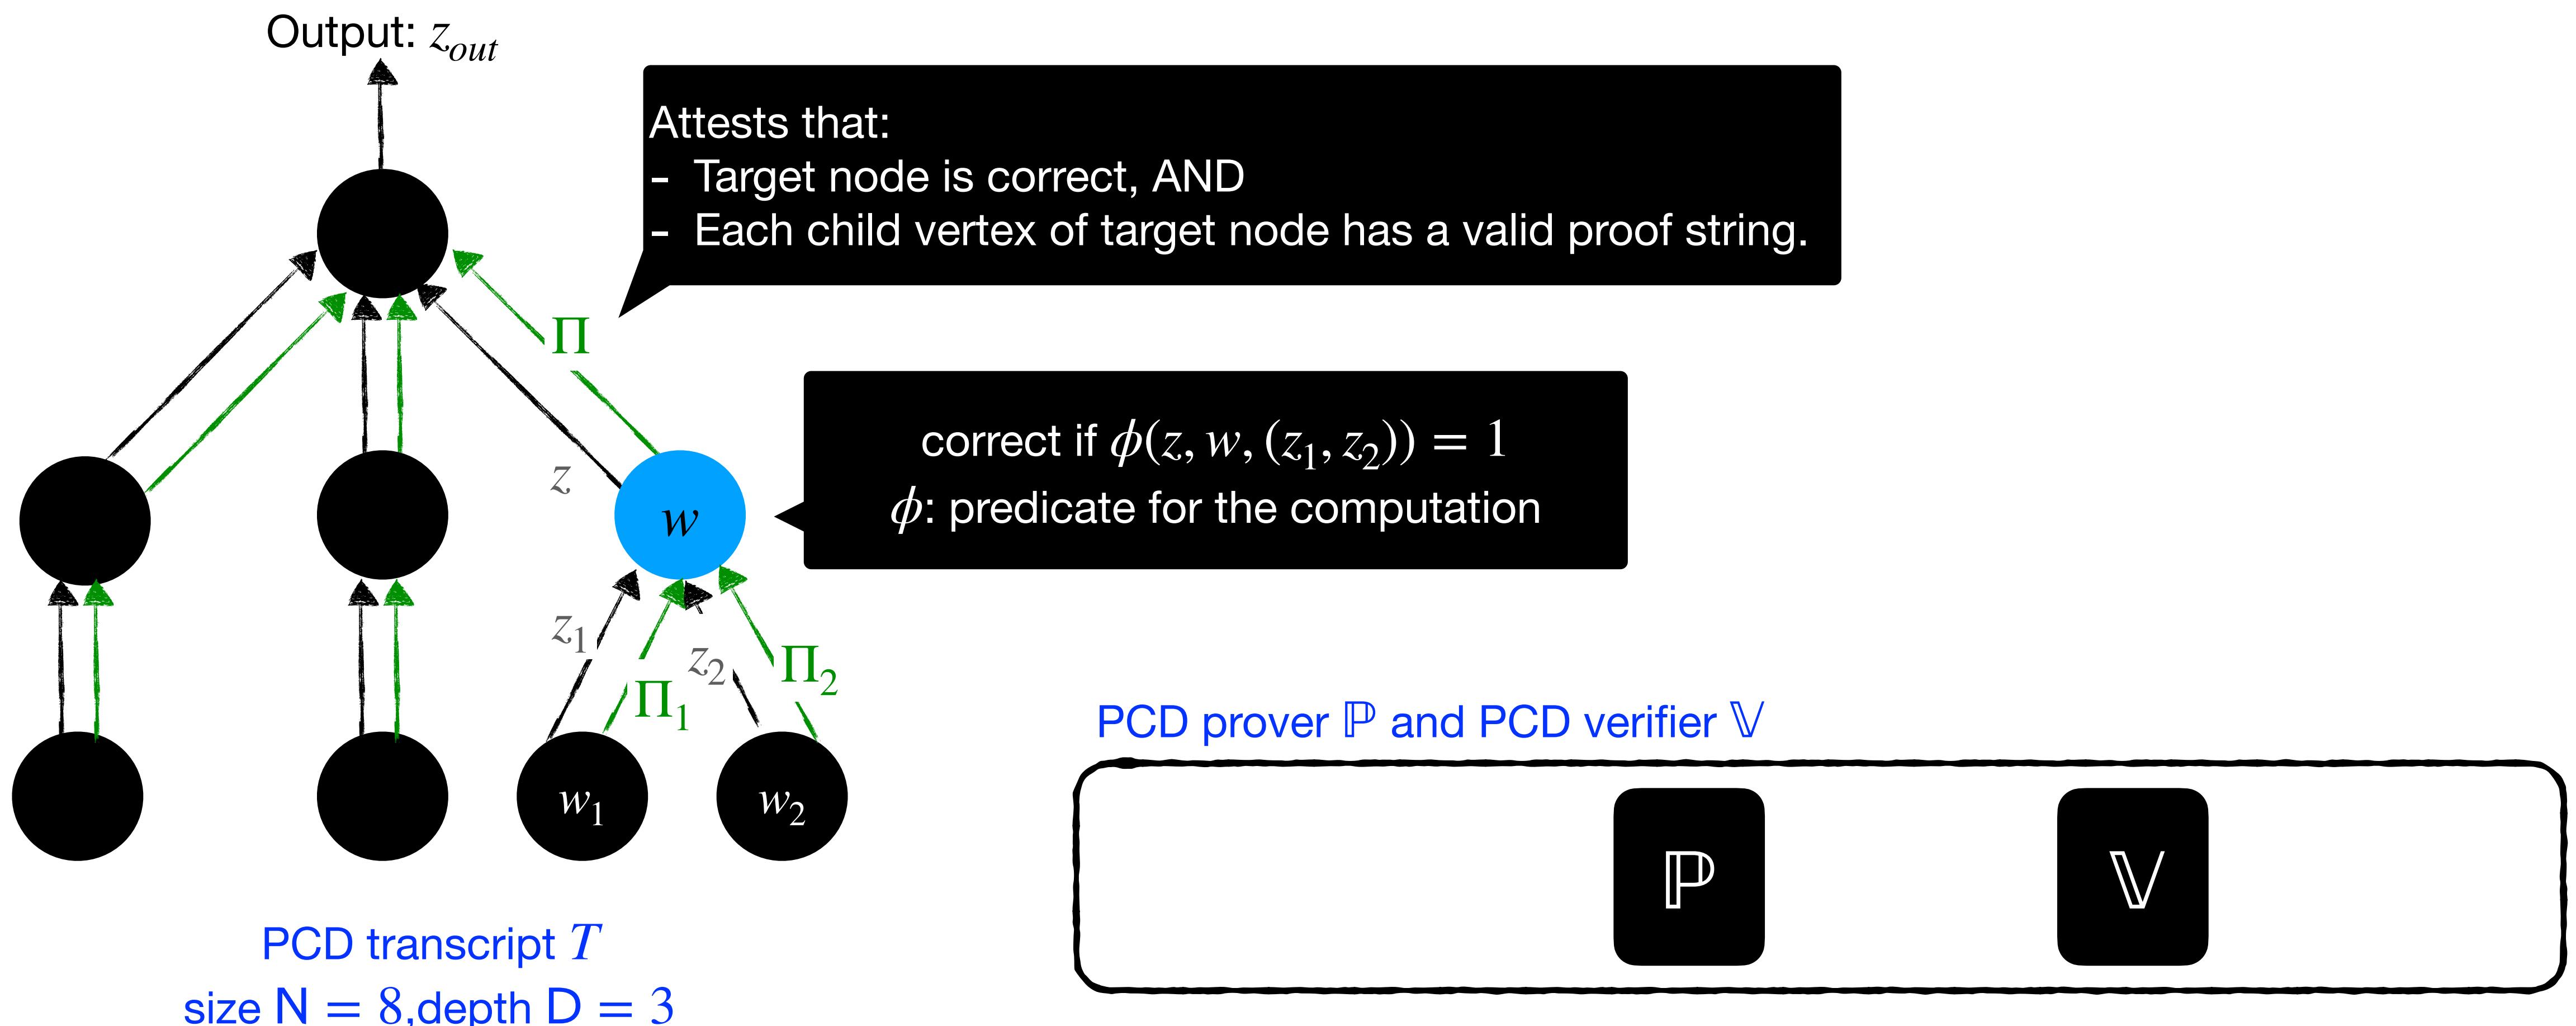
PCD transcript T

size $N = 8$, depth $D = 3$

What is proof-carrying data (PCD)? [2/2]

Proof-carrying data (PCD)

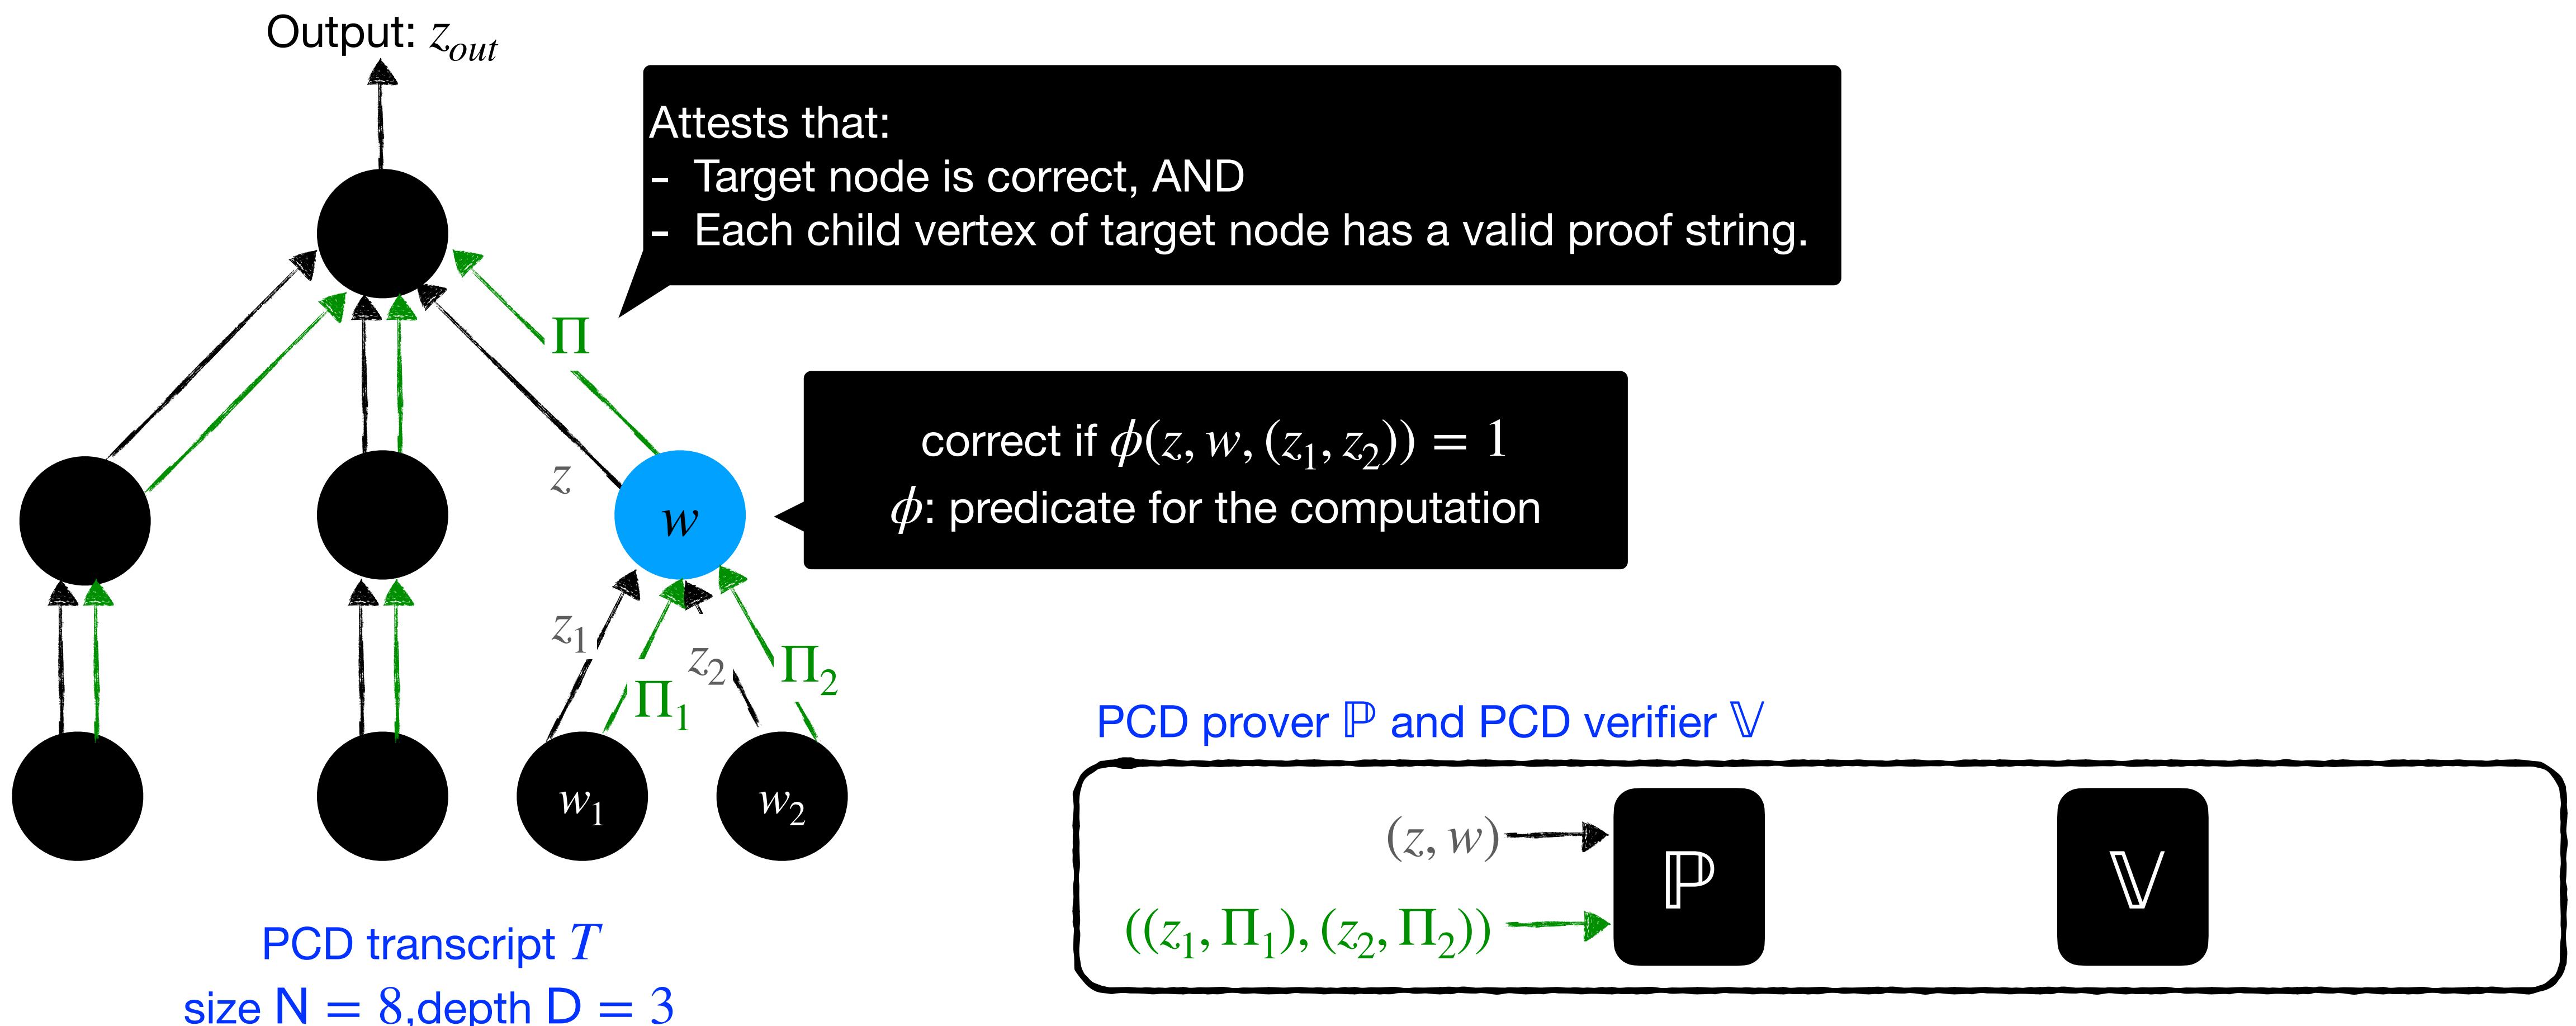
- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**



What is proof-carrying data (PCD)? [2/2]

Proof-carrying data (PCD)

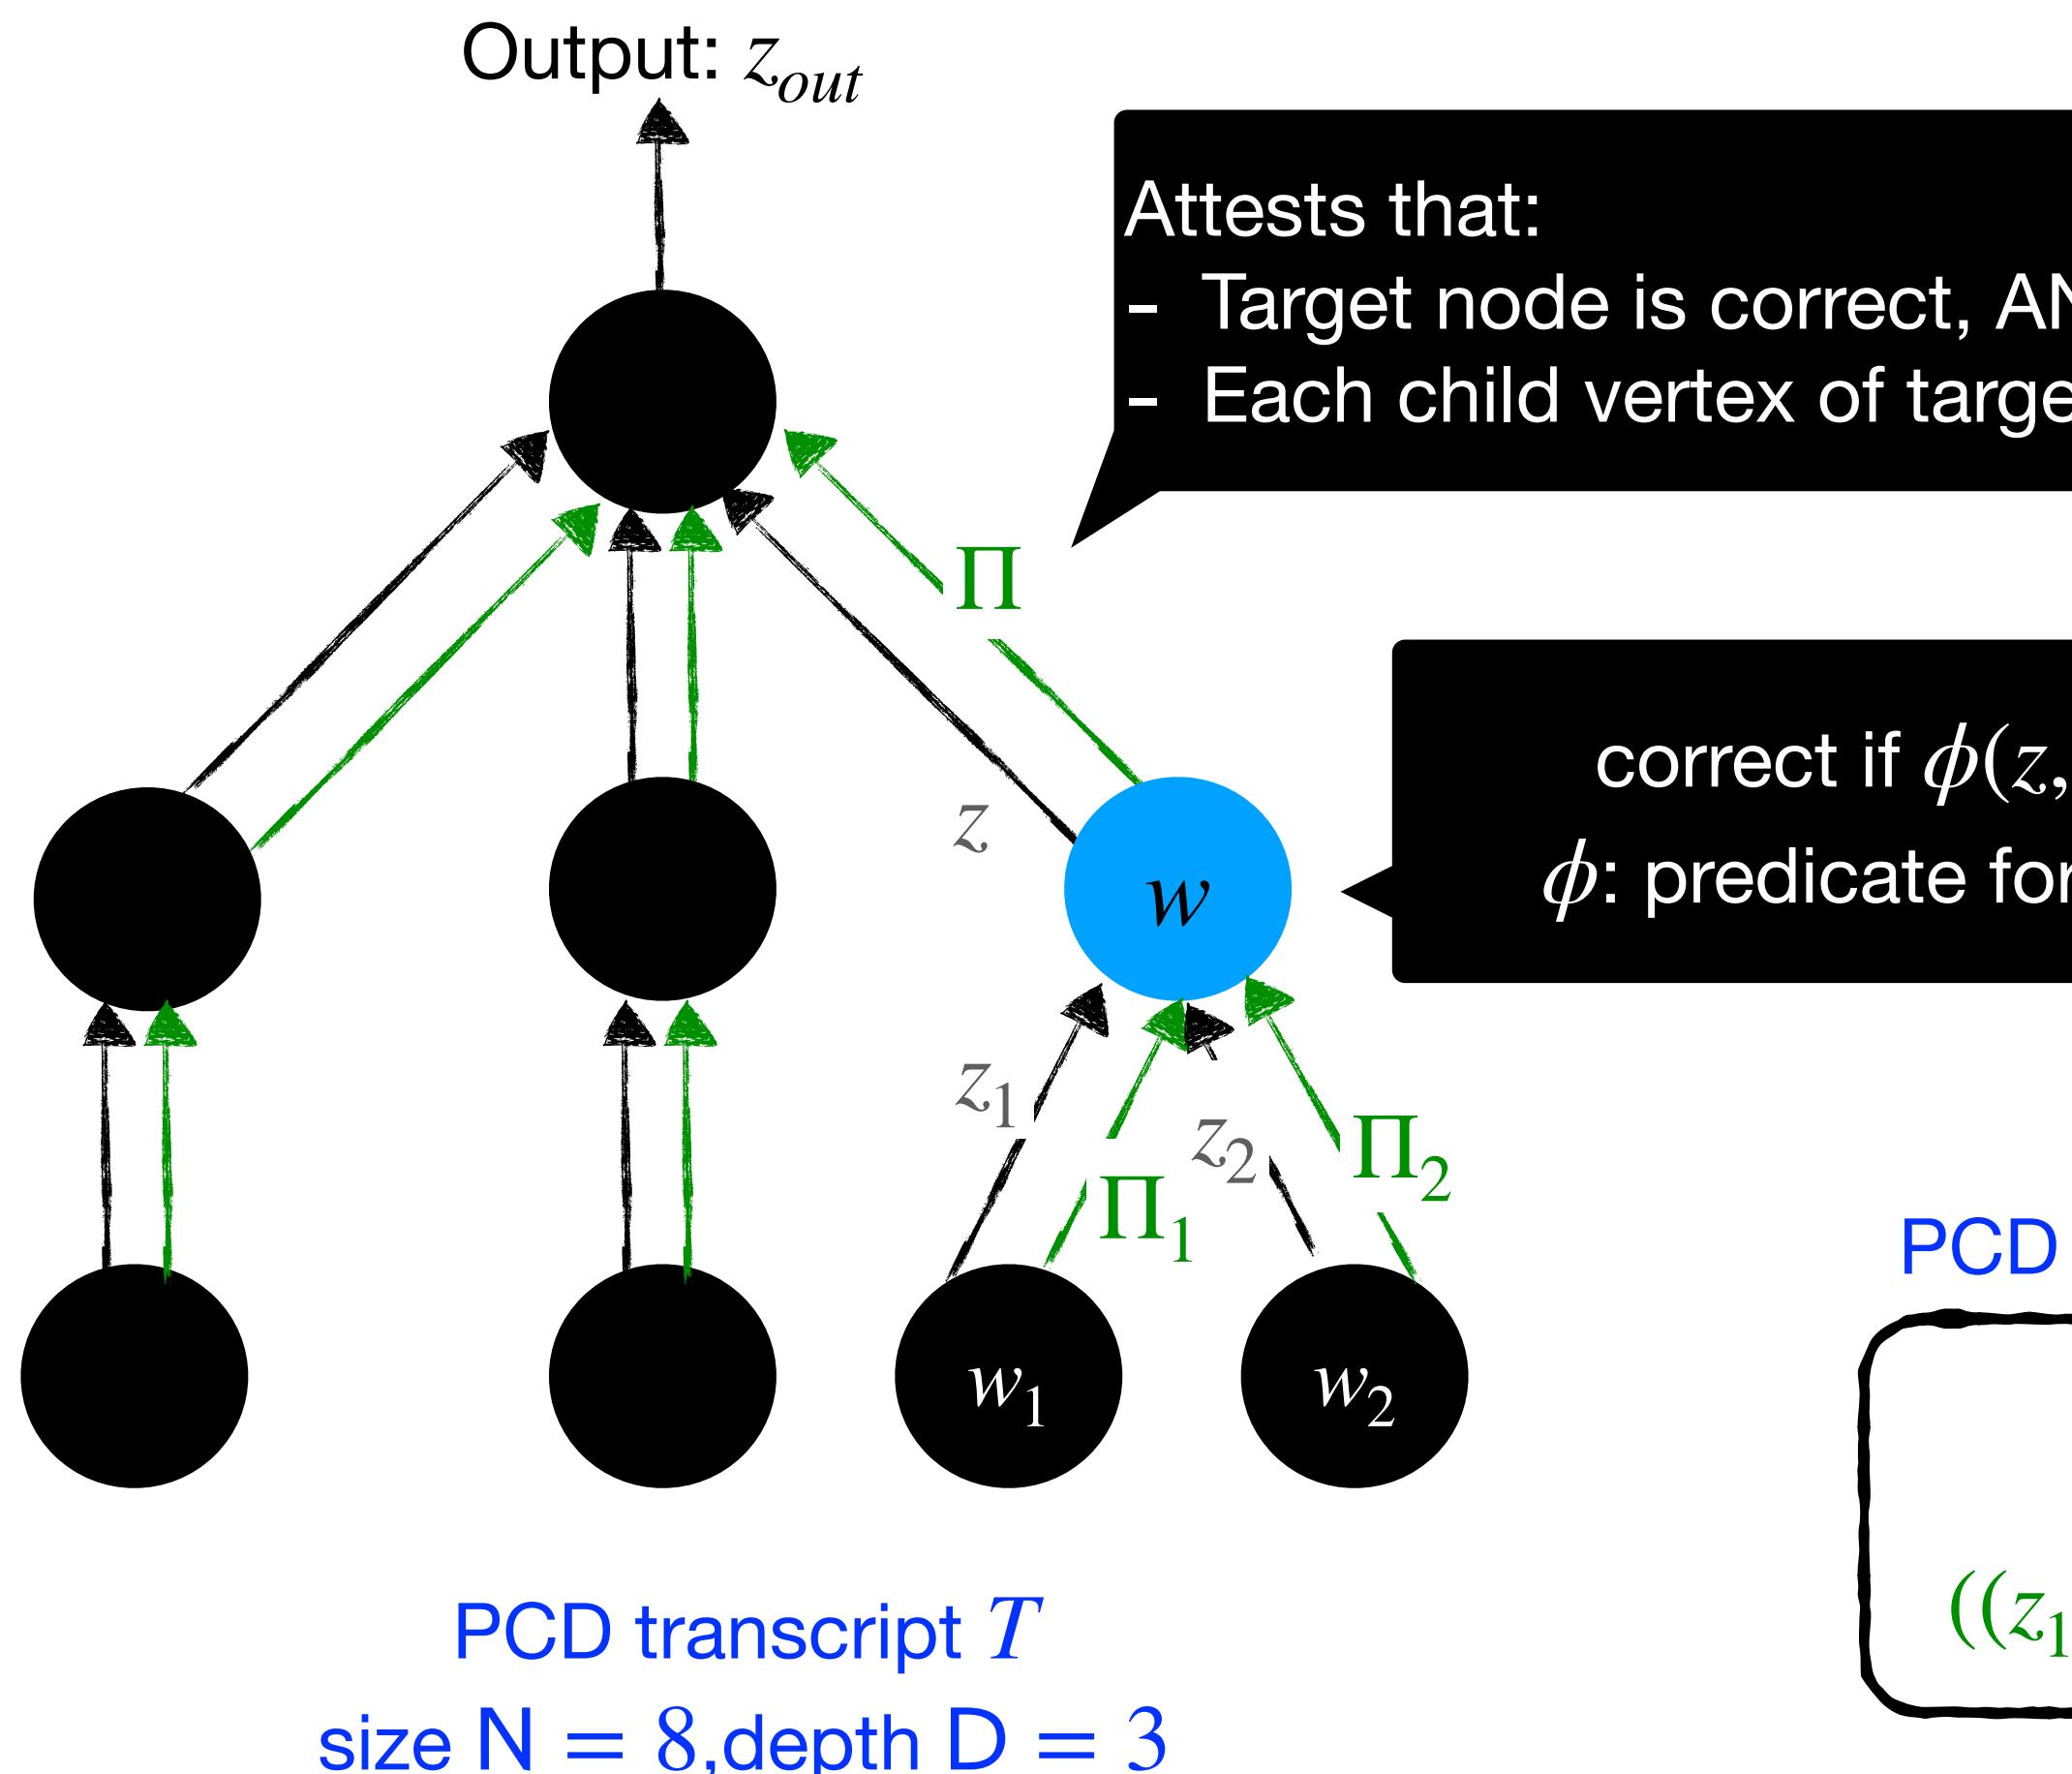
- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**



What is proof-carrying data (PCD)? [2/2]

Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**



Attests that:

- Target node is correct, AND
- Each child vertex of target node has a valid proof string.

correct if $\phi(z, w, (z_1, z_2)) = 1$
 ϕ : predicate for the computation

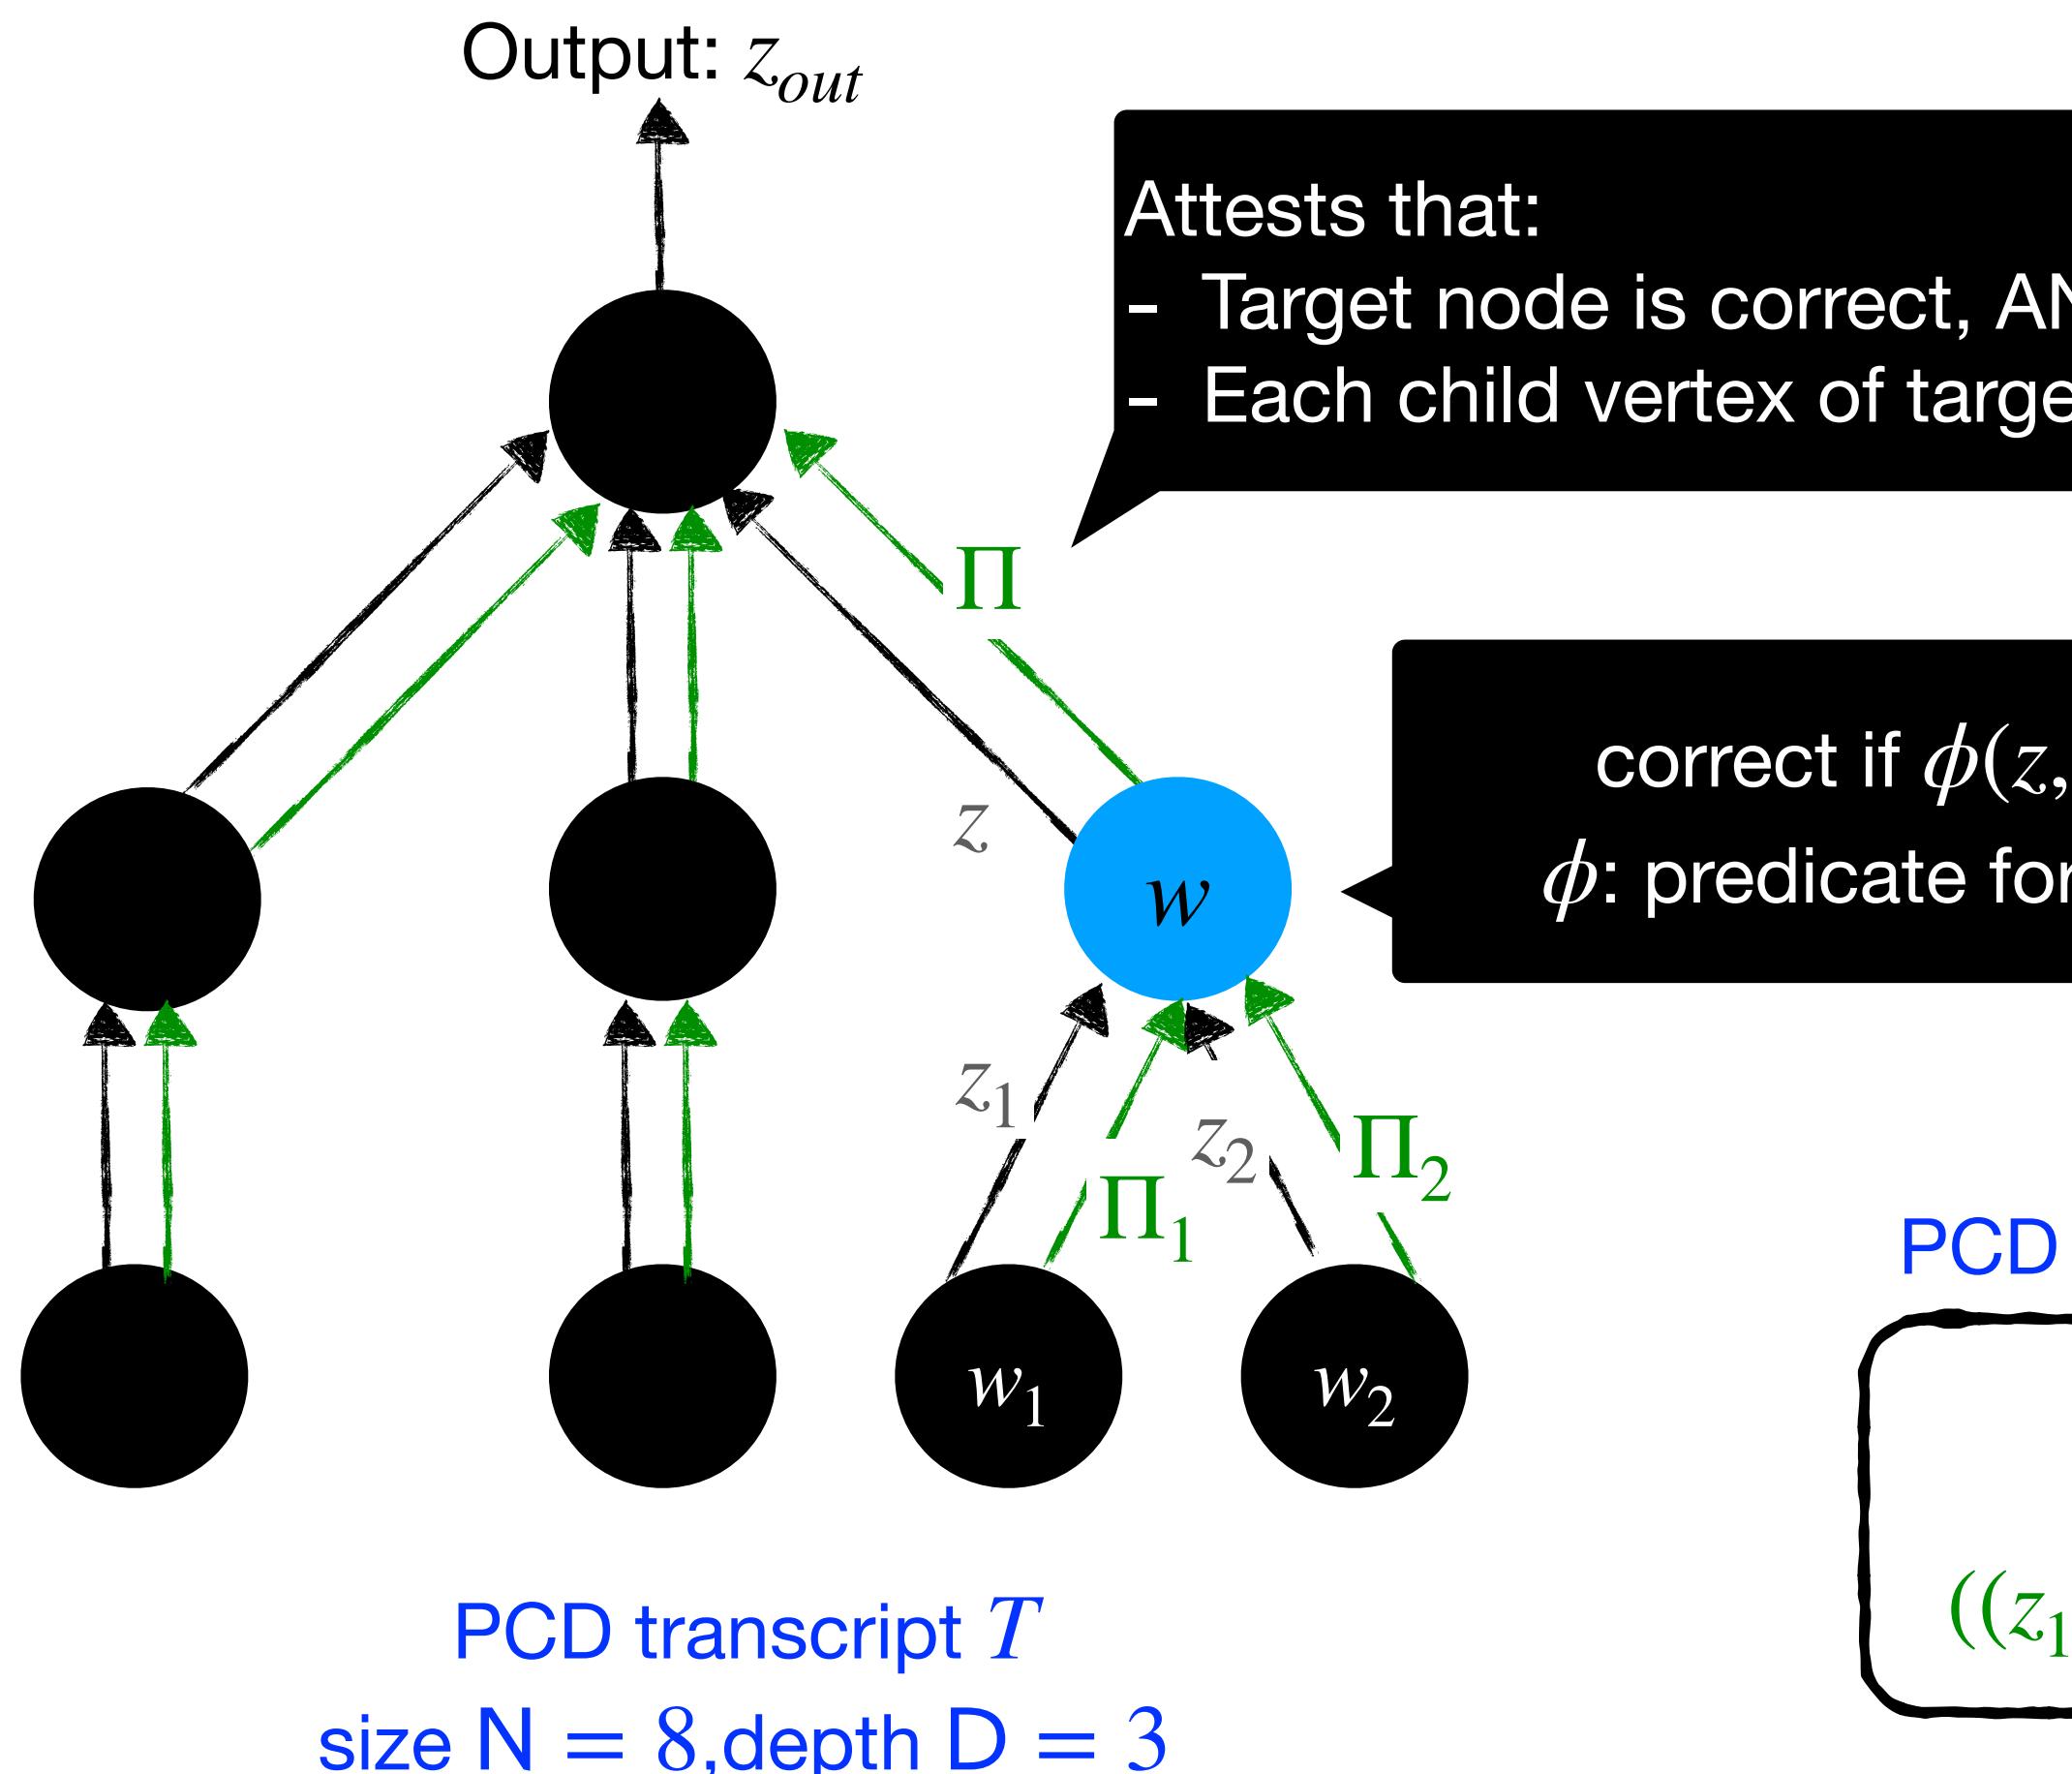
PCD prover \mathbb{P} and PCD verifier \mathbb{V}



What is proof-carrying data (PCD)? [2/2]

Proof-carrying data (PCD)

- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**

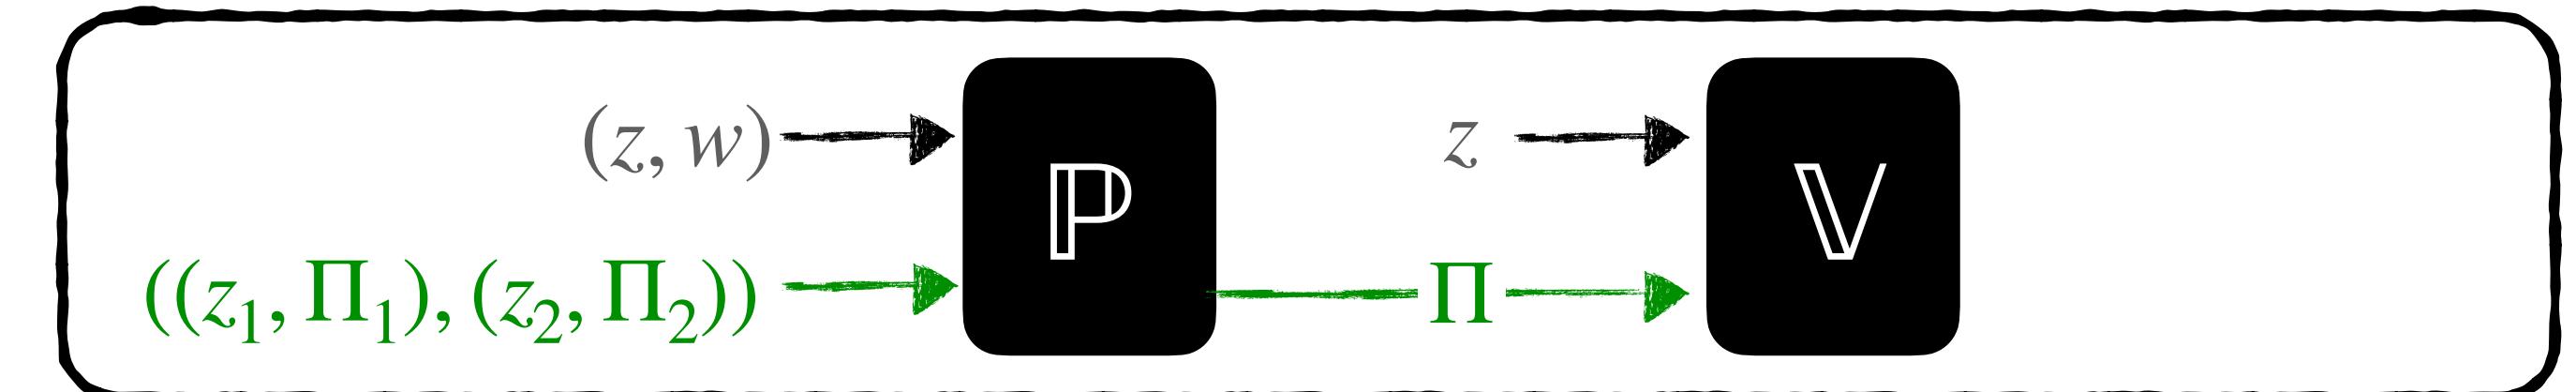


Attests that:

- Target node is correct, AND
- Each child vertex of target node has a valid proof string.

correct if $\phi(z, w, (z_1, z_2)) = 1$
 ϕ : predicate for the computation

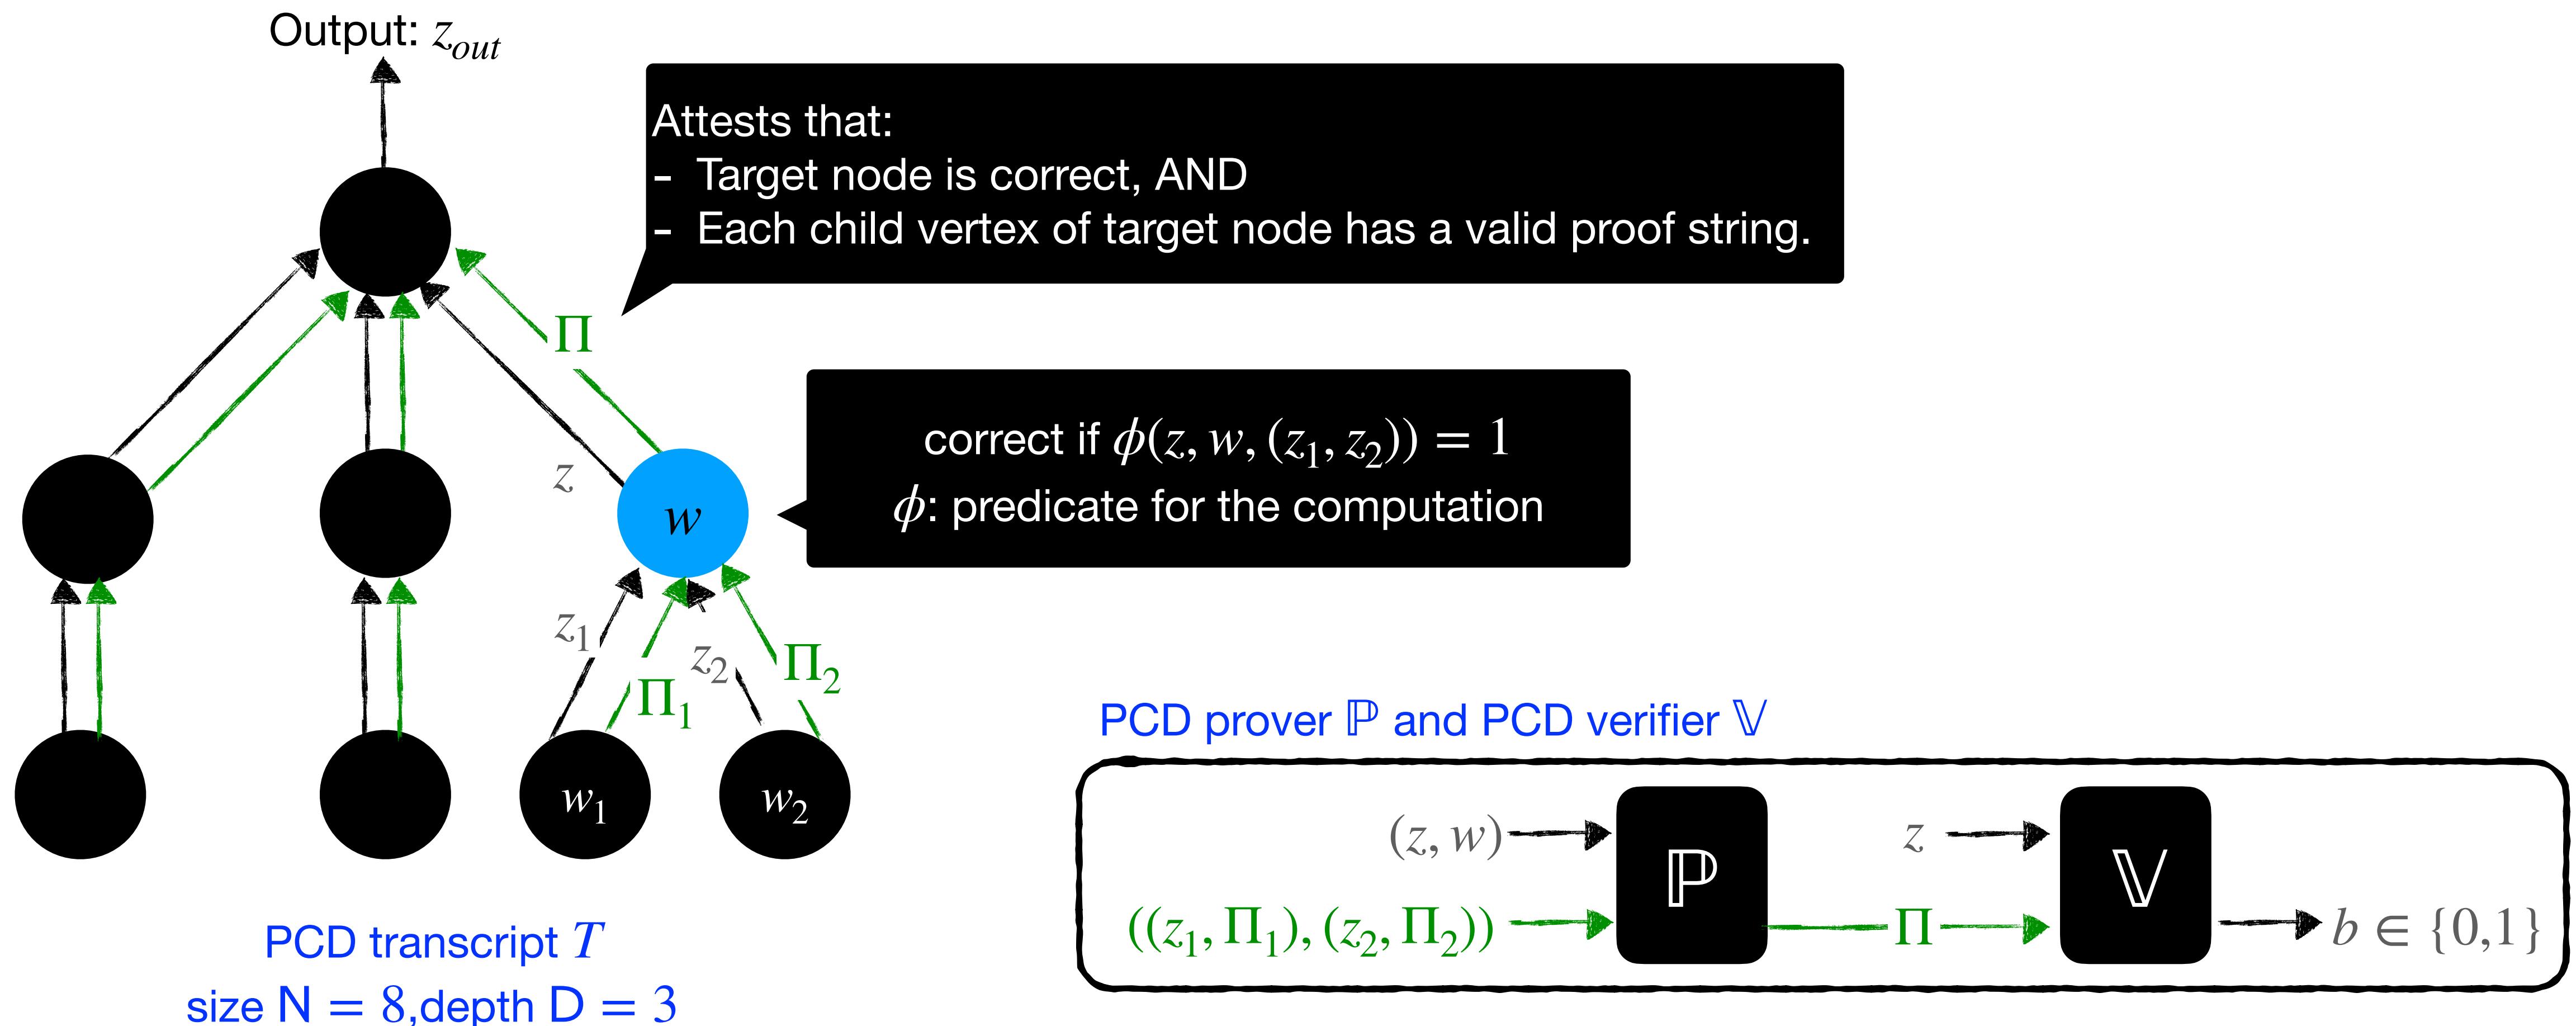
PCD prover \mathbb{P} and PCD verifier \mathbb{V}



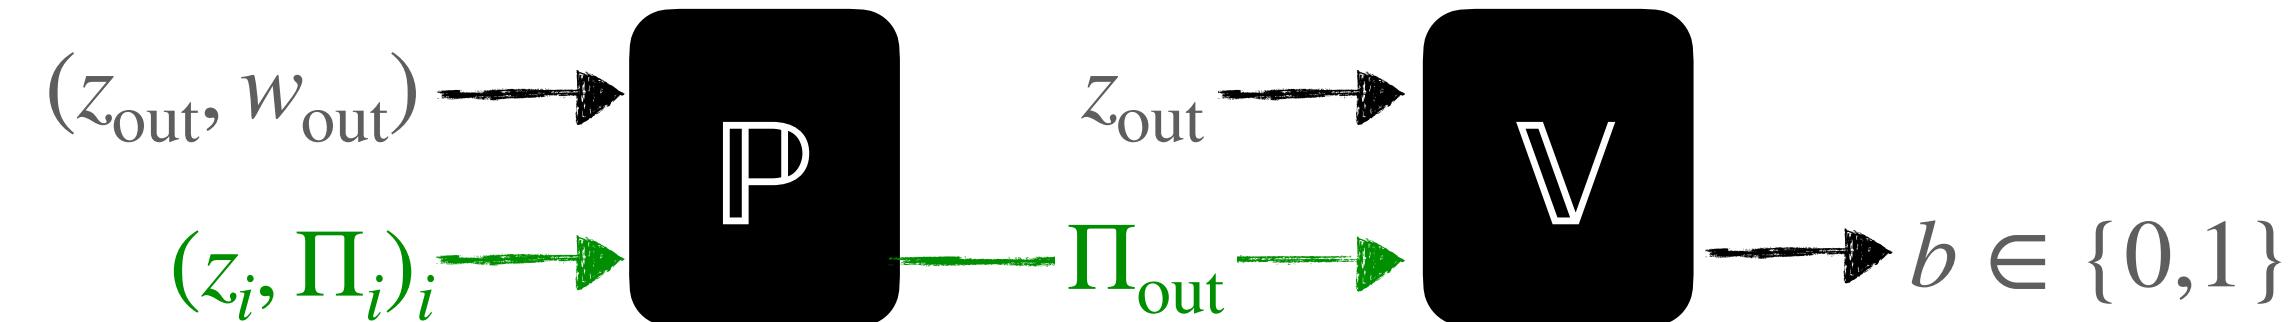
What is proof-carrying data (PCD)? [2/2]

Proof-carrying data (PCD)

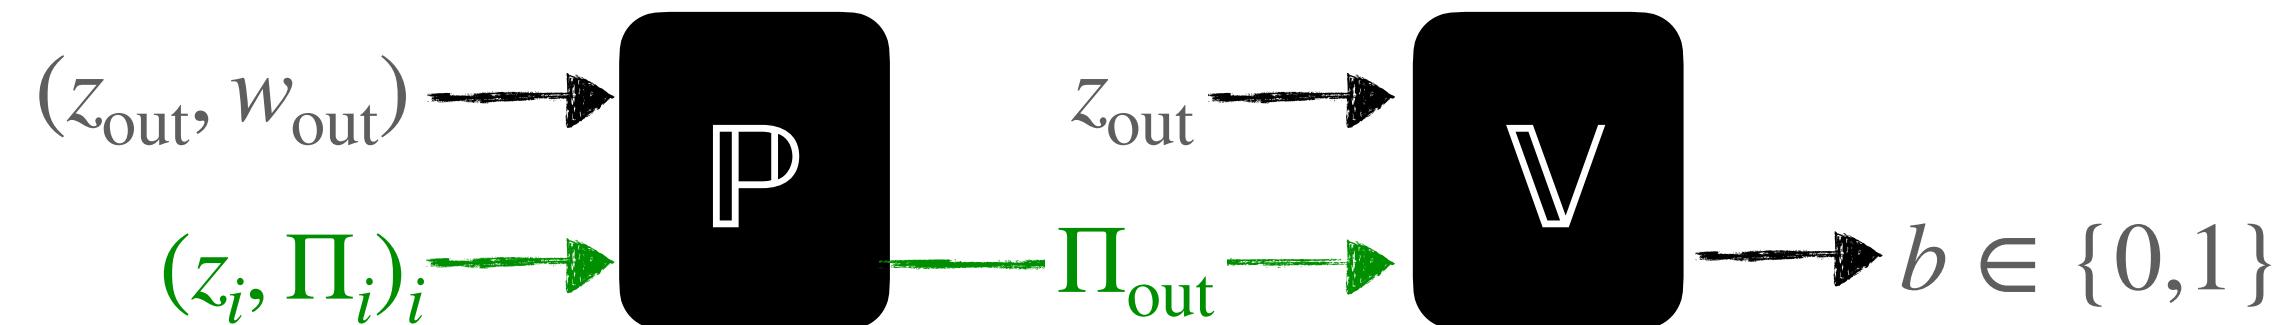
- Enables mutually distrustful parties to perform a distributed computation
- The correctness of each step can be **verified efficiently**



Security guarantee of PCD



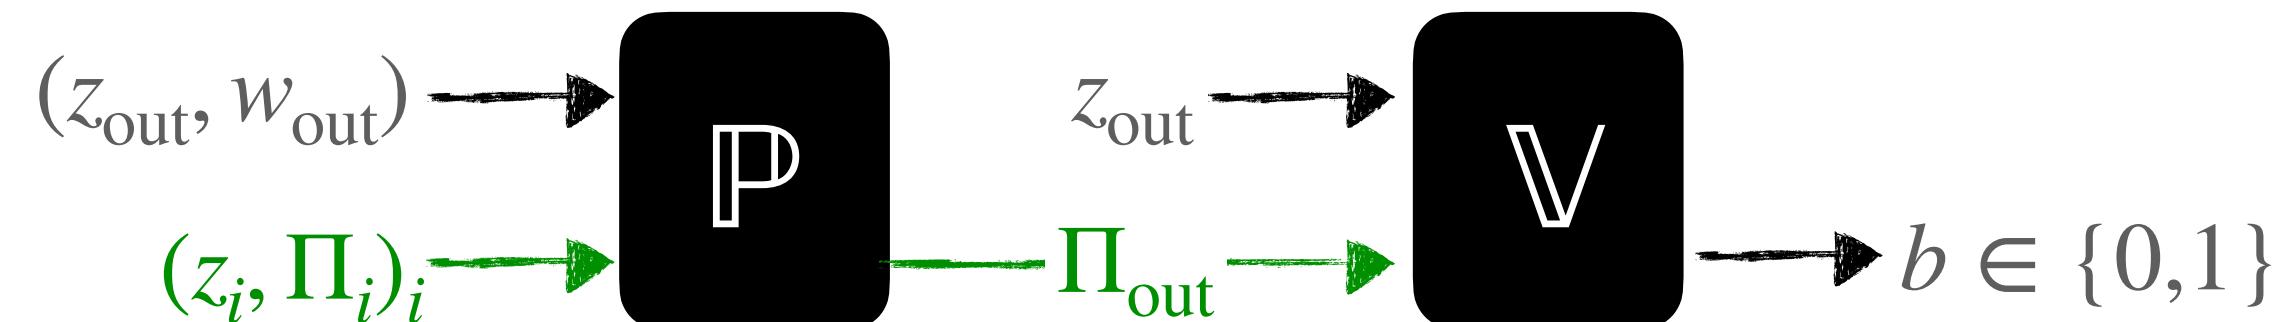
Security guarantee of PCD



T : computation transcript
 λ : security parameter
 D : maximum transcript depth
 N : maximum transcript size

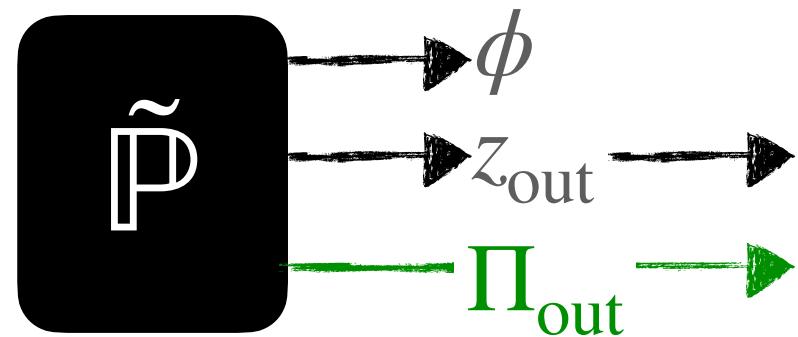
Knowledge soundness: \forall bounded $\tilde{\mathbb{P}}$, \exists an efficient extractor $\mathbb{E}_{\tilde{\mathbb{P}}}$ who outputs T such that the following happens with probability at most $\kappa(\lambda, D, N)$:

Security guarantee of PCD

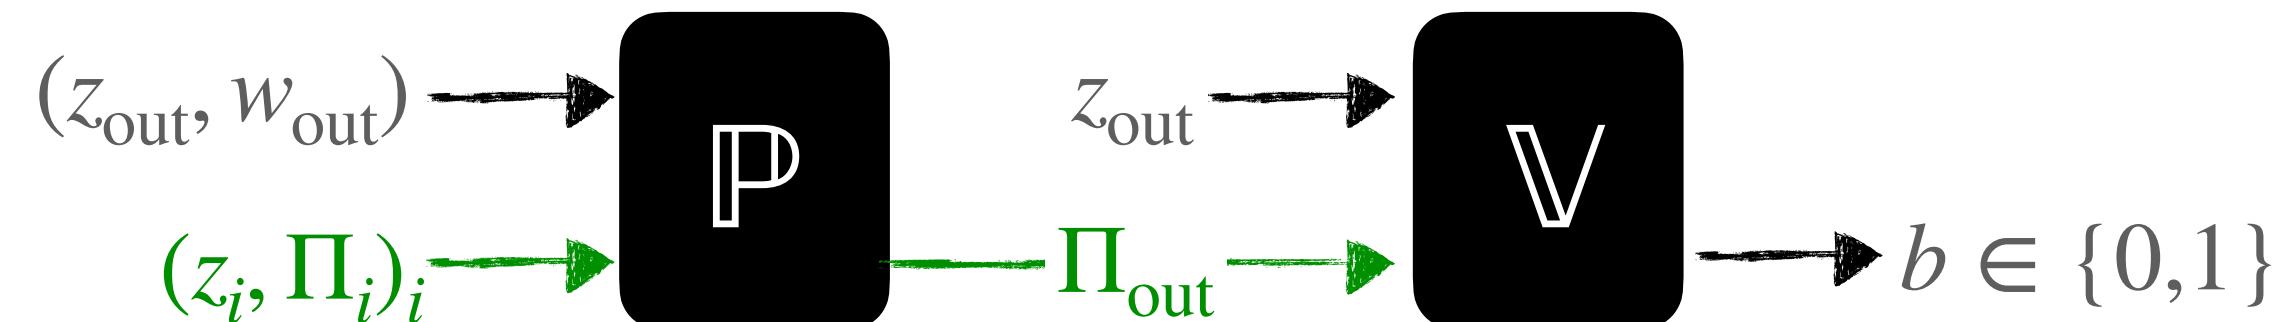


T : computation transcript
 λ : security parameter
 D : maximum transcript depth
 N : maximum transcript size

Knowledge soundness: \forall bounded $\tilde{\mathbb{P}}$, \exists an efficient extractor $\mathbb{E}_{\tilde{\mathbb{P}}}$ who outputs T such that the following happens with probability at most $\kappa(\lambda, D, N)$:

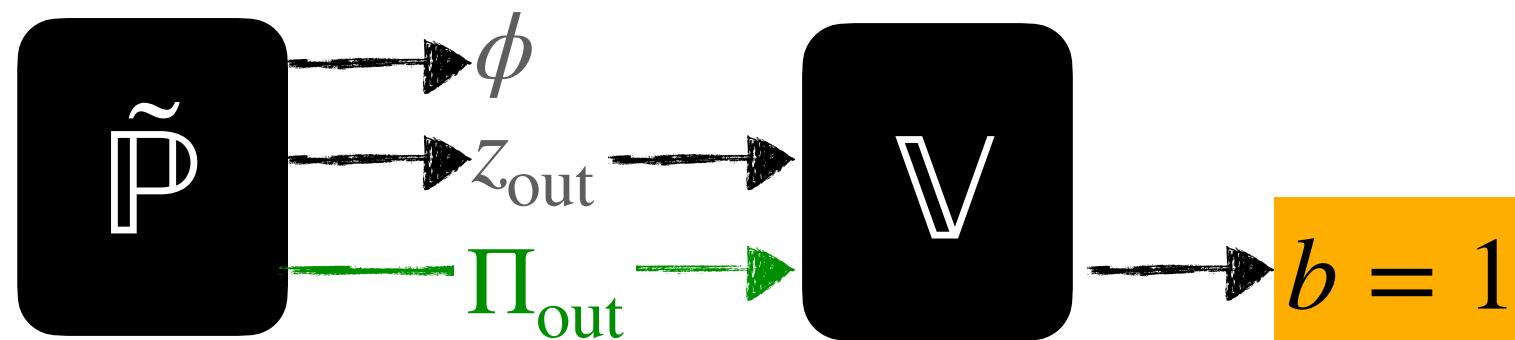


Security guarantee of PCD

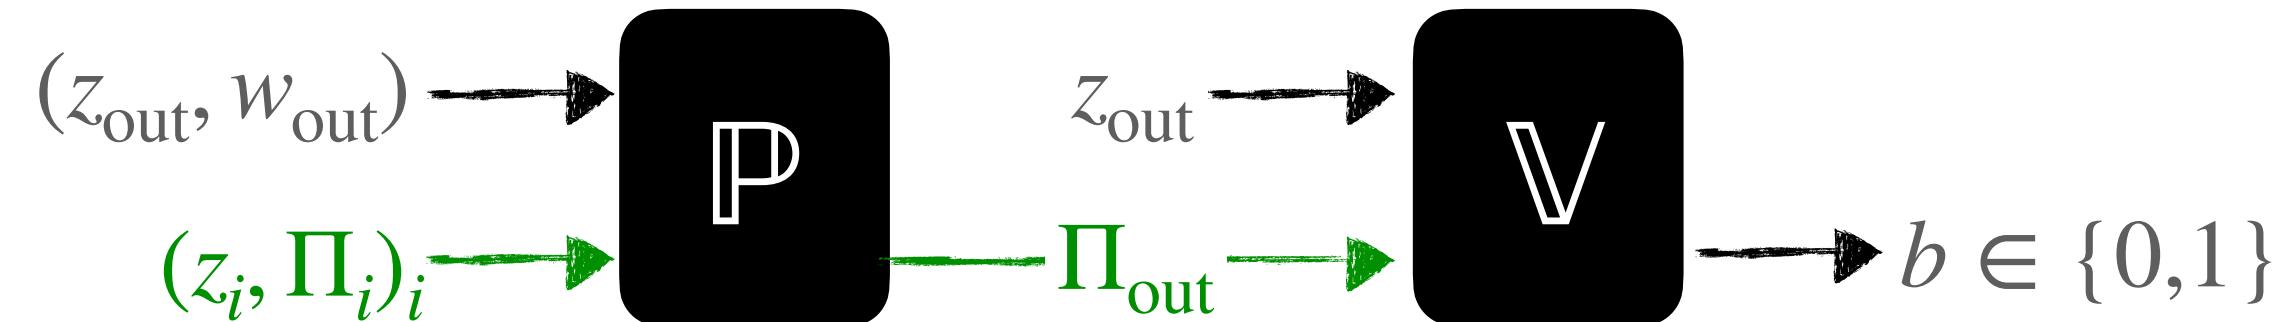


T : computation transcript
 λ : security parameter
 D : maximum transcript depth
 N : maximum transcript size

Knowledge soundness: \forall bounded \tilde{P} , \exists an efficient extractor $E_{\tilde{P}}$ who outputs T such that the following happens with probability at most $\kappa(\lambda, D, N)$:

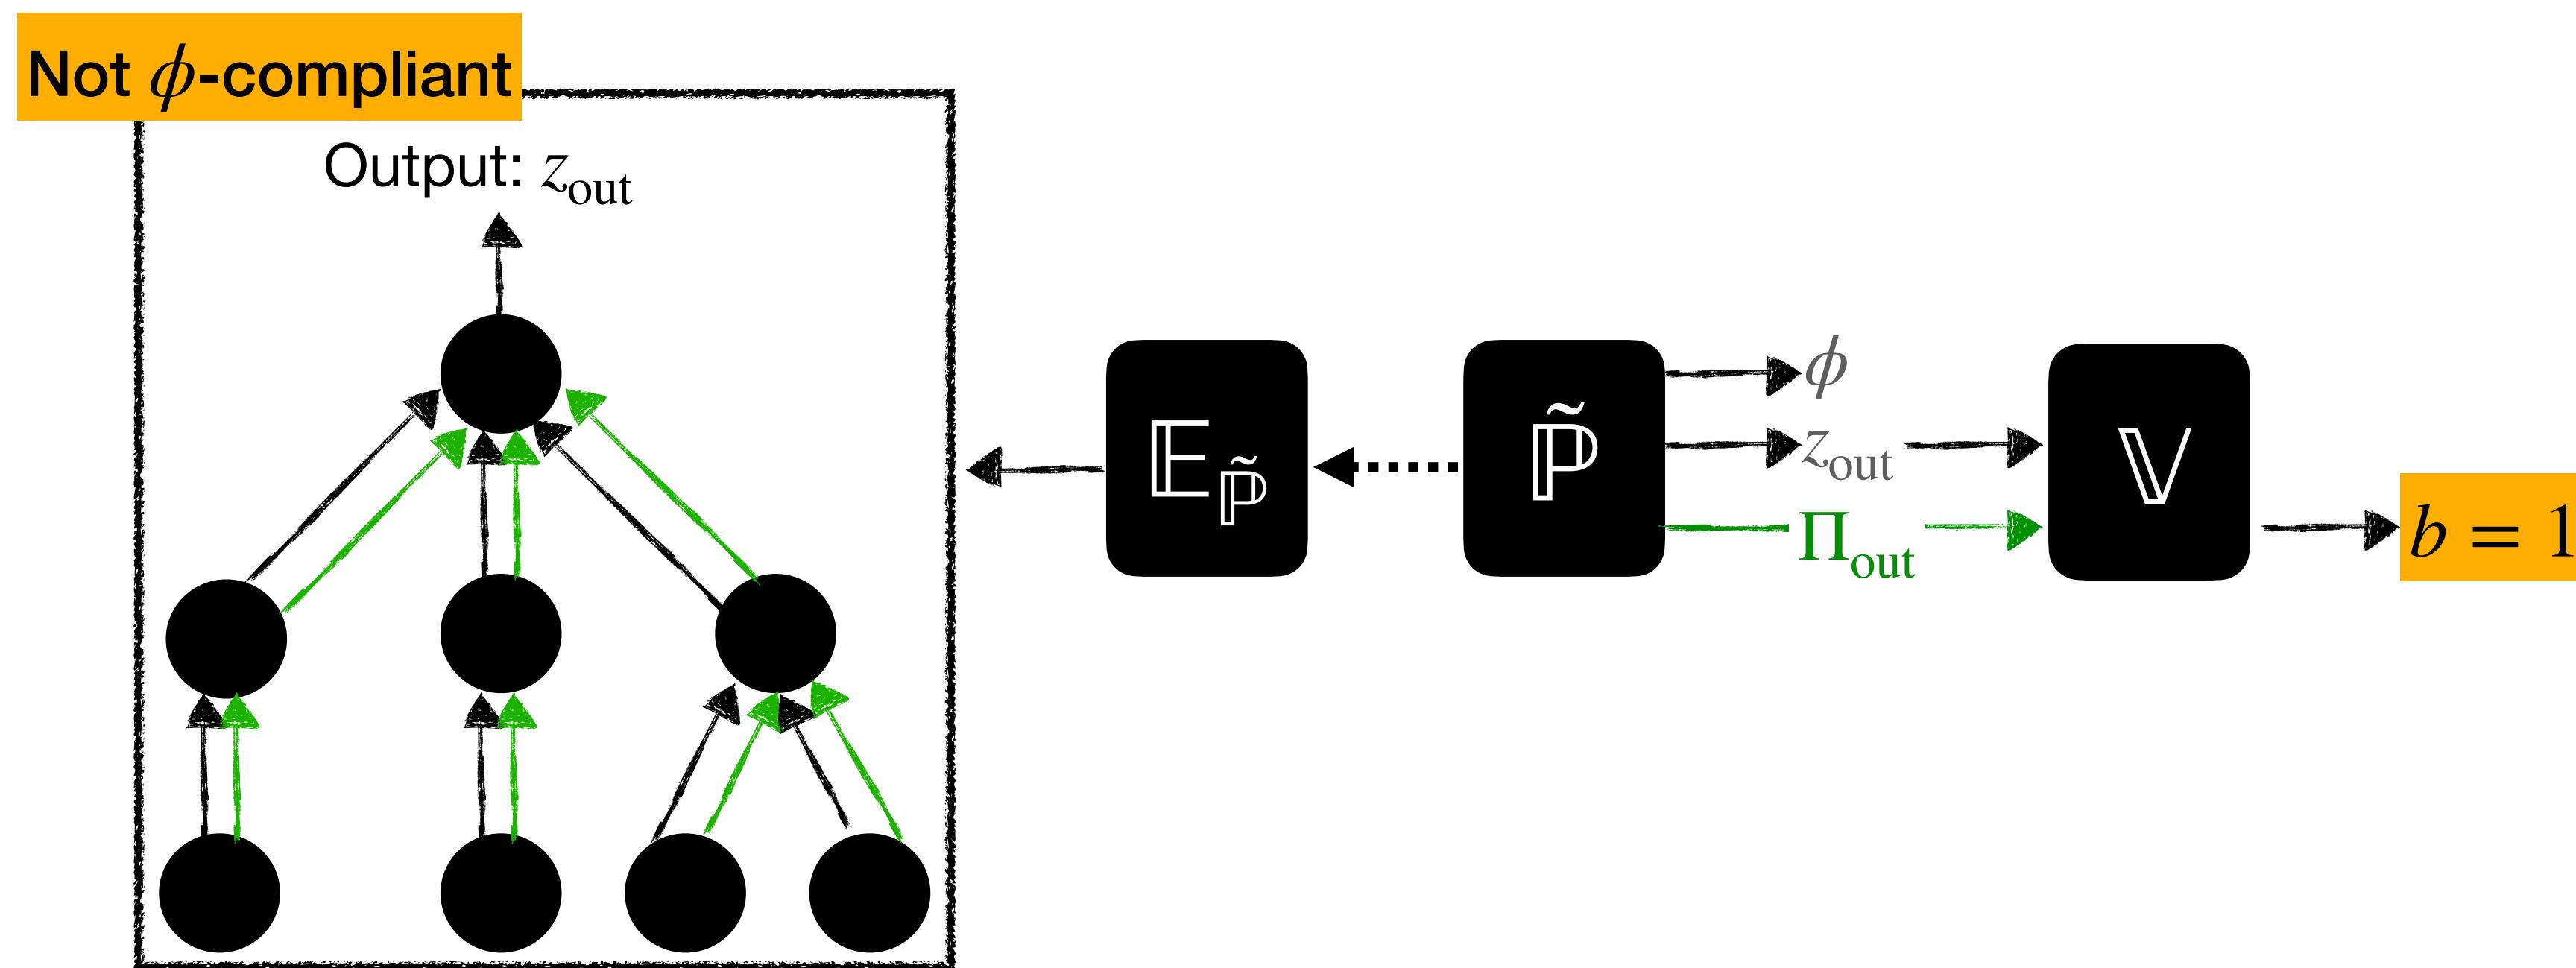


Security guarantee of PCD



T : computation transcript
 λ : security parameter
 D : maximum transcript depth
 N : maximum transcript size

Knowledge soundness: \forall bounded \tilde{P} , \exists an efficient extractor $E_{\tilde{P}}$ who outputs T such that the following happens with probability at most $\kappa(\lambda, D, N)$:



Review: SNARK

Review: SNARK

PCD can be constructed from a SNARK (e.g., for CSAT).

Review: SNARK

PCD can be constructed from a SNARK (e.g., for CSAT).

$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$

Review: SNARK

PCD can be constructed from a SNARK (e.g., for CSAT).

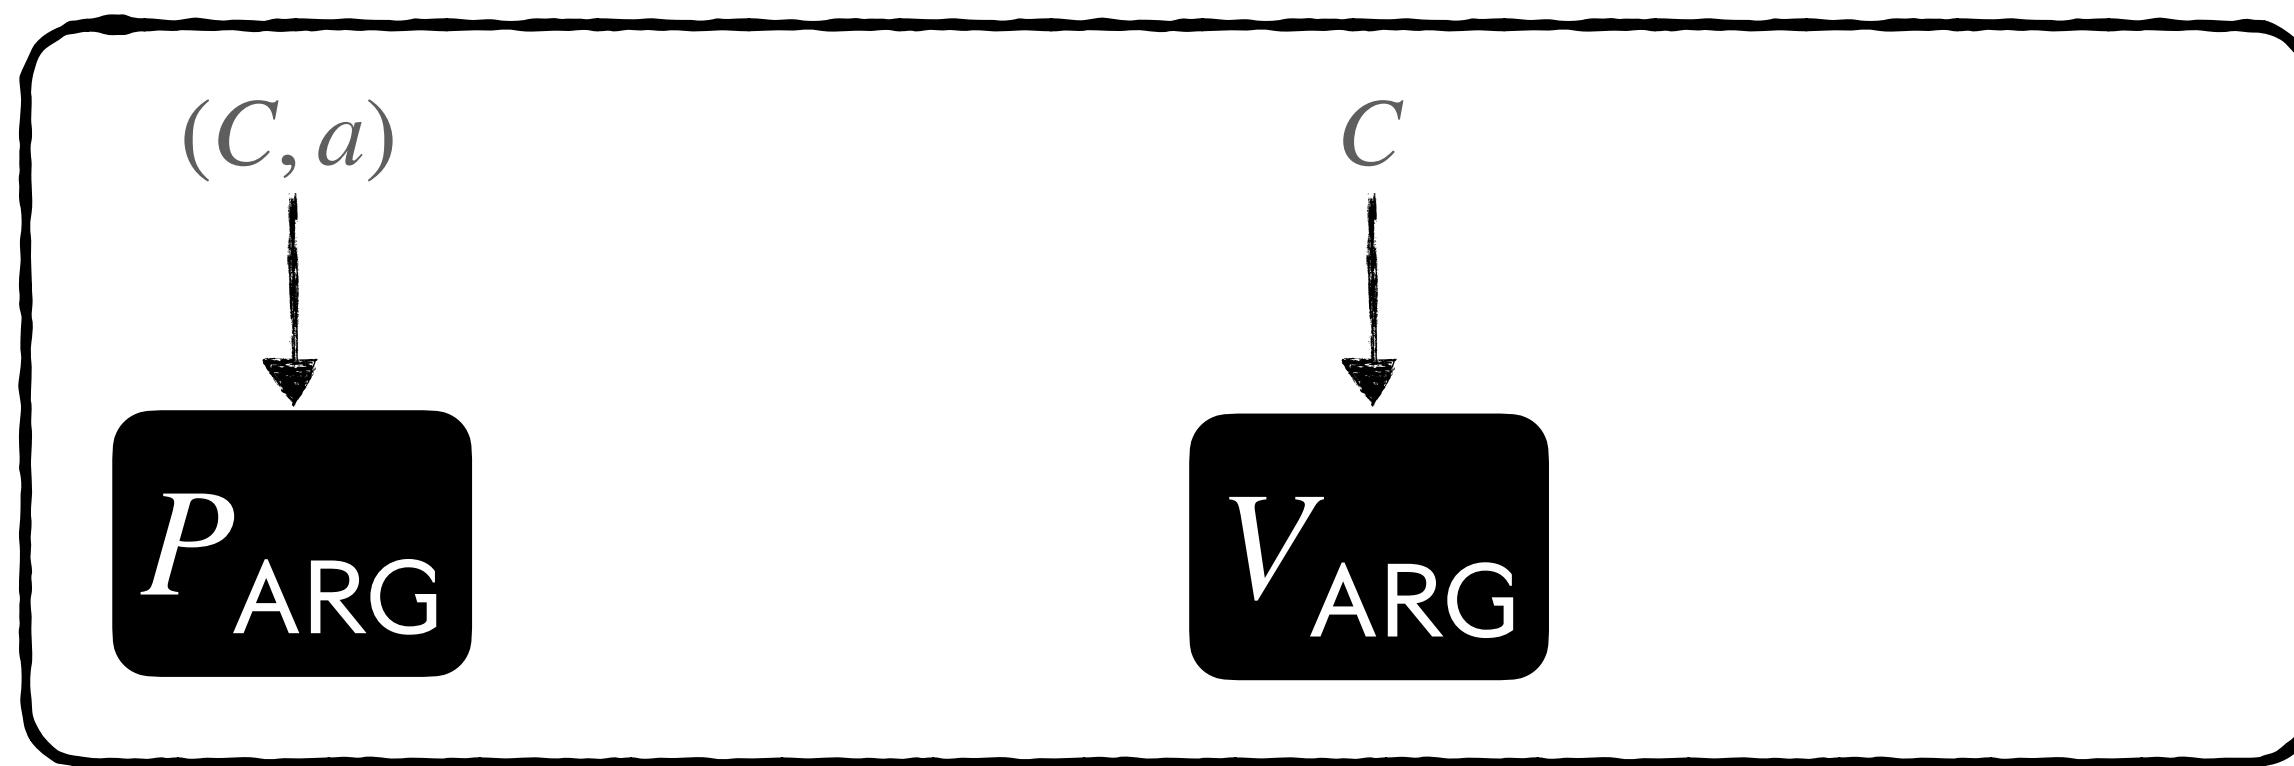
$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$



Review: SNARK

PCD can be constructed from a SNARK (e.g., for CSAT).

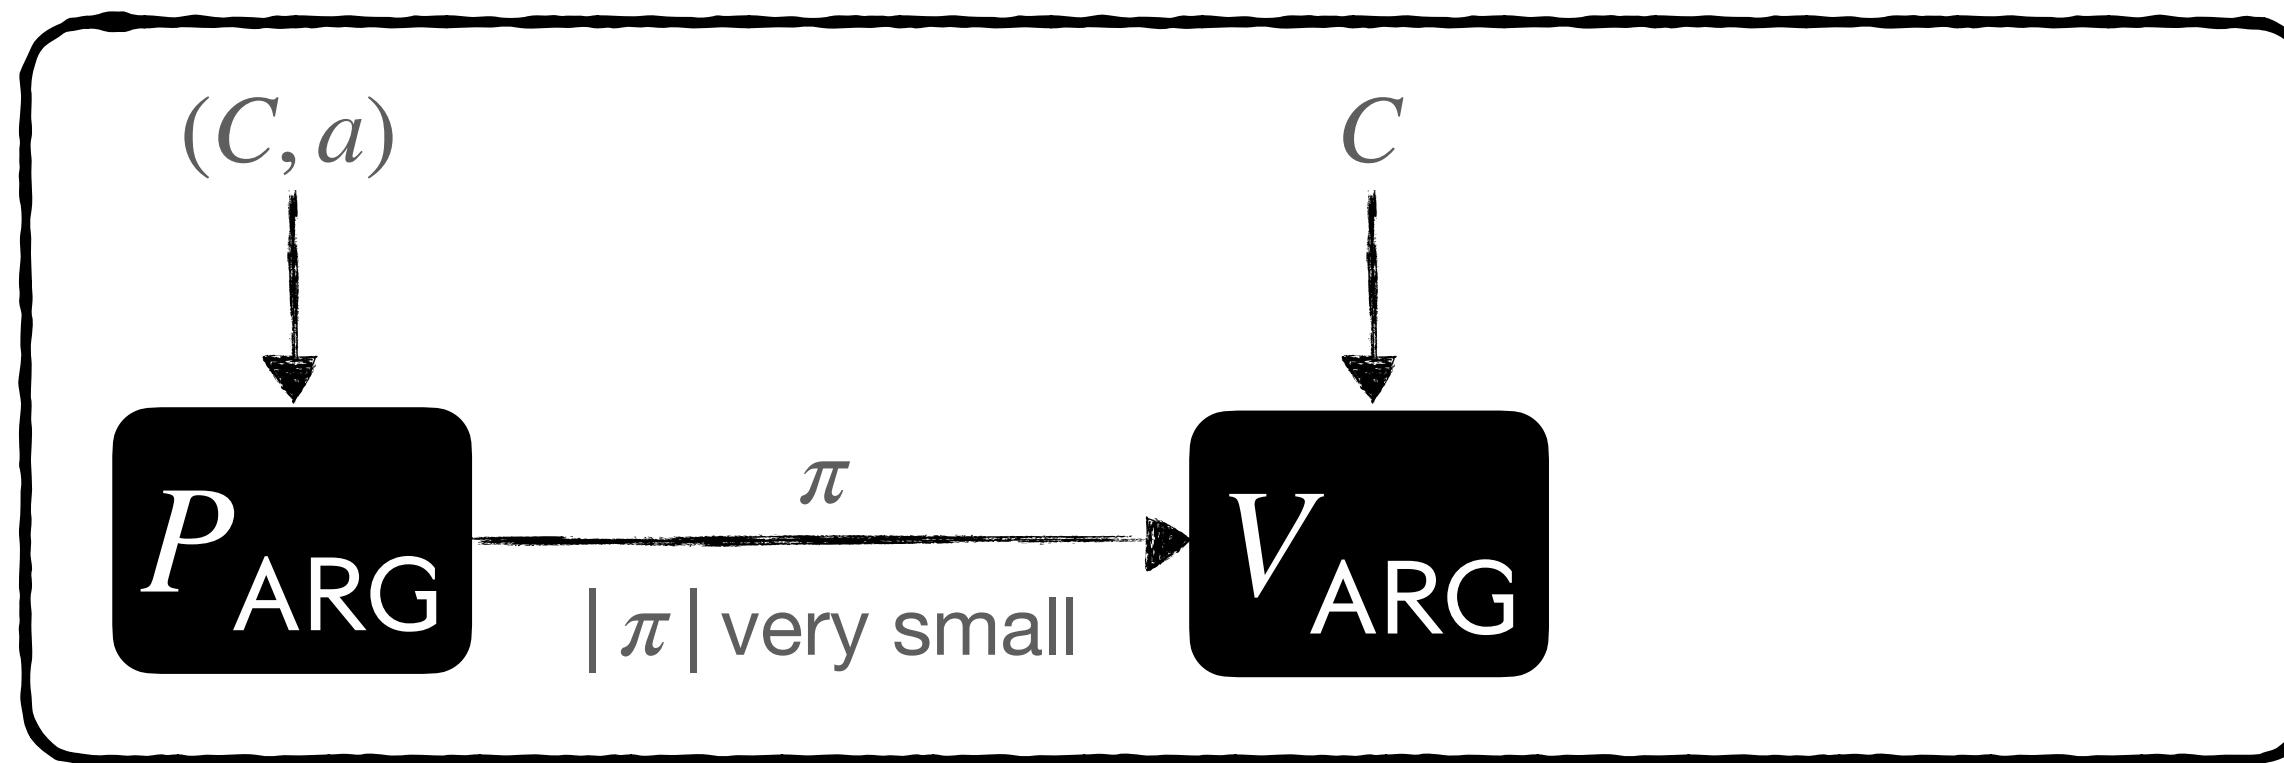
$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$



Review: SNARK

PCD can be constructed from a SNARK (e.g., for CSAT).

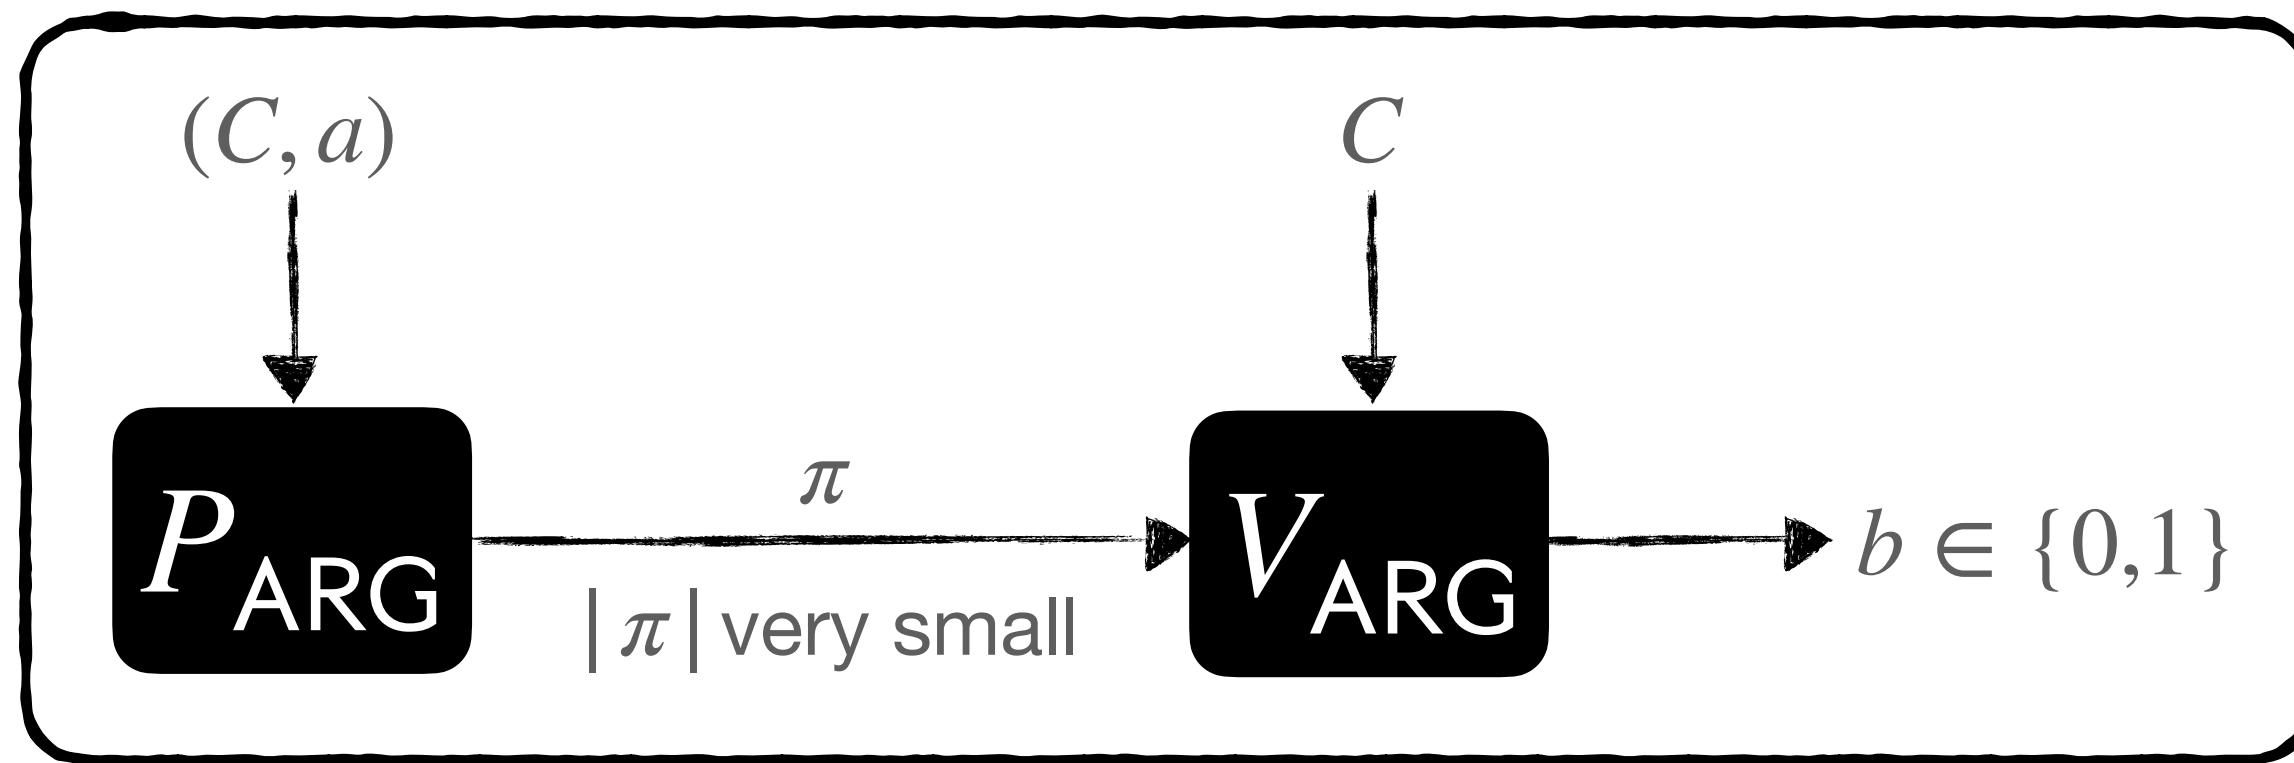
$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$



Review: SNARK

PCD can be constructed from a SNARK (e.g., for CSAT).

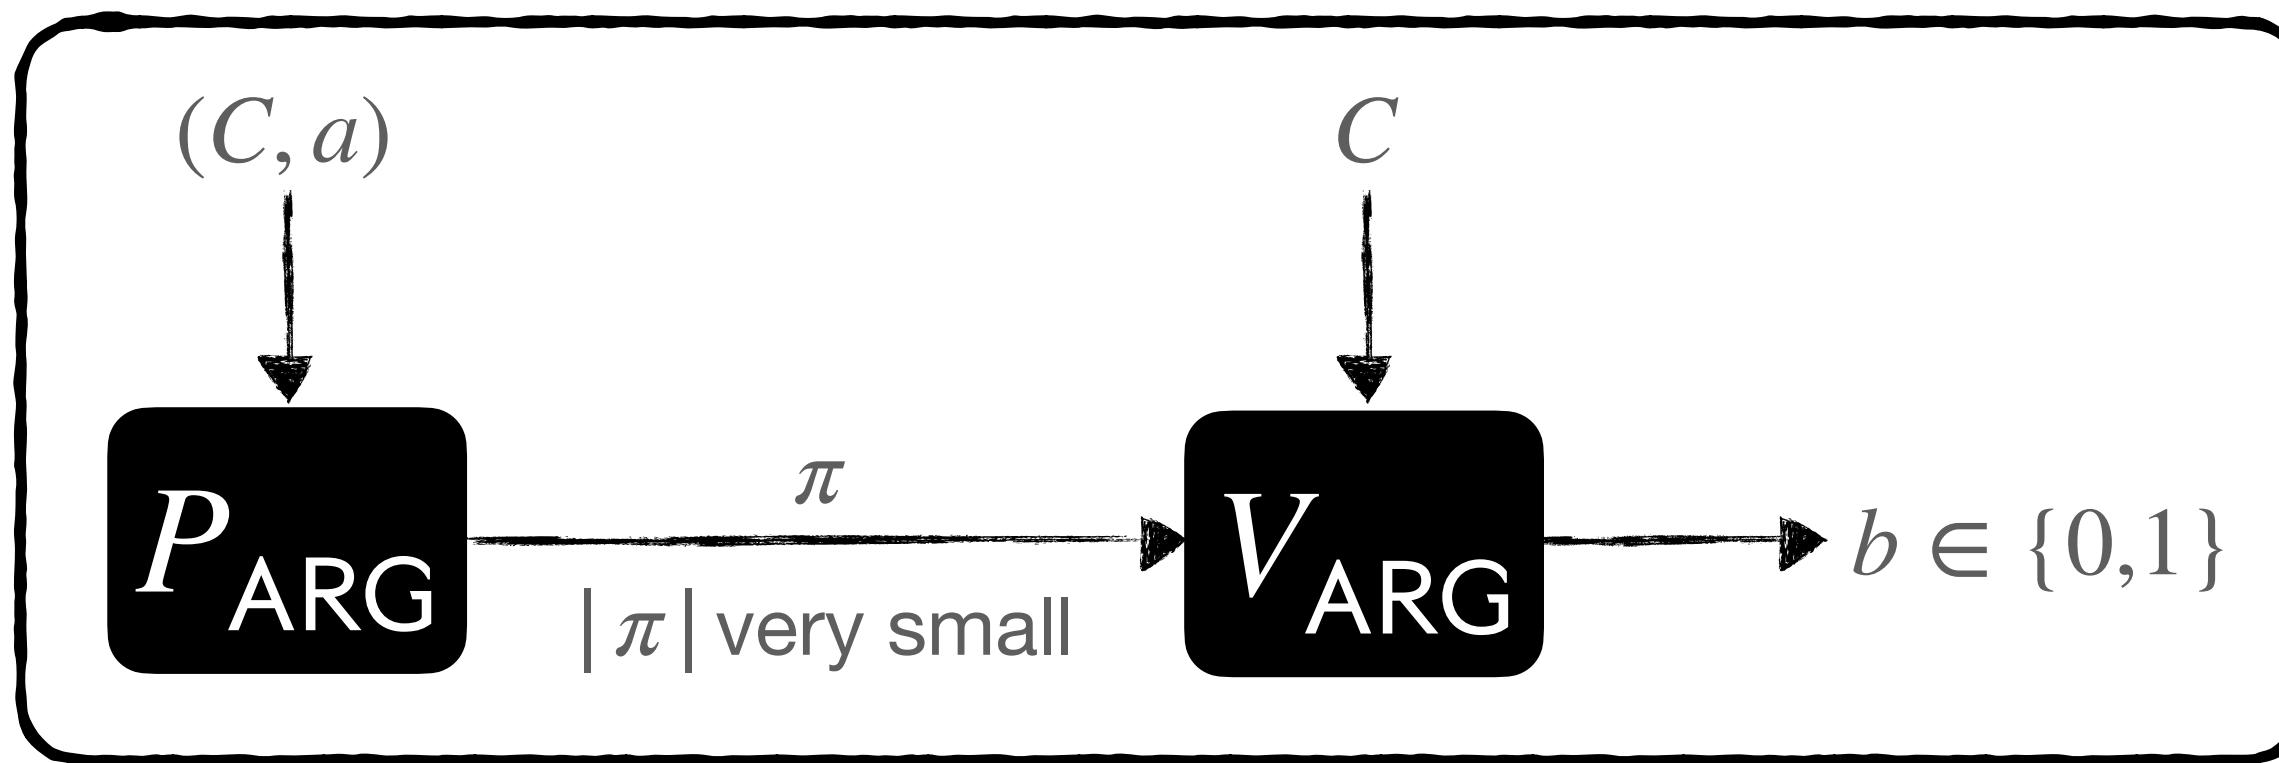
$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$



Review: SNARK

PCD can be constructed from a SNARK (e.g., for CSAT).

$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$



λ : security parameter

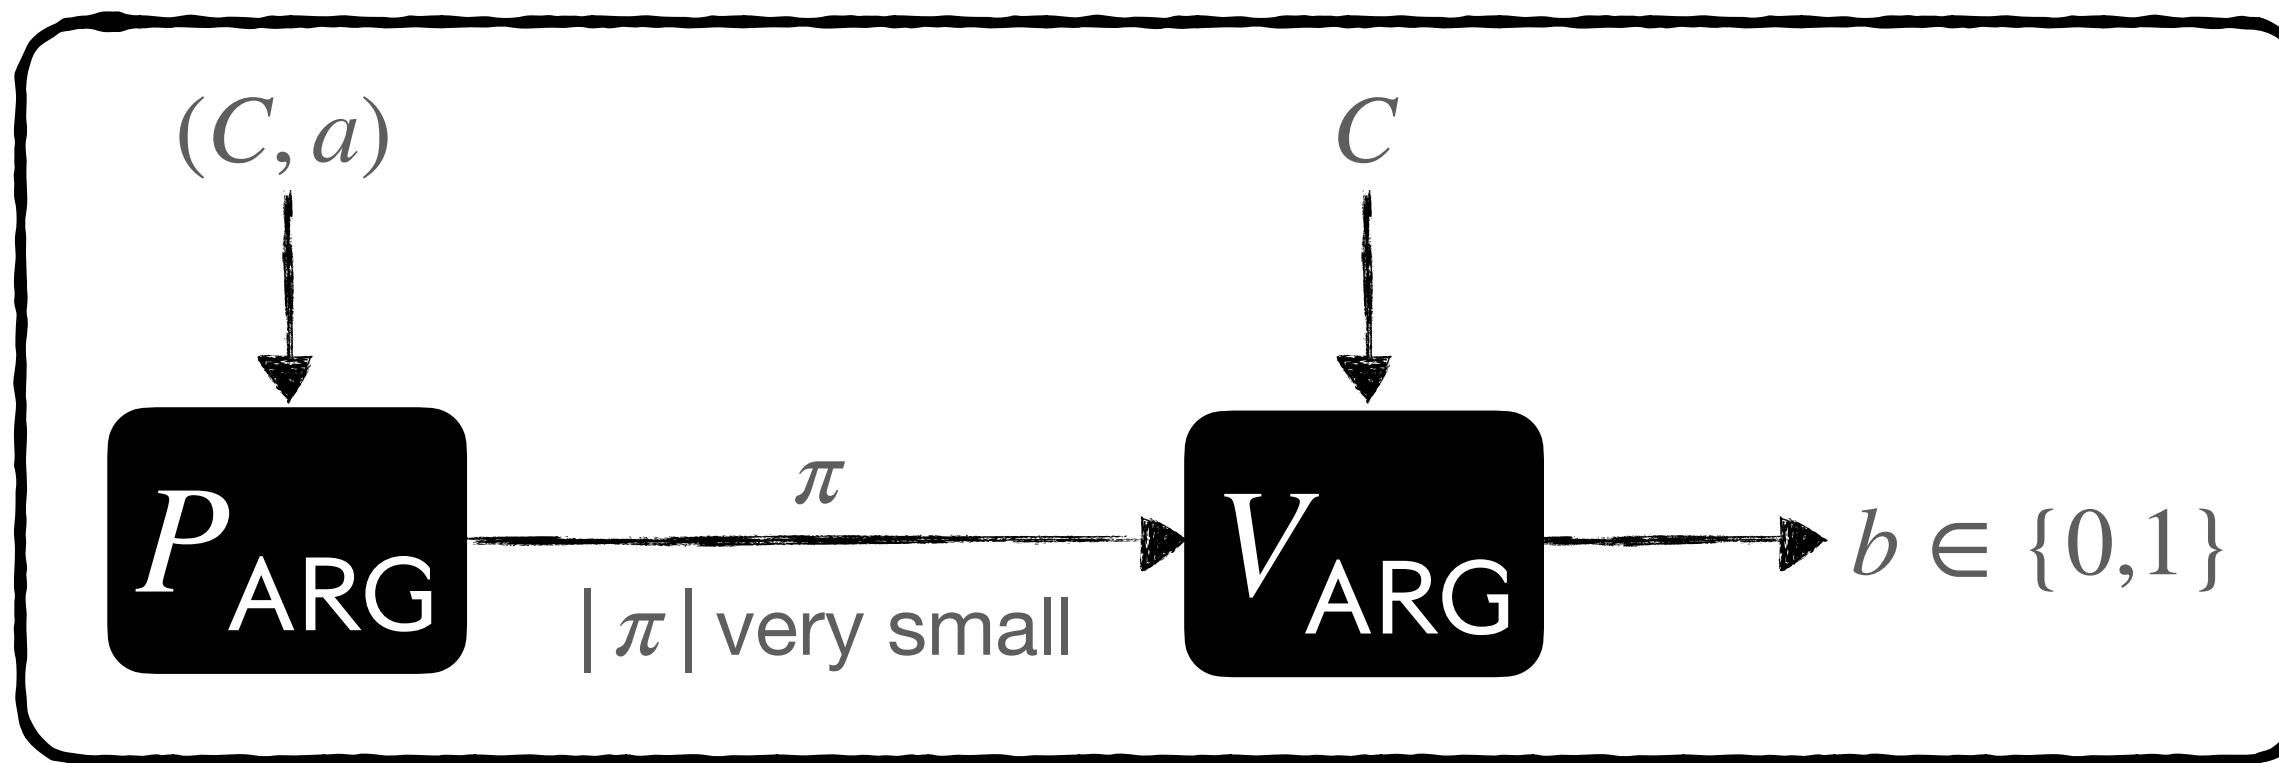
n : upper bound on size of C

Knowledge soundness: \forall bounded \tilde{P}_{ARG} , \exists an efficient extractor $E_{\tilde{P}_{\text{ARG}}}$ who outputs an assignment a such that the following happens with probability at most $\kappa_{\text{ARG}}(\lambda, n)$:

Review: SNARK

PCD can be constructed from a SNARK (e.g., for CSAT).

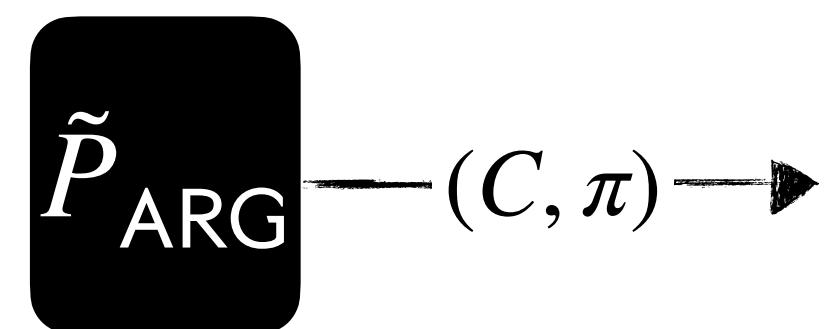
$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$



λ : security parameter

n : upper bound on size of C

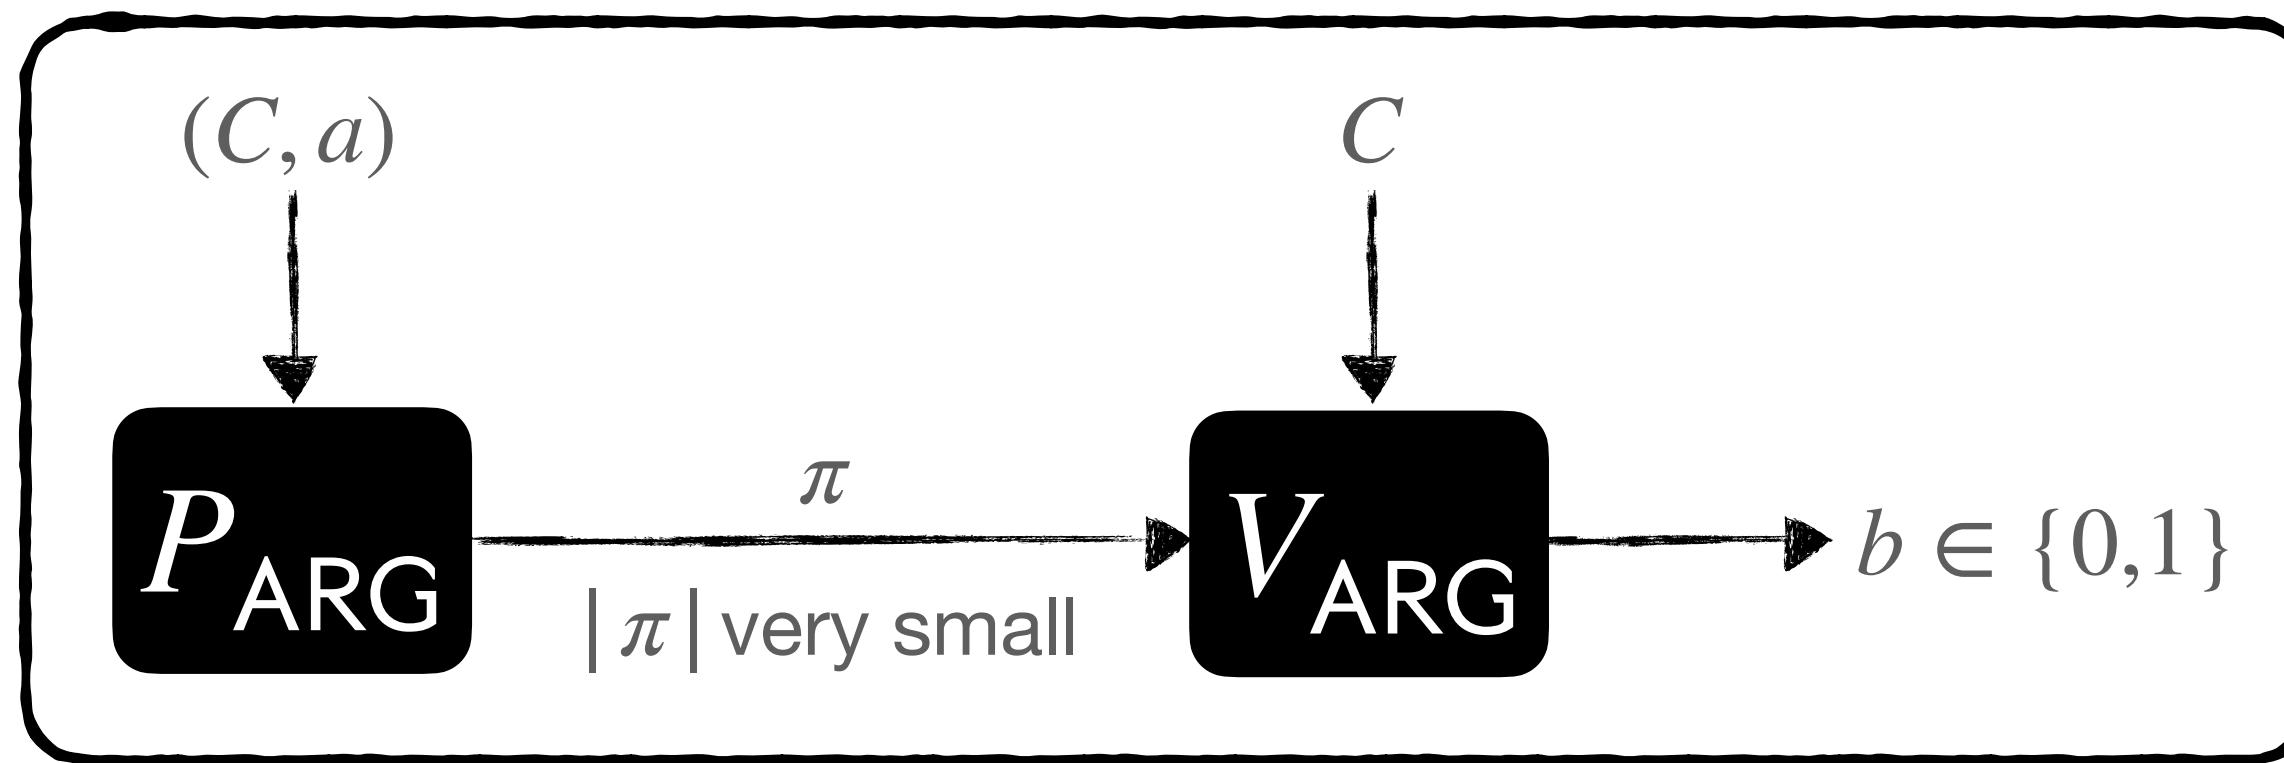
Knowledge soundness: \forall bounded \tilde{P}_{ARG} , \exists an efficient extractor $E_{\tilde{P}_{\text{ARG}}}$ who outputs an assignment a such that the following happens with probability at most $\kappa_{\text{ARG}}(\lambda, n)$:



Review: SNARK

PCD can be constructed from a SNARK (e.g., for CSAT).

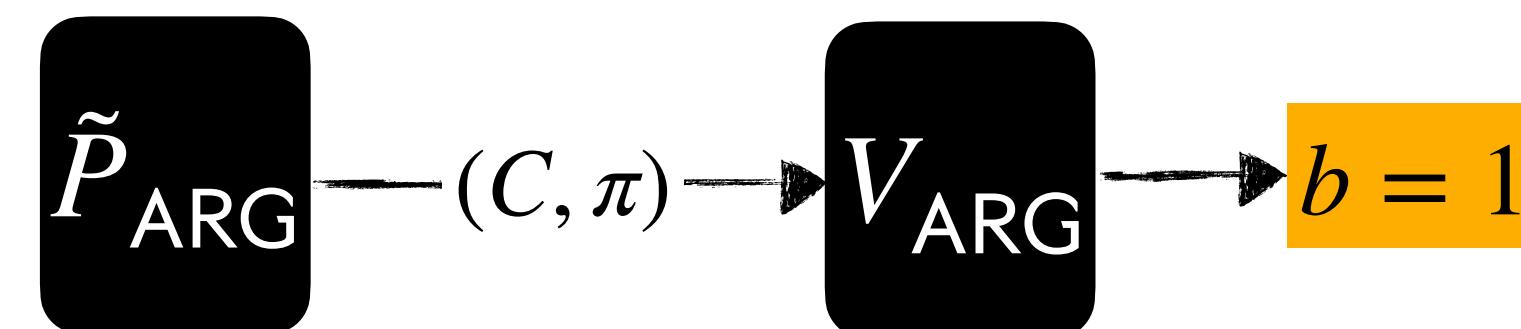
$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$



λ : security parameter

n : upper bound on size of C

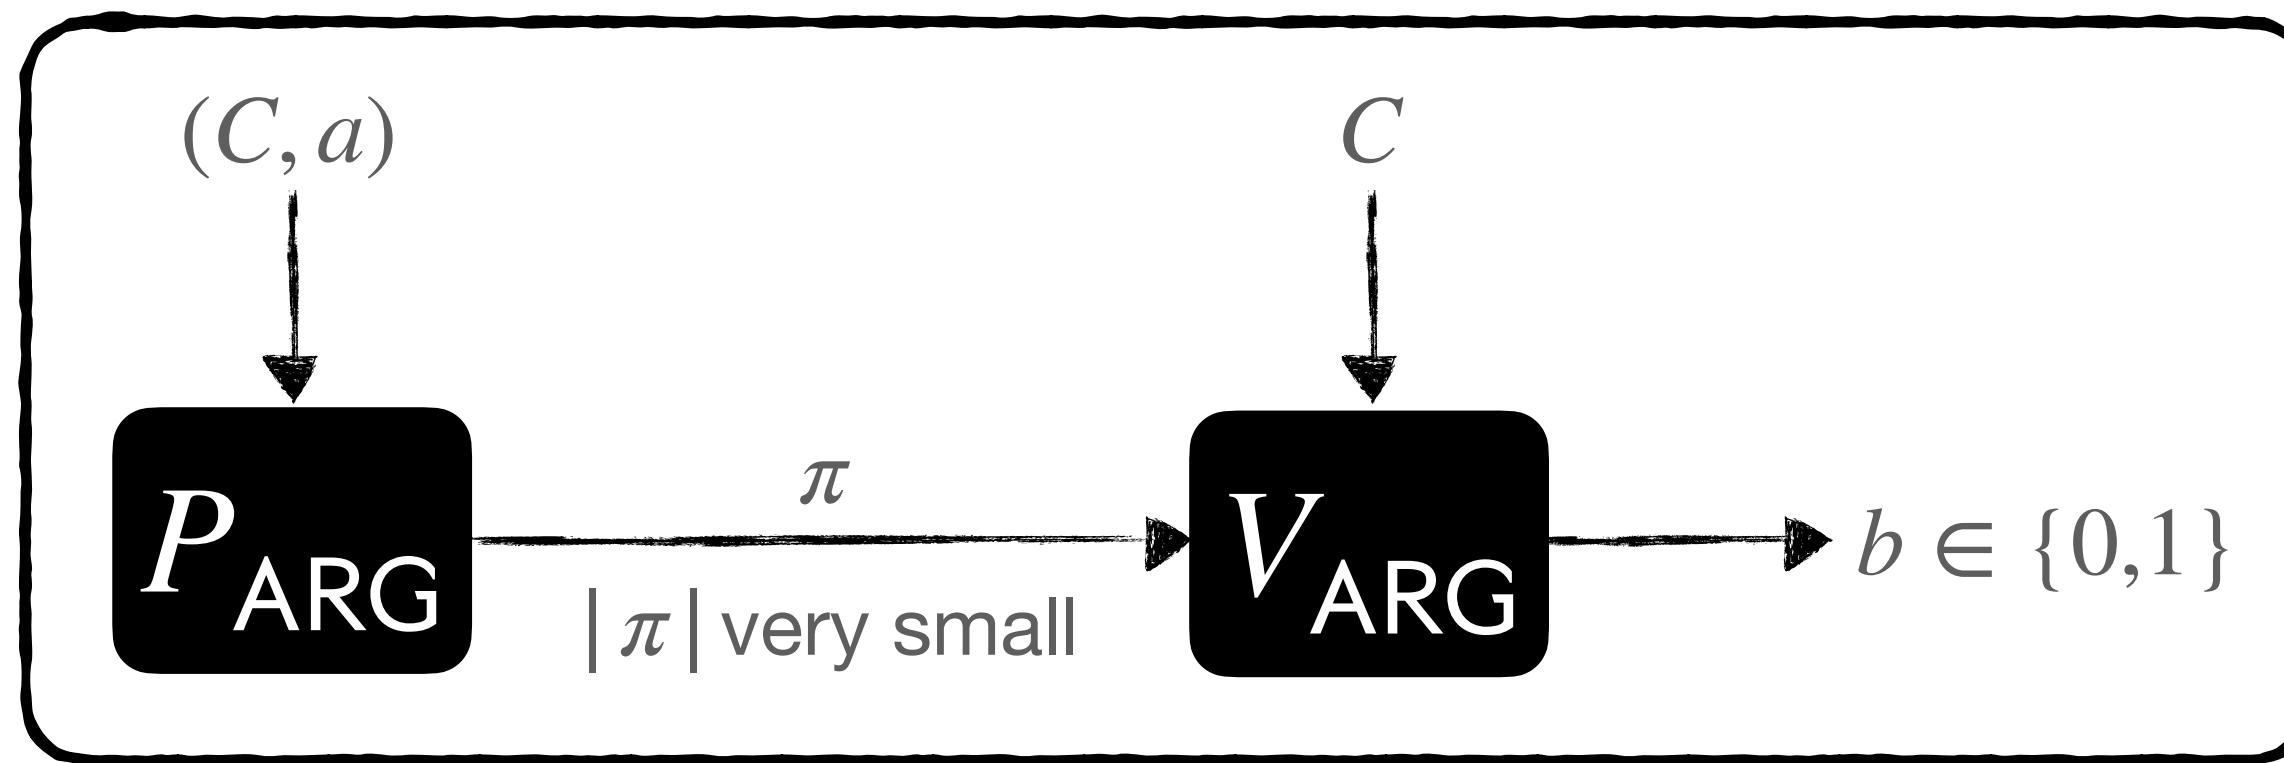
Knowledge soundness: \forall bounded \tilde{P}_{ARG} , \exists an efficient extractor $E_{\tilde{P}_{\text{ARG}}}$ who outputs an assignment a such that the following happens with probability at most $\kappa_{\text{ARG}}(\lambda, n)$:



Review: SNARK

PCD can be constructed from a SNARK (e.g., for CSAT).

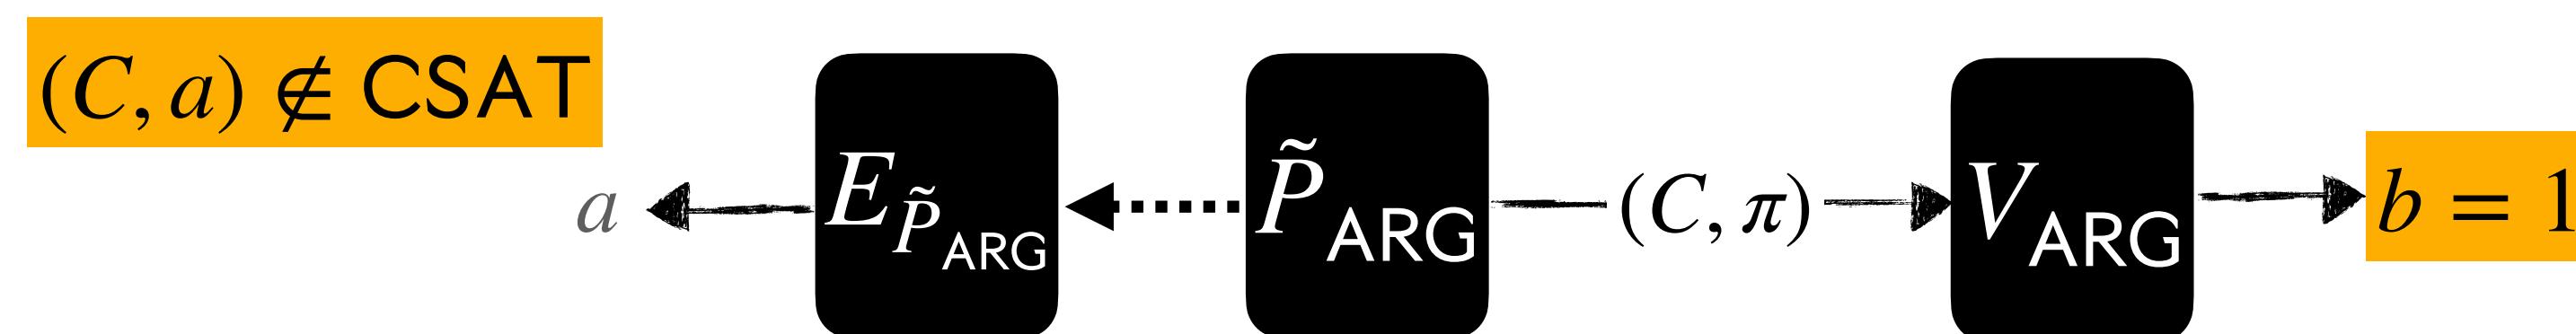
$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$



λ : security parameter

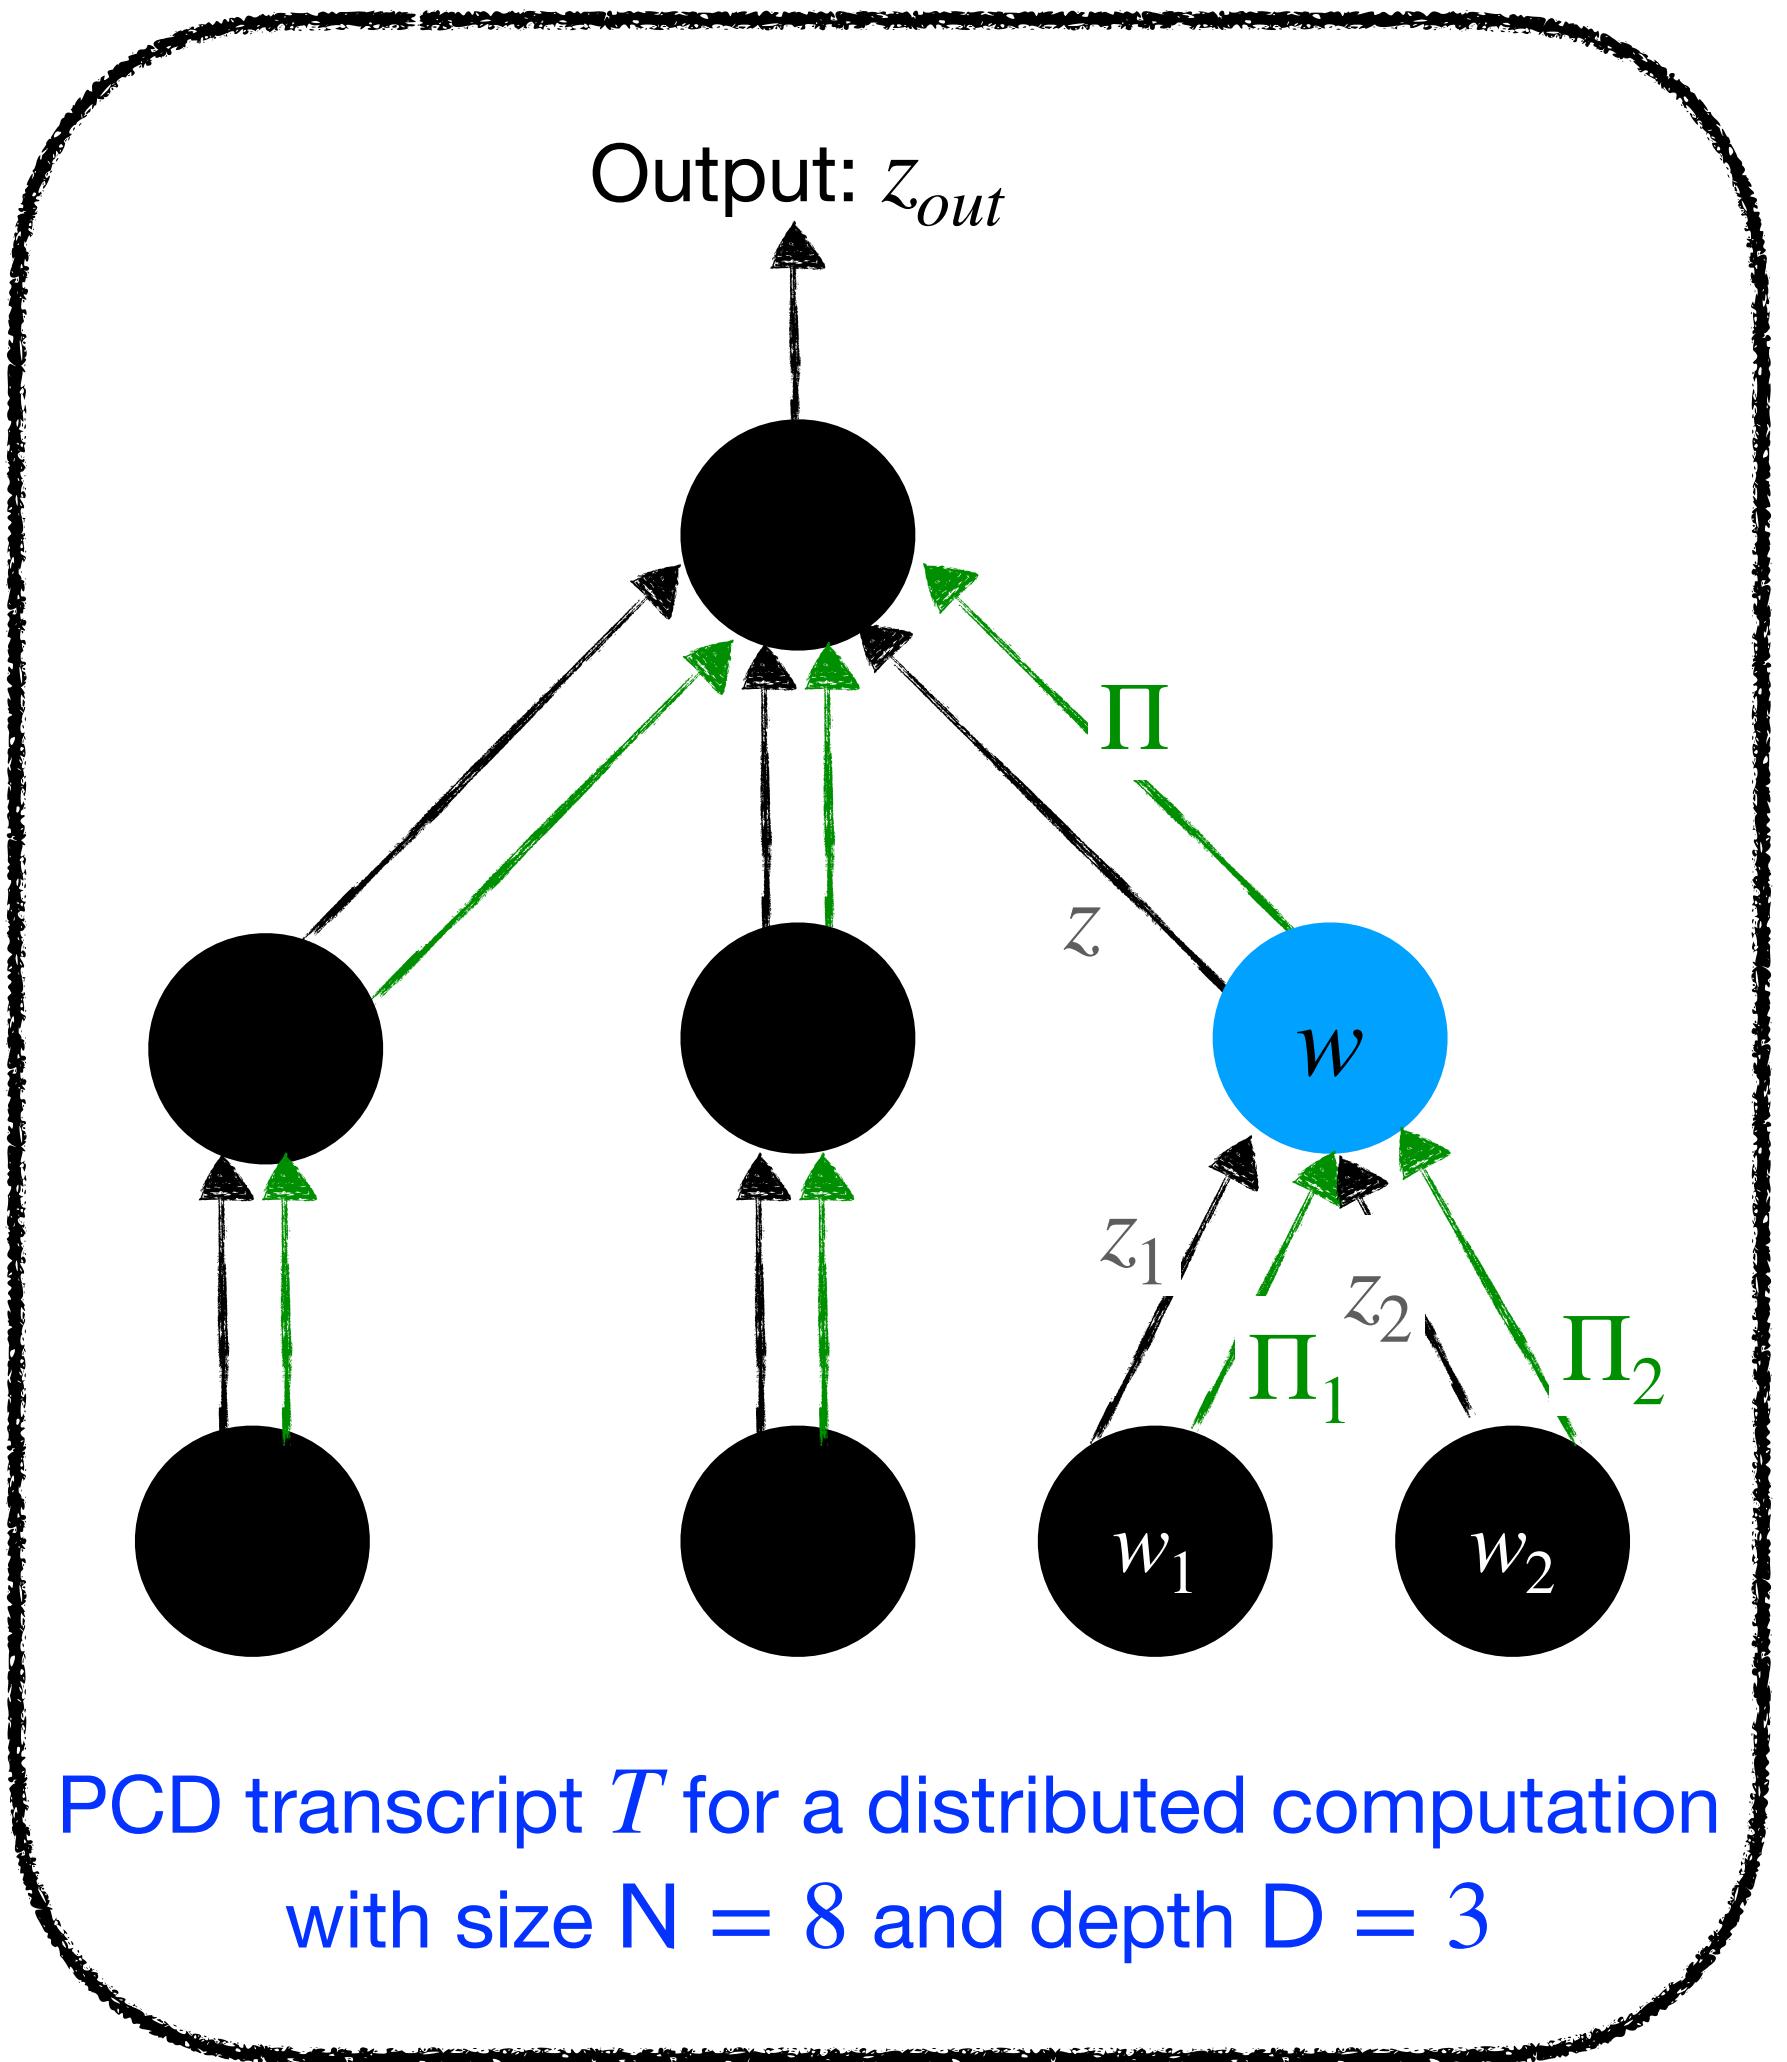
n : upper bound on size of C

Knowledge soundness: \forall bounded \tilde{P}_{ARG} , \exists an efficient extractor $E_{\tilde{P}_{\text{ARG}}}$ who outputs an assignment a such that the following happens with probability at most $\kappa_{\text{ARG}}(\lambda, n)$:



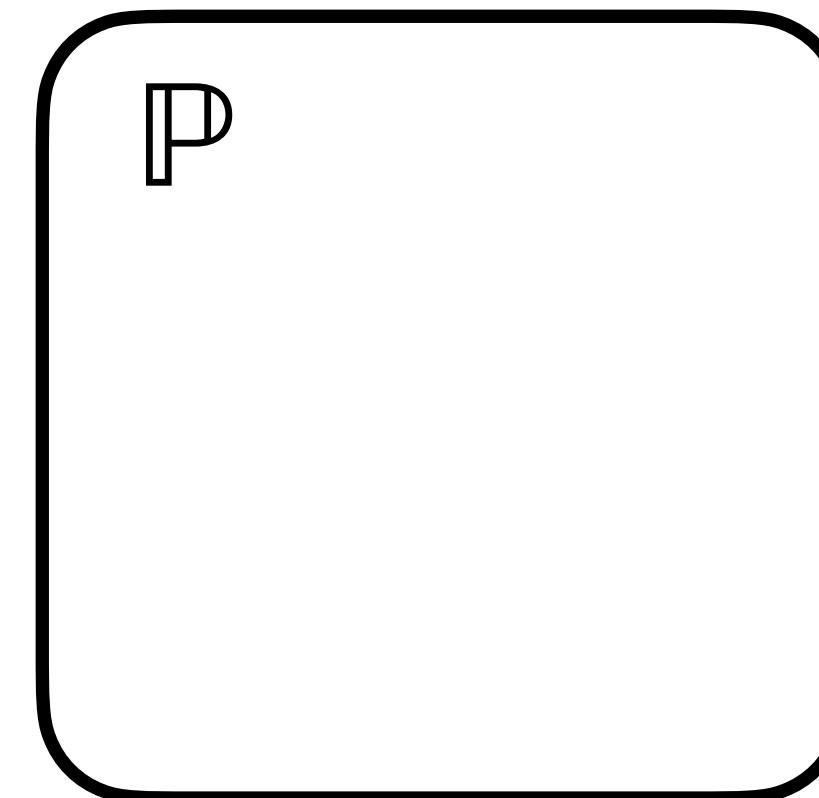
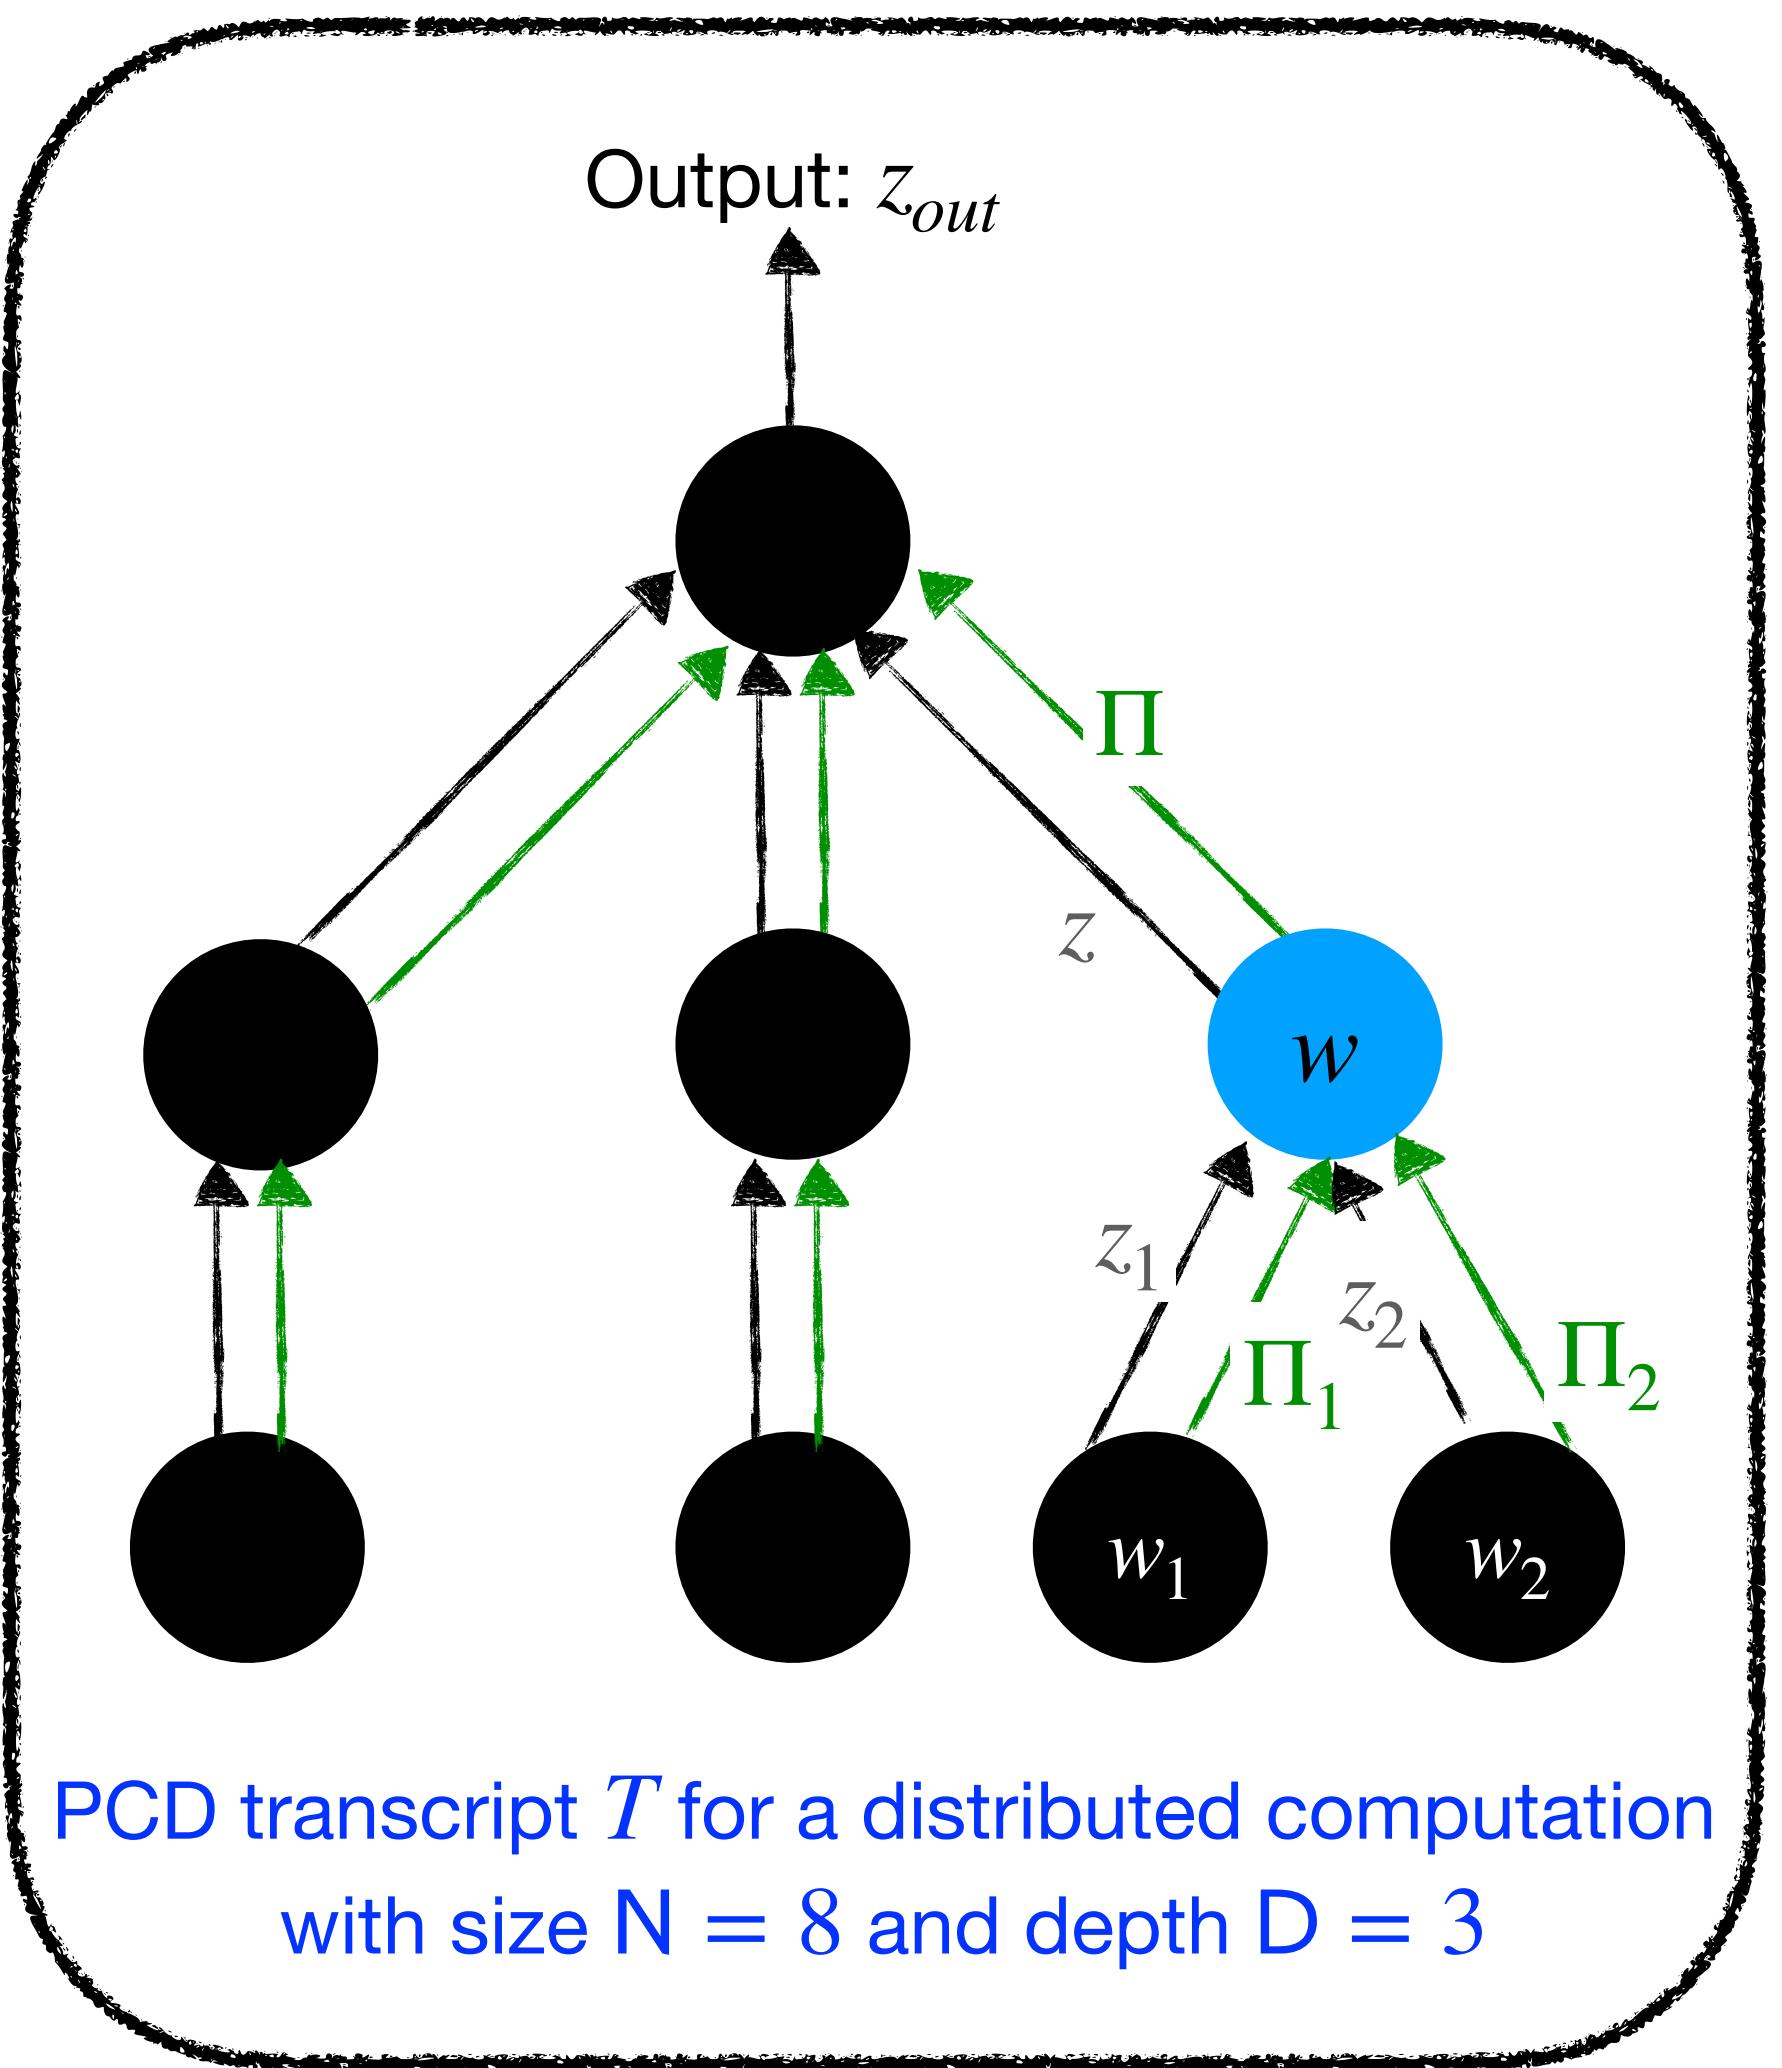
Naive PCD construction

Concatenate SNARK proofs



Naive PCD construction

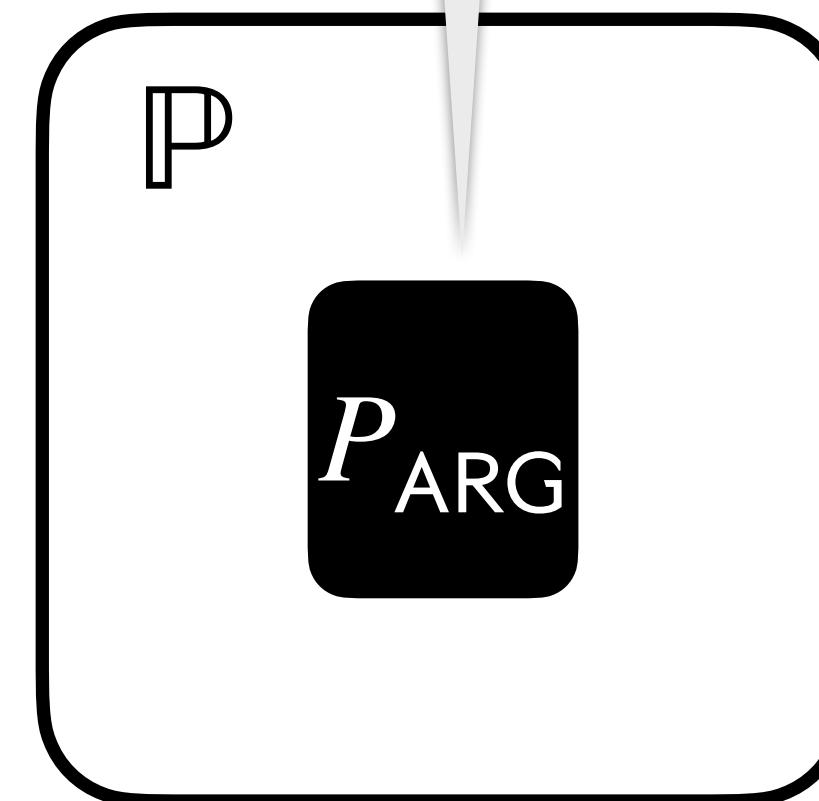
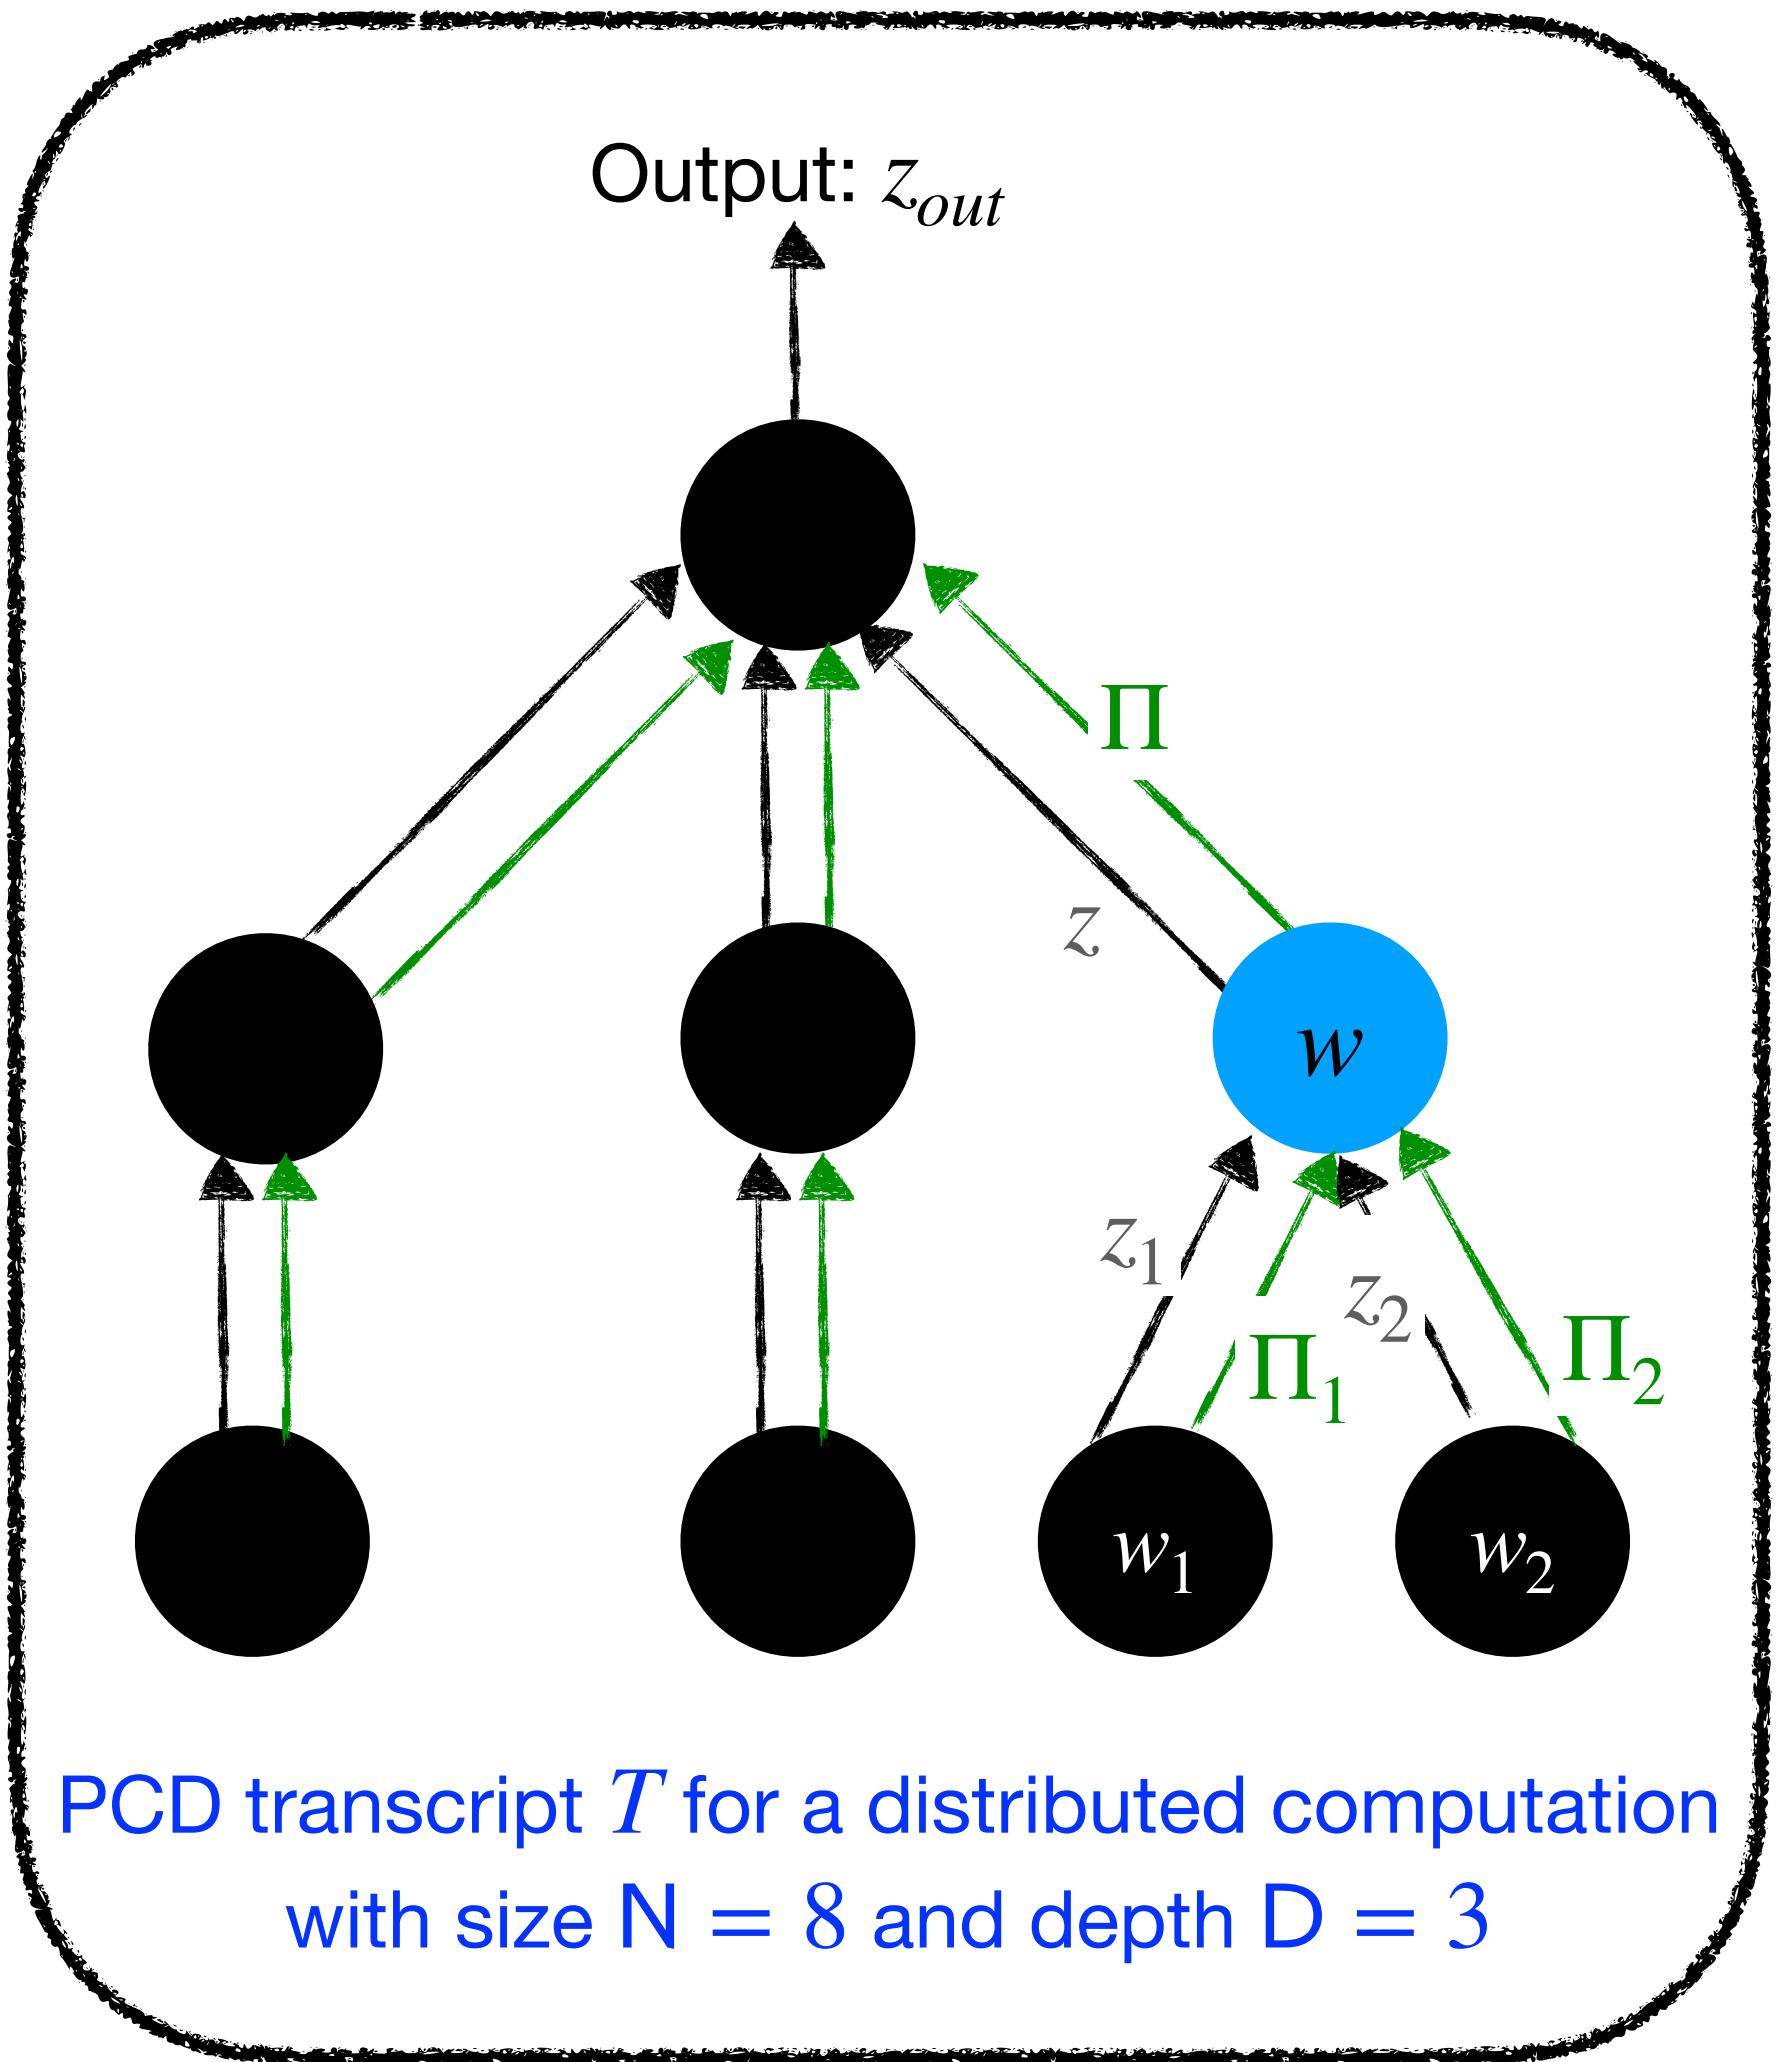
Concatenate SNARK proofs



Naive PCD construction

Concatenate SNARK proofs

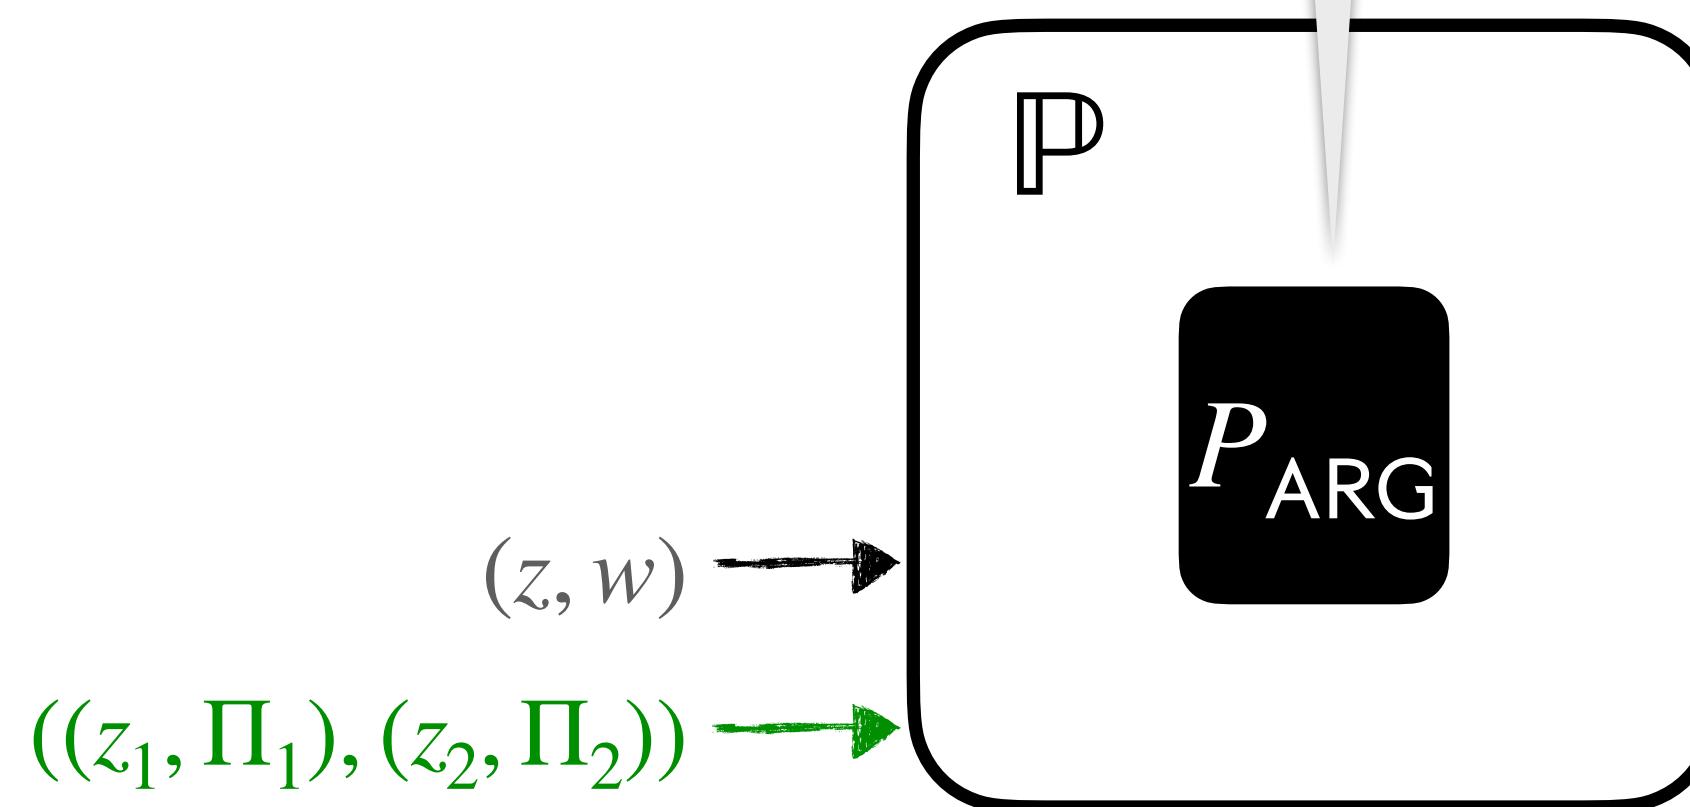
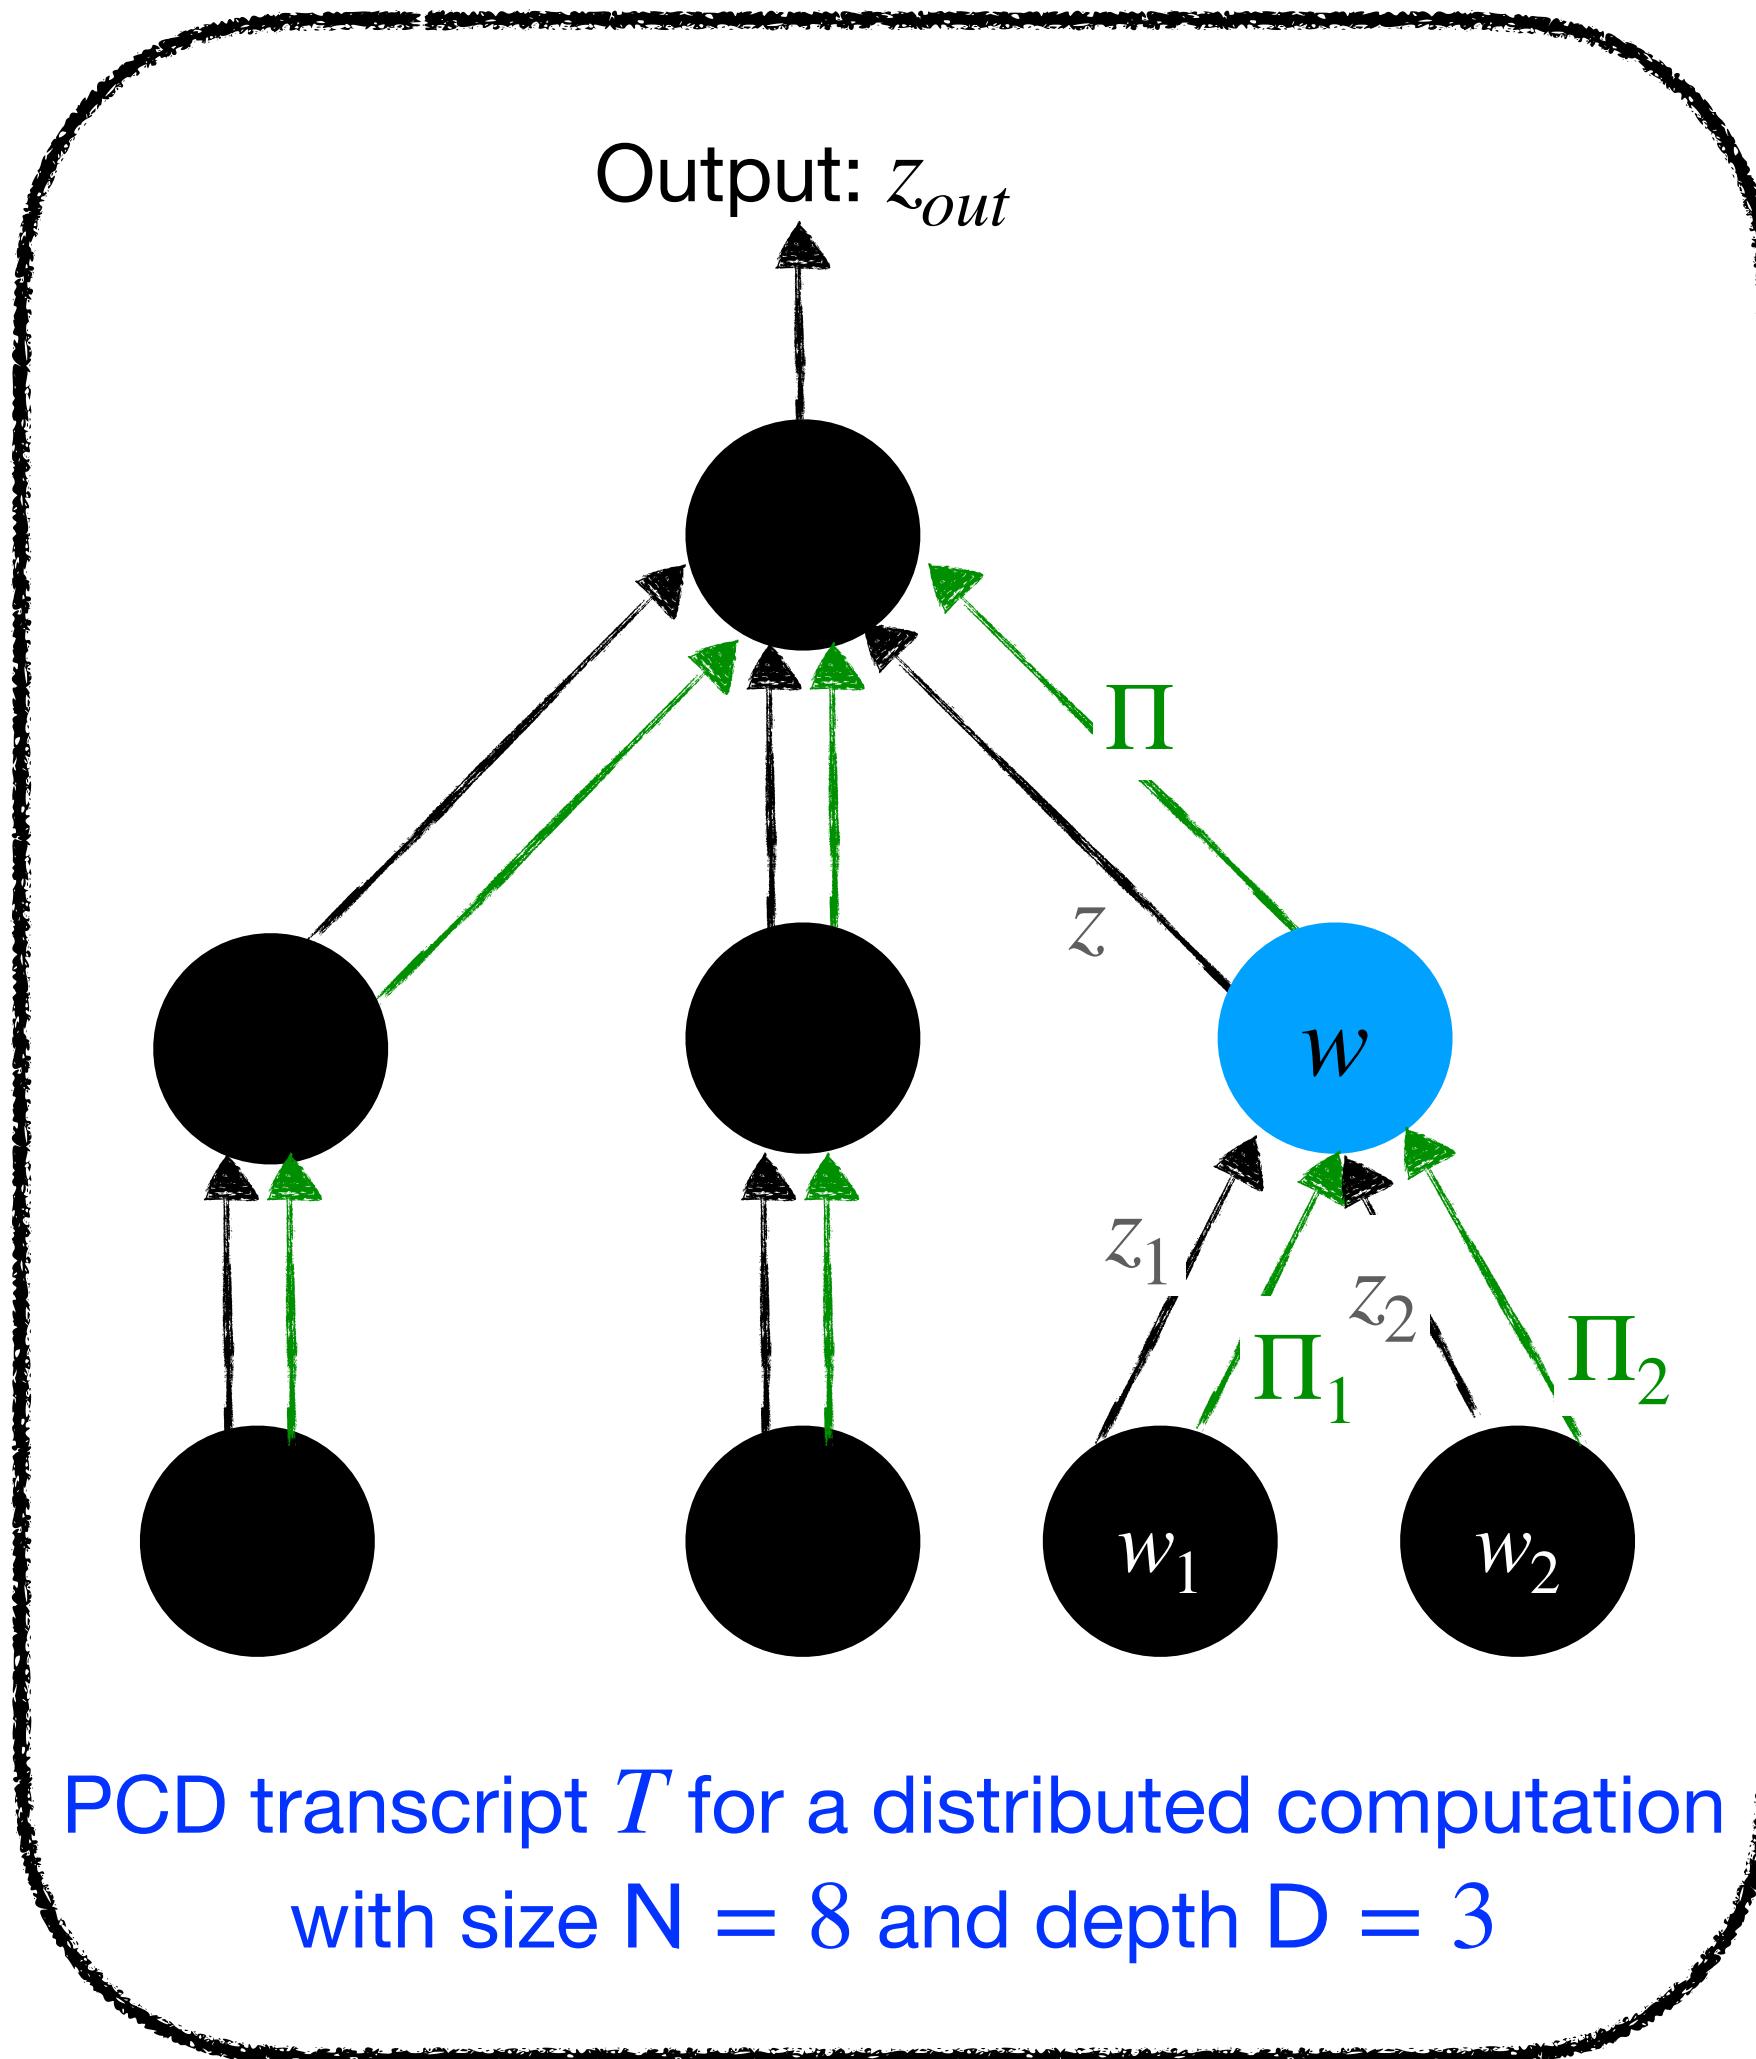
SNARK prover for compliance predicate ϕ



Naive PCD construction

Concatenate SNARK proofs

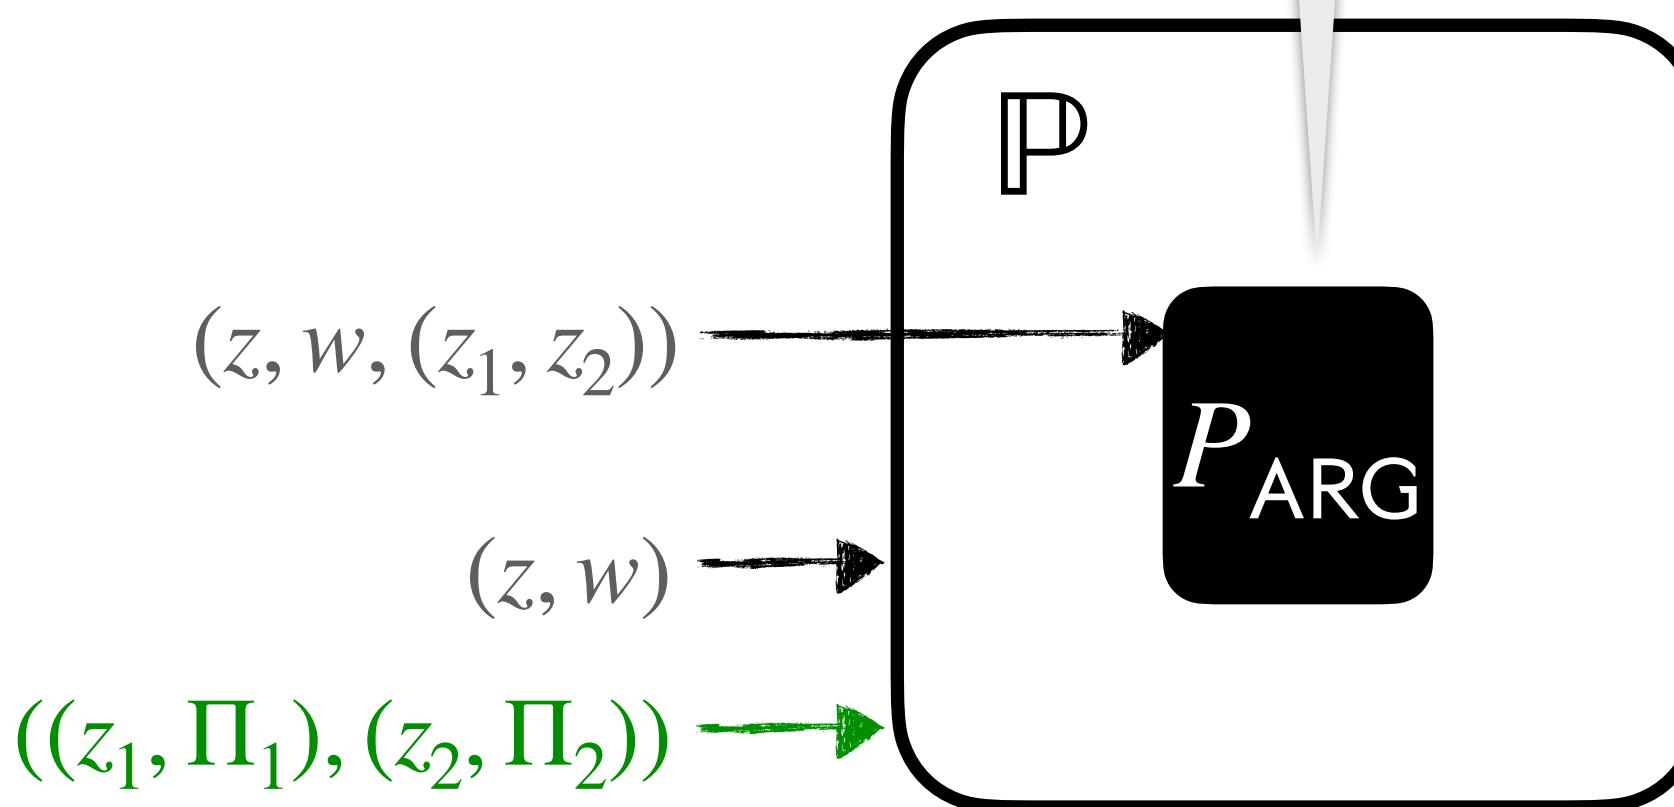
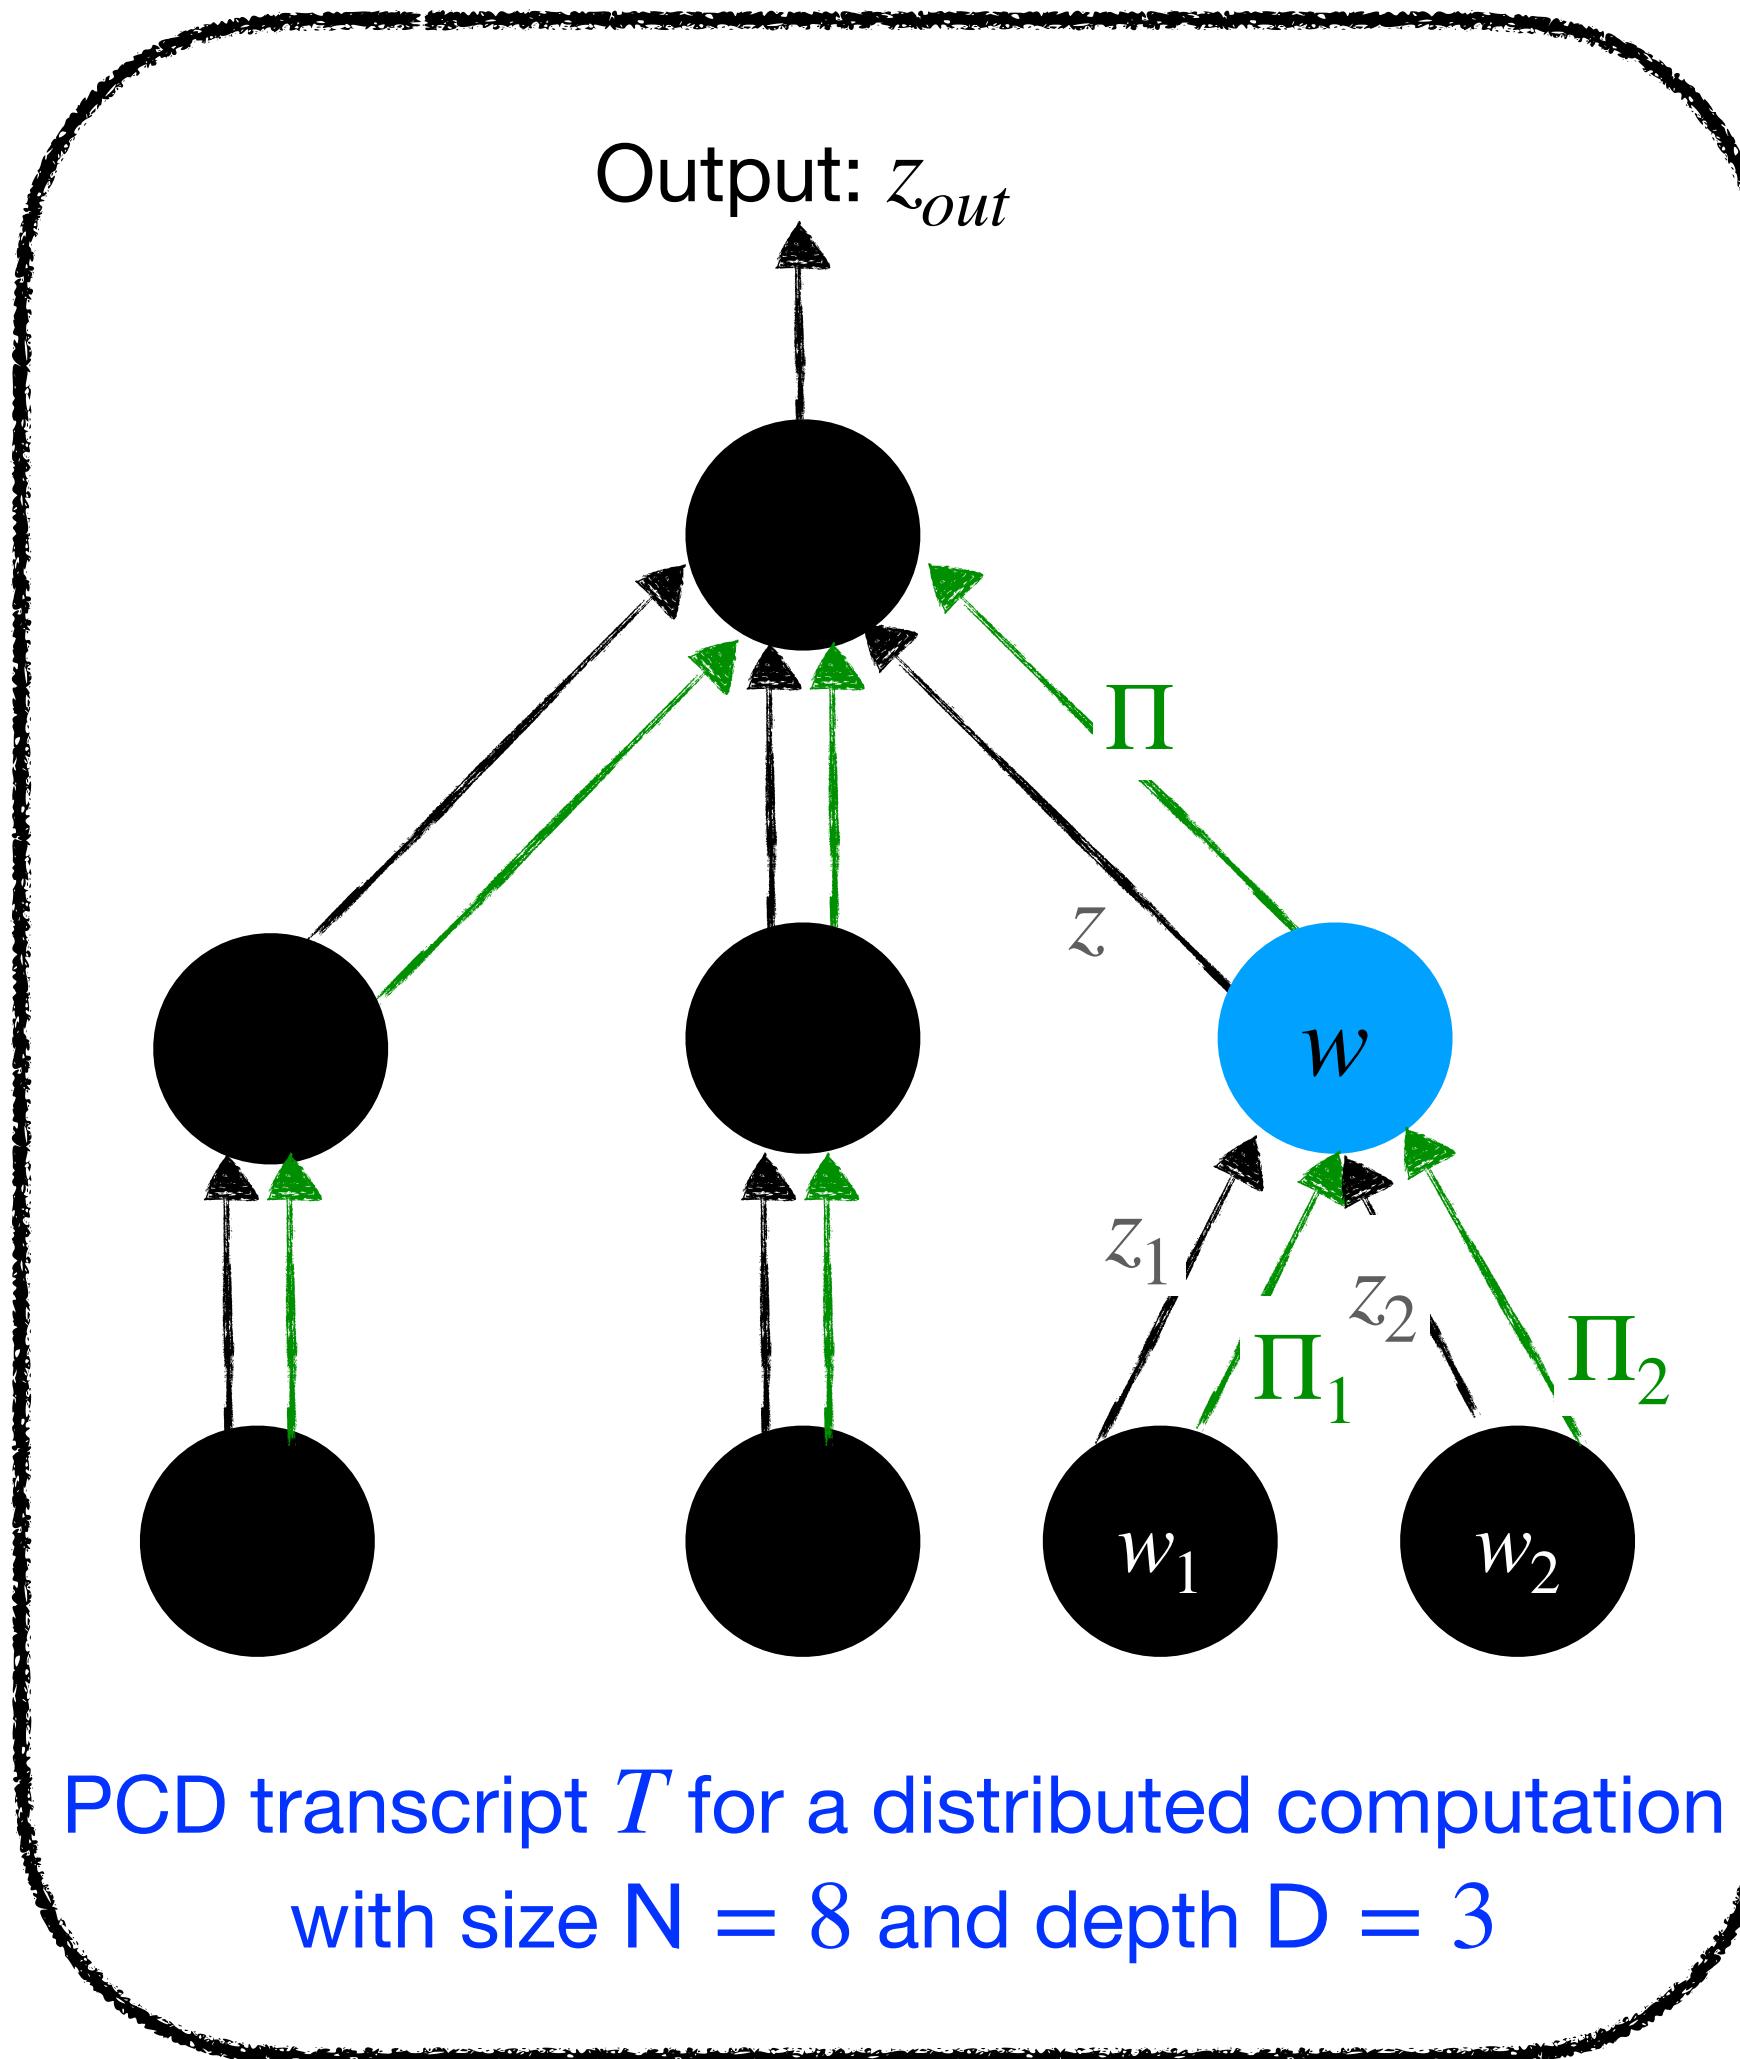
SNARK prover for compliance predicate ϕ



Naive PCD construction

Concatenate SNARK proofs

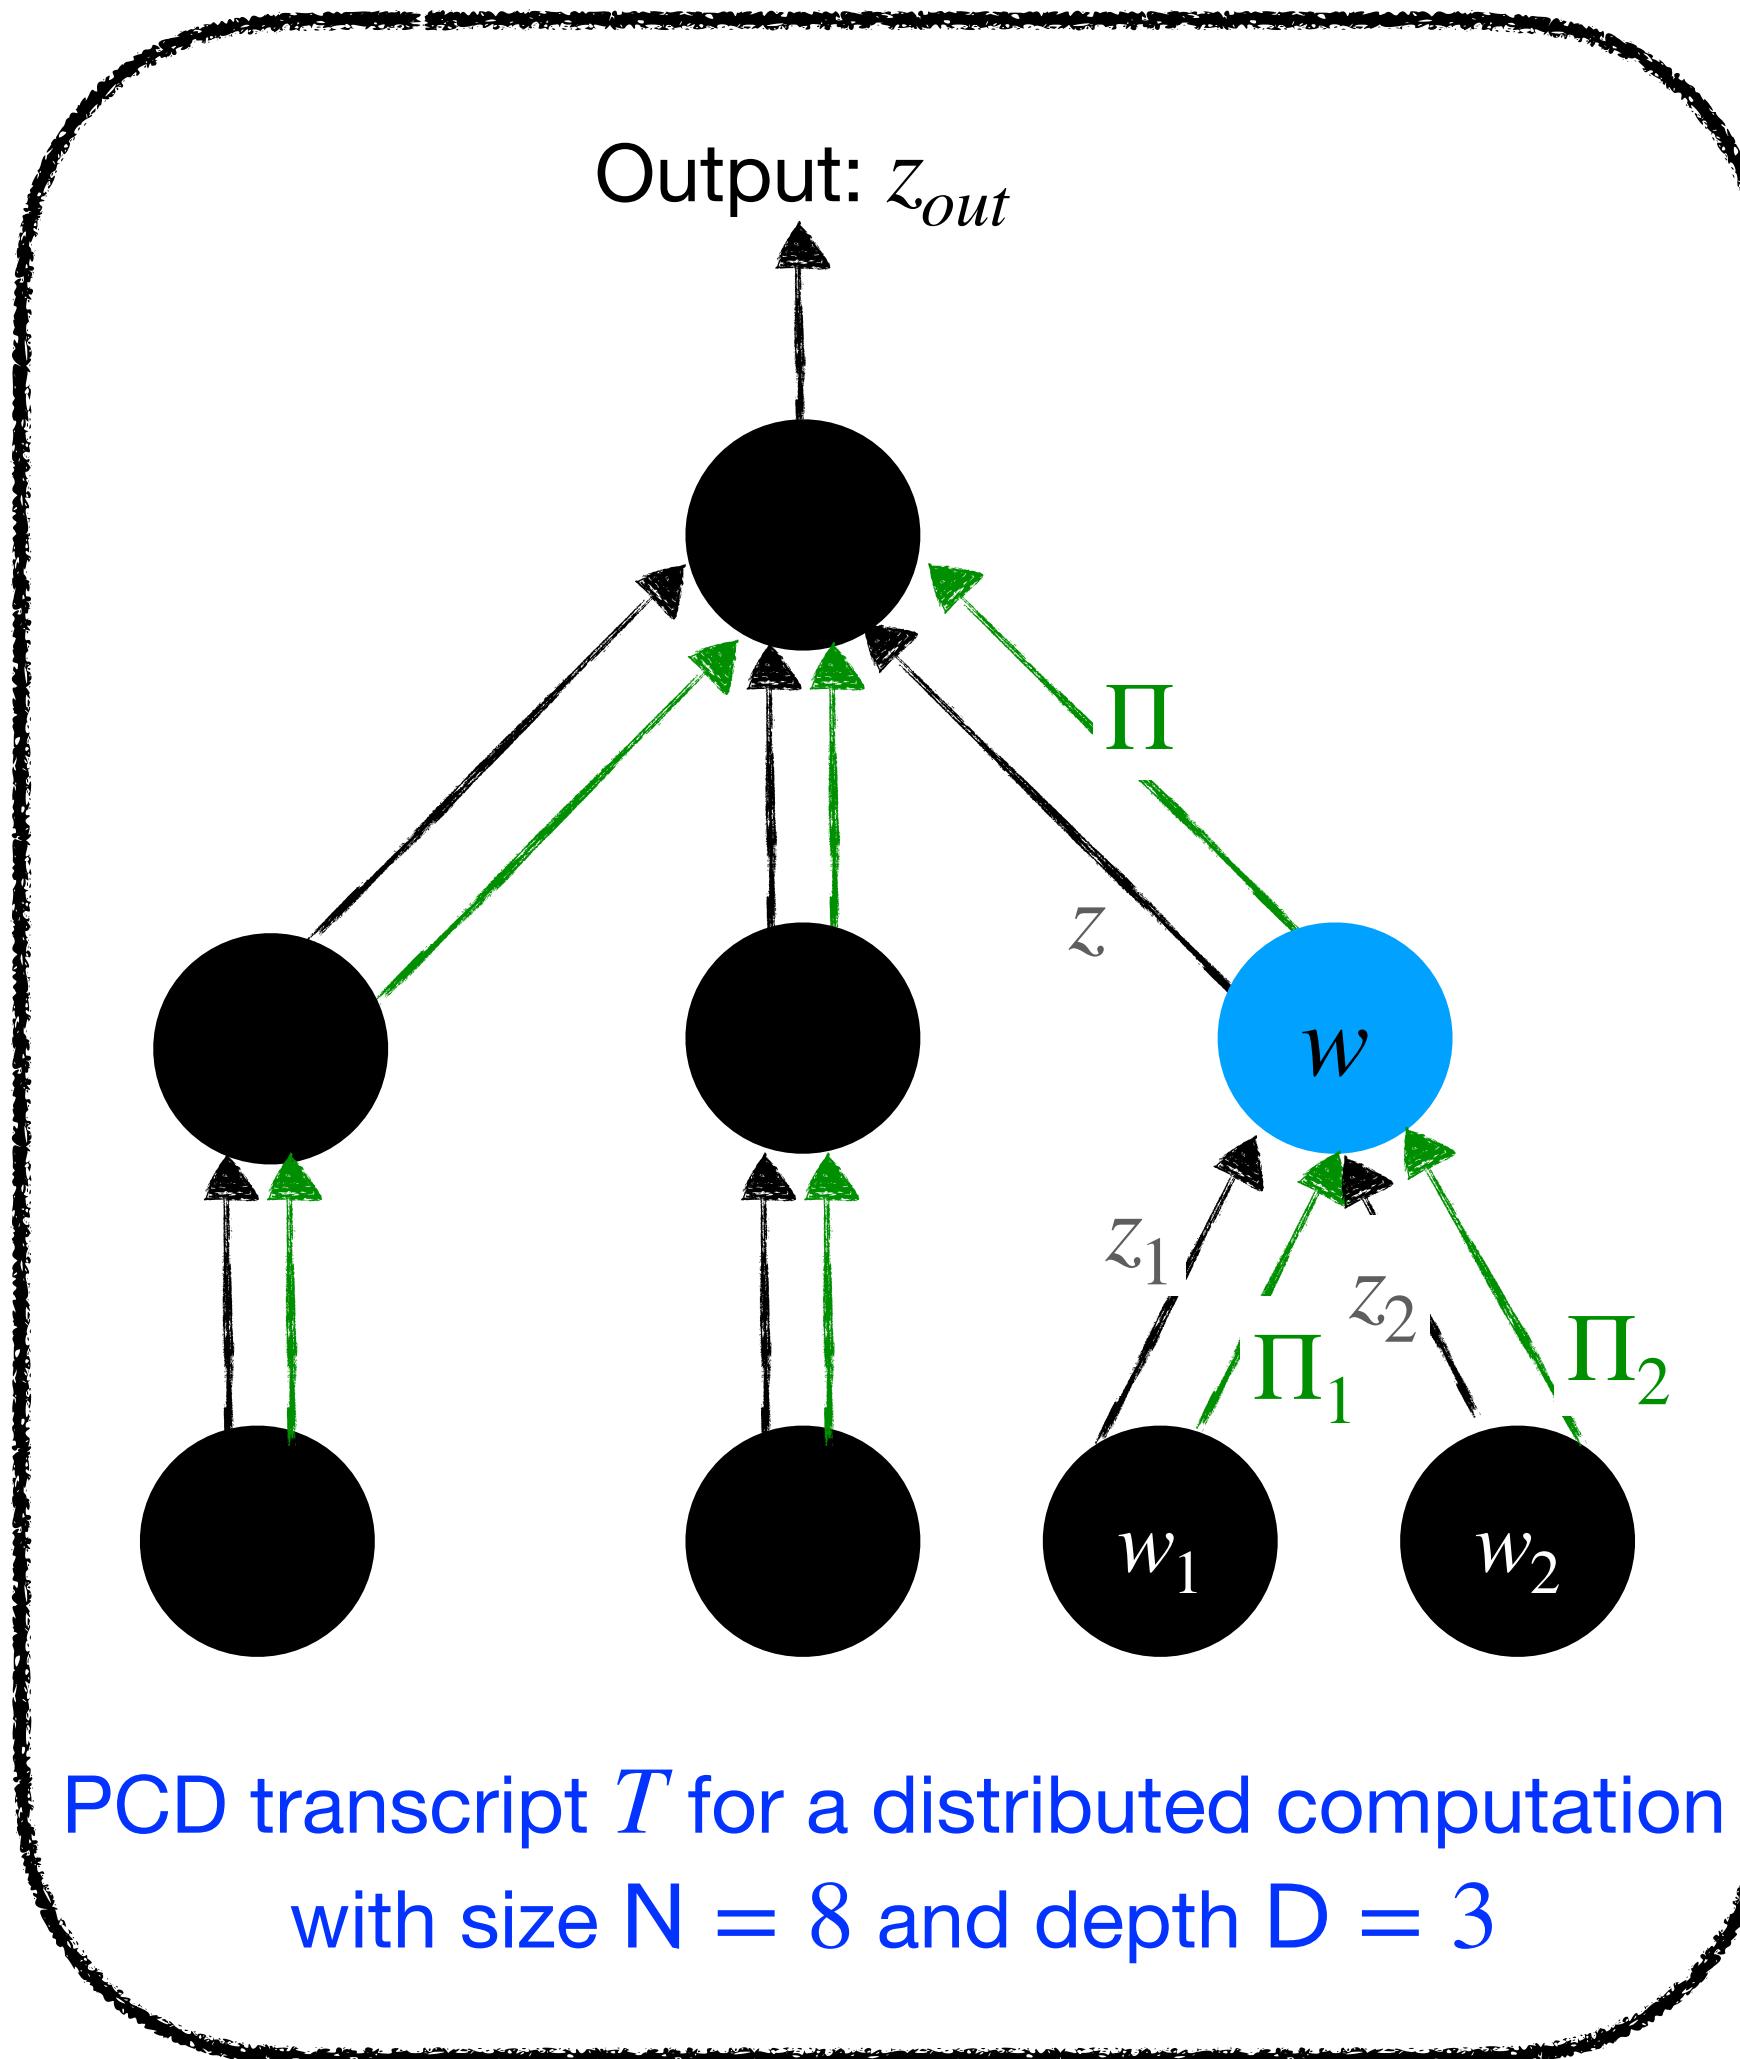
SNARK prover for compliance predicate ϕ



Naive PCD construction

Concatenate SNARK proofs

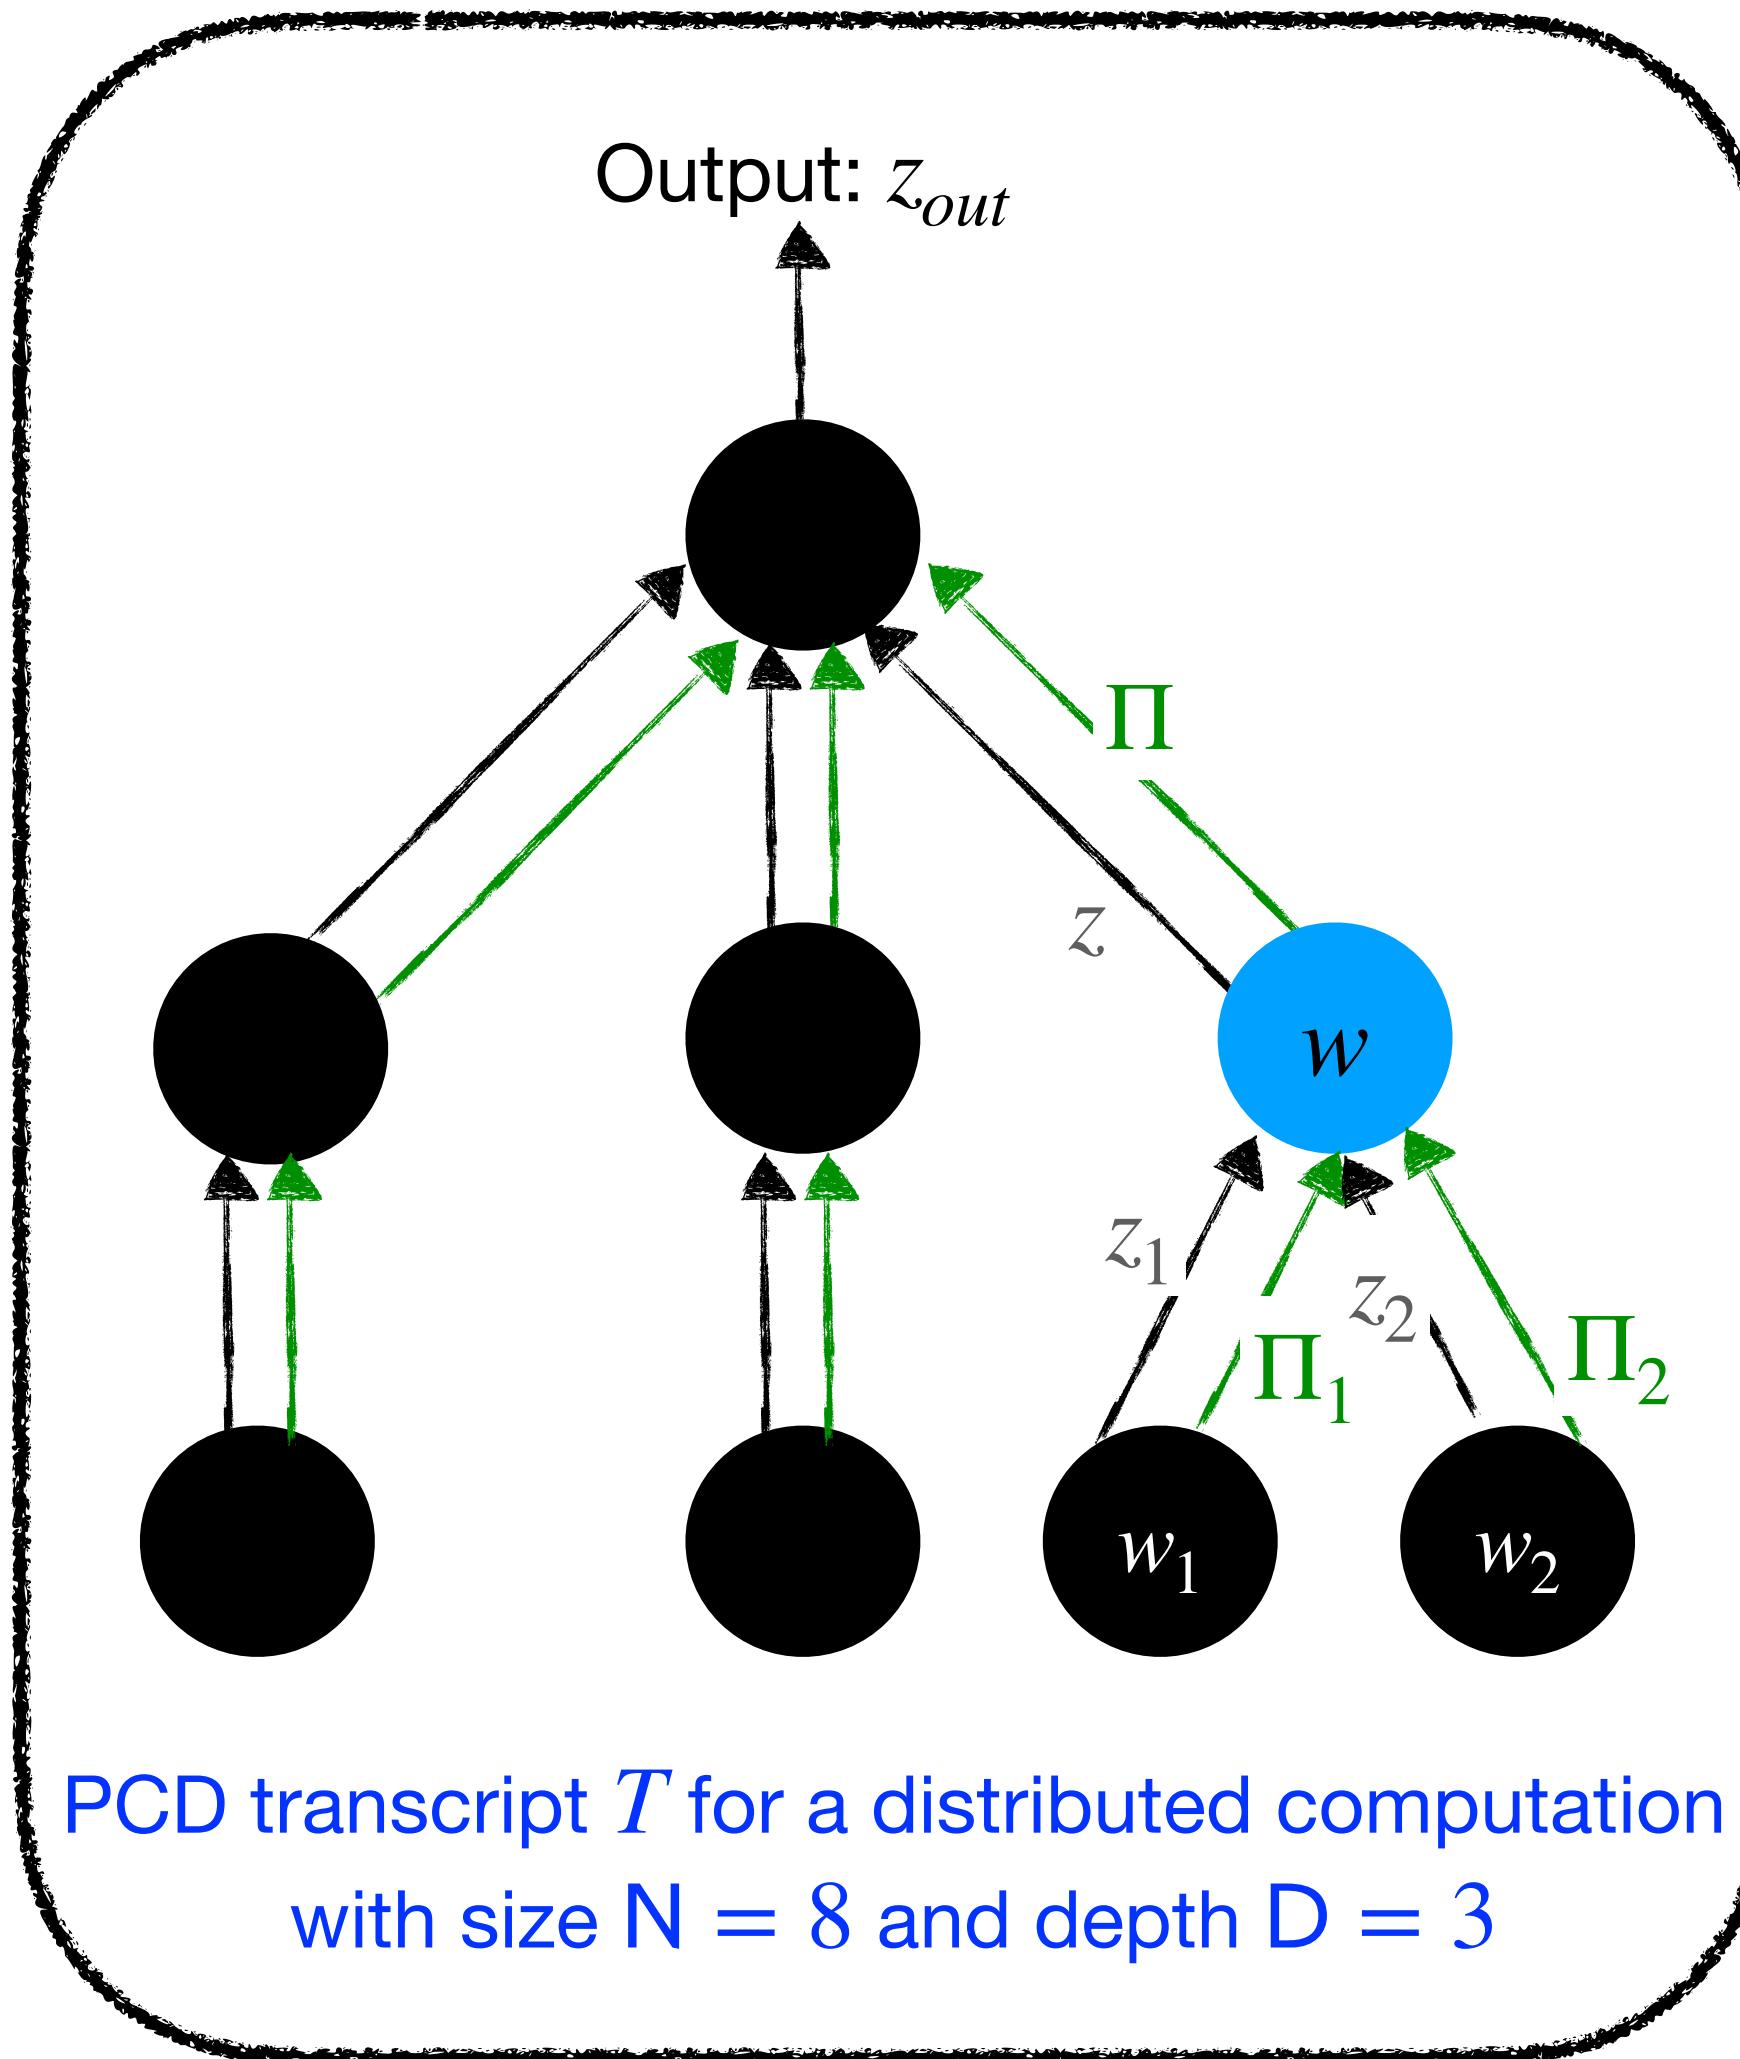
SNARK prover for compliance predicate ϕ



Naive PCD construction

Concatenate SNARK proofs

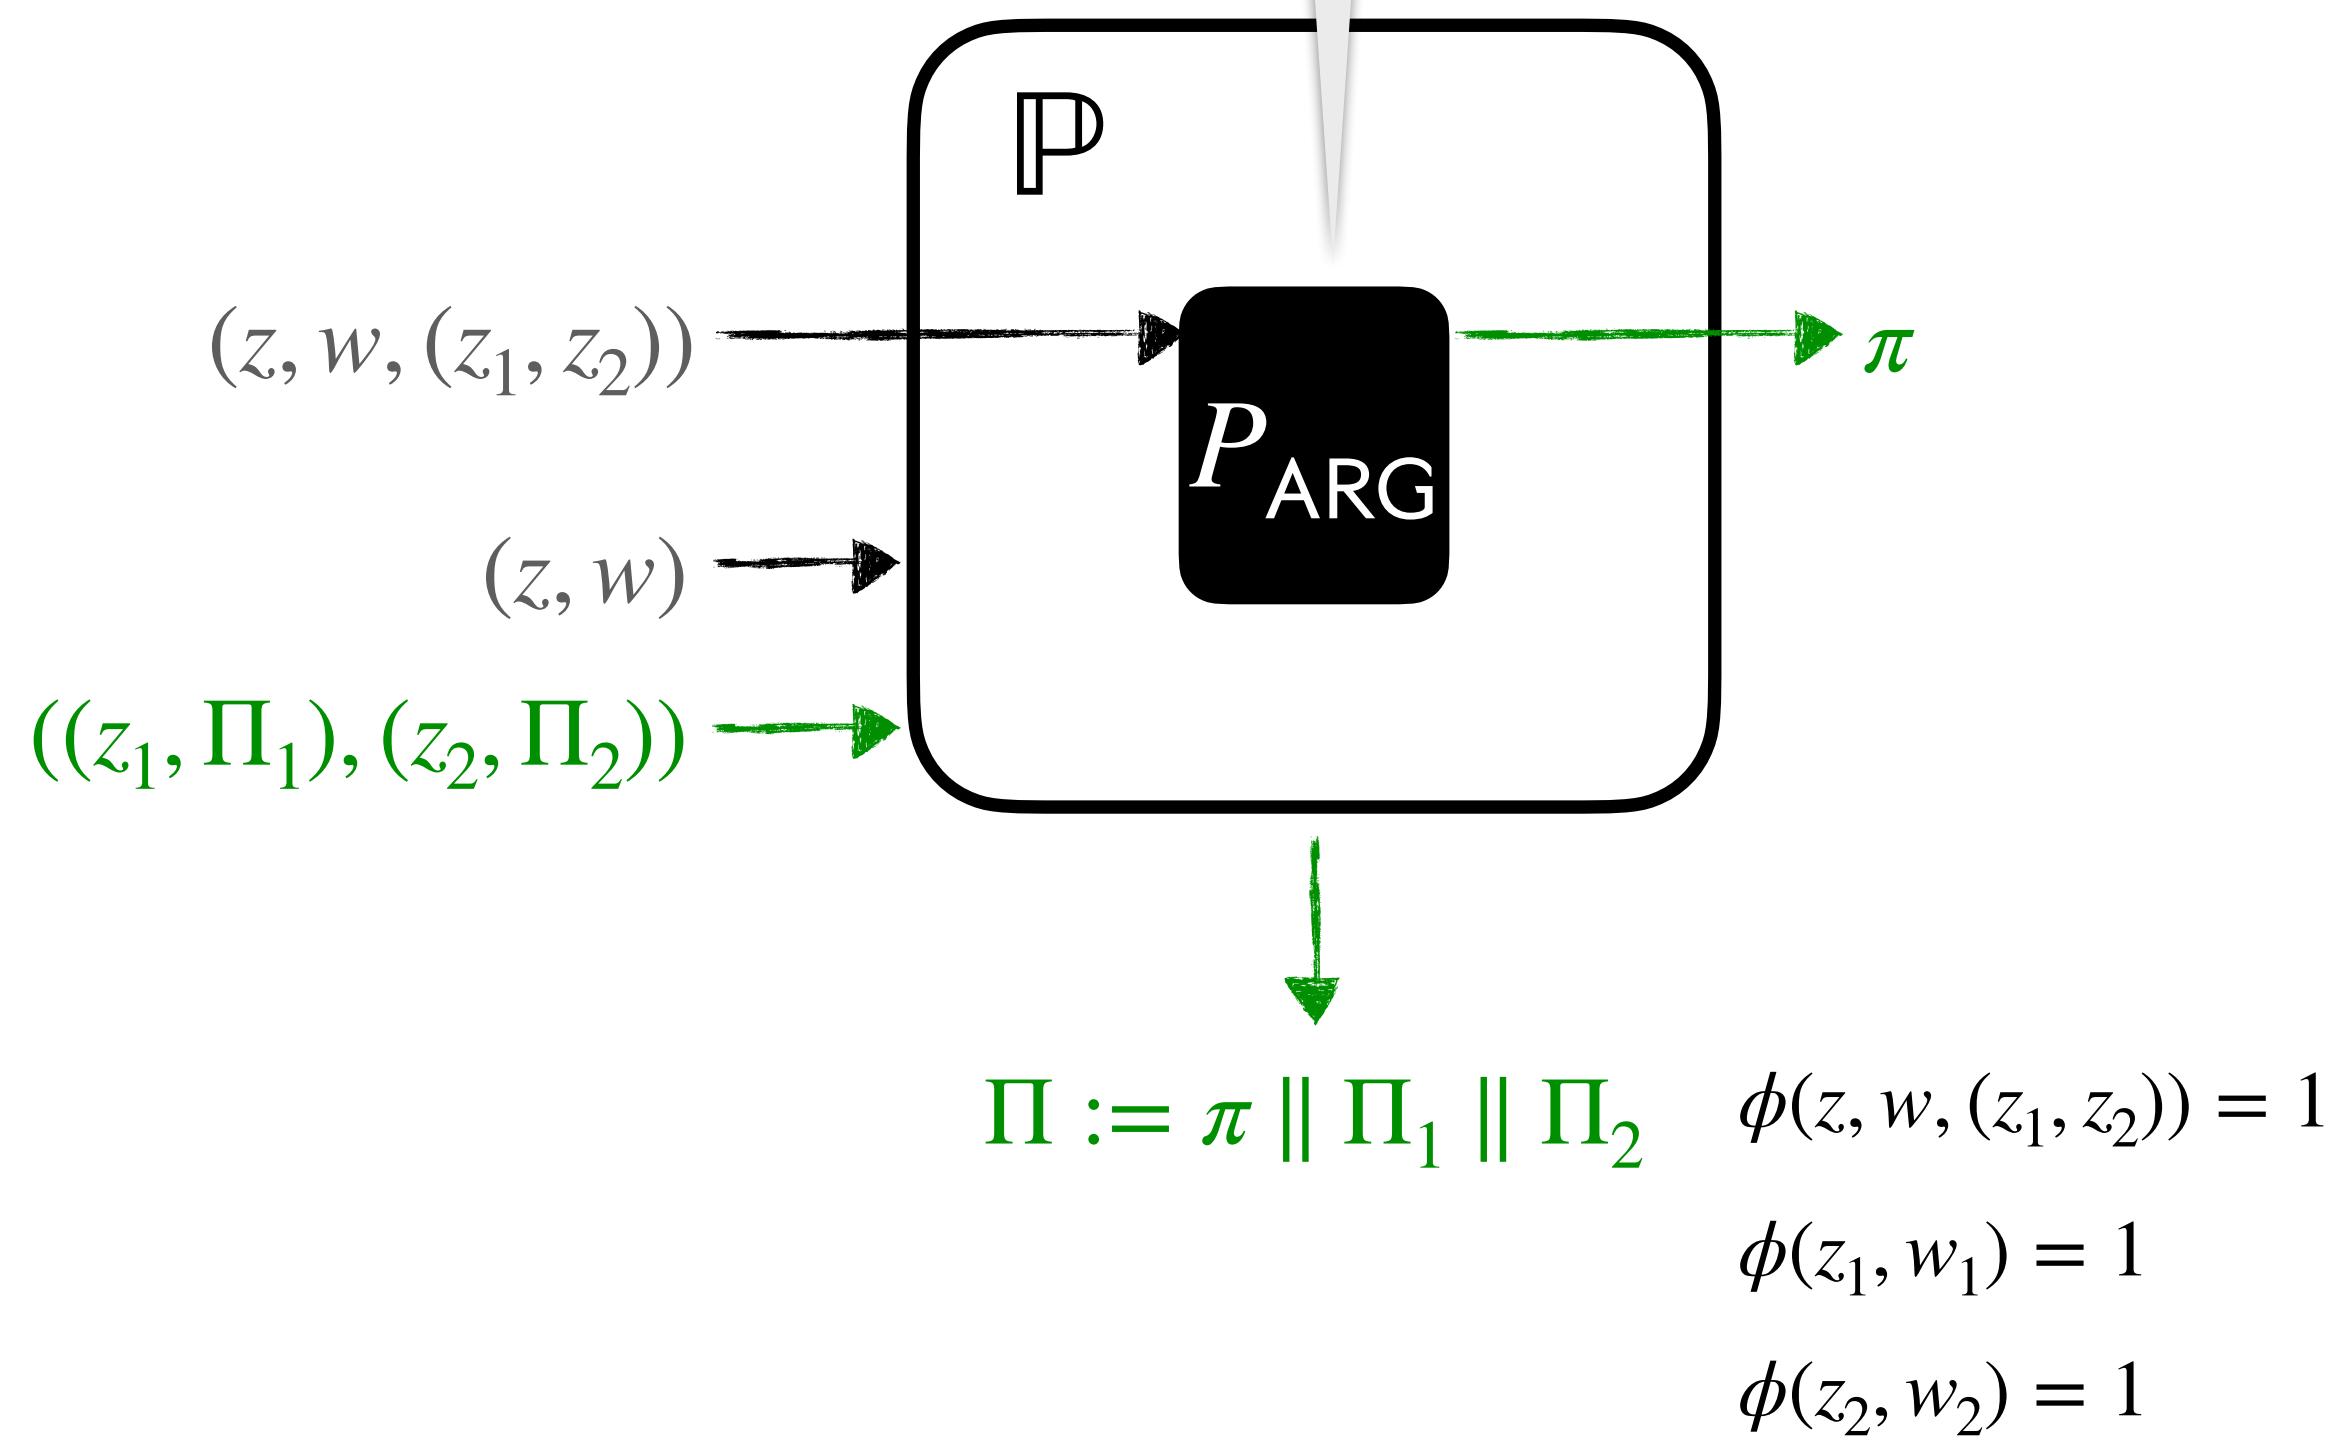
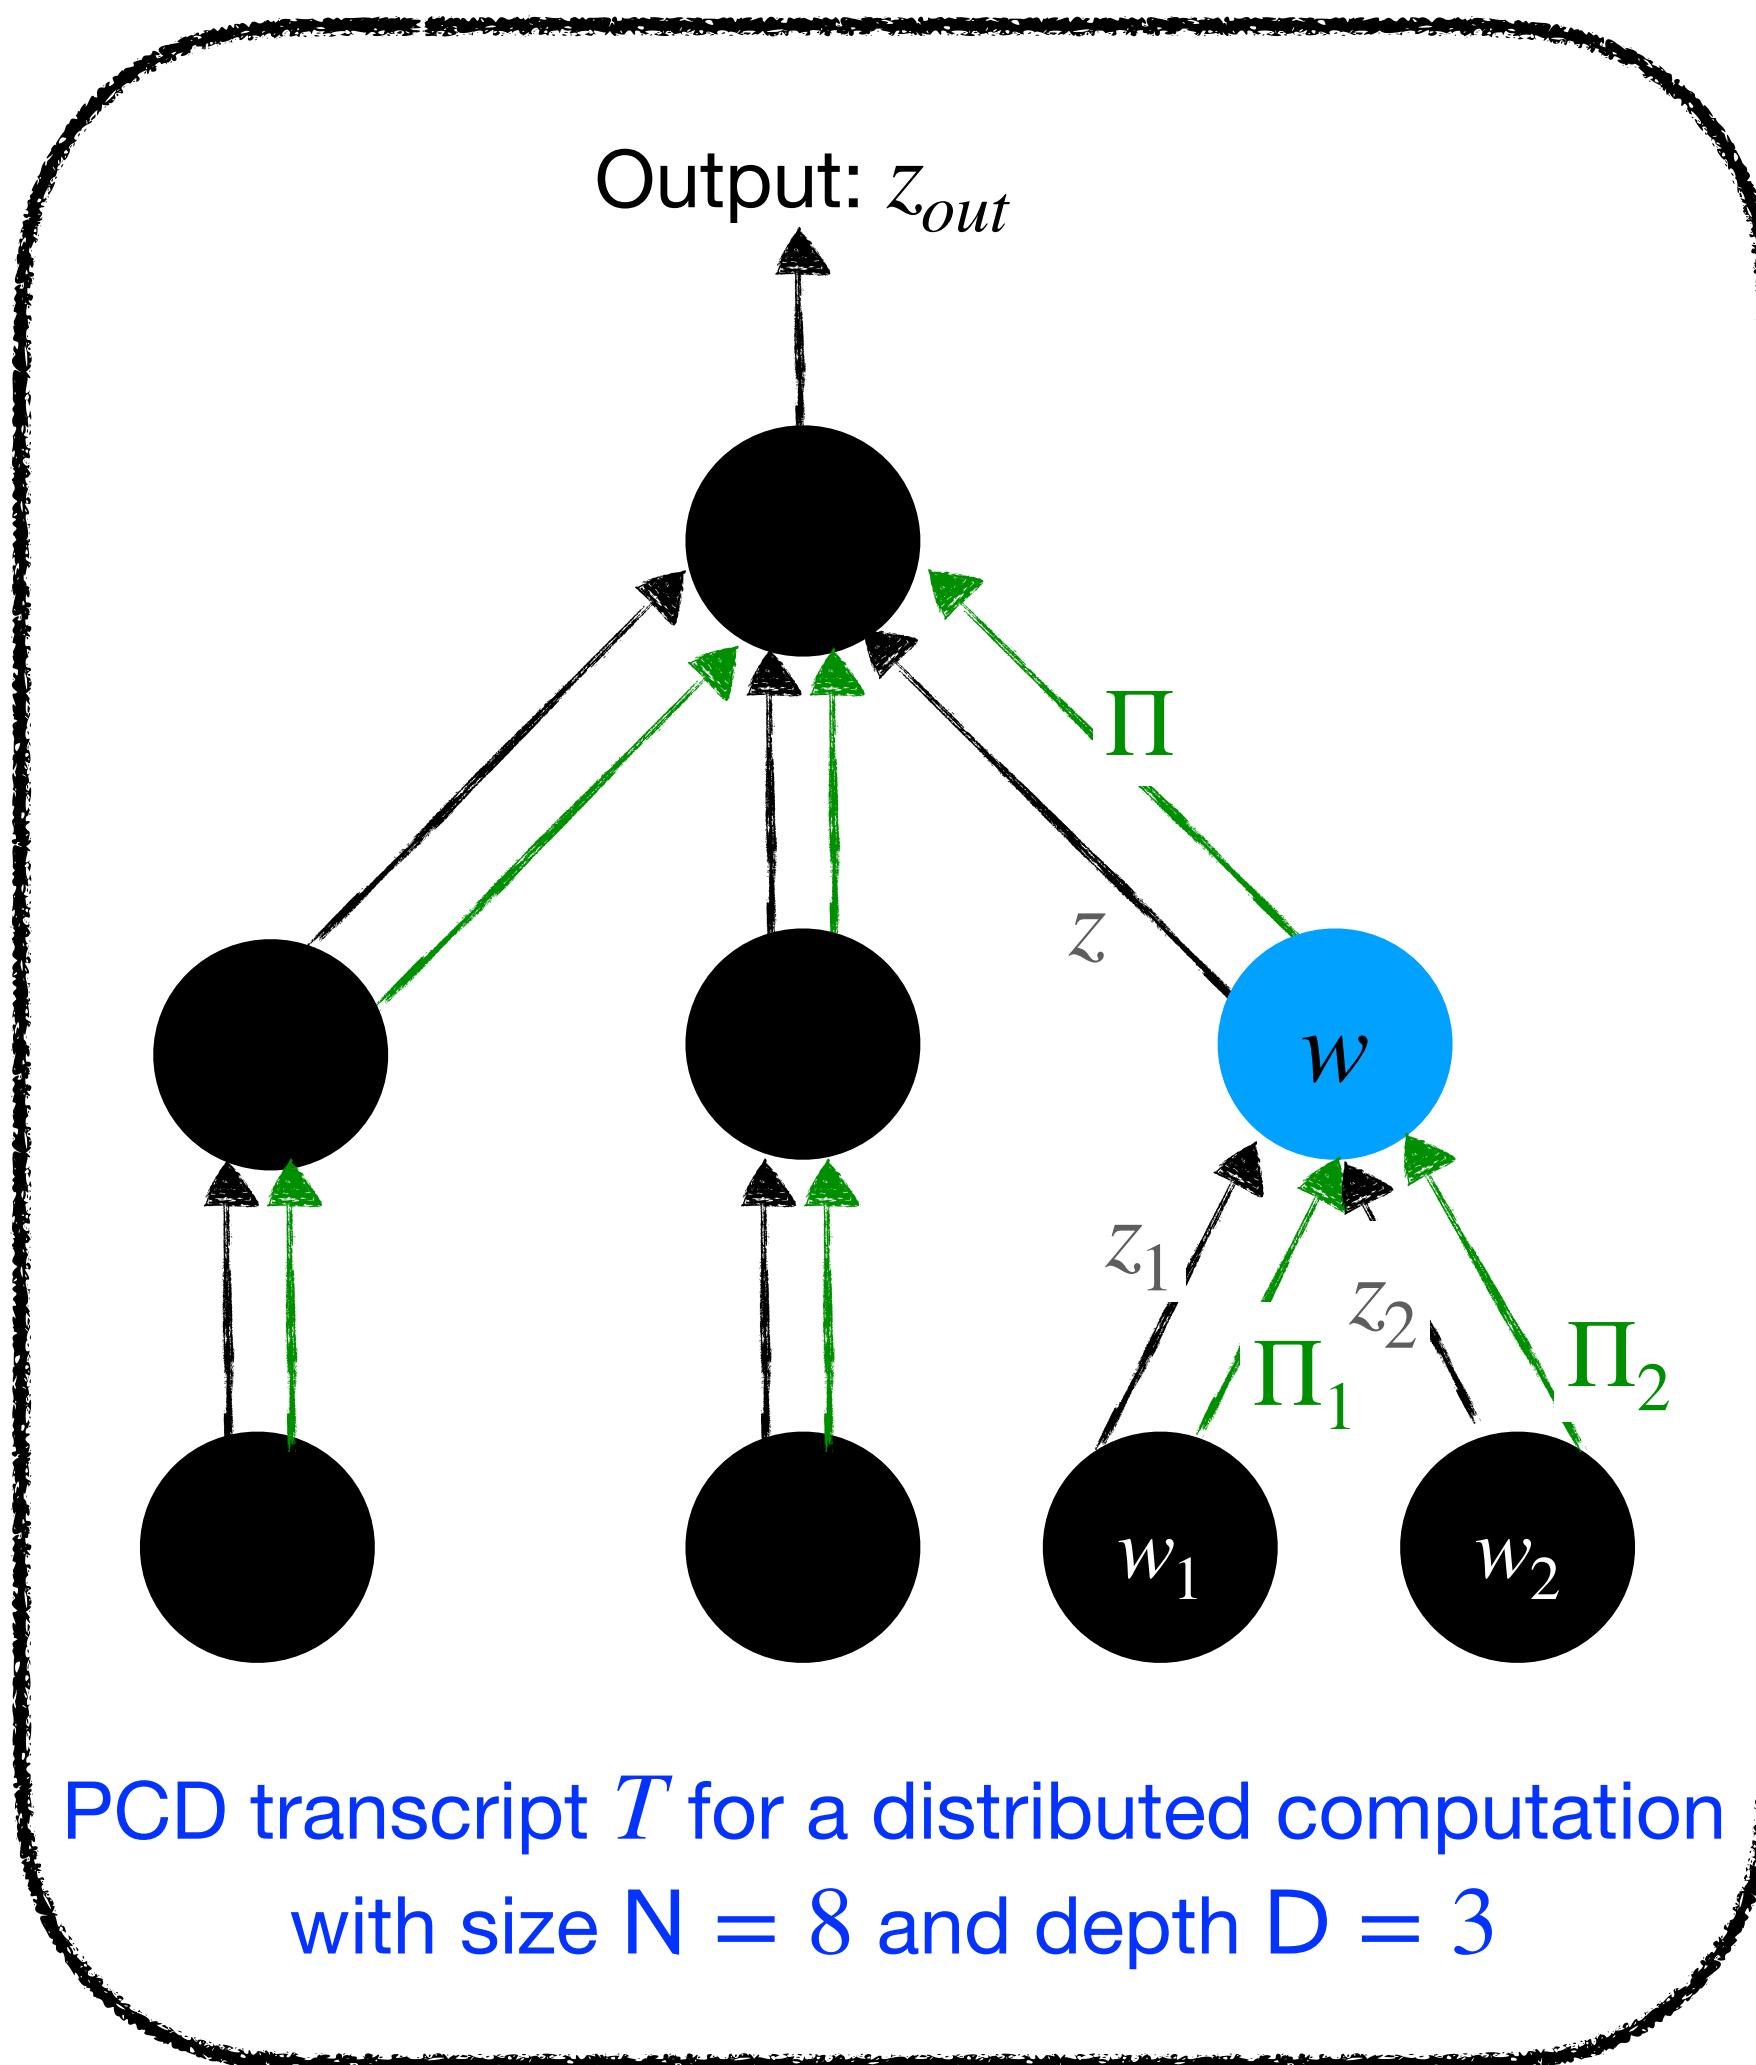
SNARK prover for compliance predicate ϕ



Naive PCD construction

Concatenate SNARK proofs

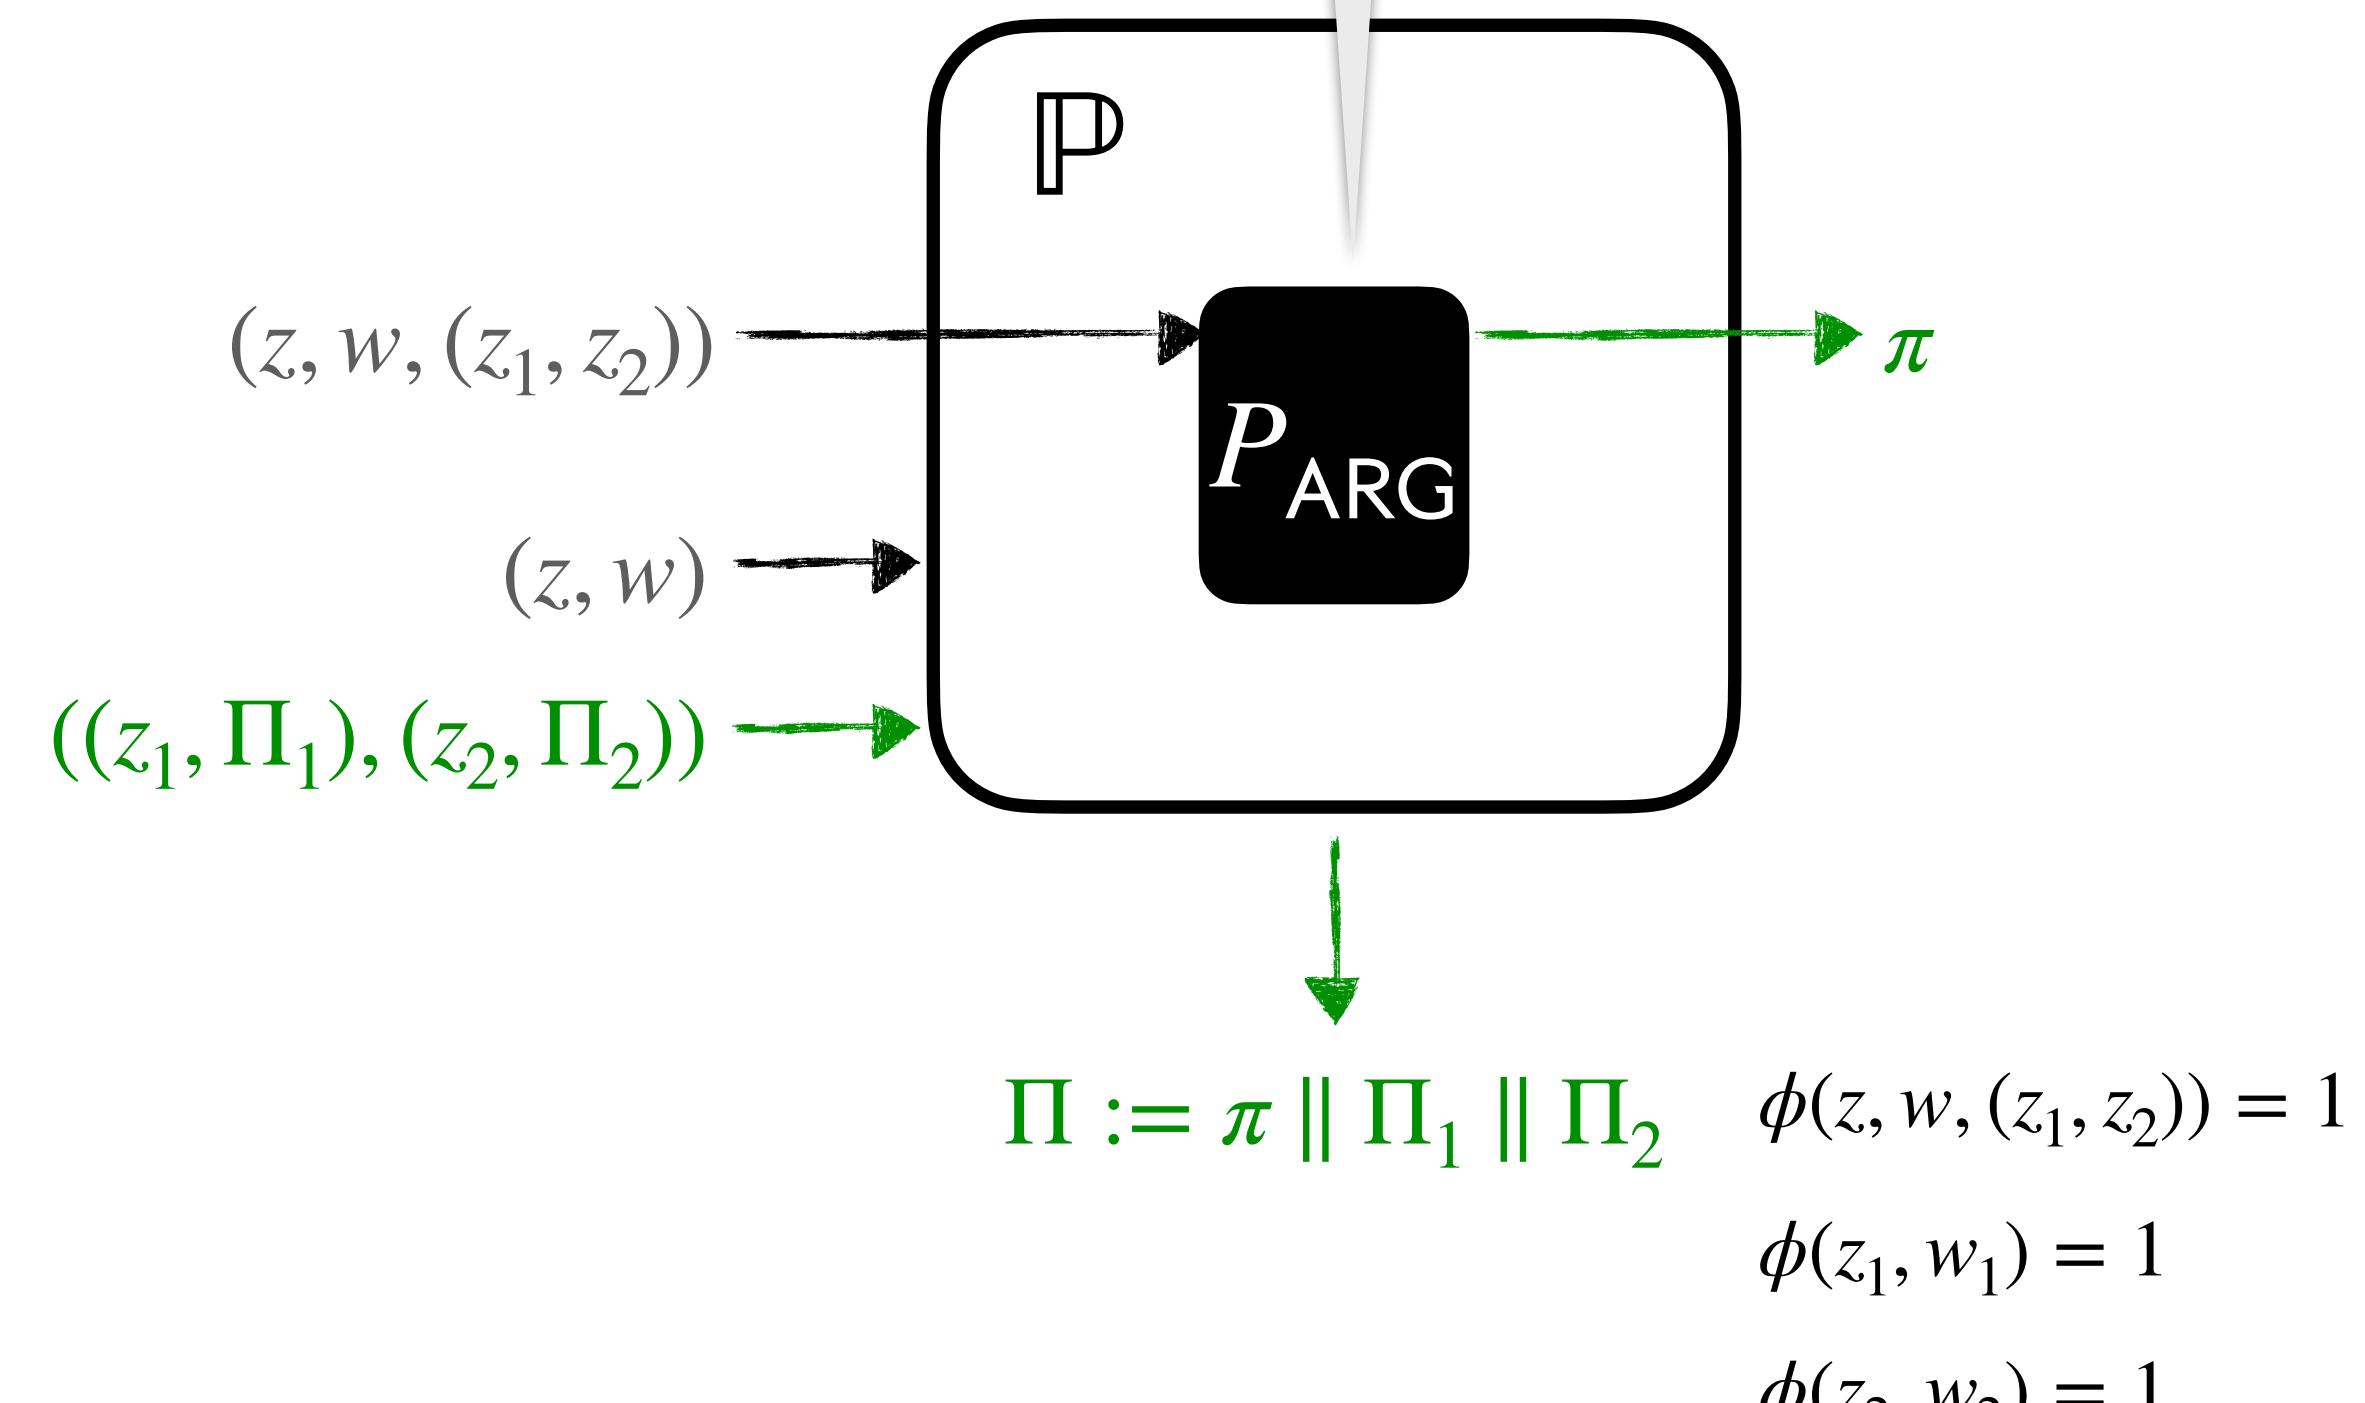
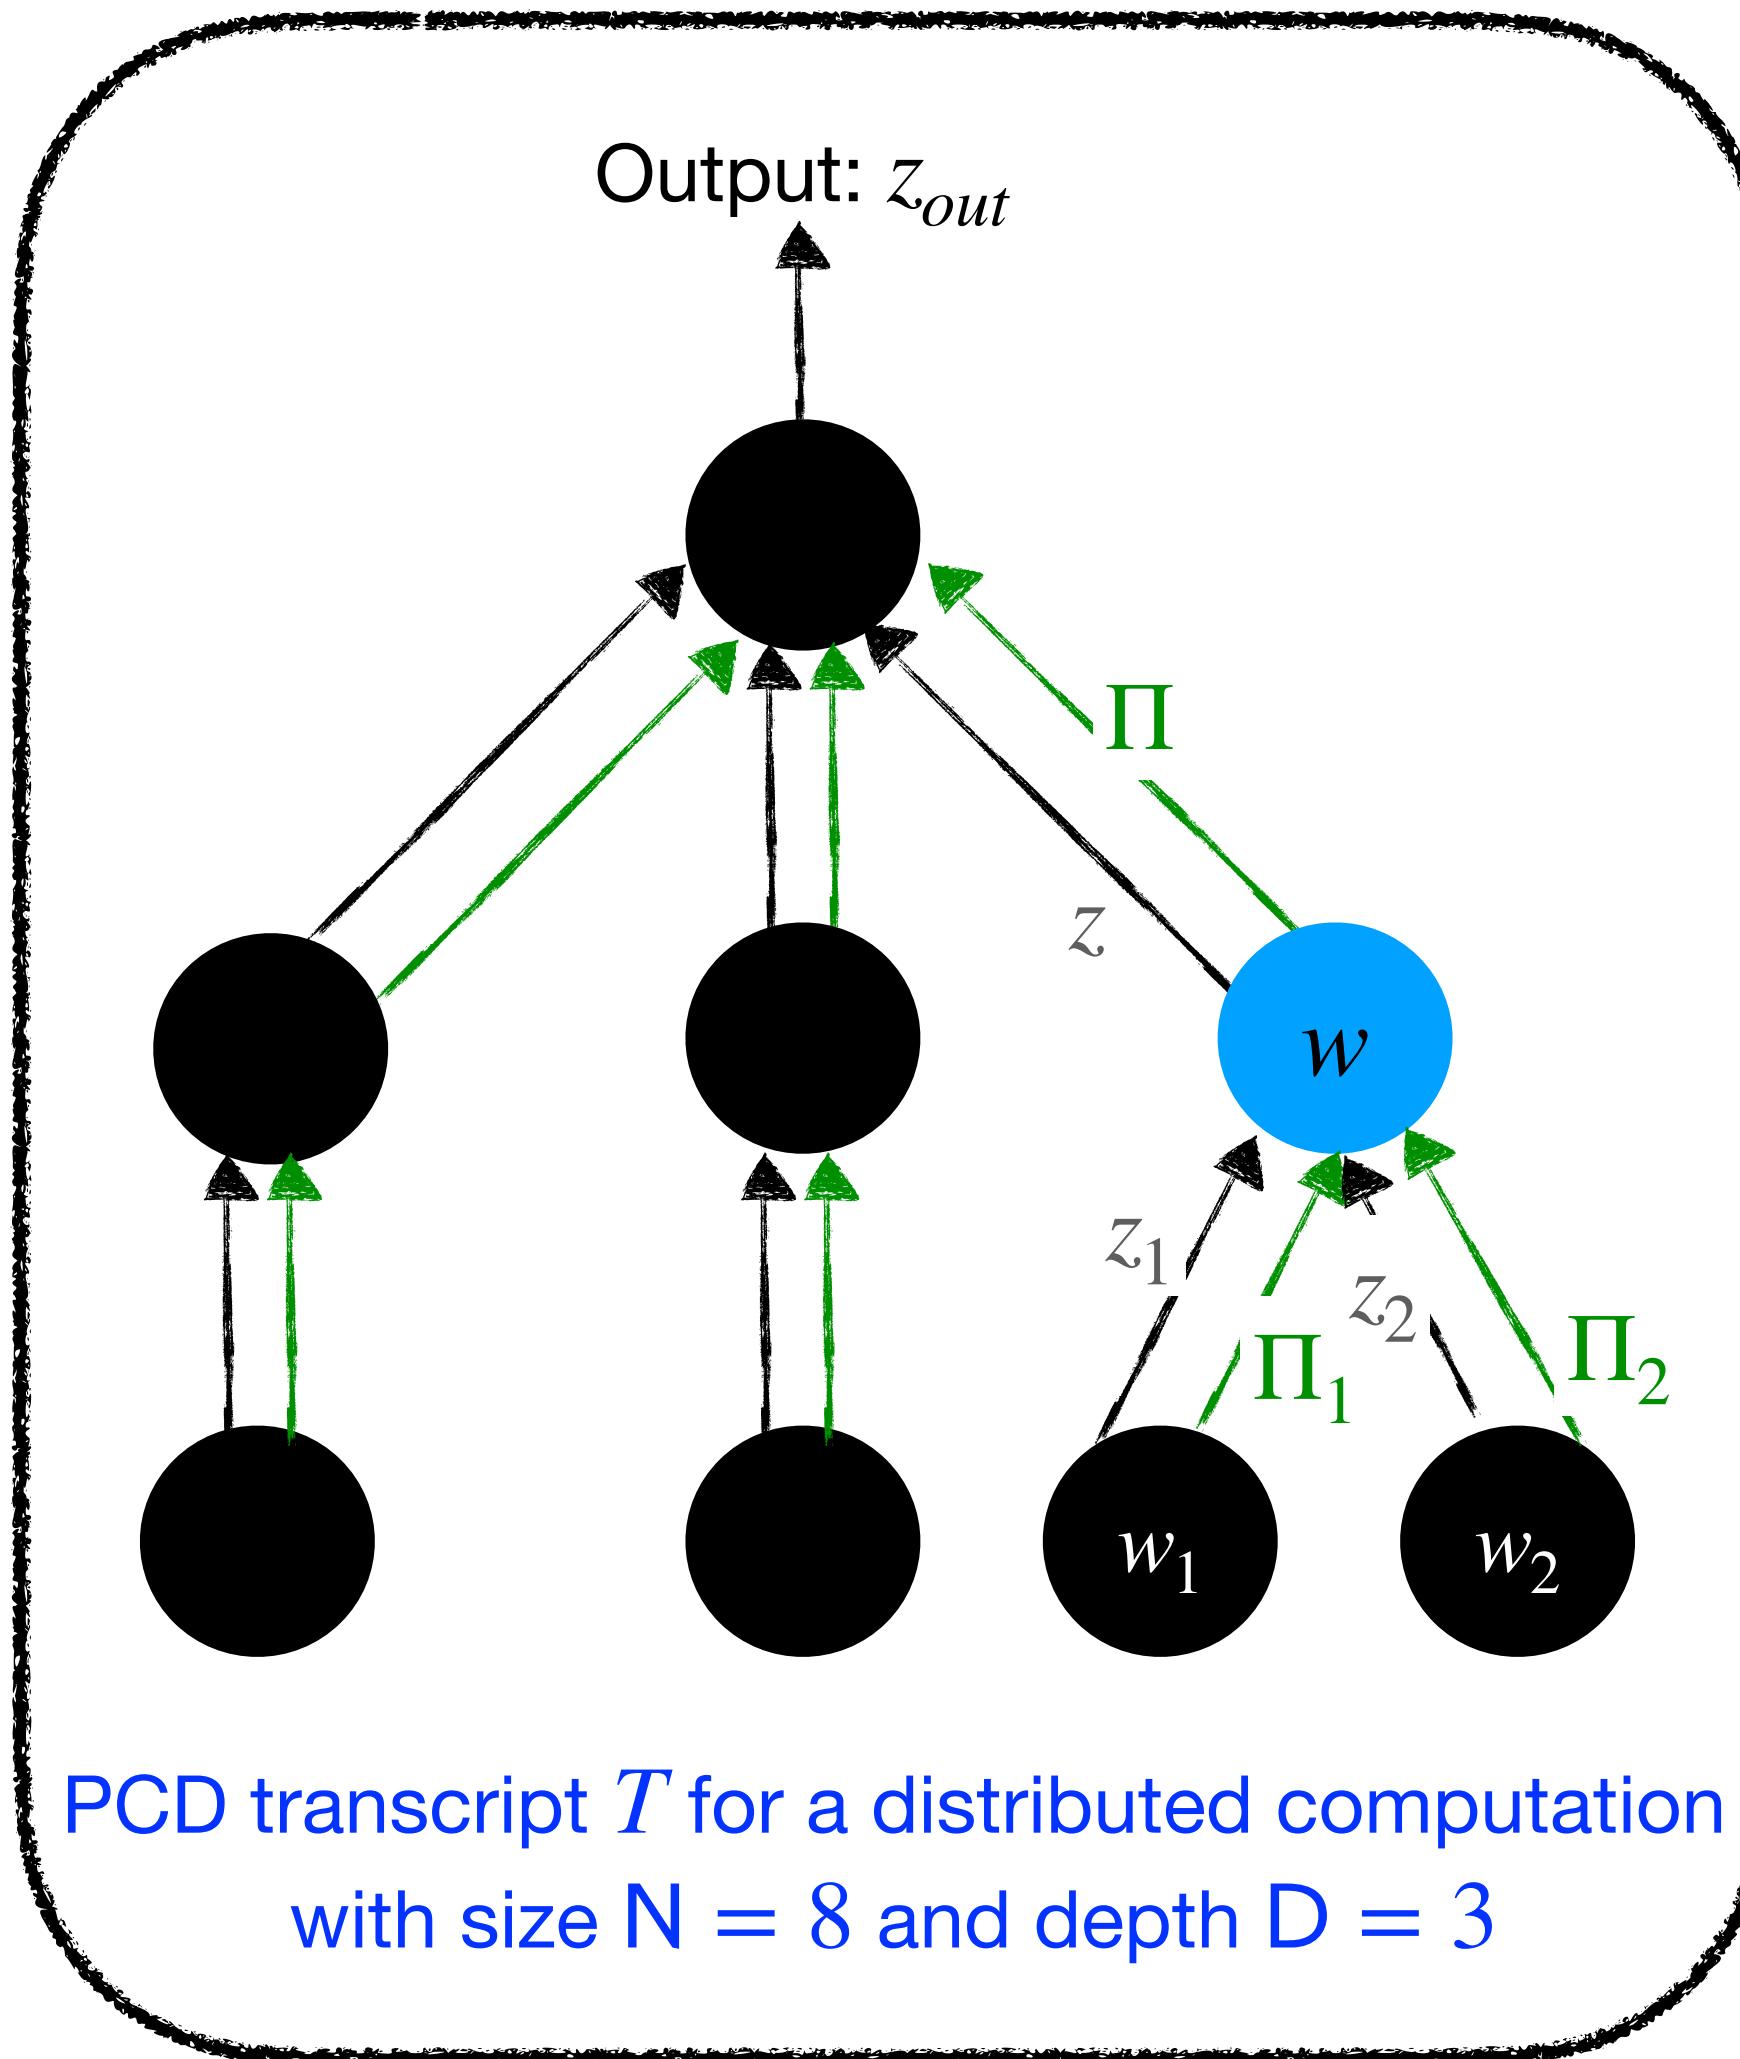
SNARK prover for compliance predicate ϕ



Naive PCD construction

Concatenate SNARK proofs

SNARK prover for compliance predicate ϕ

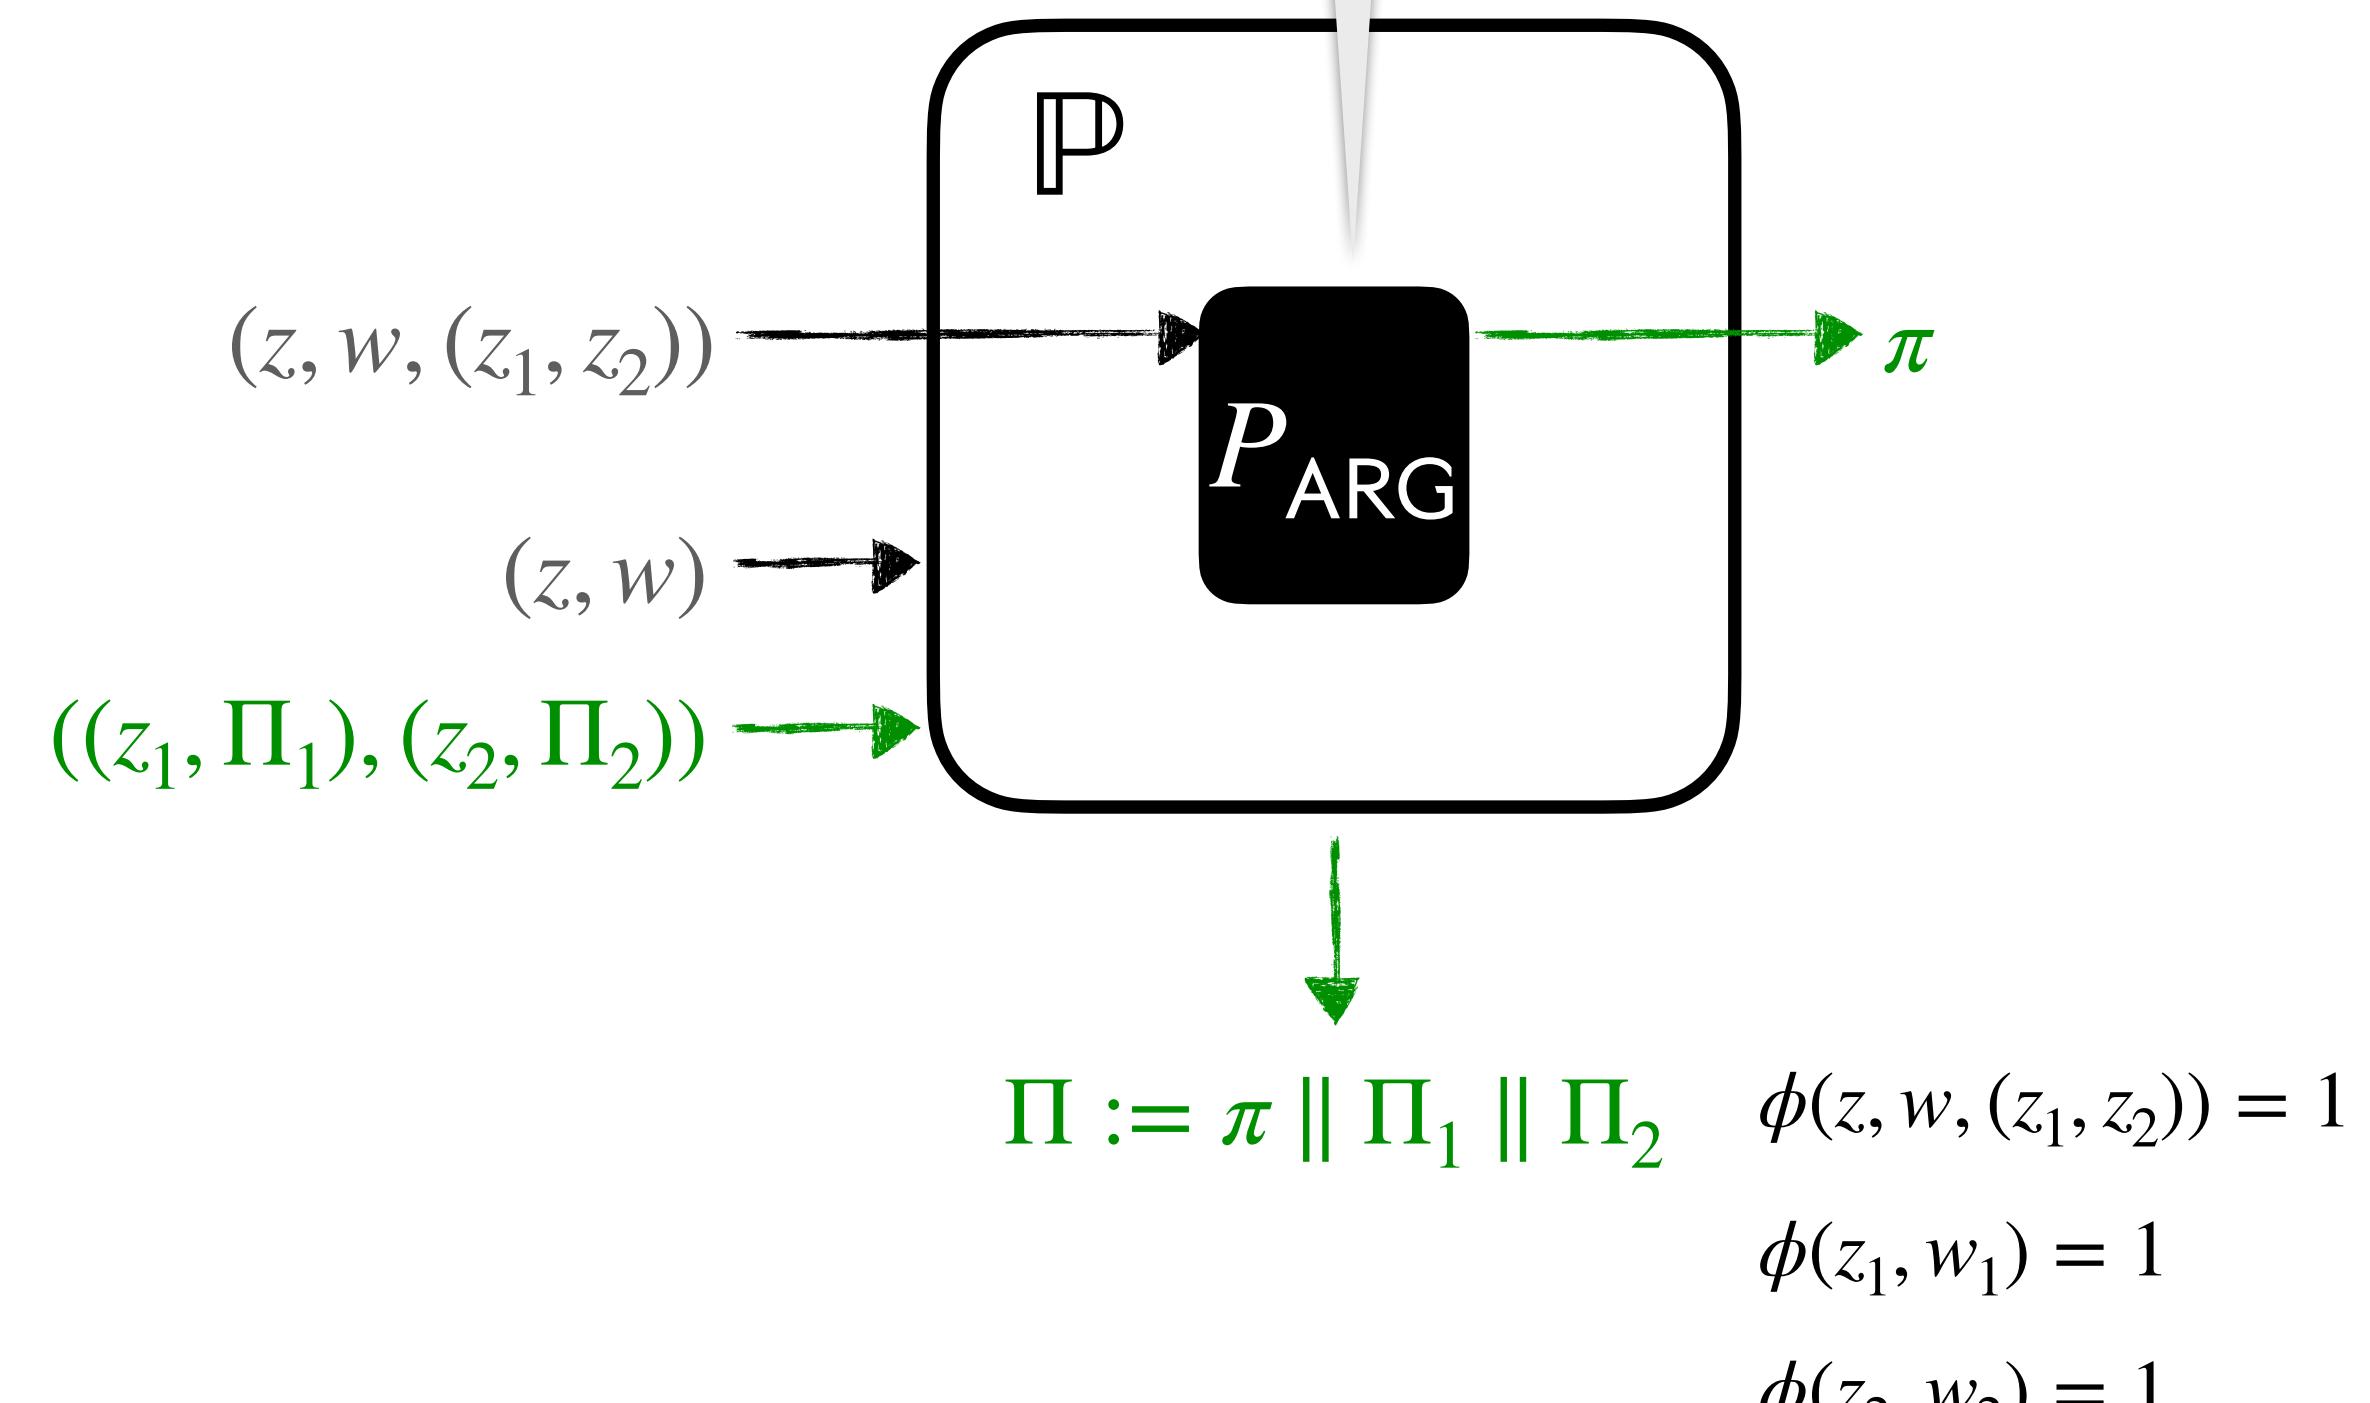
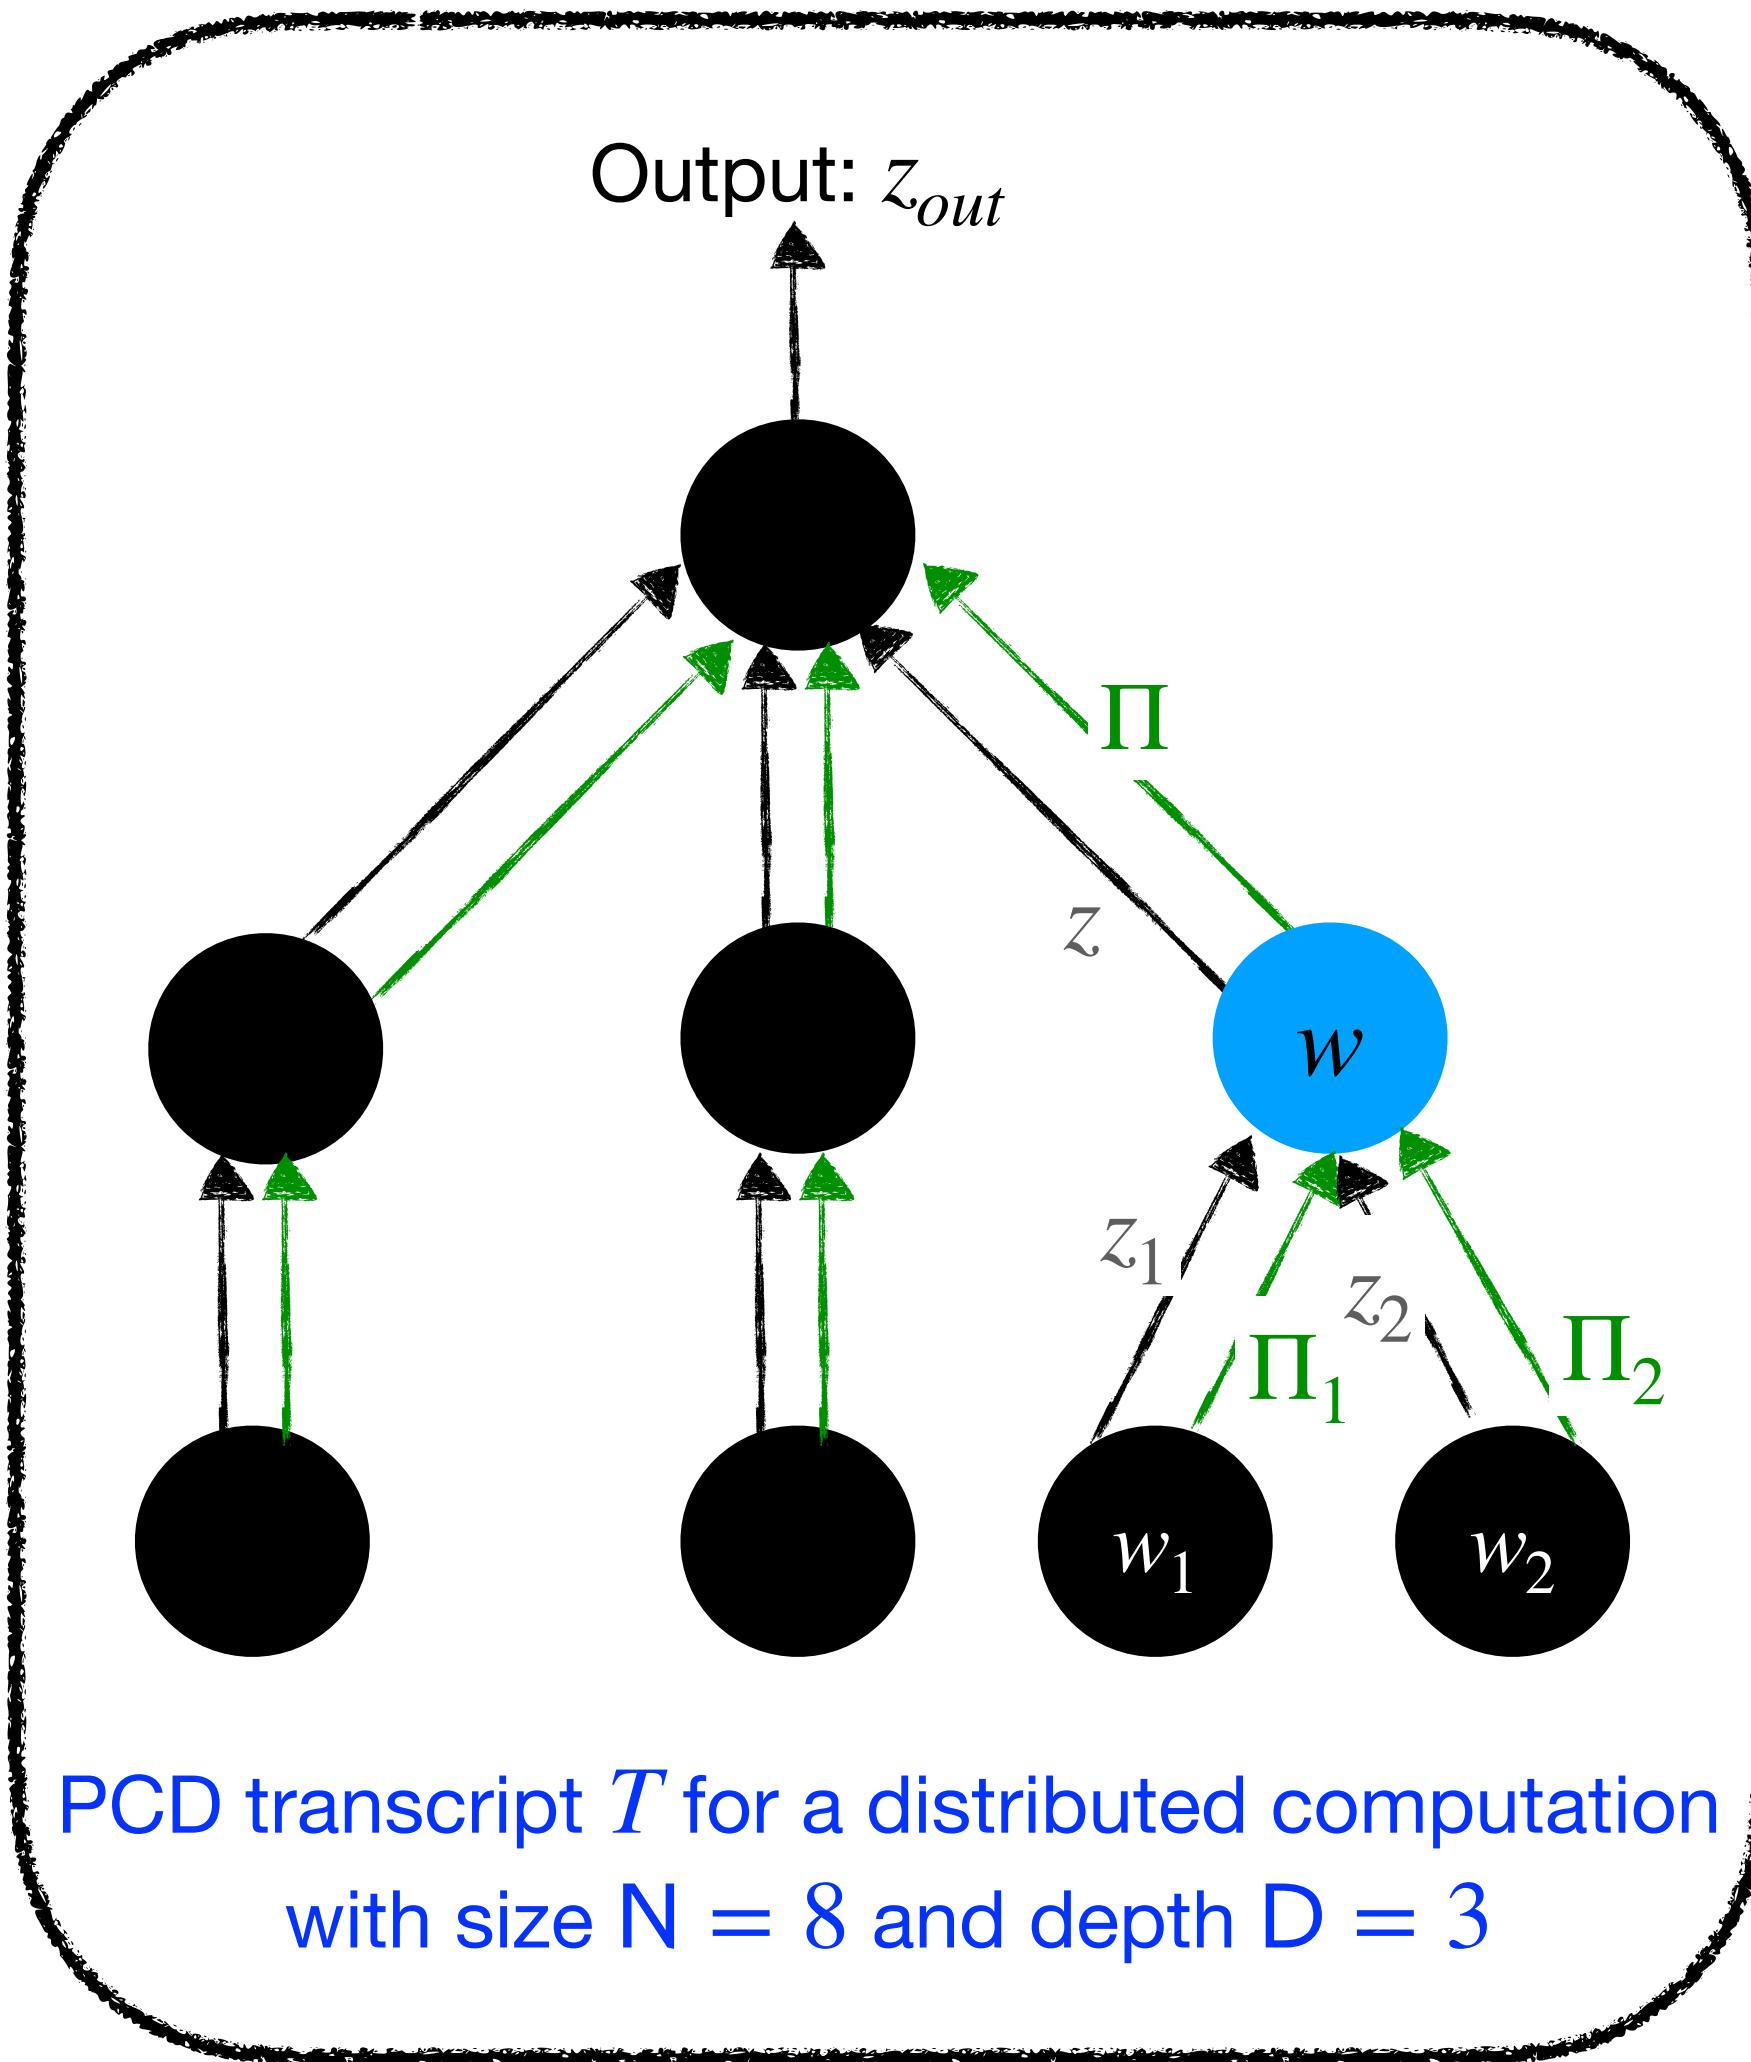


Pro: as secure as underlying SNARKs!

Naive PCD construction

Concatenate SNARK proofs

SNARK prover for compliance predicate ϕ



Pro: as secure as underlying SNARKs!

Con: Π is NOT succinct (linear in number of vertices)

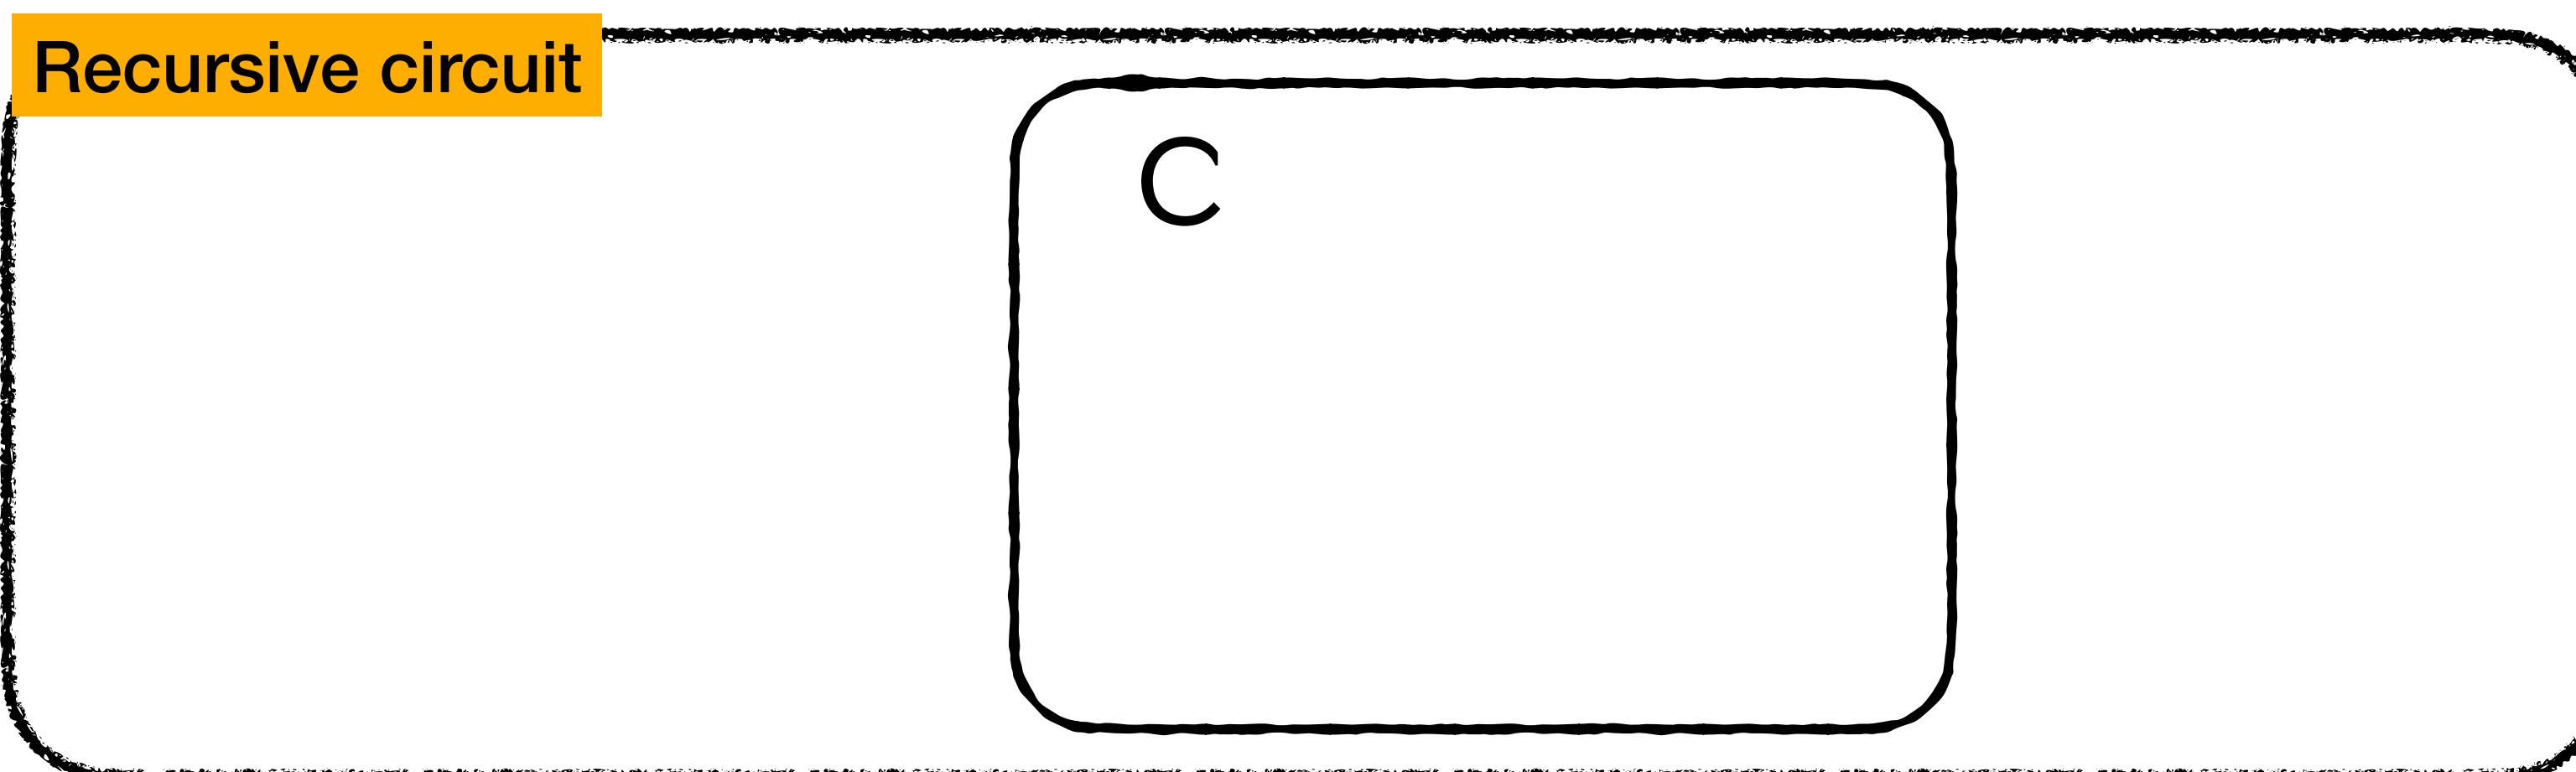
Recursively compose the SNARK proofs

Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.

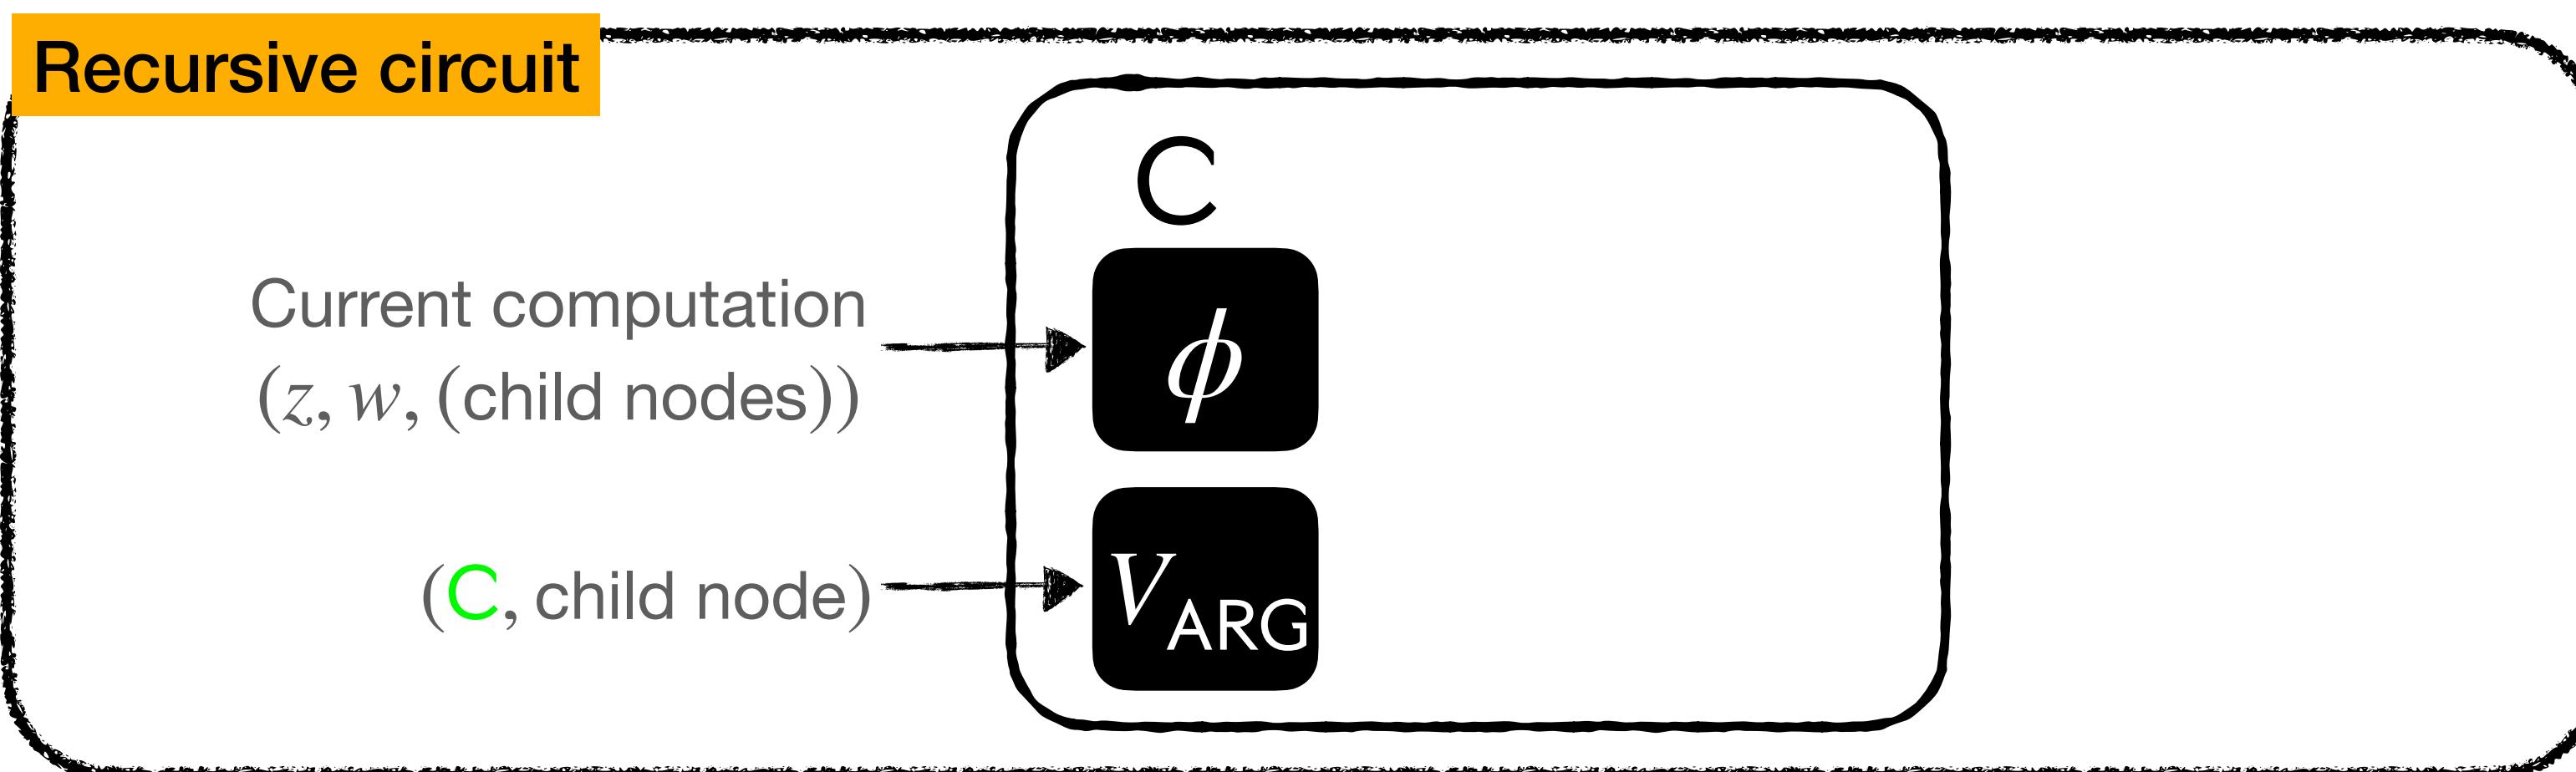
Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



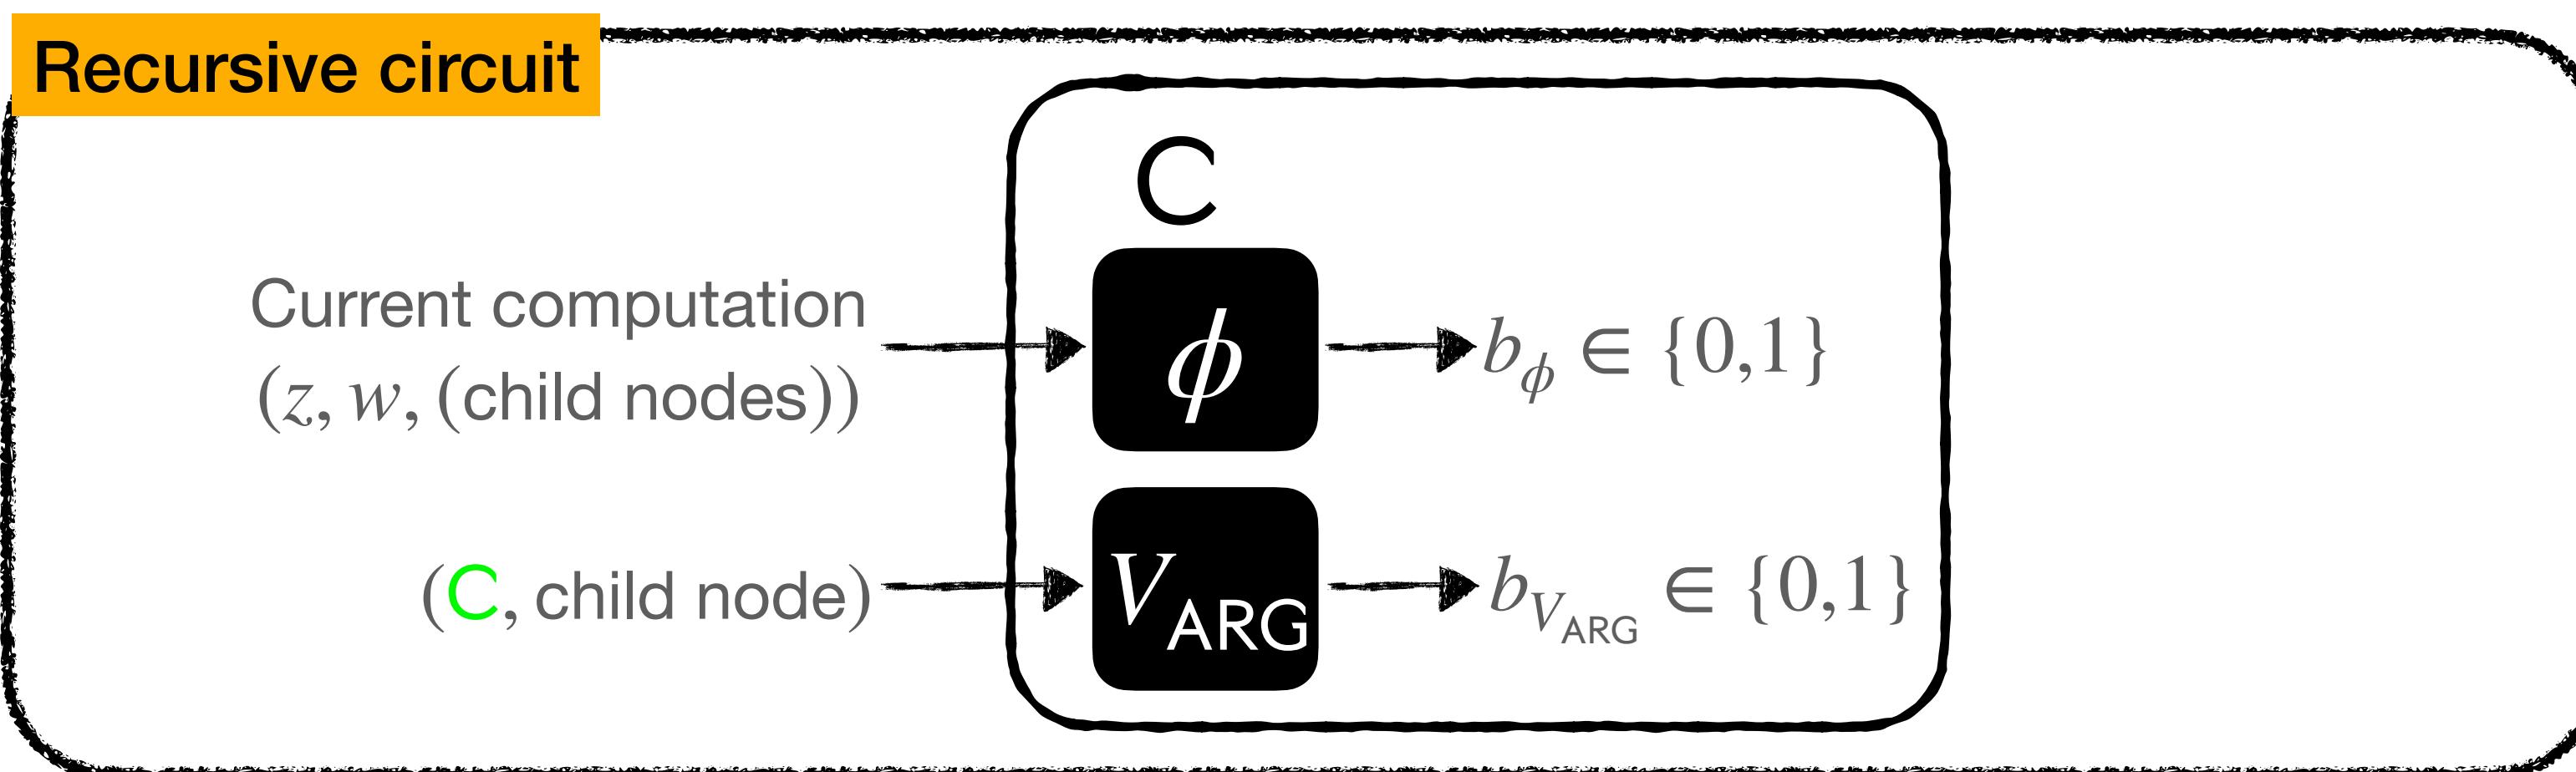
Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



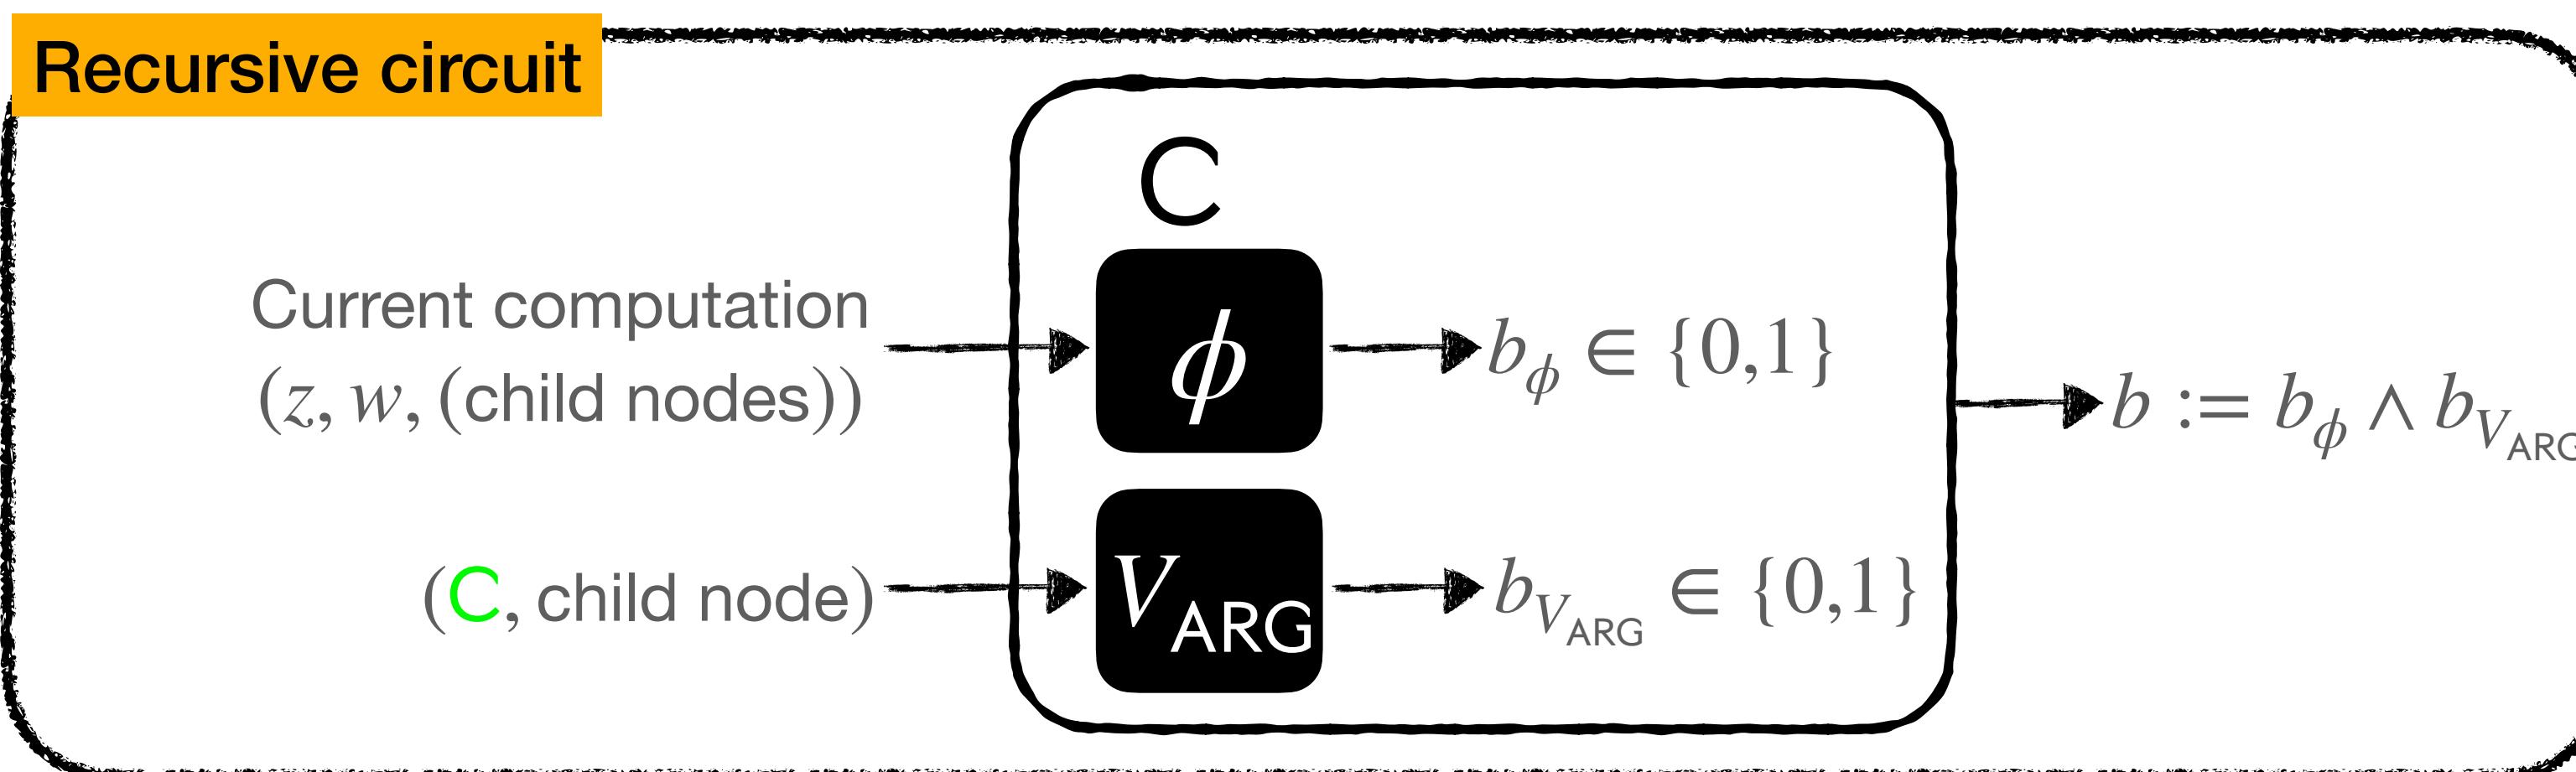
Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



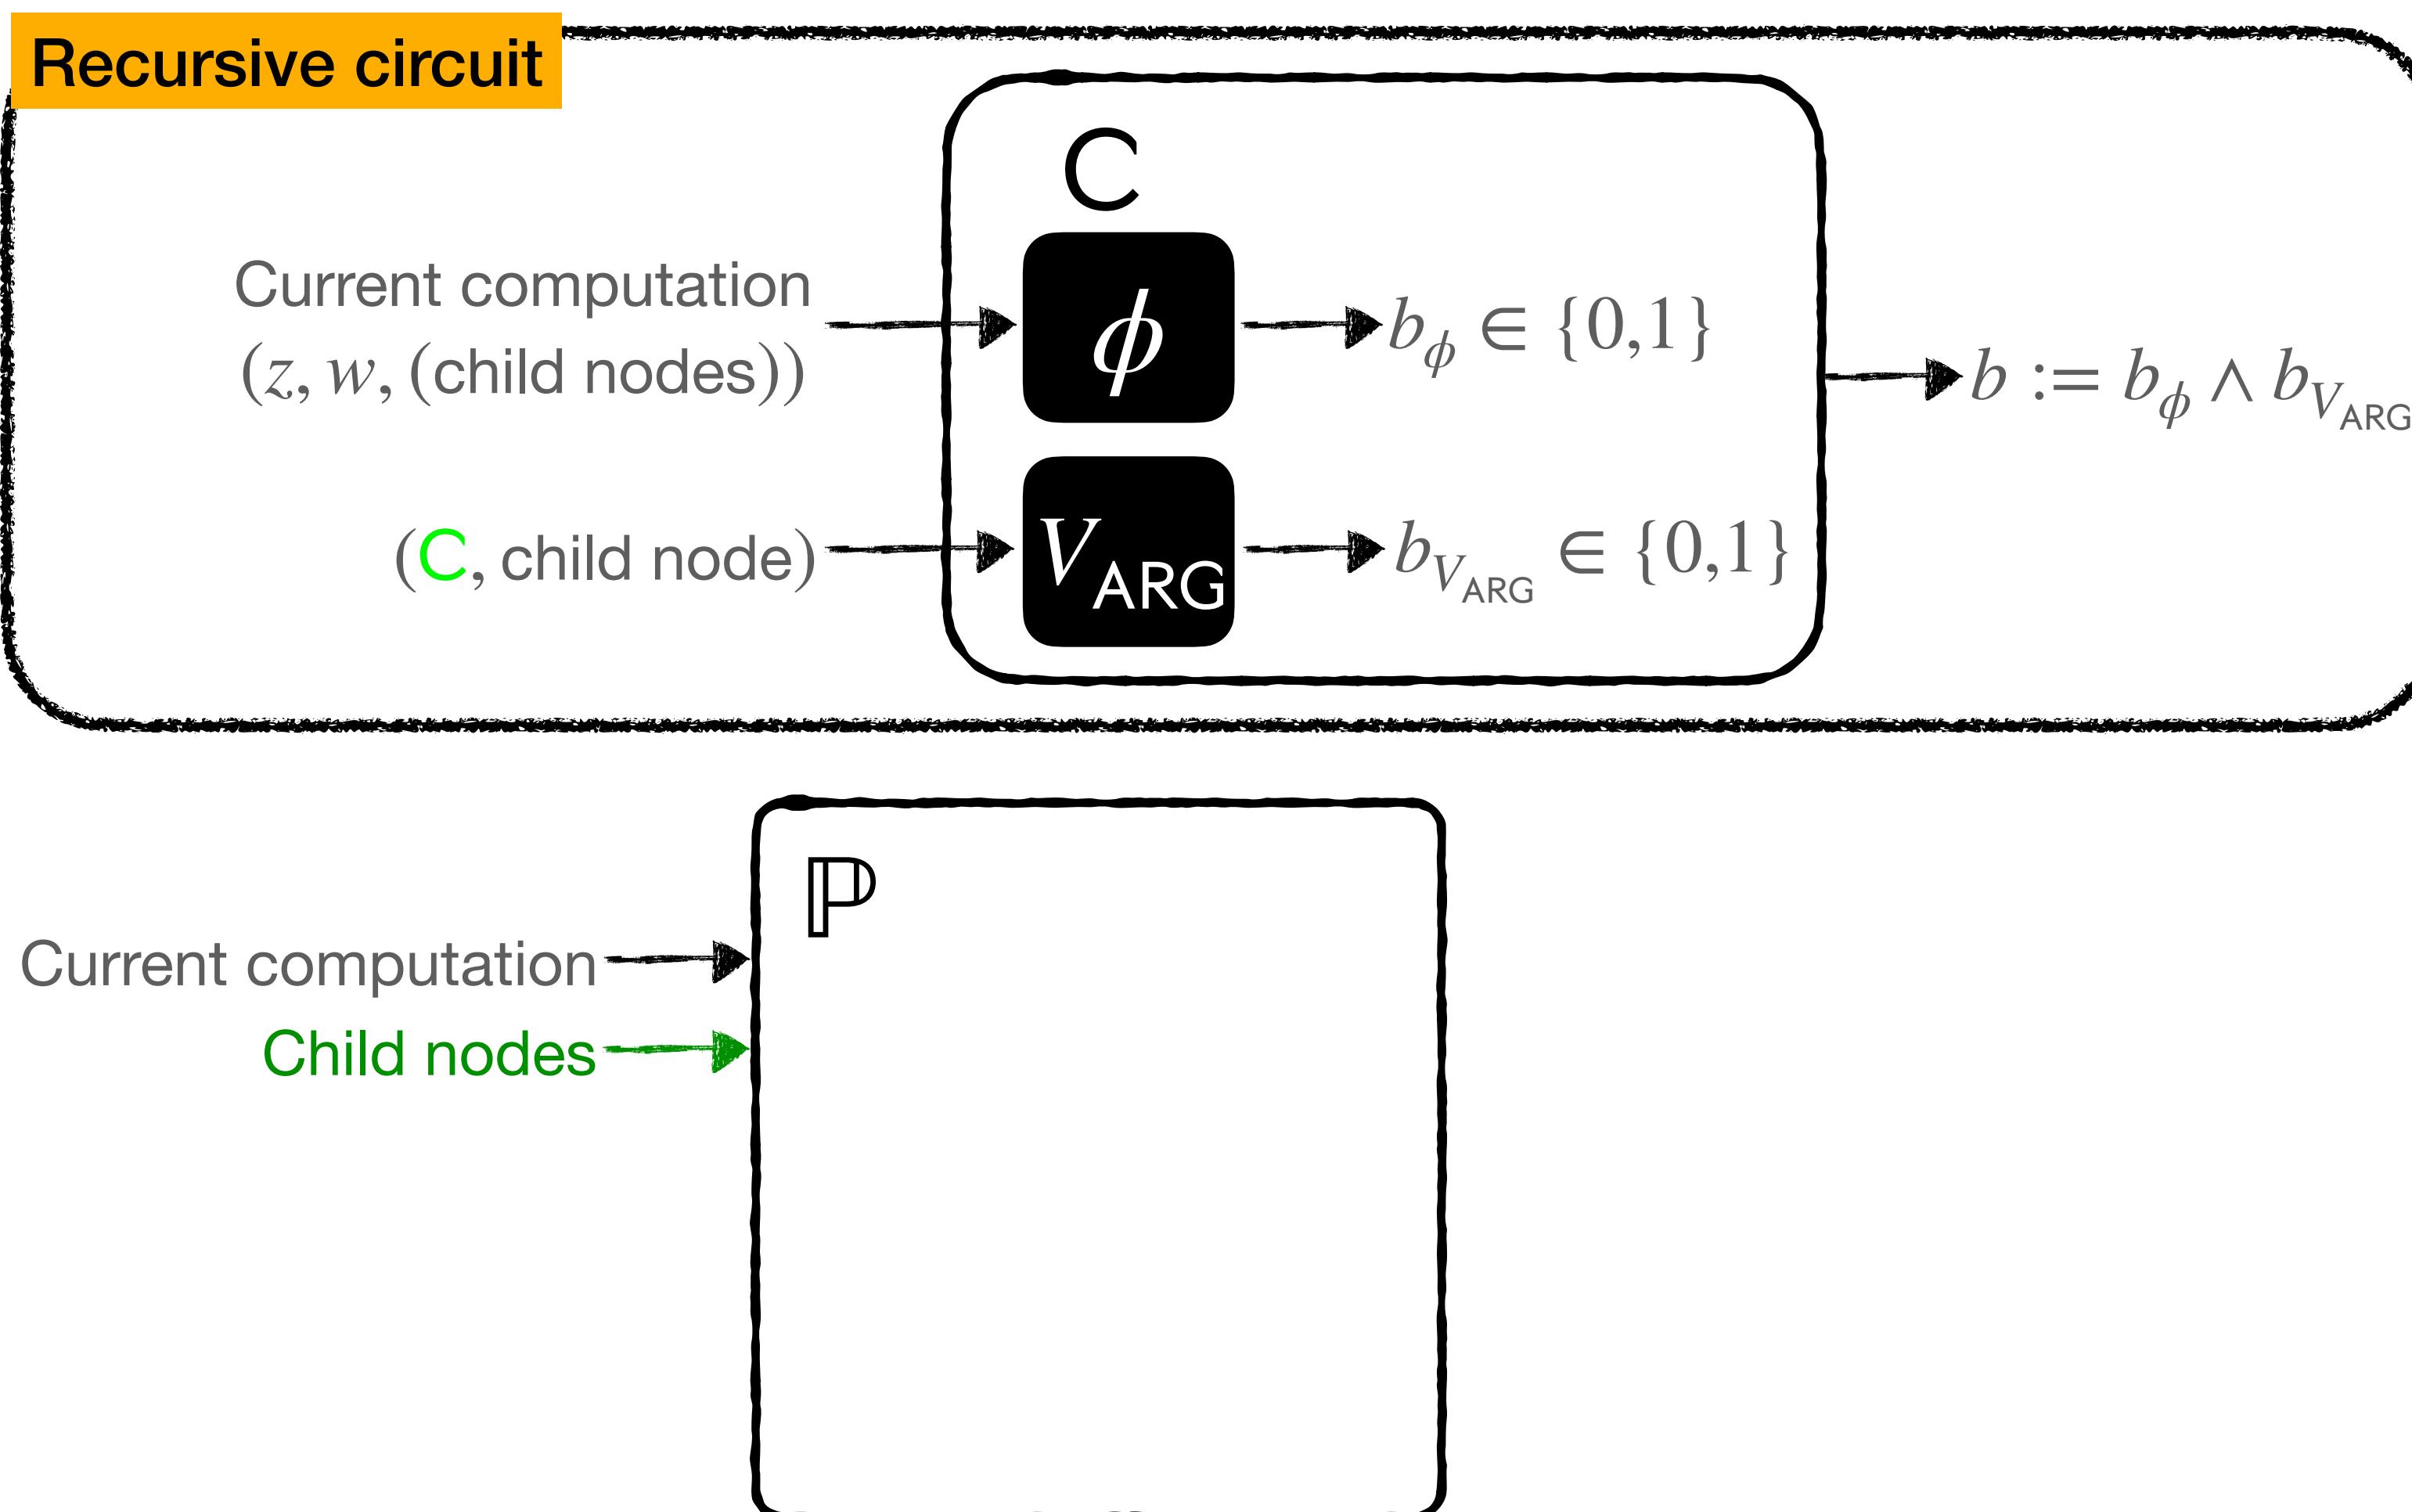
Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



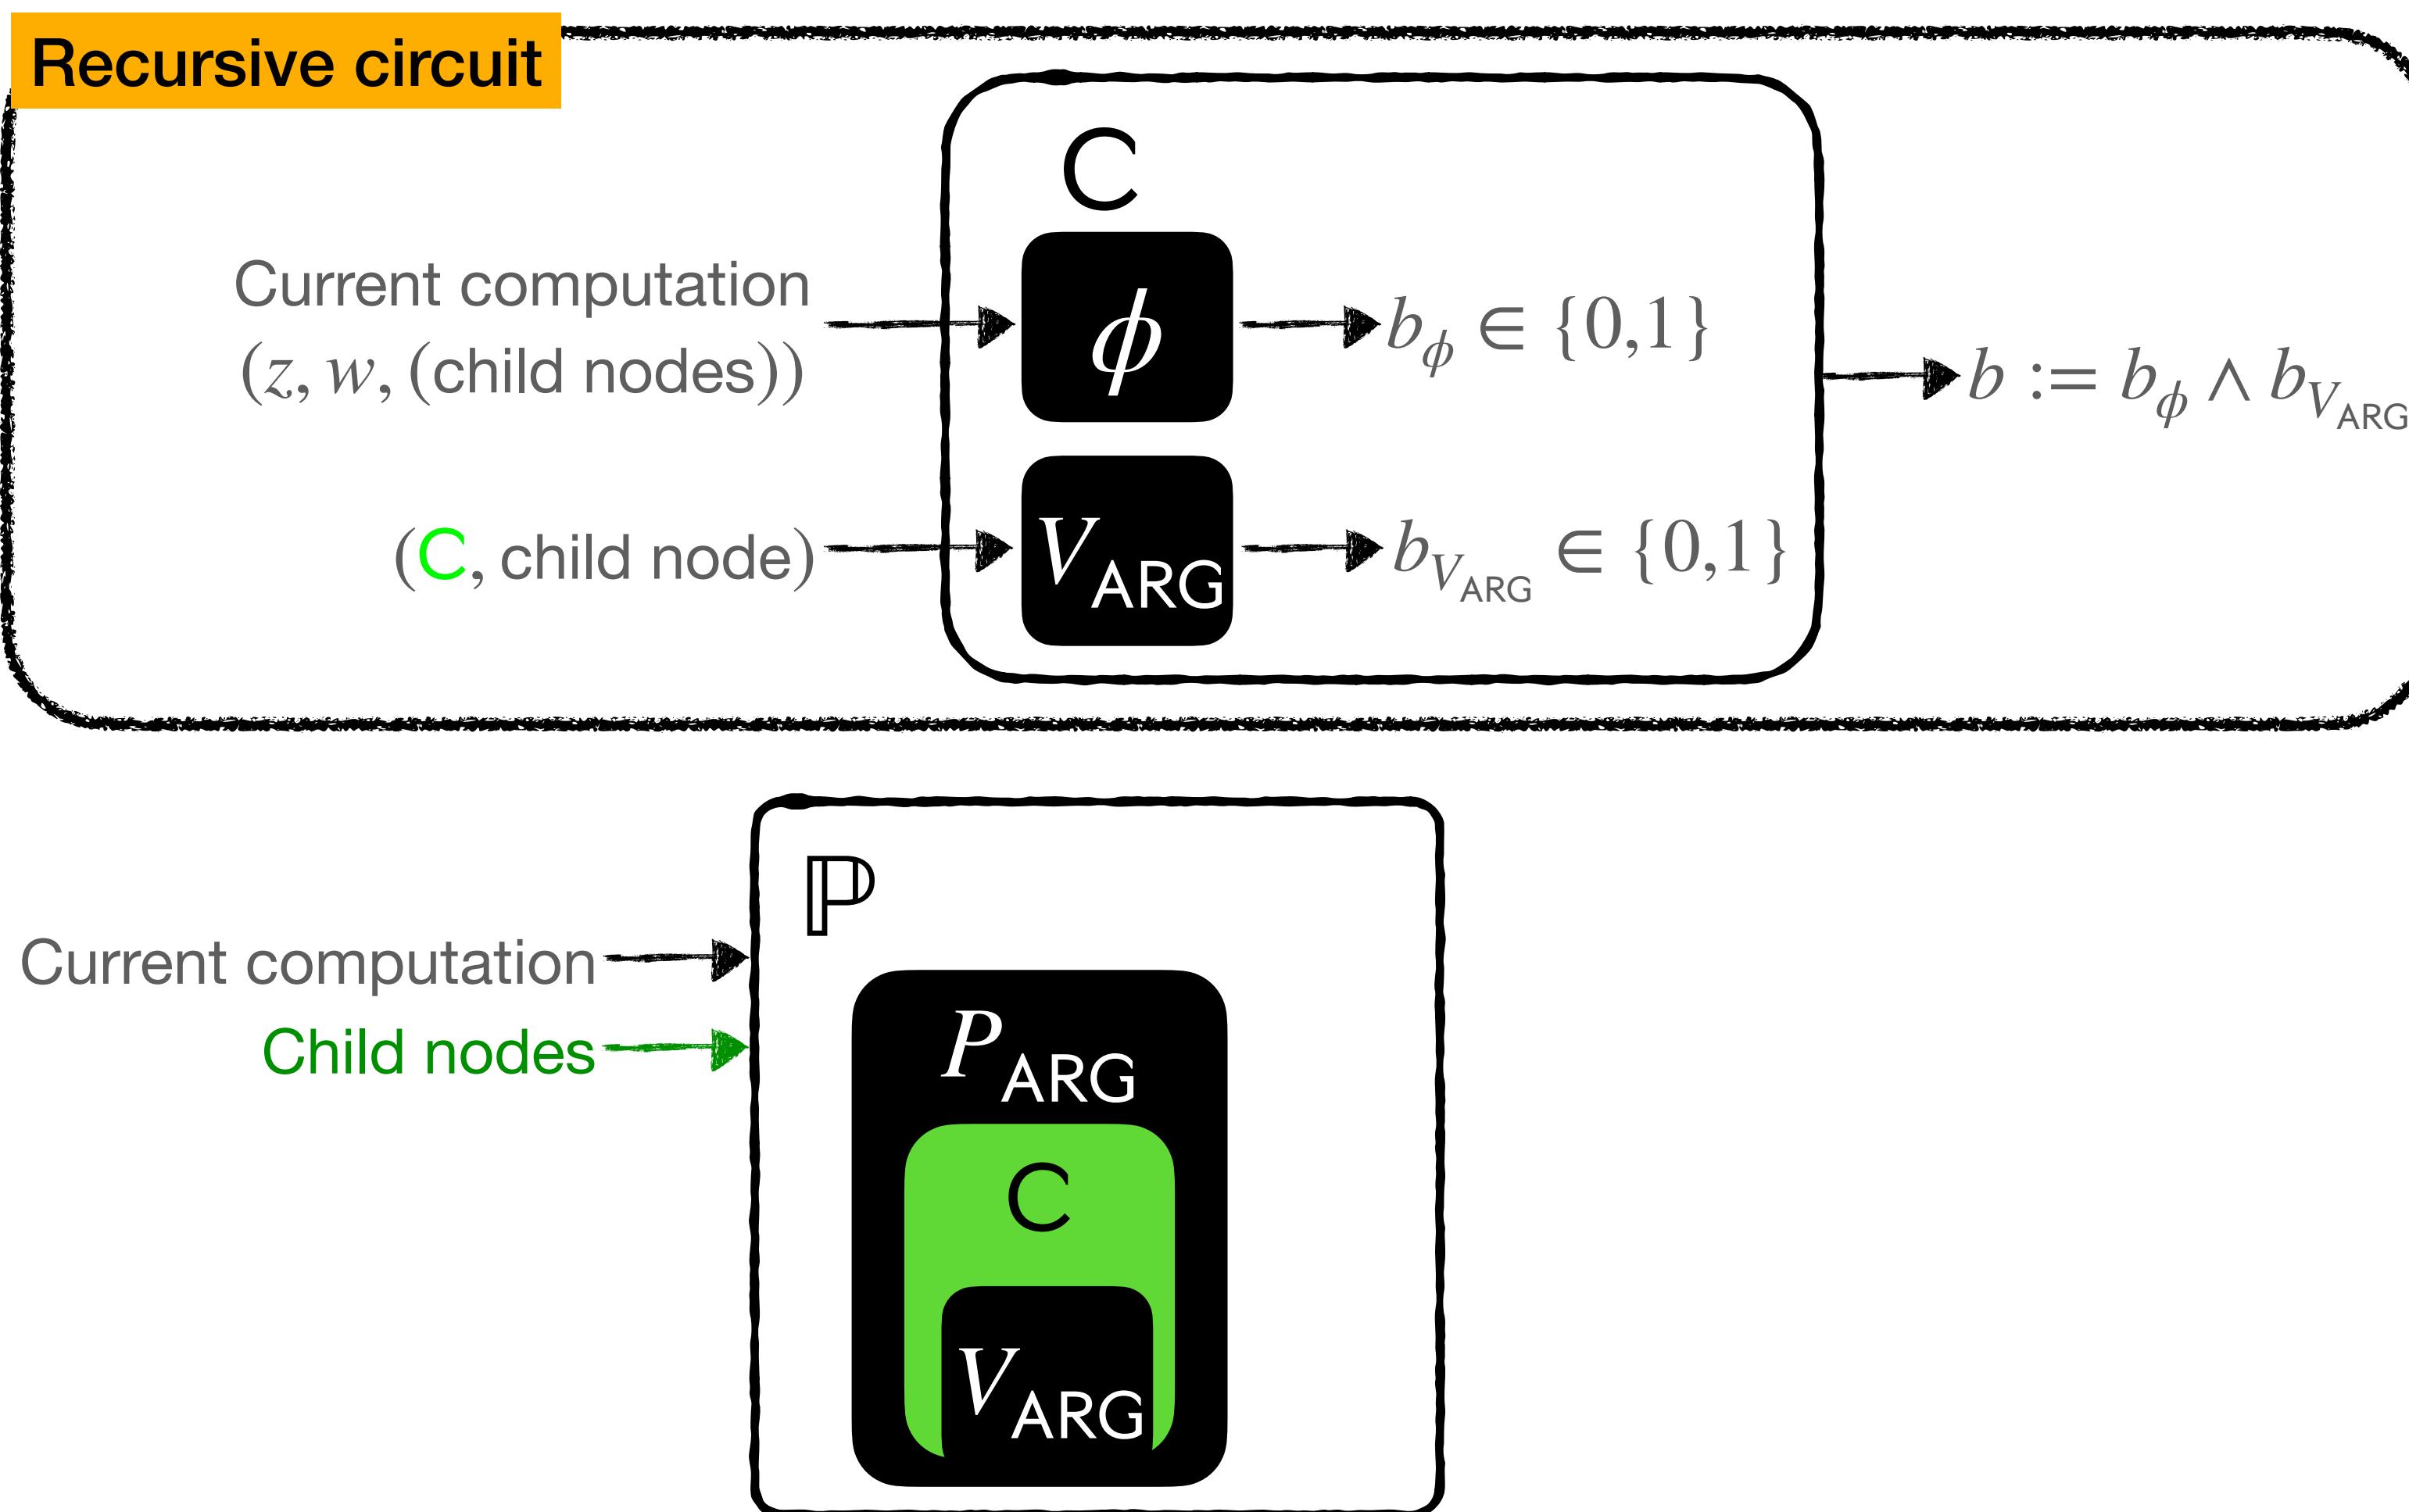
Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



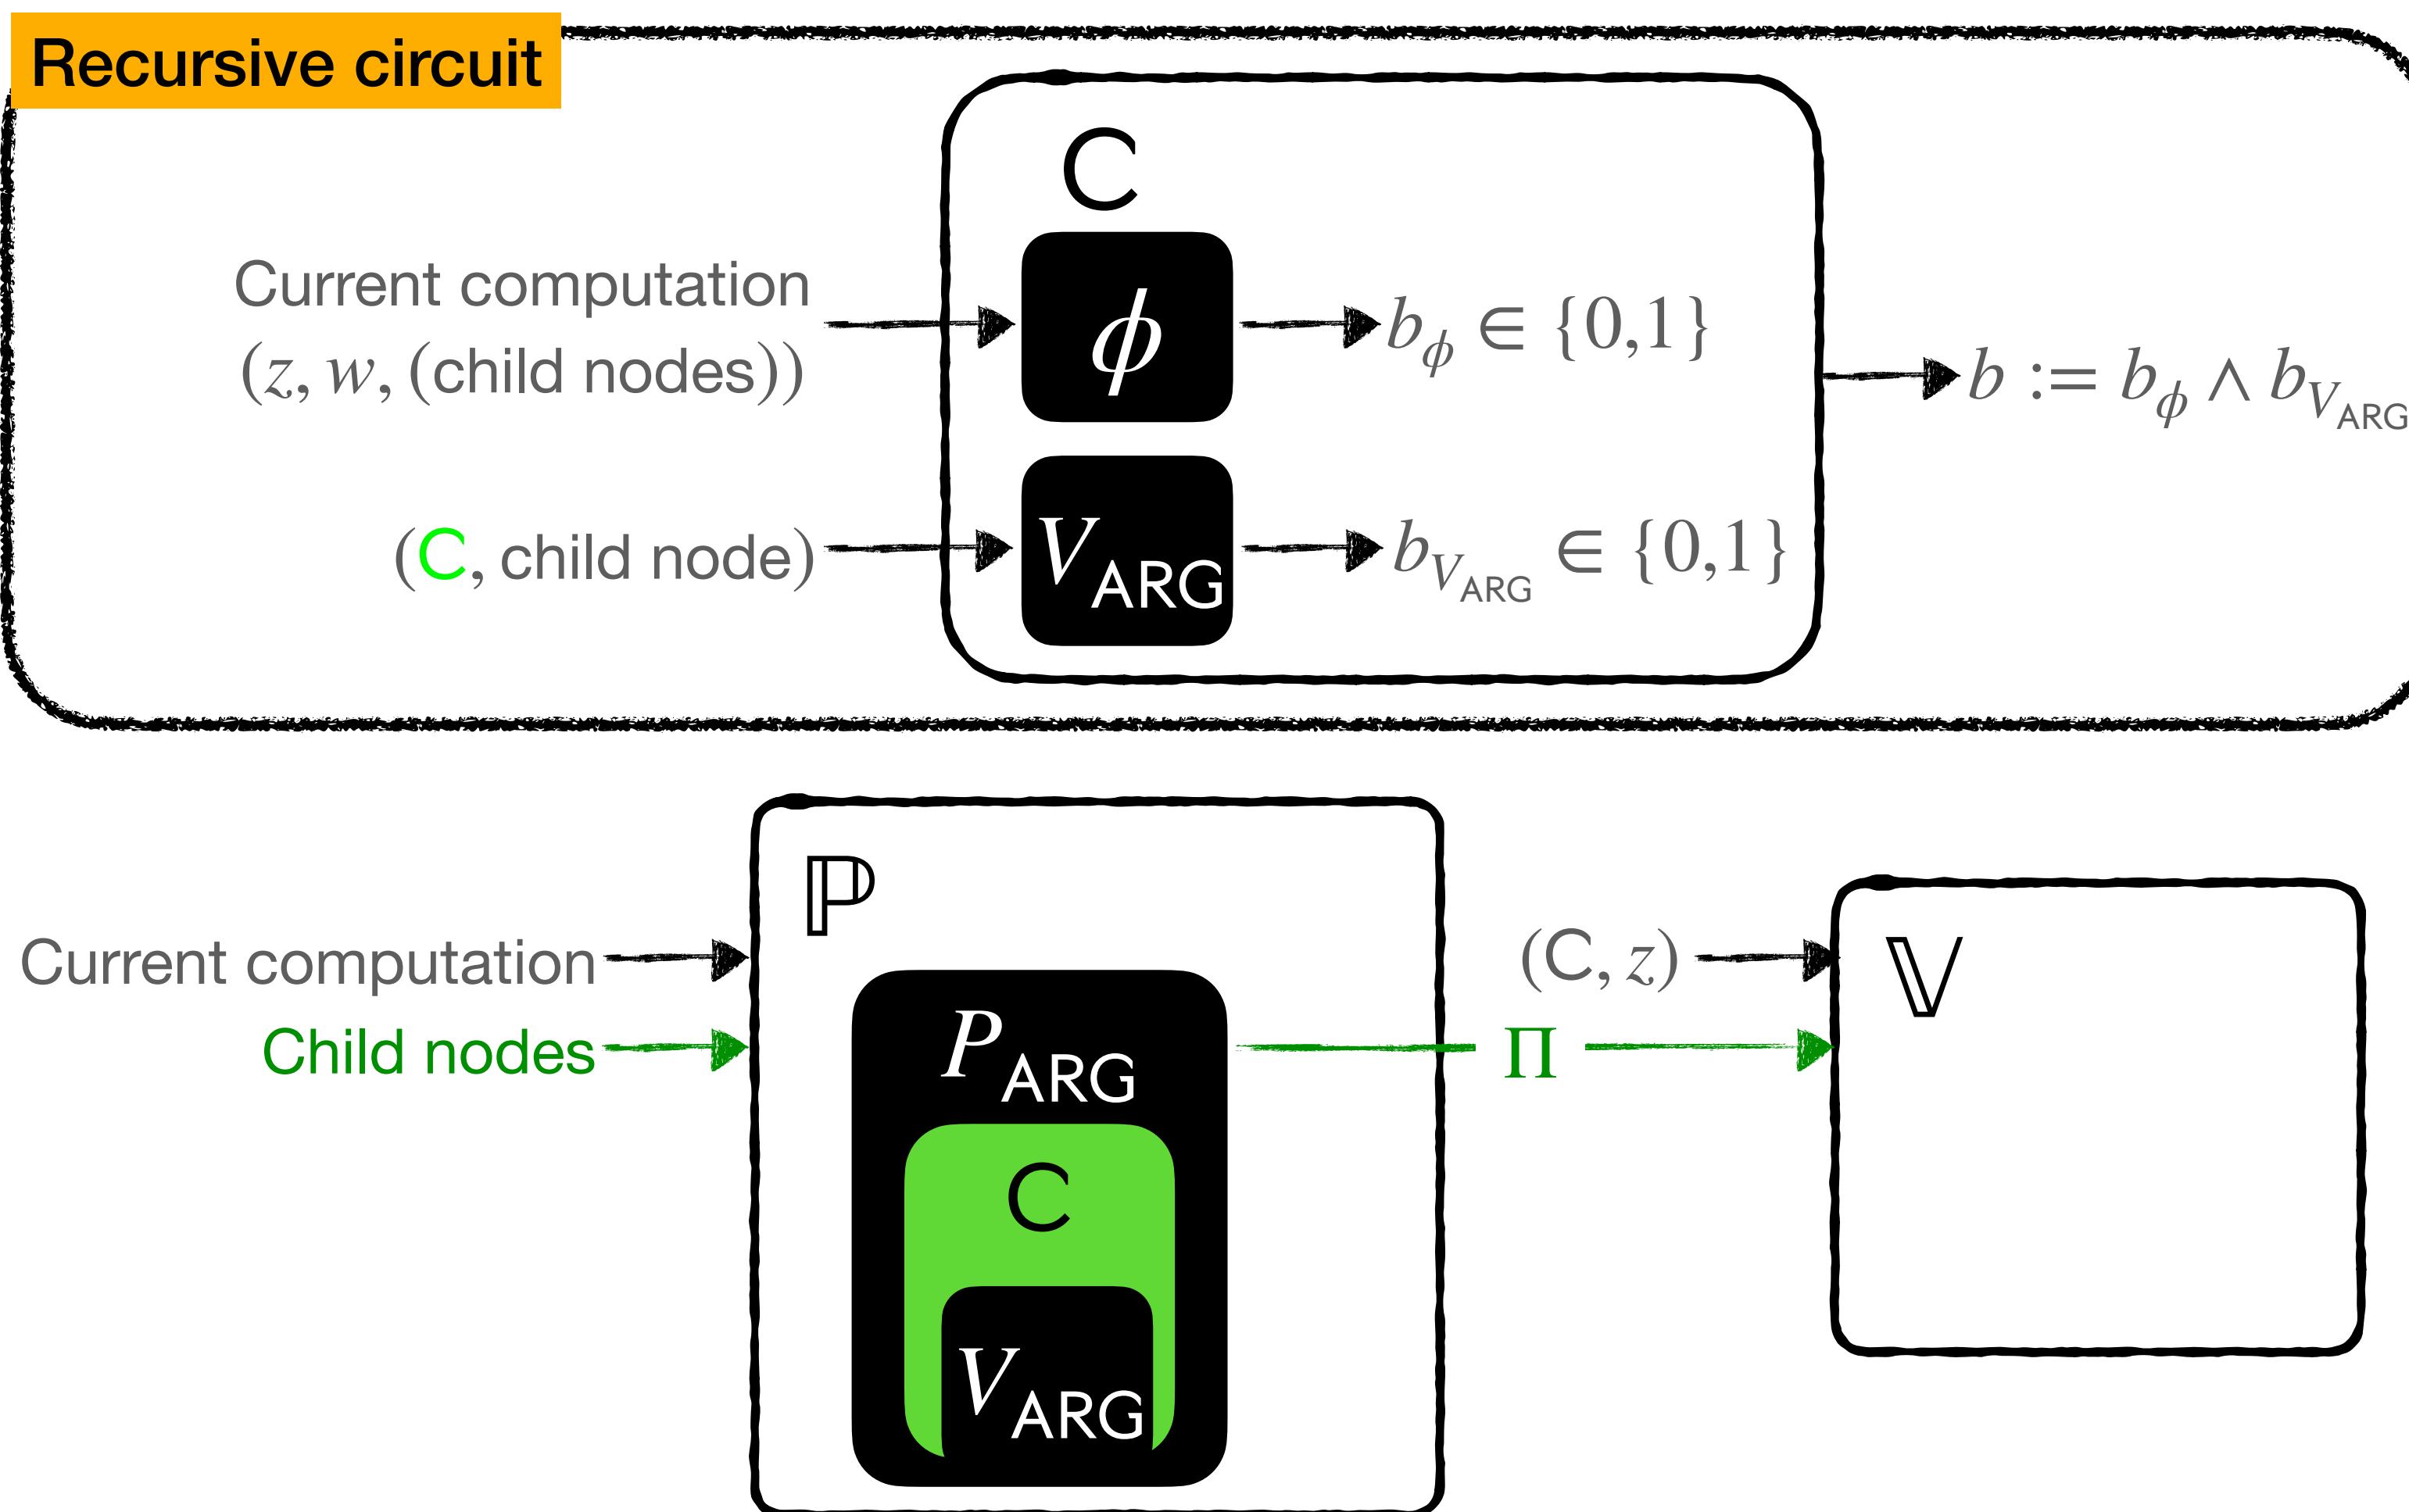
Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



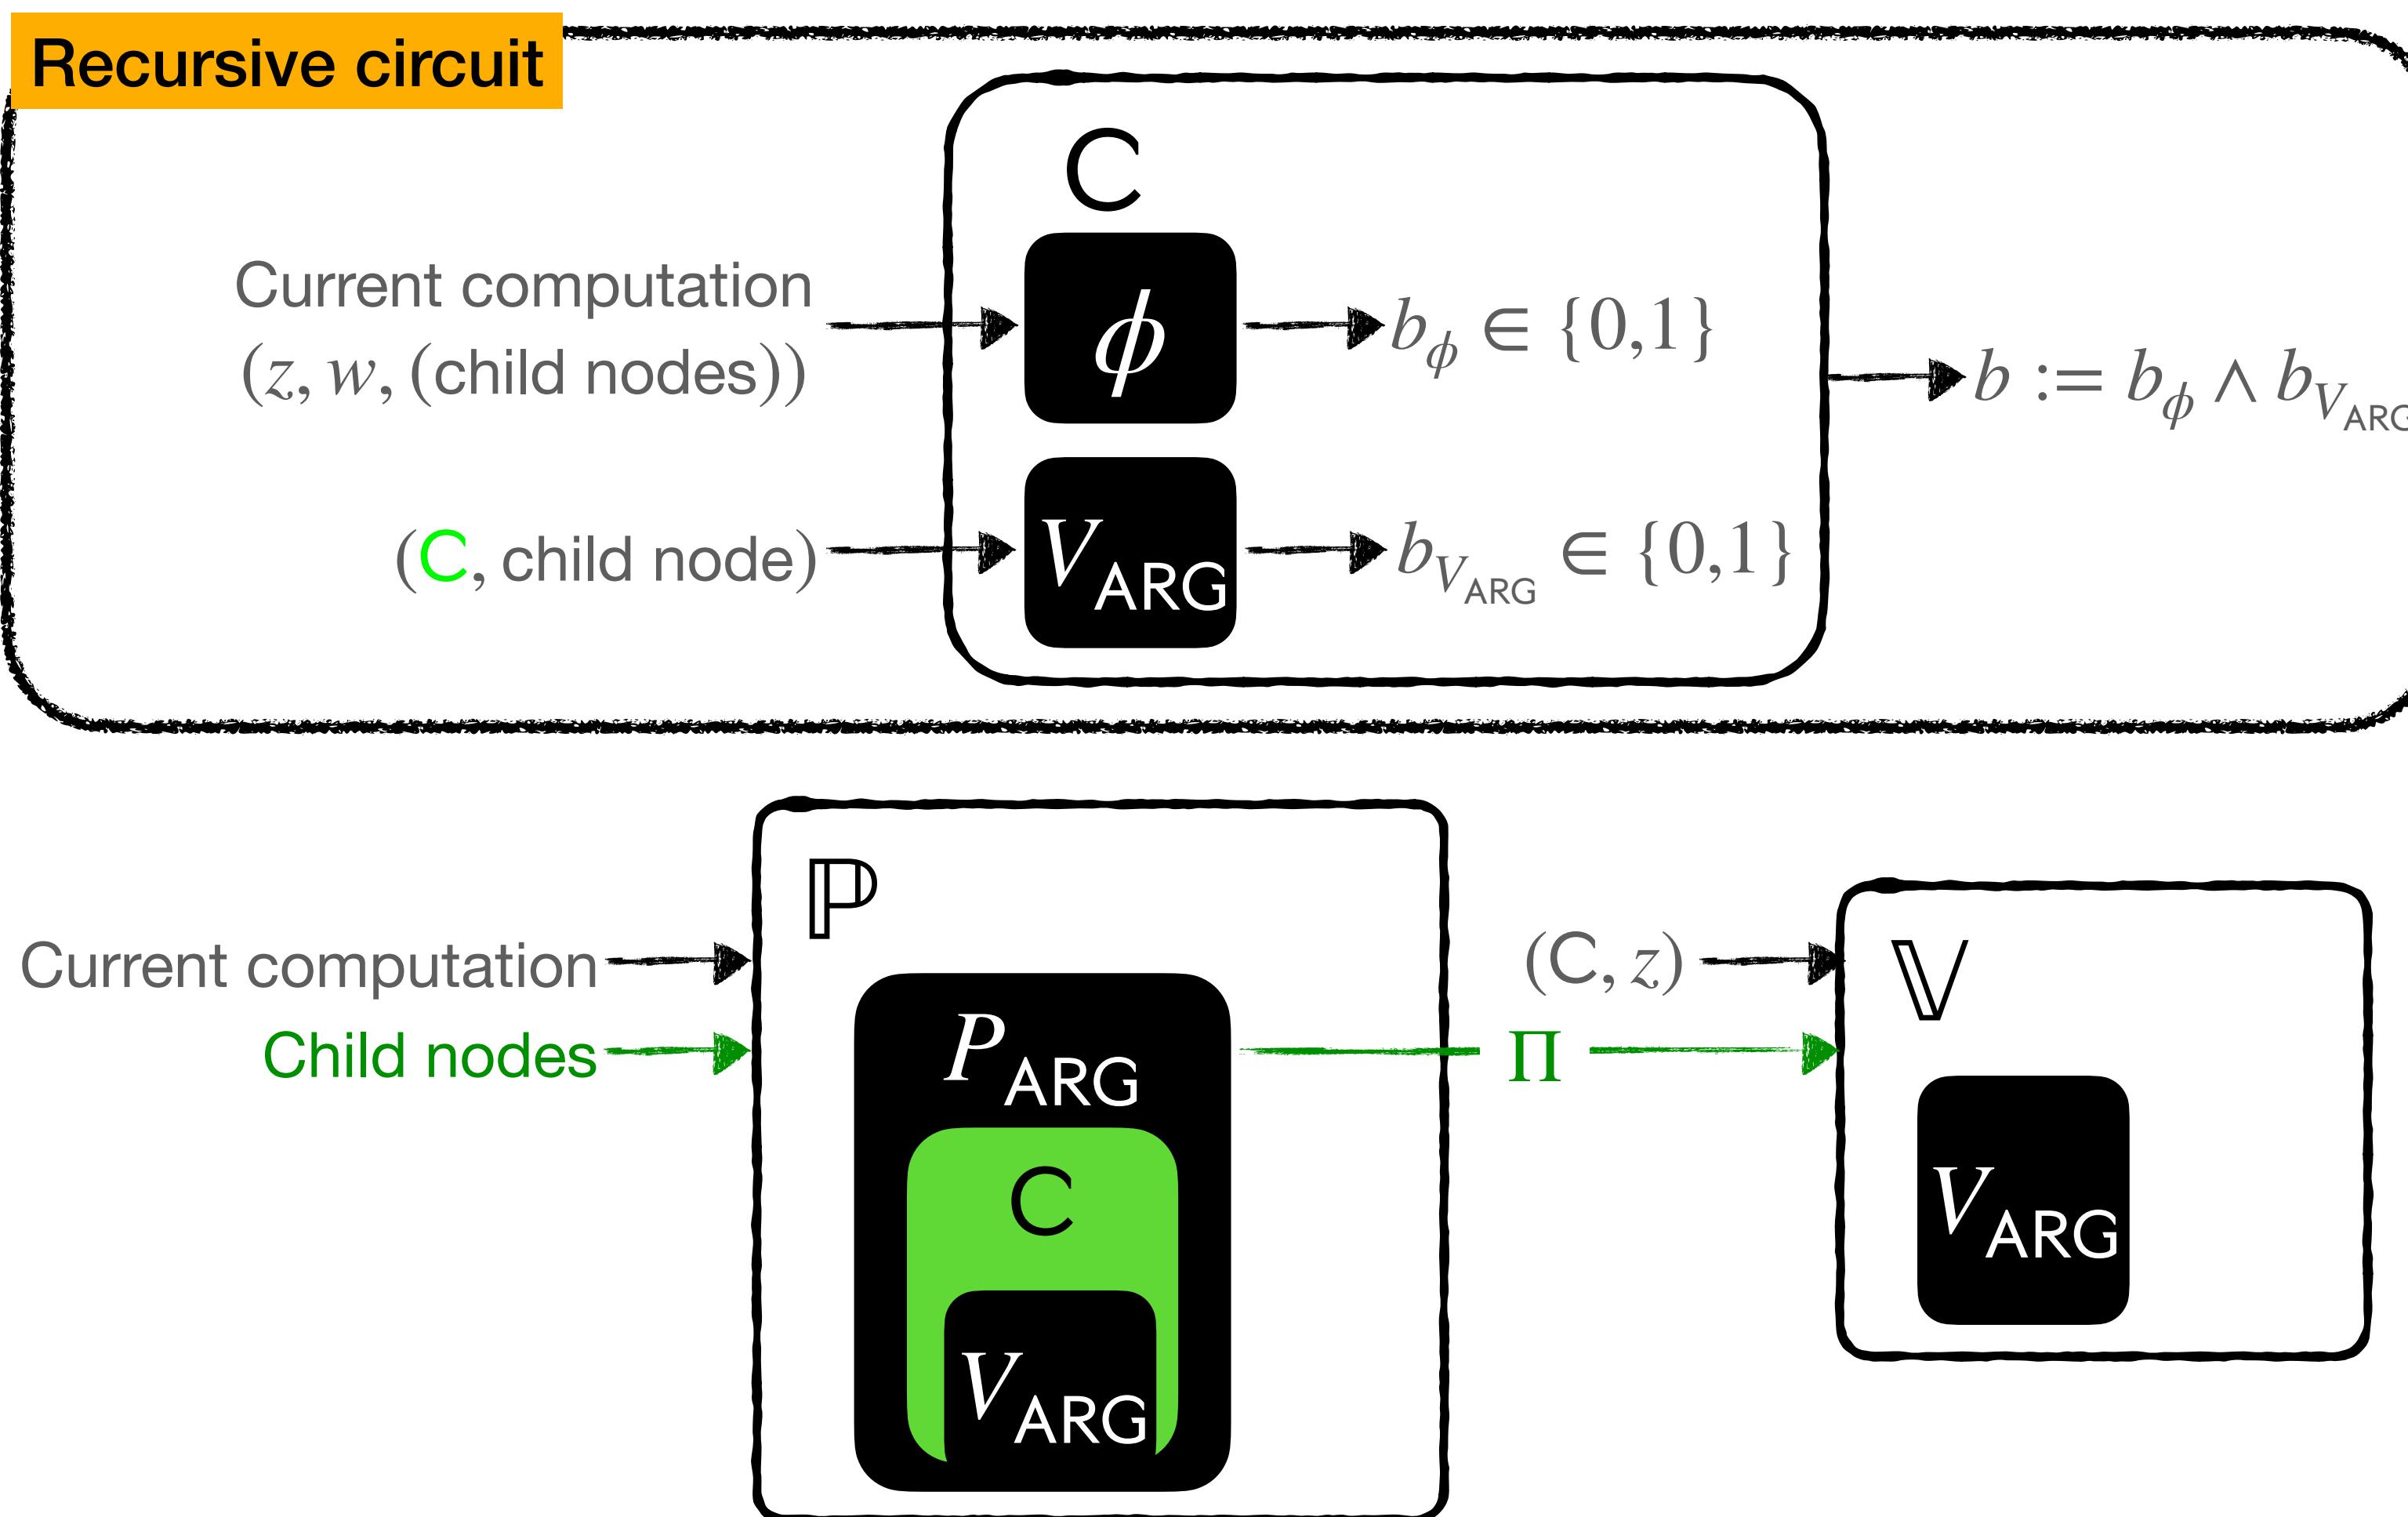
Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



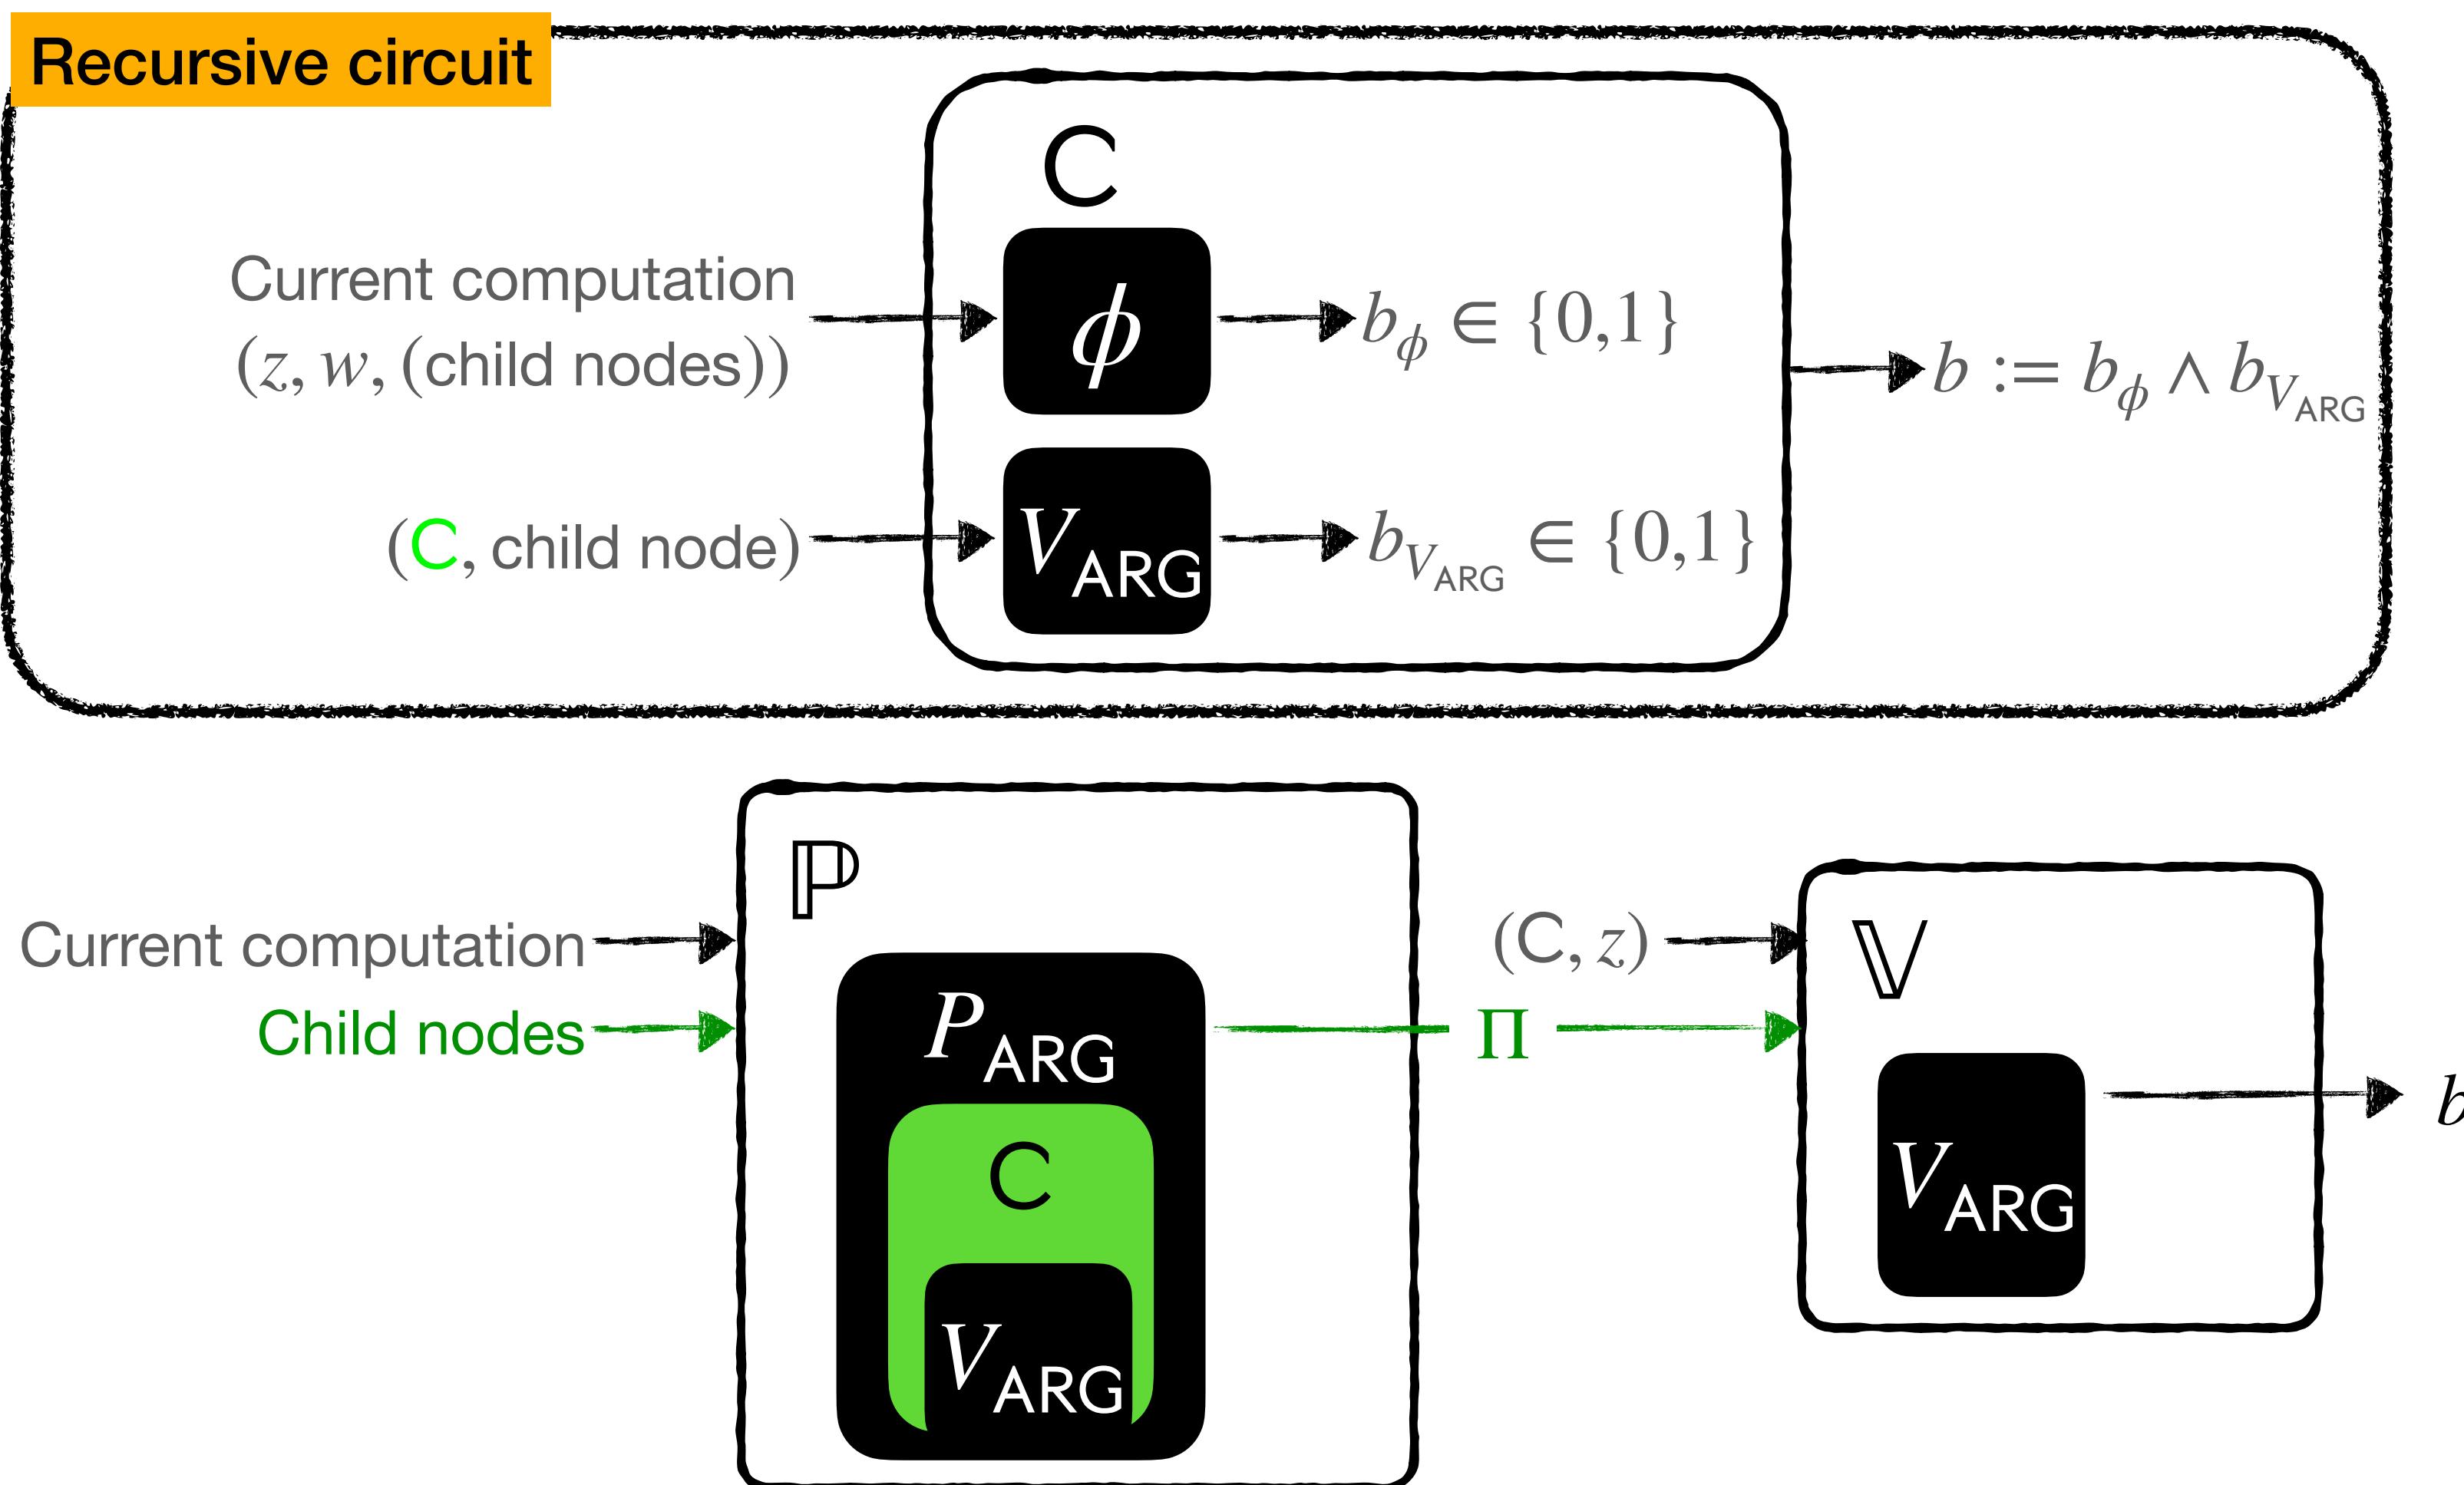
Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



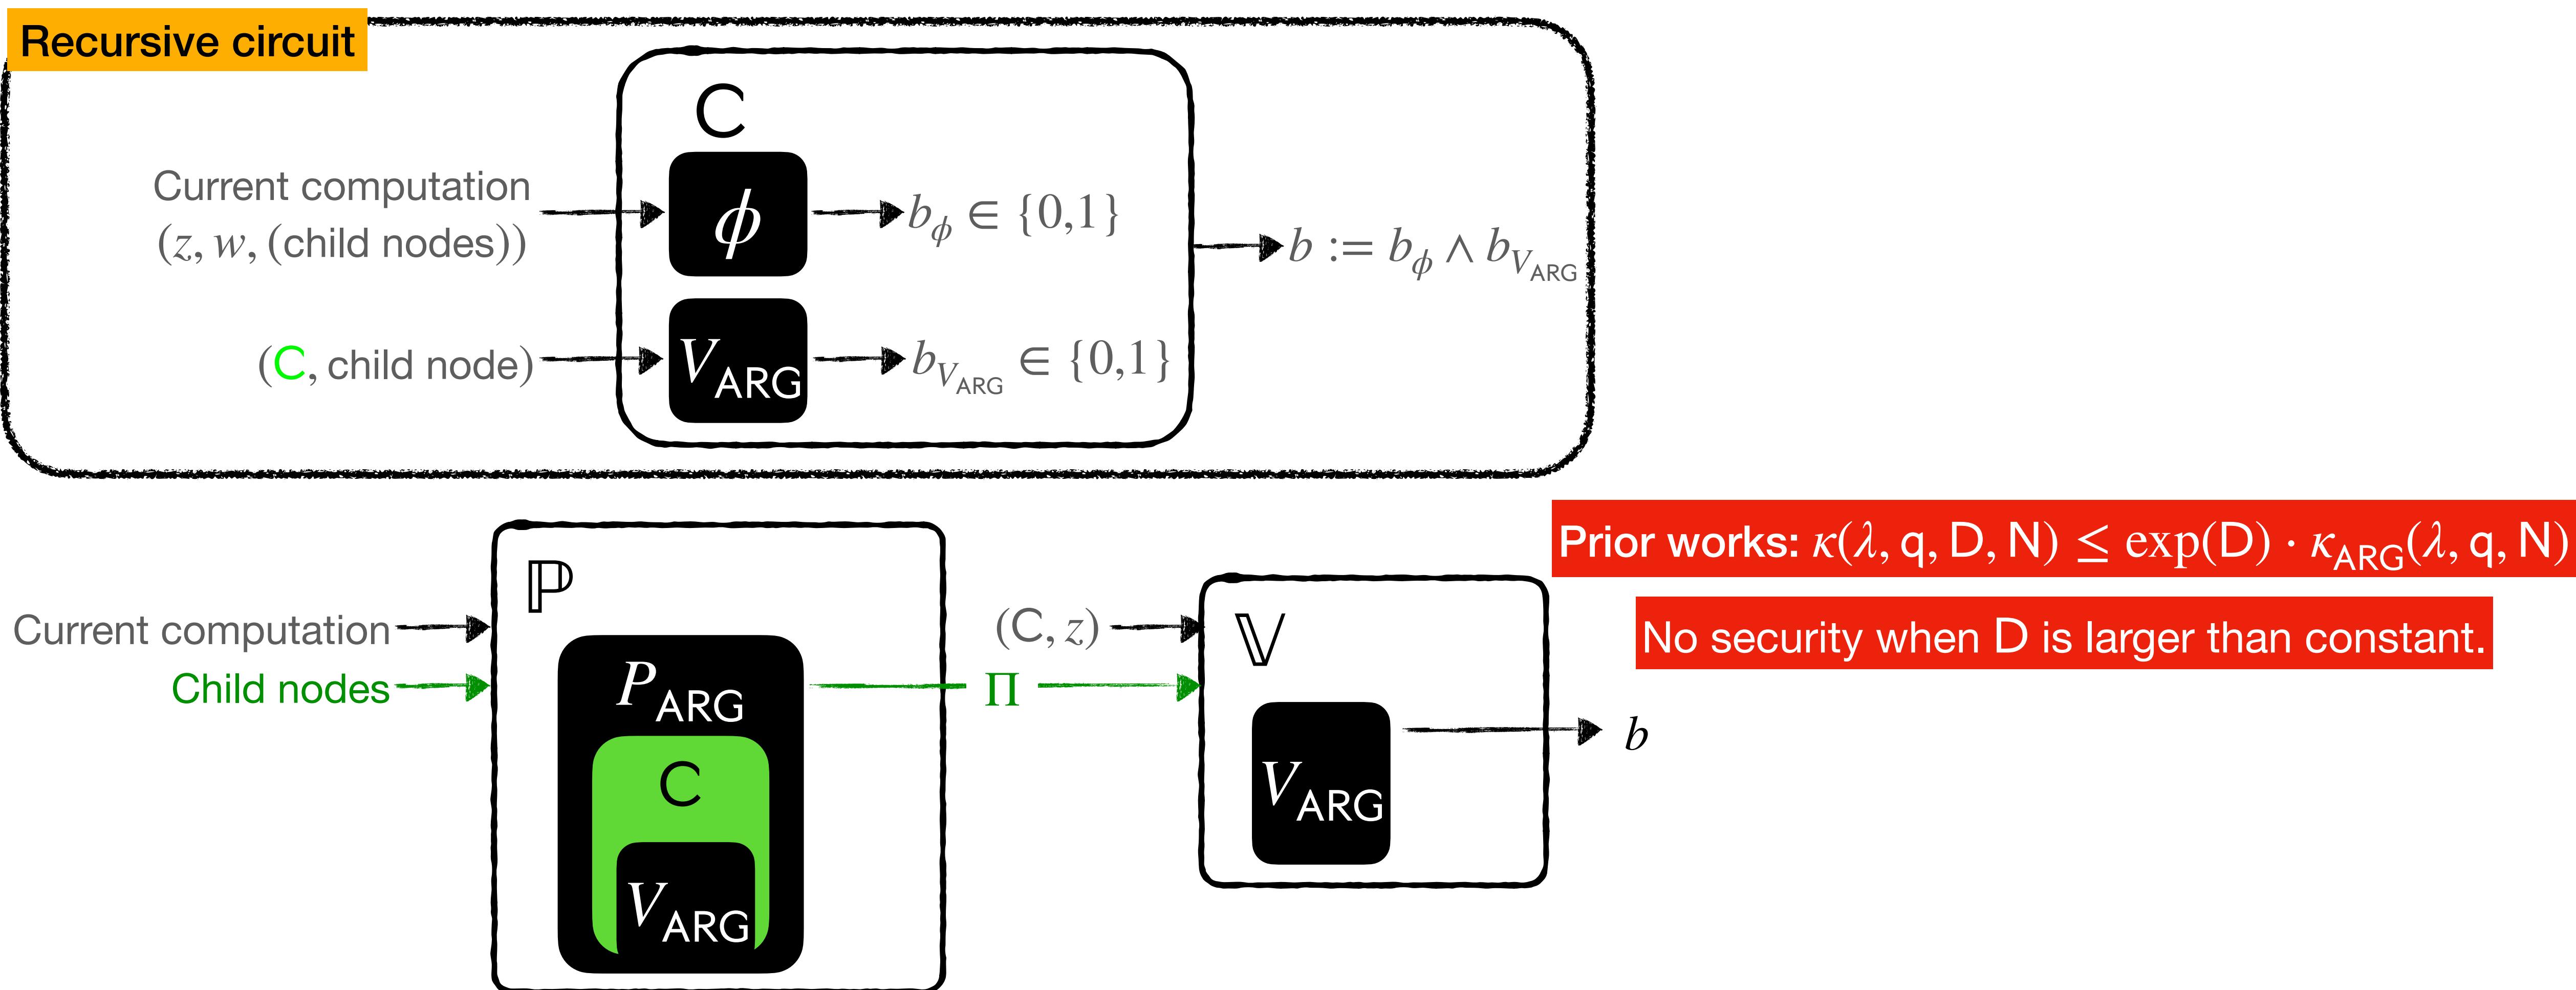
Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



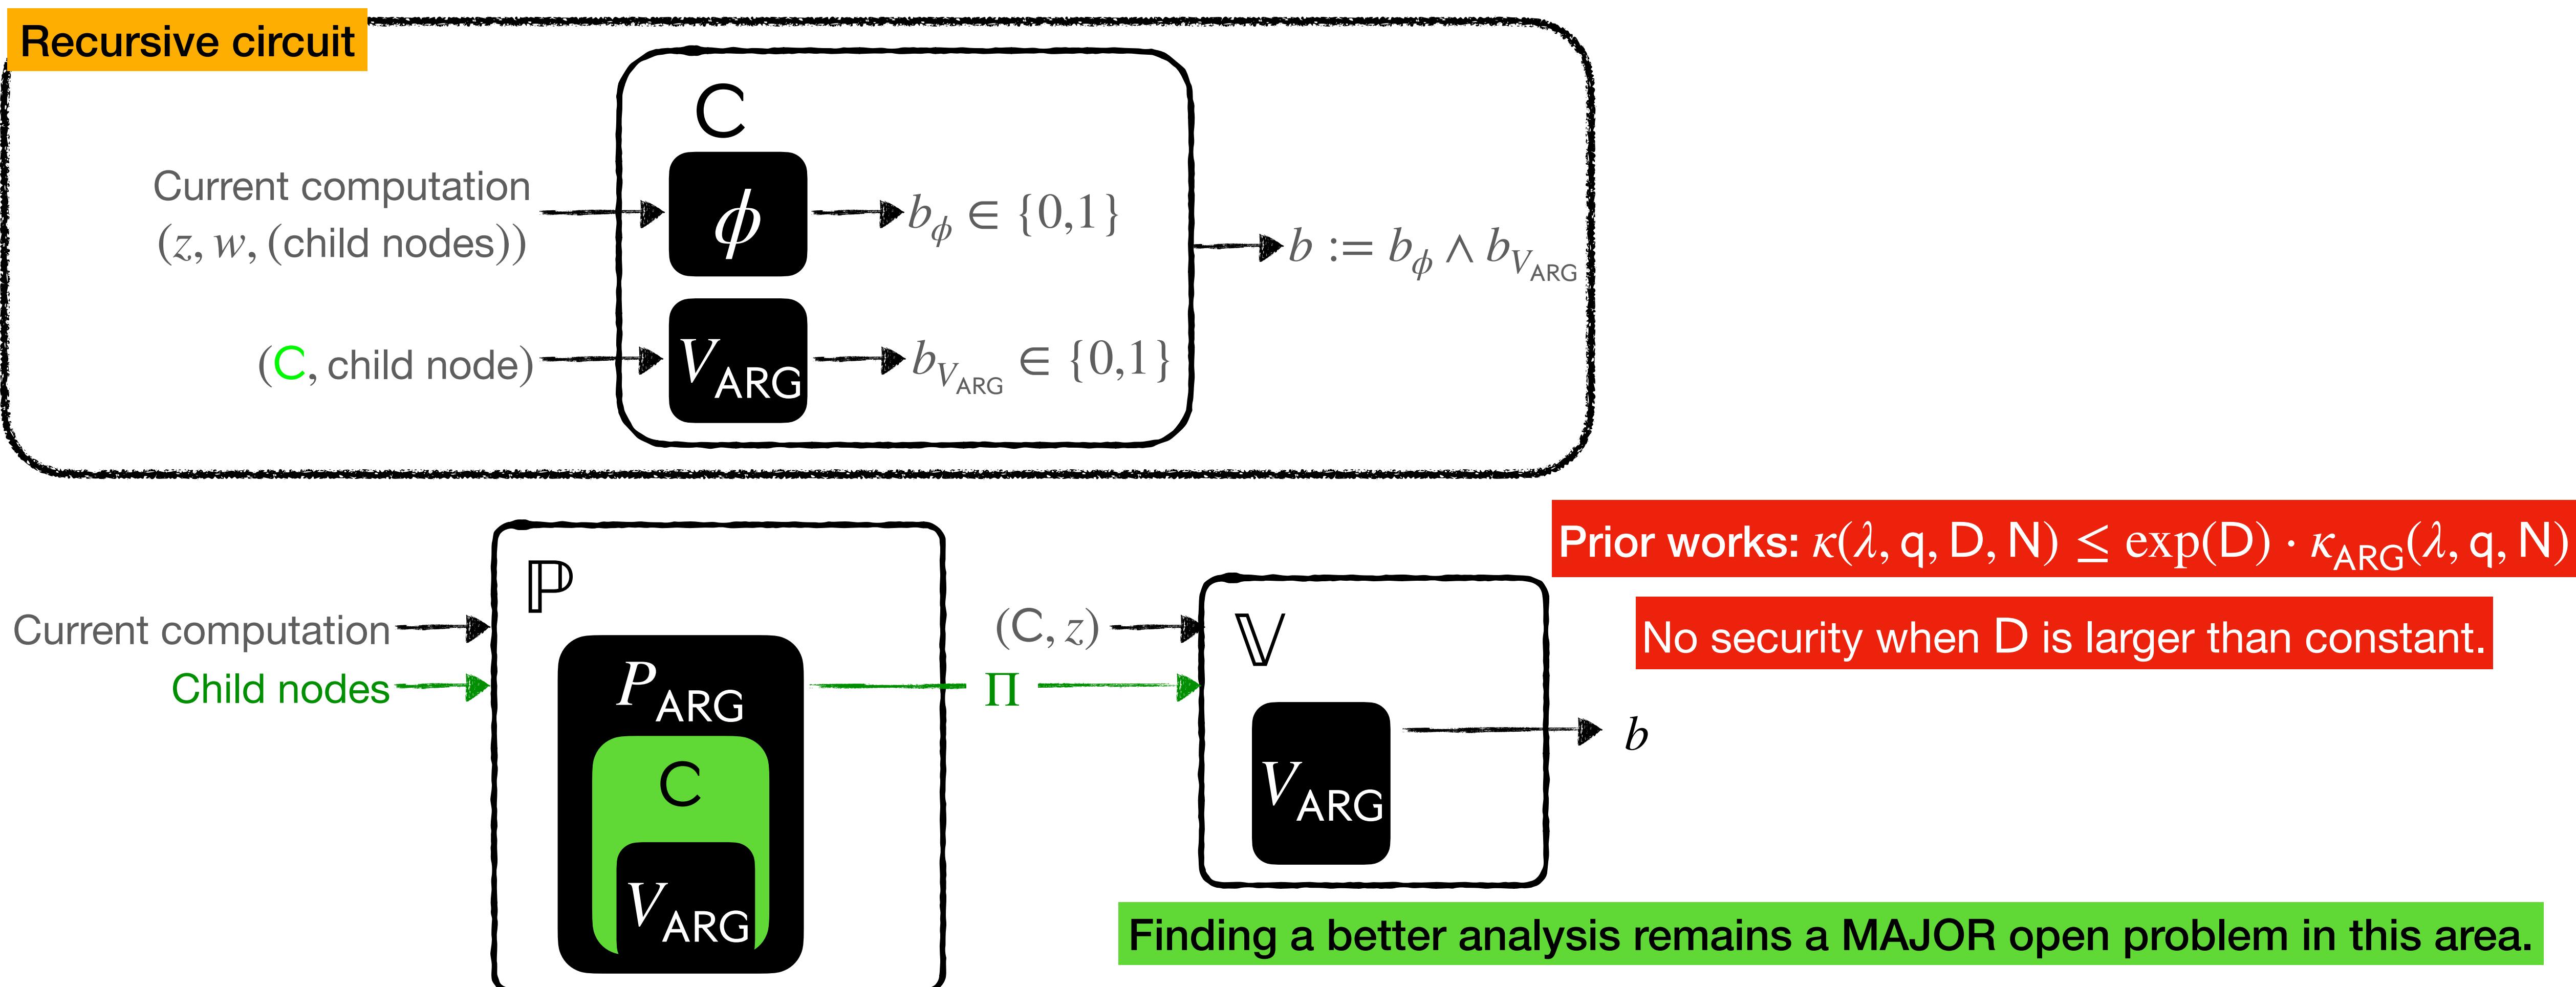
Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



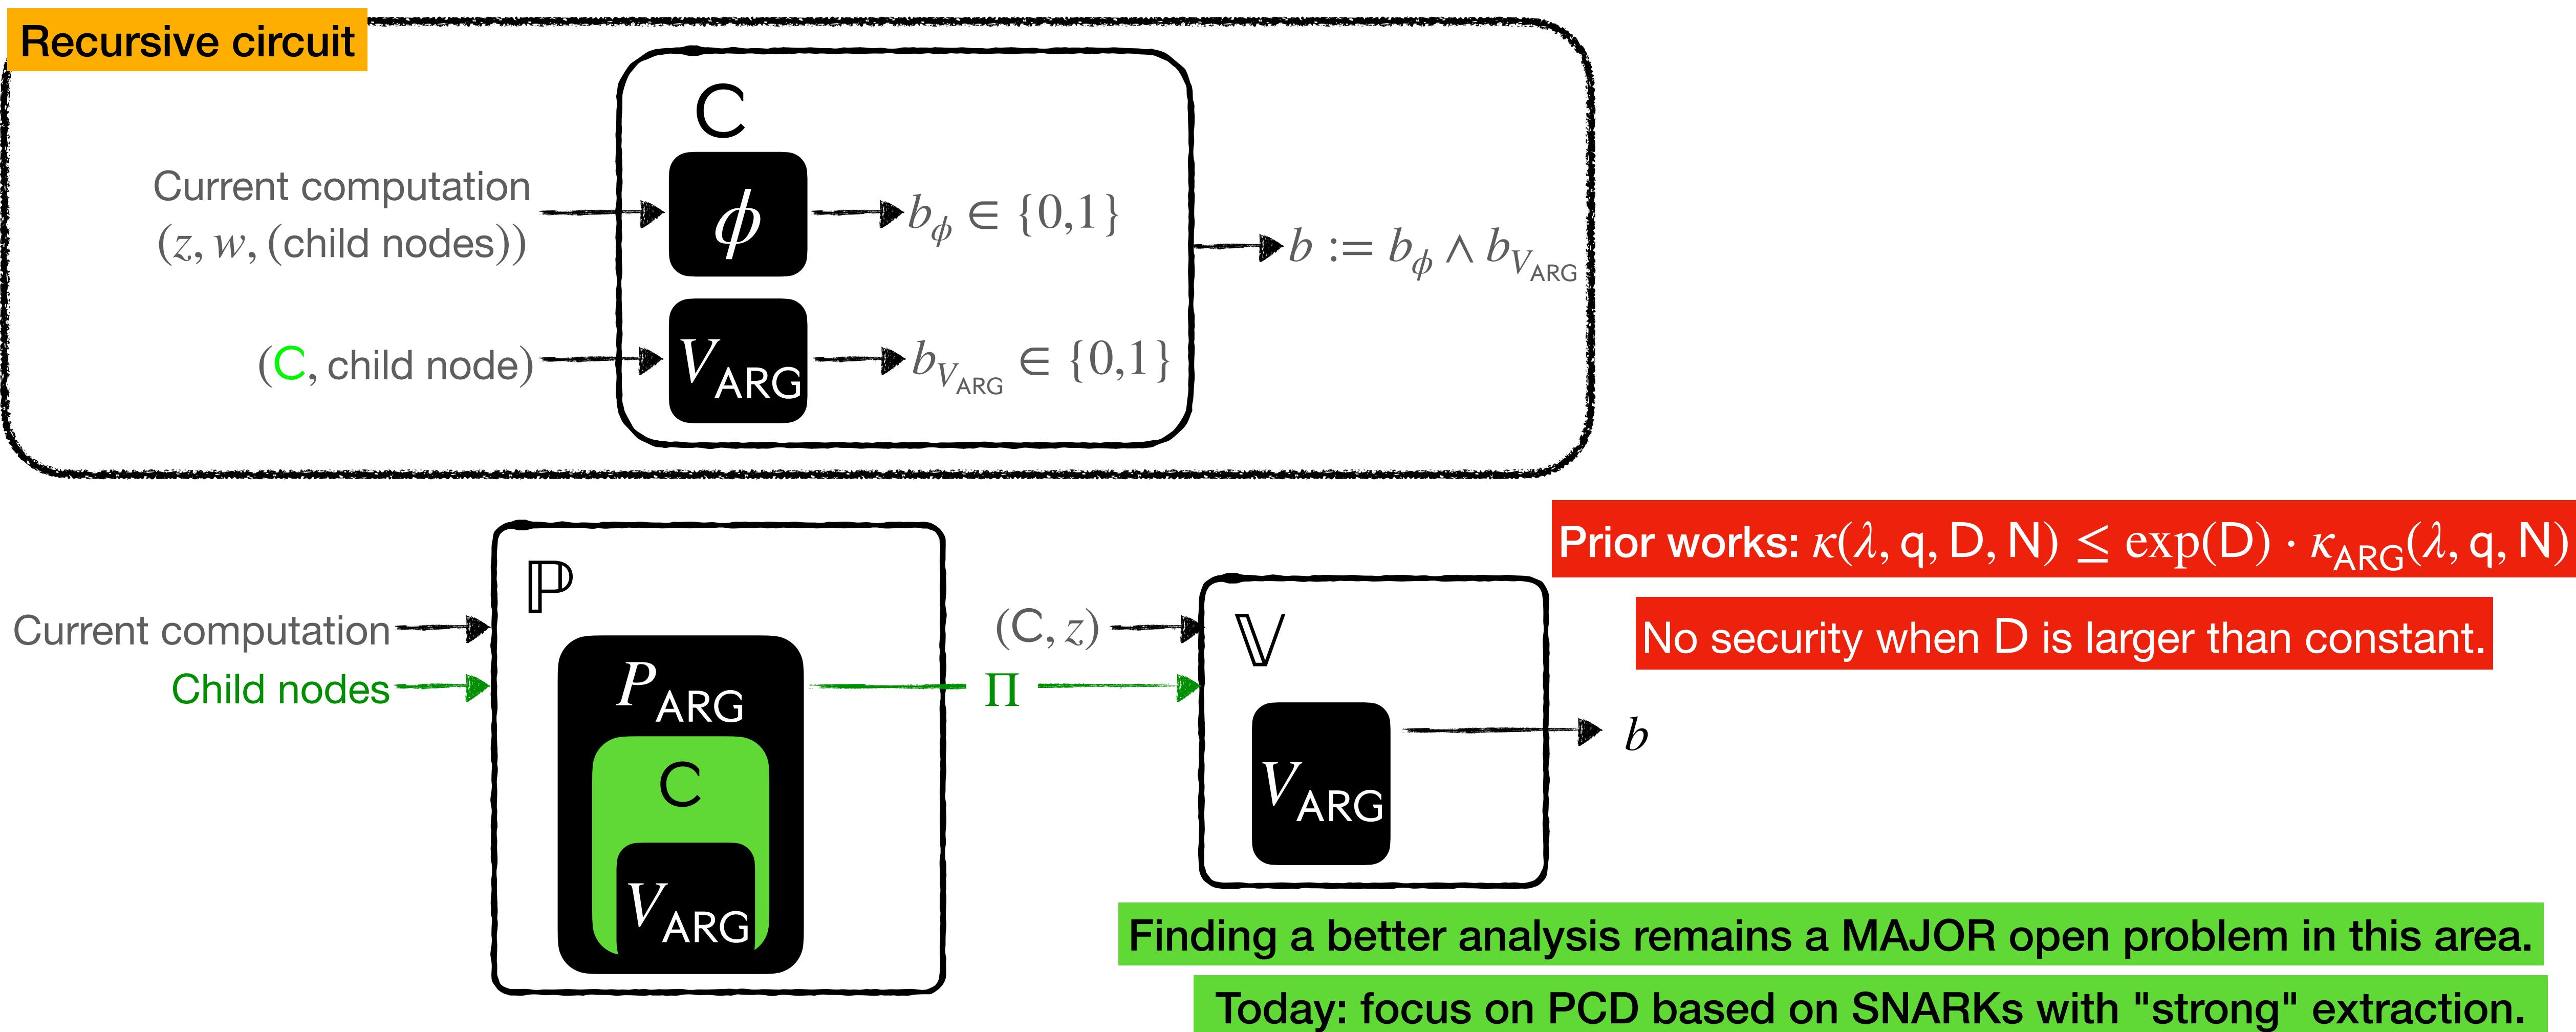
Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



Recursively compose the SNARK proofs

PCD = recursive proof composition of a SNARK: PCD invokes SNARK (for CSAT) for the recursive circuit C.



Our result

Our result

Theorem.

Our result

Theorem. We prove a significantly improved security bound for PCD based on SNARKs with **straightline extraction**:

Our result

Theorem. We prove a significantly improved security bound for PCD based on SNARKs with **straightline extraction**:

SNARK for CSAT
with straightline extraction

Recursive proof composition

PCD with straightline extraction

Our result

Theorem. We prove a significantly improved security bound for PCD based on SNARKs with **straightline extraction**:

SNARK for CSAT
with straightline extraction

Recursive proof composition

PCD with straightline extraction

$$\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N)$$

Our result

Theorem. We prove a significantly improved security bound for PCD based on SNARKs with **straightline extraction**:

SNARK for CSAT
with straightline extraction

Recursive proof composition

PCD with straightline extraction

$$\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N)$$

Significant improvement:
 κ doesn't grow exponentially with D

Our result

Theorem. We prove a significantly improved security bound for PCD based on SNARKs with **straightline extraction**:

SNARK for CSAT
with straightline extraction

Recursive proof composition

PCD with straightline extraction

$$\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N)$$

Significant improvement:
 κ doesn't grow exponentially with D

In practice, SNARKs have non-black-box knowledge soundness.

Our result

Theorem. We prove a significantly improved security bound for PCD based on SNARKs with **straightline extraction**:

SNARK for CSAT
with straightline extraction

Recursive proof composition

PCD with straightline extraction

$$\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N)$$

Significant improvement:

κ doesn't grow exponentially with D

In practice, SNARKs have non-black-box knowledge soundness.
Straightline extraction only exists in idealized models.

Our result

Theorem. We prove a significantly improved security bound for PCD based on SNARKs with **straightline extraction**:

SNARK for CSAT
with straightline extraction

Recursive proof composition

PCD with straightline extraction

$$\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N)$$

Significant improvement:
 κ doesn't grow exponentially with D

In practice, SNARKs have non-black-box knowledge soundness.
Straightline extraction only exists in idealized models.
How can we apply our theorem in practice then?

Applications

Applications

Application 1 [main].

Applications

Application 1 [main].

- We propose a new idealization of hash-based PCD used in practice as a “PCD” in the ROM.

Applications

Application 1 [main].

- We propose a new idealization of hash-based PCD used in practice as a “PCD” in the ROM.
- We apply our theorem: $\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N) = \kappa_{\text{ARG}}(\lambda)$.

Applications

Application 1 [main].

- We propose a new idealization of hash-based PCD used in practice as a “PCD” in the ROM.
- We apply our theorem: $\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N) = \kappa_{\text{ARG}}(\lambda)$.
- First justification for current choice of parameters of hash-based PCD in practice! [Polygon, Sharp]

Applications

Application 1 [main].

- We propose a new idealization of hash-based PCD used in practice as a “PCD” in the ROM.
- We apply our theorem: $\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N) = \kappa_{\text{ARG}}(\lambda)$.
- First justification for current choice of parameters of hash-based PCD in practice! [Polygon, Sharp]

Application 2.

Applications

Application 1 [main].

- We propose a new idealization of hash-based PCD used in practice as a “PCD” in the ROM.
- We apply our theorem: $\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N) = \kappa_{\text{ARG}}(\lambda)$.
- First justification for current choice of parameters of hash-based PCD in practice! [Polygon, Sharp]

Application 2.

- [CT10]: SNARK with straightline extraction in the SROM (*signed random oracle model*).

Applications

Application 1 [main].

- We propose a new idealization of hash-based PCD used in practice as a “PCD” in the ROM.
- We apply our theorem: $\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N) = \kappa_{\text{ARG}}(\lambda)$.
- First justification for current choice of parameters of hash-based PCD in practice! [Polygon, Sharp]

Application 2.

- [CT10]: SNARK with straightline extraction in the SROM (*signed random oracle model*).
- Their bound: $\kappa(\lambda, D, N, t) \leq \textcolor{red}{N} \cdot \kappa_{\text{ARG}}(\lambda, N, N \cdot t)$.

Applications

Application 1 [main].

- We propose a new idealization of hash-based PCD used in practice as a “PCD” in the ROM.
- We apply our theorem: $\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N) = \kappa_{\text{ARG}}(\lambda)$.
- First justification for current choice of parameters of hash-based PCD in practice! [Polygon, Sharp]

Application 2.

- [CT10]: SNARK with straightline extraction in the SROM (*signed random oracle model*).
- Their bound: $\kappa(\lambda, D, N, t) \leq \mathbf{N} \cdot \kappa_{\text{ARG}}(\lambda, N, N \cdot t)$.
- Our bound: $\kappa(\lambda, D, N, t) \leq \kappa_{\text{ARG}}(\lambda, N, t + N)$.

Applications

Application 1 [main].

- We propose a new idealization of hash-based PCD used in practice as a “PCD” in the ROM.
- We apply our theorem: $\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N) = \kappa_{\text{ARG}}(\lambda)$.
- First justification for current choice of parameters of hash-based PCD in practice! [Polygon, Sharp]

Application 2.

- [CT10]: SNARK with straightline extraction in the SROM (*signed random oracle model*).
- Their bound: $\kappa(\lambda, D, N, t) \leq \mathbf{N} \cdot \kappa_{\text{ARG}}(\lambda, N, N \cdot t)$.
- Our bound: $\kappa(\lambda, D, N, t) \leq \kappa_{\text{ARG}}(\lambda, N, t + N)$.

Application 3.

Applications

Application 1 [main].

- We propose a new idealization of hash-based PCD used in practice as a “PCD” in the ROM.
- We apply our theorem: $\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N) = \kappa_{\text{ARG}}(\lambda)$.
- First justification for current choice of parameters of hash-based PCD in practice! [Polygon, Sharp]

Application 2.

- [CT10]: SNARK with straightline extraction in the SROM (*signed random oracle model*).
- Their bound: $\kappa(\lambda, D, N, t) \leq \mathbf{N} \cdot \kappa_{\text{ARG}}(\lambda, N, N \cdot t)$.
- Our bound: $\kappa(\lambda, D, N, t) \leq \kappa_{\text{ARG}}(\lambda, N, t + N)$.

Application 3.

- [CCGOS23]: SNARK with straightline extraction in the AROM (*arithmetized random oracle model*).

Applications

Application 1 [main].

- We propose a new idealization of hash-based PCD used in practice as a “PCD” in the ROM.
- We apply our theorem: $\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N) = \kappa_{\text{ARG}}(\lambda)$.
- First justification for current choice of parameters of hash-based PCD in practice! [Polygon, Sharp]

Application 2.

- [CT10]: SNARK with straightline extraction in the SROM (*signed random oracle model*).
- Their bound: $\kappa(\lambda, D, N, t) \leq N \cdot \kappa_{\text{ARG}}(\lambda, N, N \cdot t)$.
- Our bound: $\kappa(\lambda, D, N, t) \leq \kappa_{\text{ARG}}(\lambda, N, t + N)$.

Application 3.

- [CCGOS23]: SNARK with straightline extraction in the AROM (*arithmetized random oracle model*).
- Their bound: $\kappa(\lambda, D, N, t) \leq N \cdot \kappa_{\text{ARG}}(\lambda, N, N \cdot t)$.

Applications

Application 1 [main].

- We propose a new idealization of hash-based PCD used in practice as a “PCD” in the ROM.
- We apply our theorem: $\kappa(\lambda, D, N) \leq \kappa_{\text{ARG}}(\lambda, N) = \kappa_{\text{ARG}}(\lambda)$.
- First justification for current choice of parameters of hash-based PCD in practice! [Polygon, Sharp]

Application 2.

- [CT10]: SNARK with straightline extraction in the SROM (*signed random oracle model*).
- Their bound: $\kappa(\lambda, D, N, t) \leq N \cdot \kappa_{\text{ARG}}(\lambda, N, N \cdot t)$.
- Our bound: $\kappa(\lambda, D, N, t) \leq \kappa_{\text{ARG}}(\lambda, N, t + N)$.

Application 3.

- [CCGOS23]: SNARK with straightline extraction in the AROM (*arithmetized random oracle model*).
- Their bound: $\kappa(\lambda, D, N, t) \leq N \cdot \kappa_{\text{ARG}}(\lambda, N, N \cdot t)$.
- Our bound: $\kappa(\lambda, D, N, t) \leq \kappa_{\text{ARG}}(\lambda, N, t + N)$.

What is straightline extraction?

SNARKs with straightline extraction

SNARKs with straightline extraction

SNARKs in an oracle model (e.g. ROM):

SNARKs with straightline extraction

SNARKs in an oracle model (e.g. ROM):



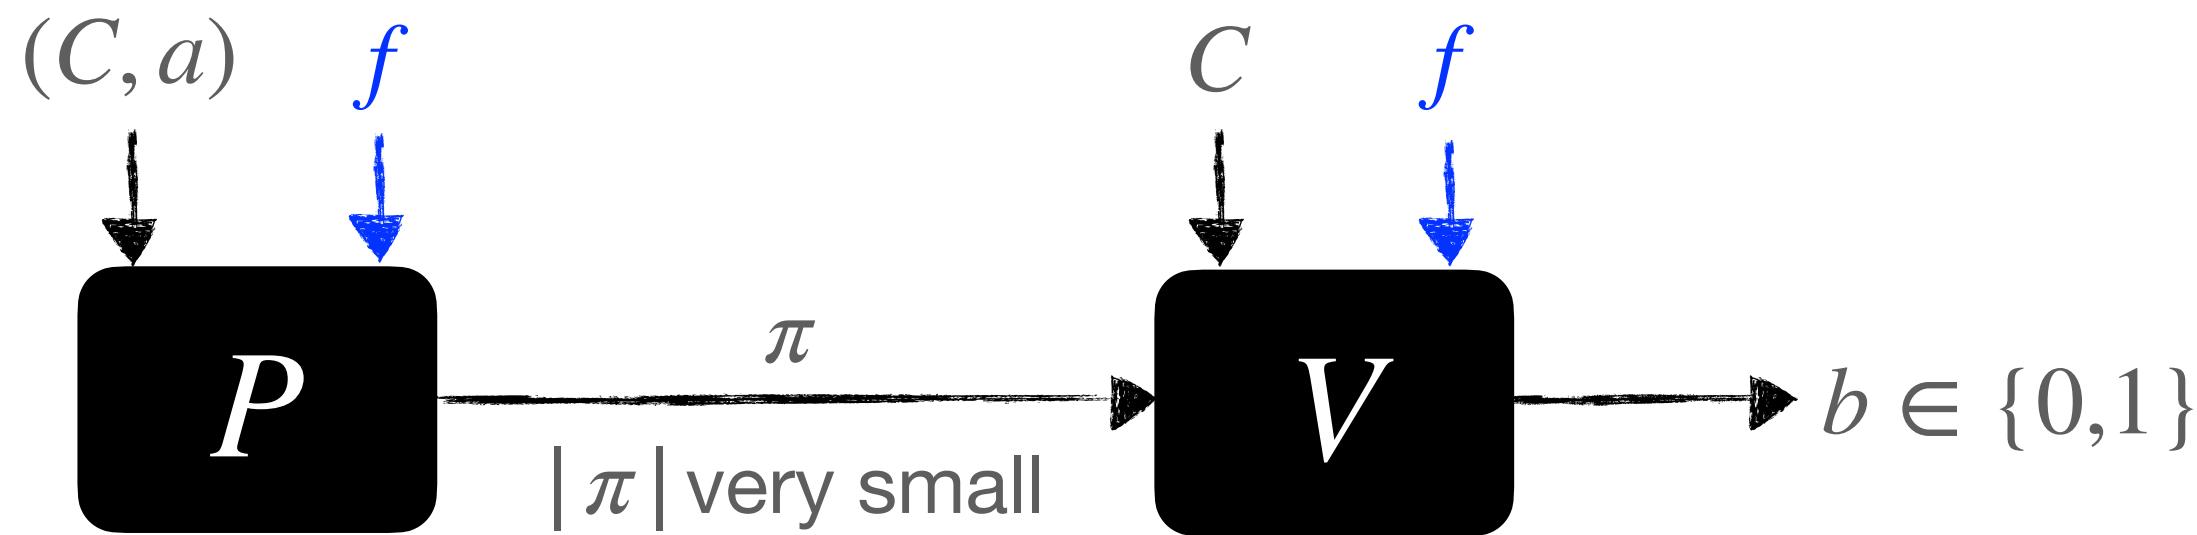
SNARKs with straightline extraction

SNARKs in an oracle model (e.g. ROM):



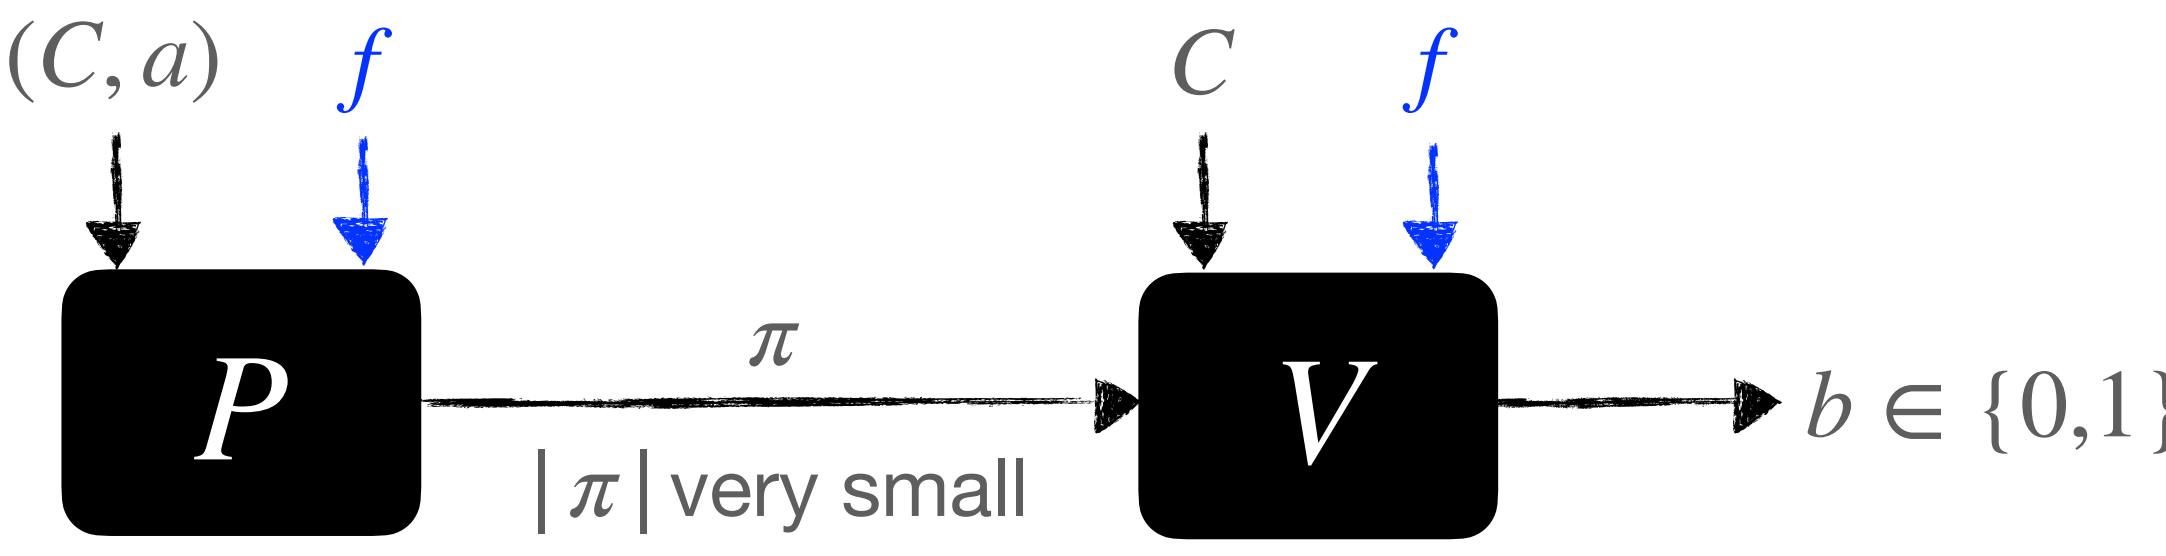
SNARKs with straightline extraction

SNARKs in an oracle model (e.g. ROM):



SNARKs with straightline extraction

SNARKs in an oracle model (e.g. ROM):

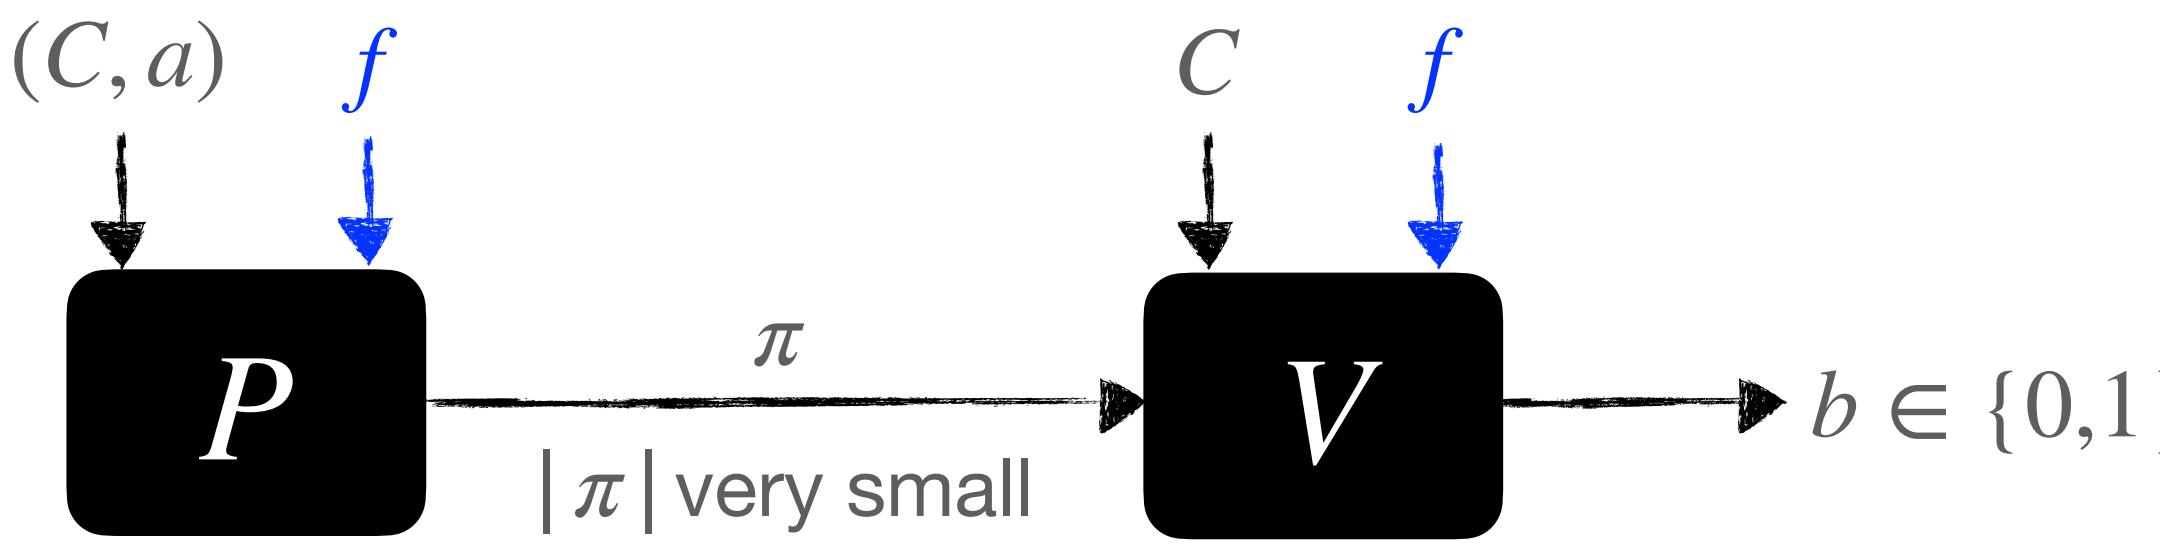


λ : security parameter
q: adversary query bound

Straightline knowledge soundness: \exists deterministic extractor E , \forall bounded adversary \tilde{P} , the following happens with probability at most $\kappa_{\text{ARG}}(\lambda, q)$:

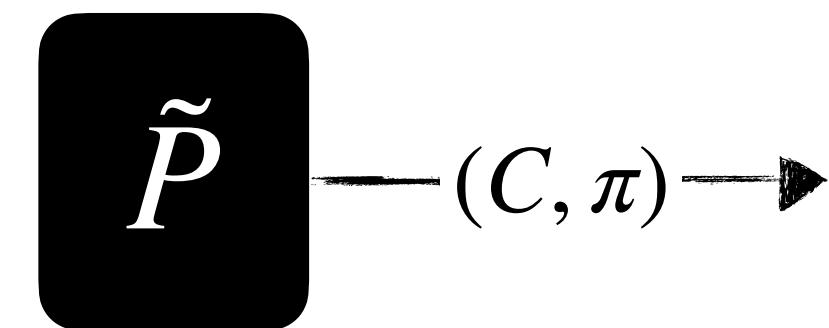
SNARKs with straightline extraction

SNARKs in an oracle model (e.g. ROM):



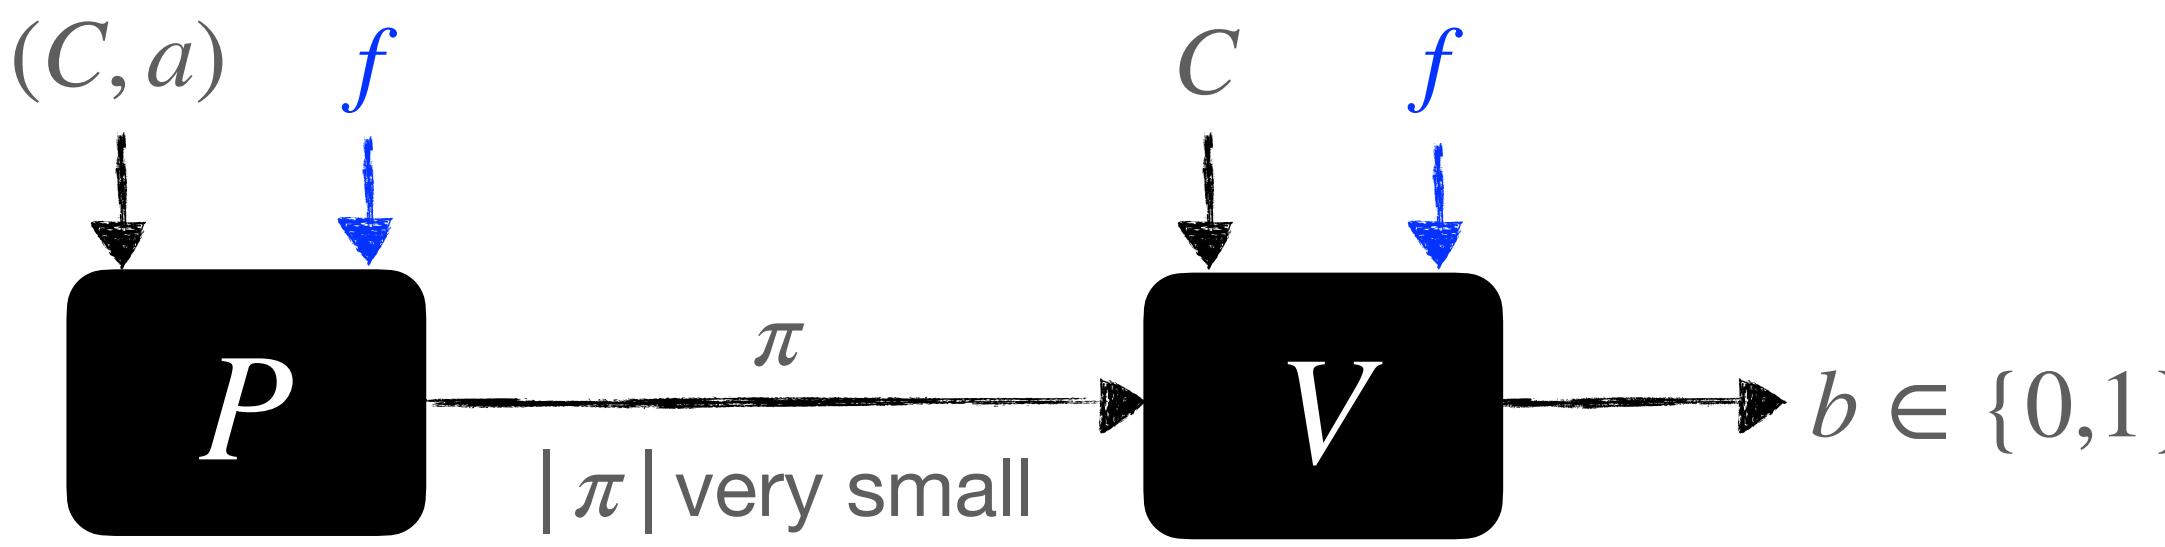
λ : security parameter
q: adversary query bound

Straightline knowledge soundness: \exists deterministic extractor E , \forall bounded adversary \tilde{P} , the following happens with probability at most $\kappa_{\text{ARG}}(\lambda, q)$:



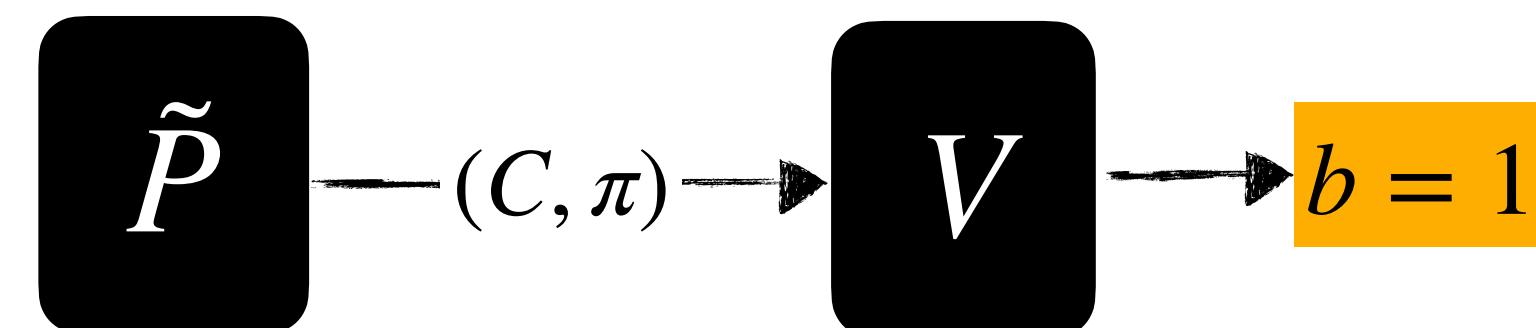
SNARKs with straightline extraction

SNARKs in an oracle model (e.g. ROM):



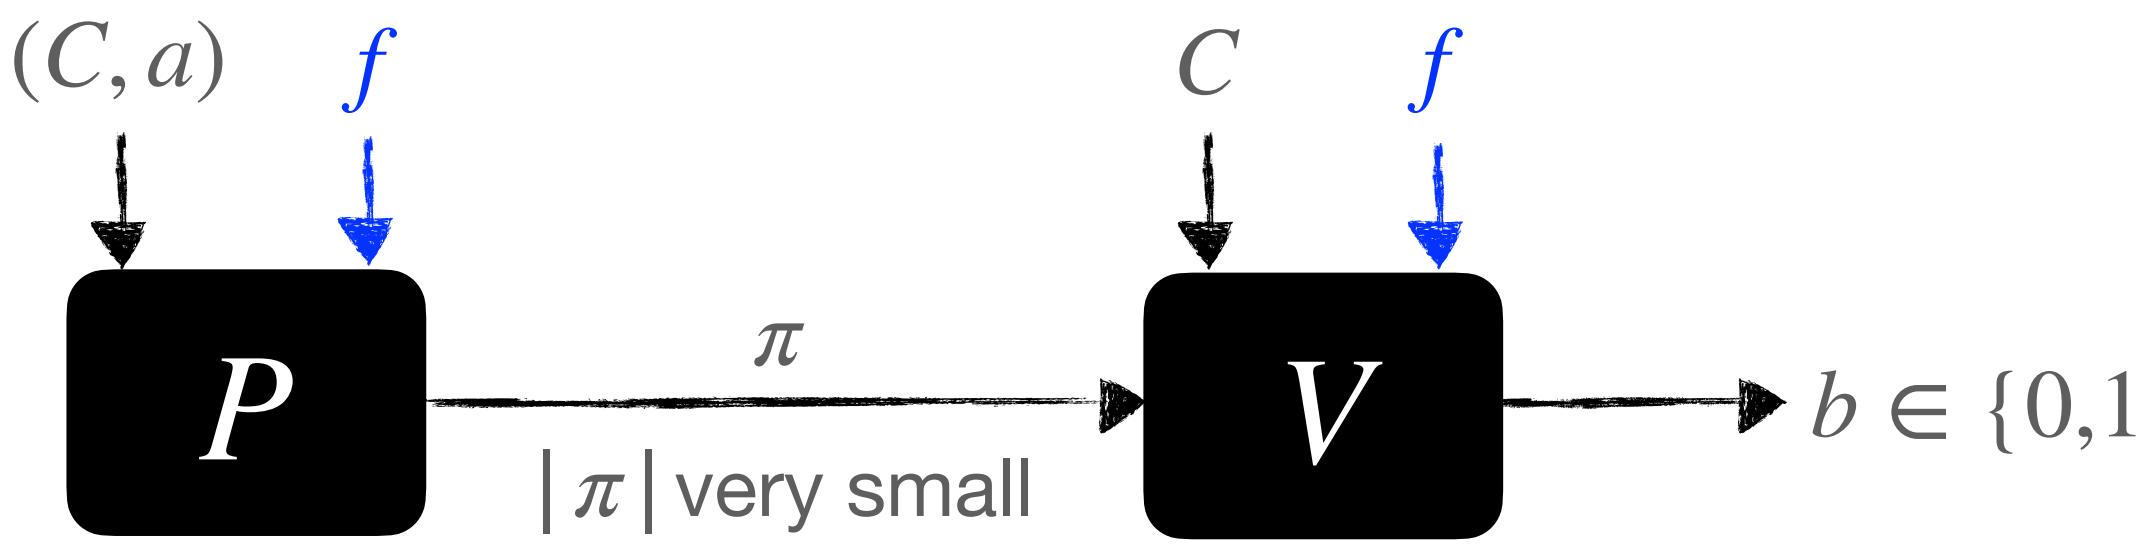
λ : security parameter
q: adversary query bound

Straightline knowledge soundness: \exists deterministic extractor E , \forall bounded adversary \tilde{P} , the following happens with probability at most $\kappa_{\text{ARG}}(\lambda, q)$:



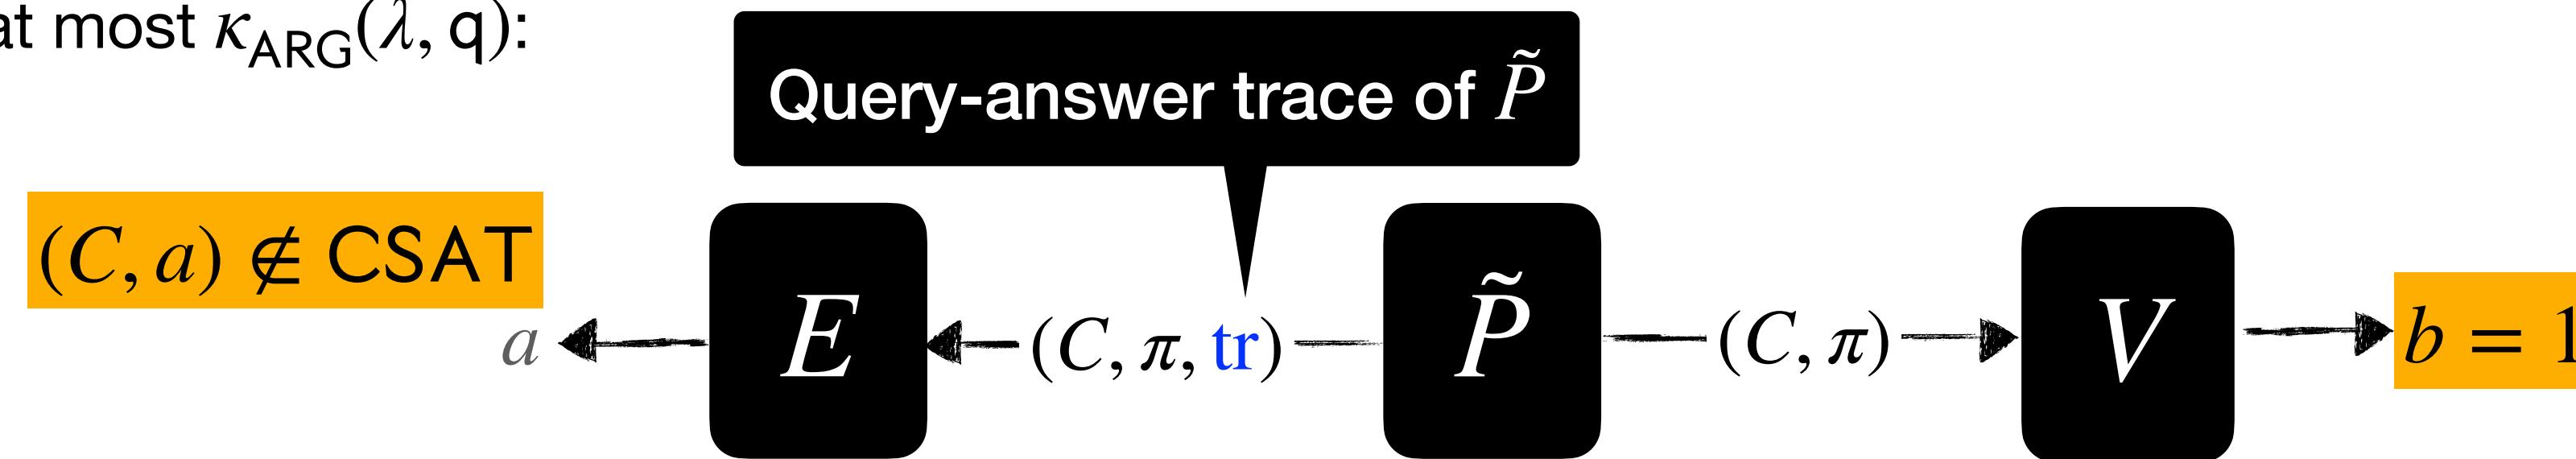
SNARKs with straightline extraction

SNARKs in an oracle model (e.g. ROM):



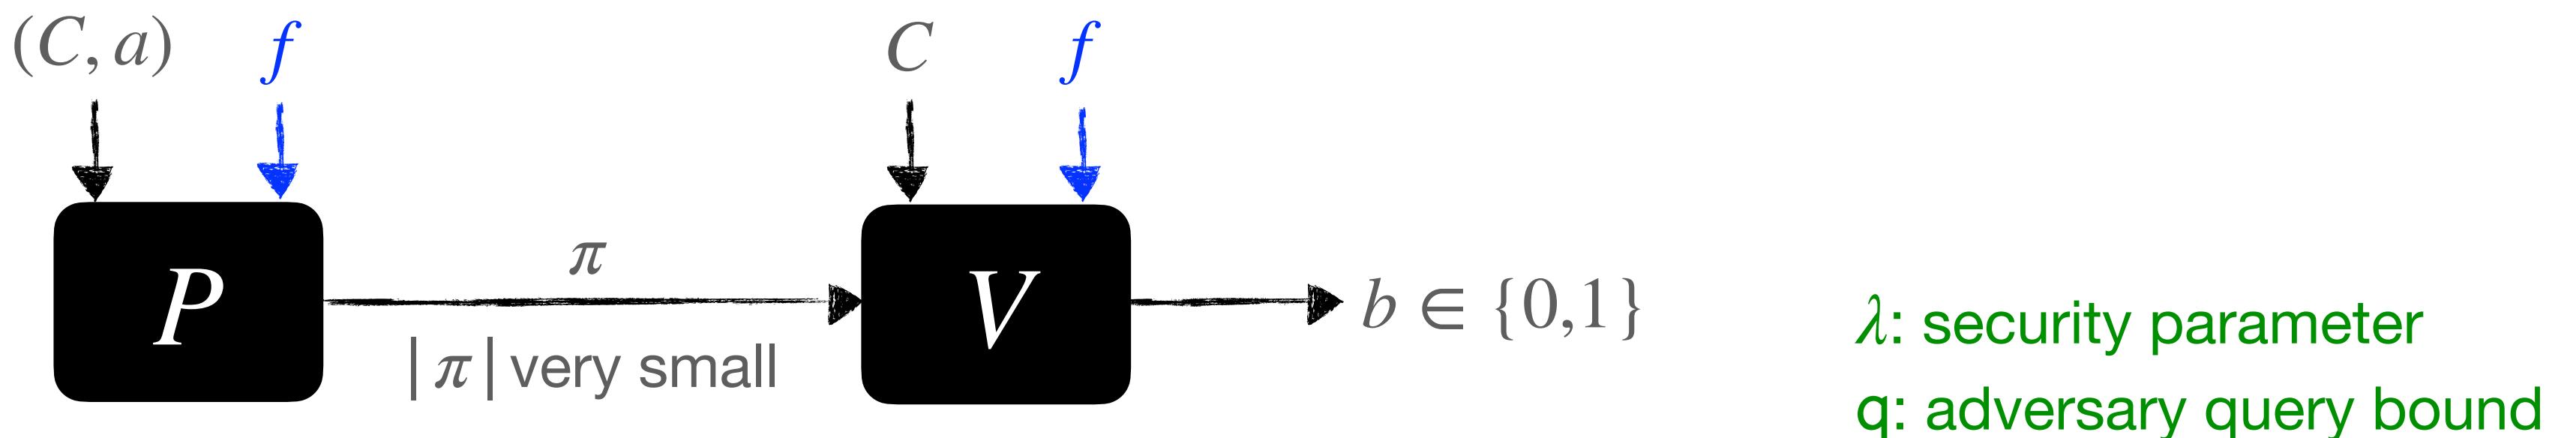
λ : security parameter
q: adversary query bound

Straightline knowledge soundness: \exists deterministic extractor E , \forall bounded adversary \tilde{P} , the following happens with probability at most $\kappa_{\text{ARG}}(\lambda, q)$:

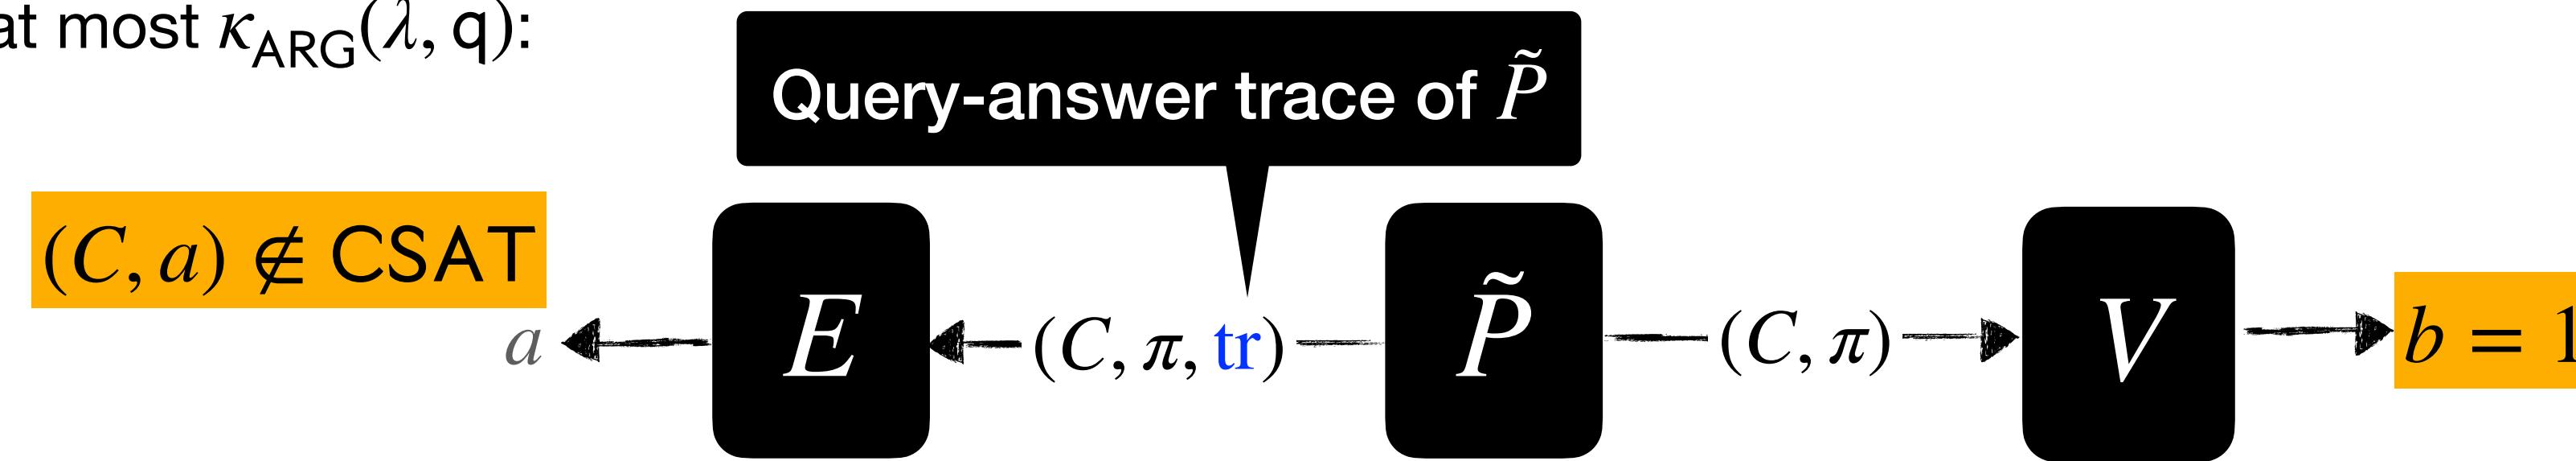


SNARKs with straightline extraction

SNARKs in an oracle model (e.g. ROM):



Straightline knowledge soundness: \exists deterministic extractor E , \forall bounded adversary \tilde{P} , the following happens with probability at most $\kappa_{\text{ARG}}(\lambda, q)$:



Wonderful Fact: in the ROM (and other interesting oracle models) there are SNARKs of interest with straightline extraction!
(E.g., the Micali SNARK and BCS SNARK and related constructions.)

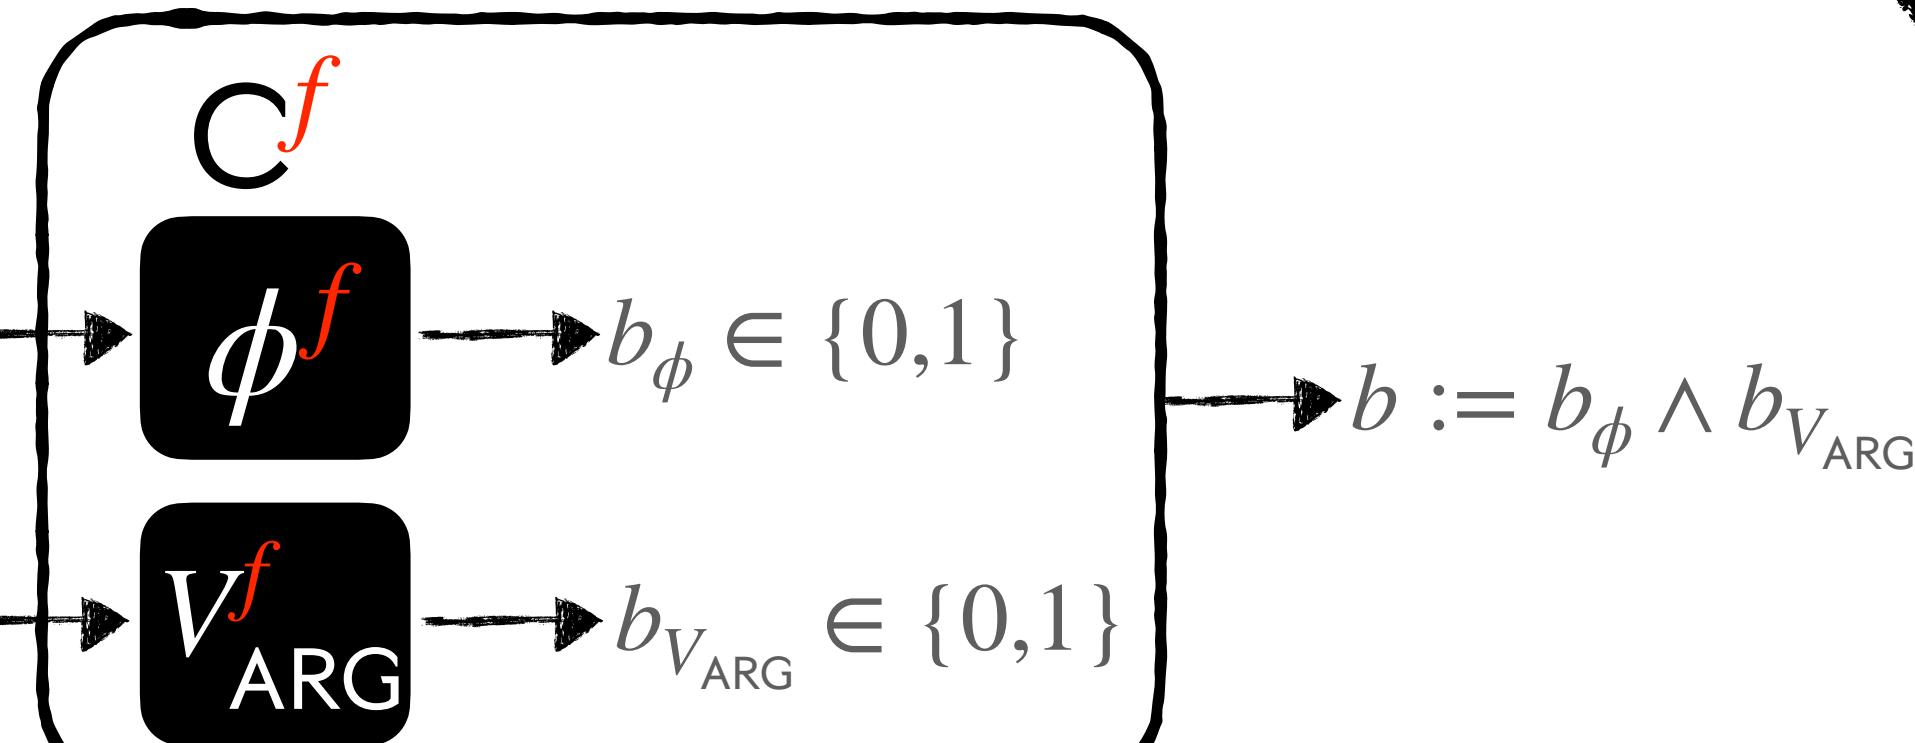
Can't we use the previous recursive composition?

Can't we use the previous recursive composition?

Recursive circuit

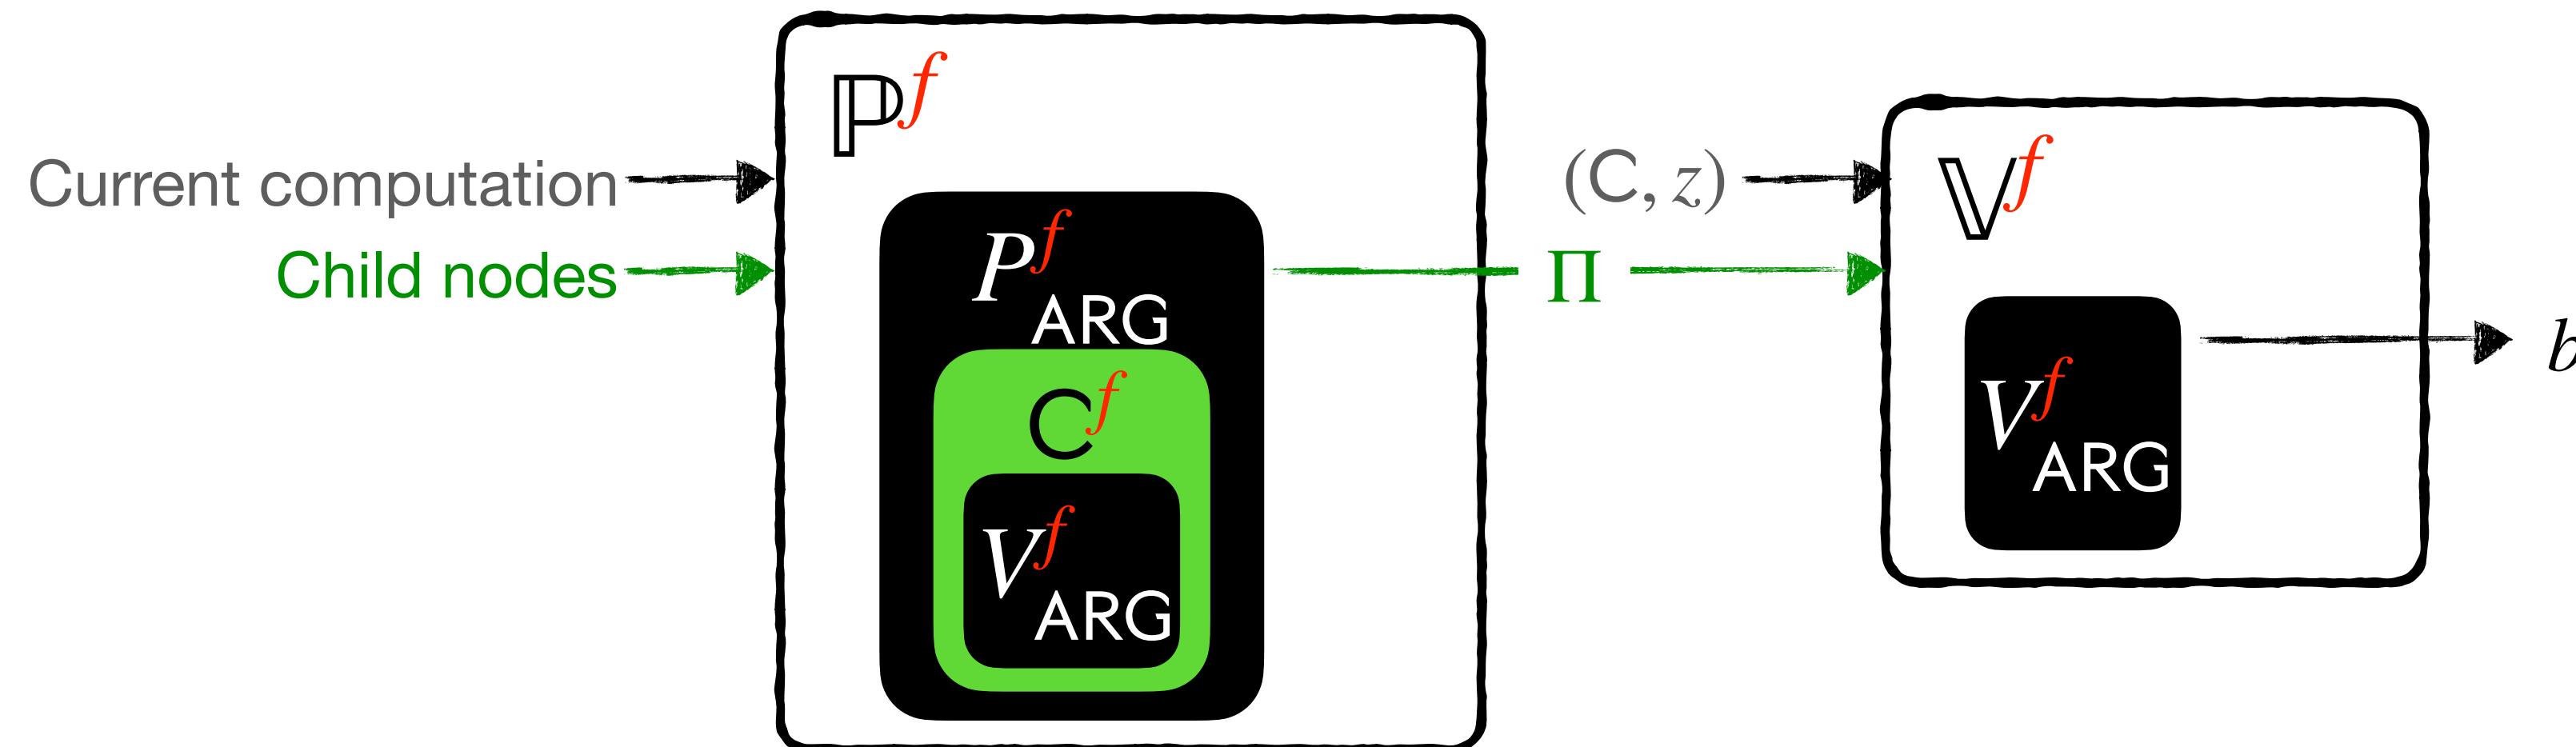
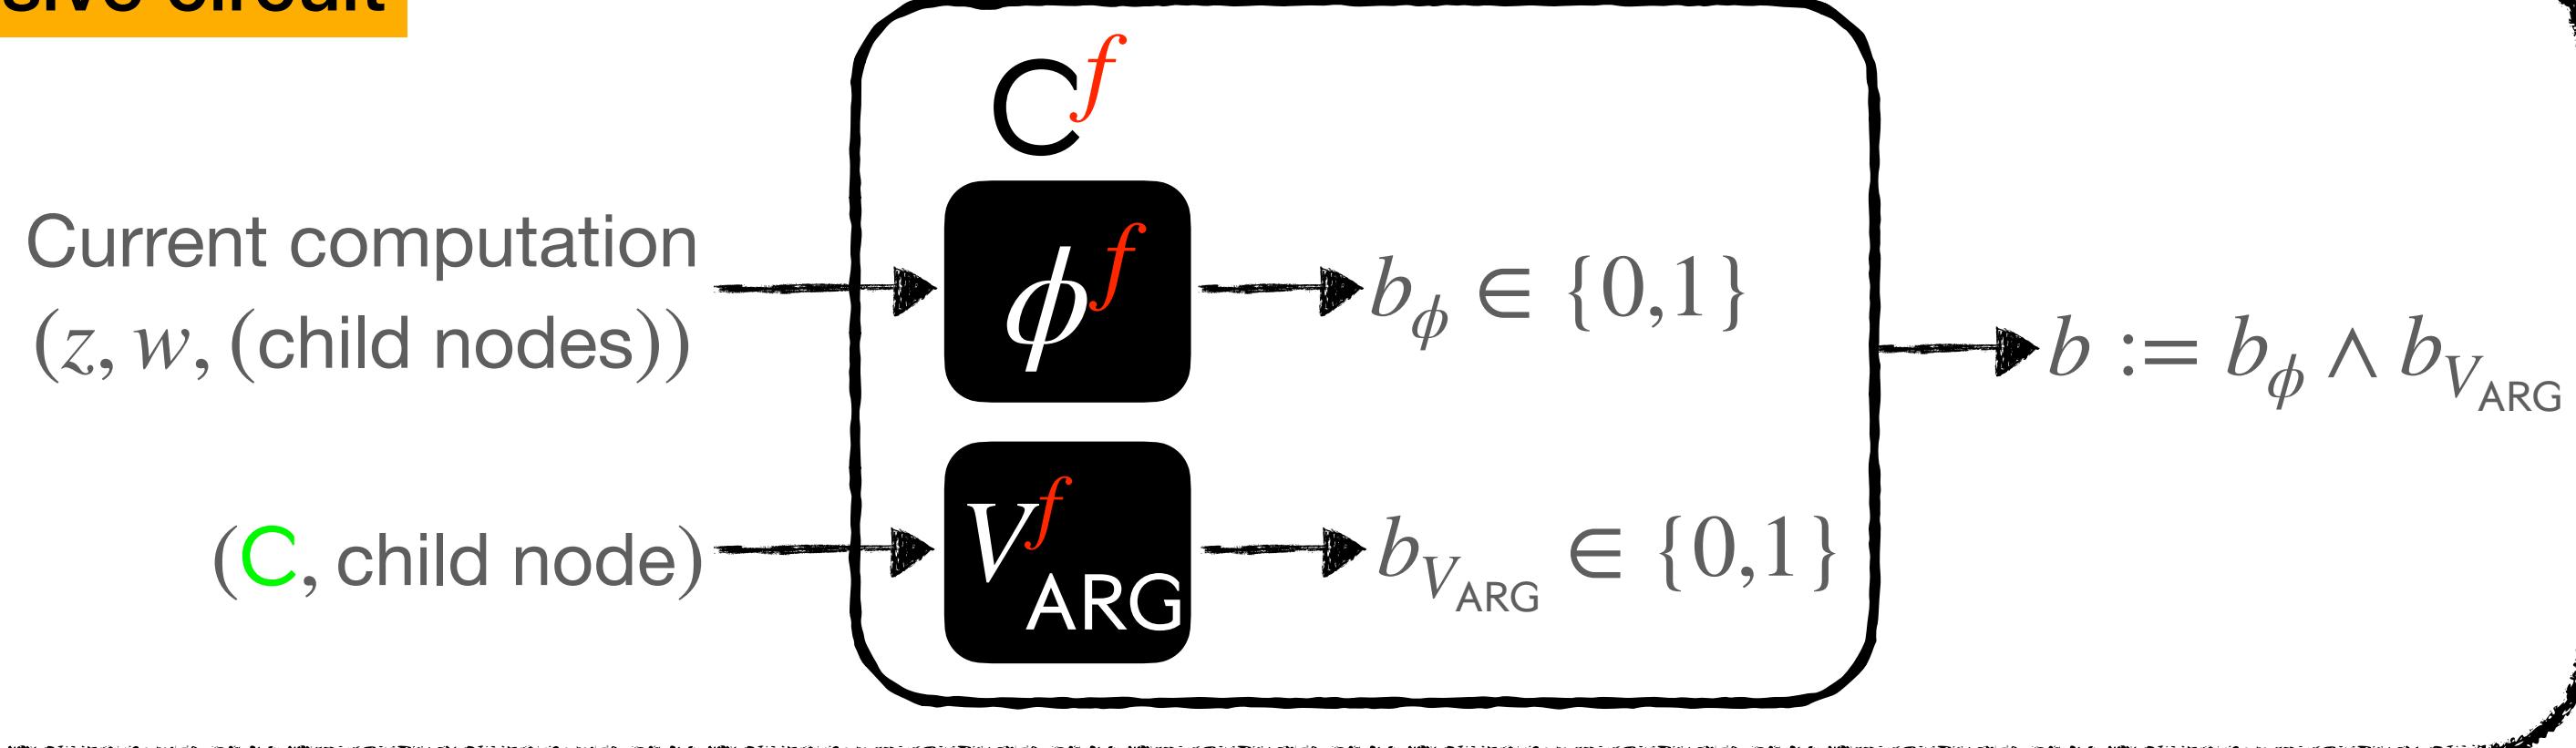
Current computation
($z, w, (\text{child nodes})$)

(C , child node)



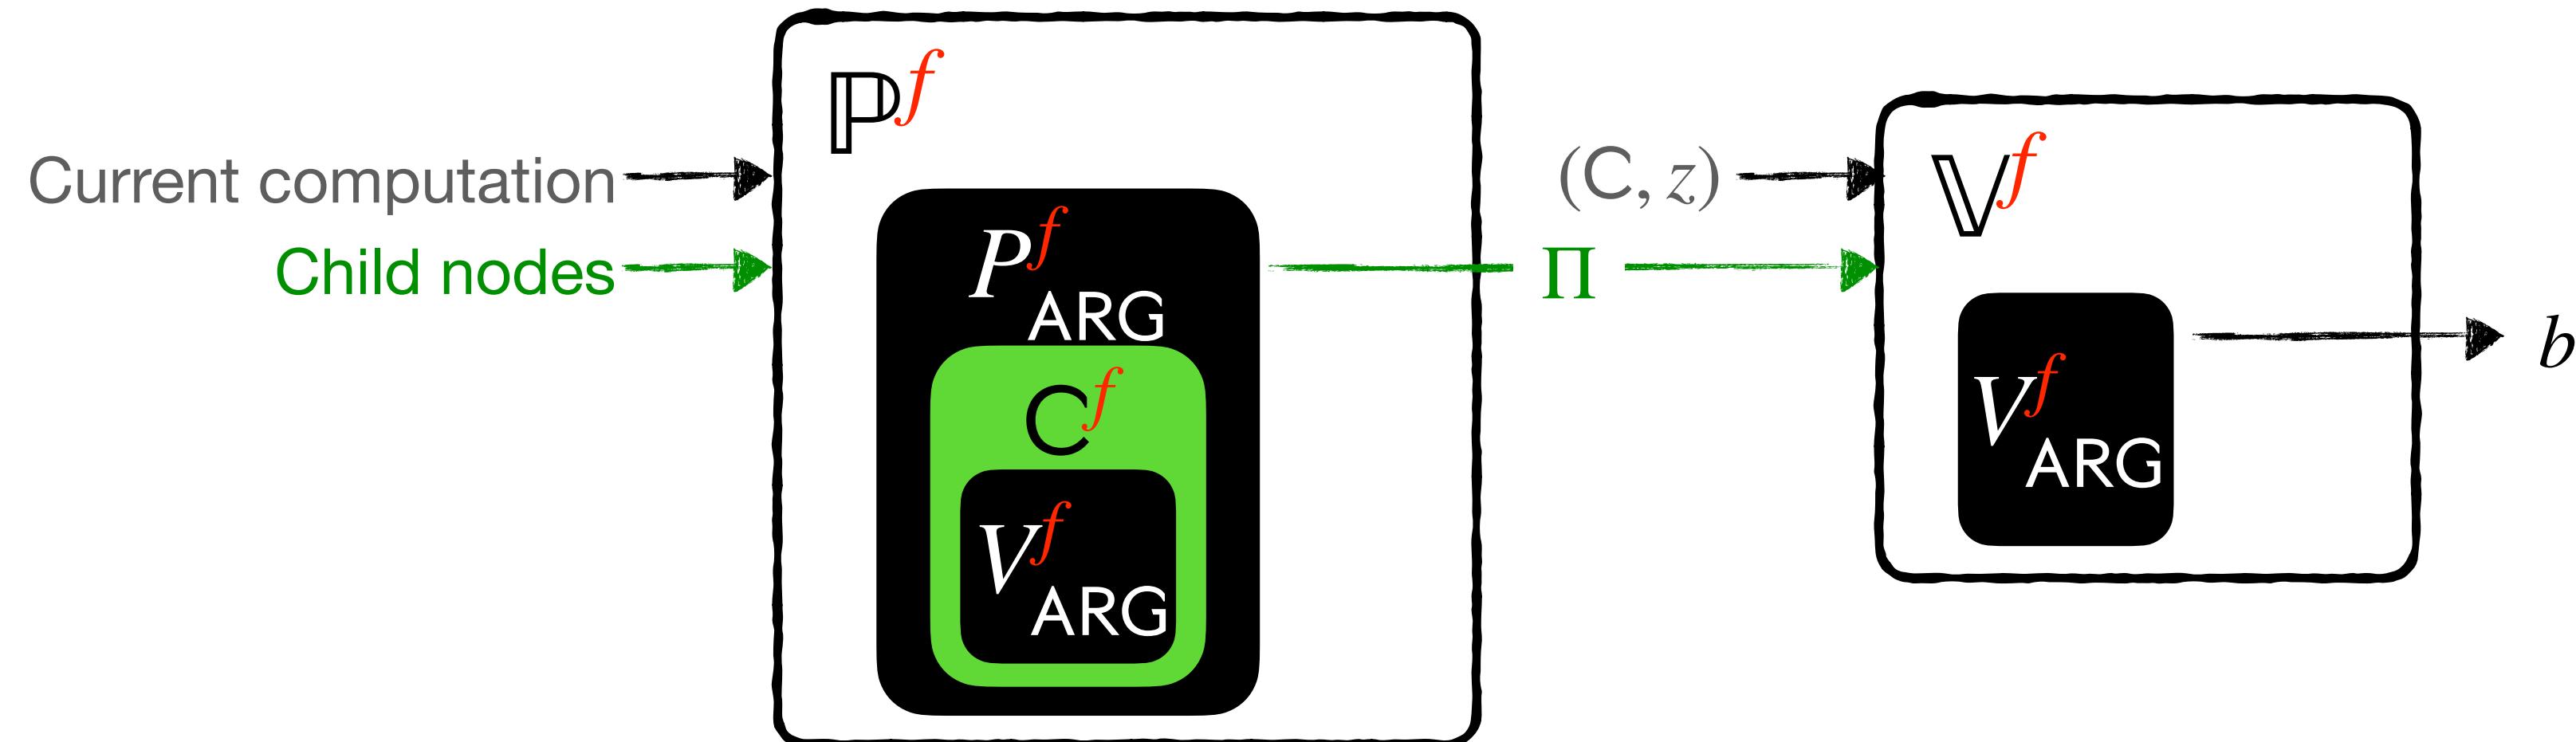
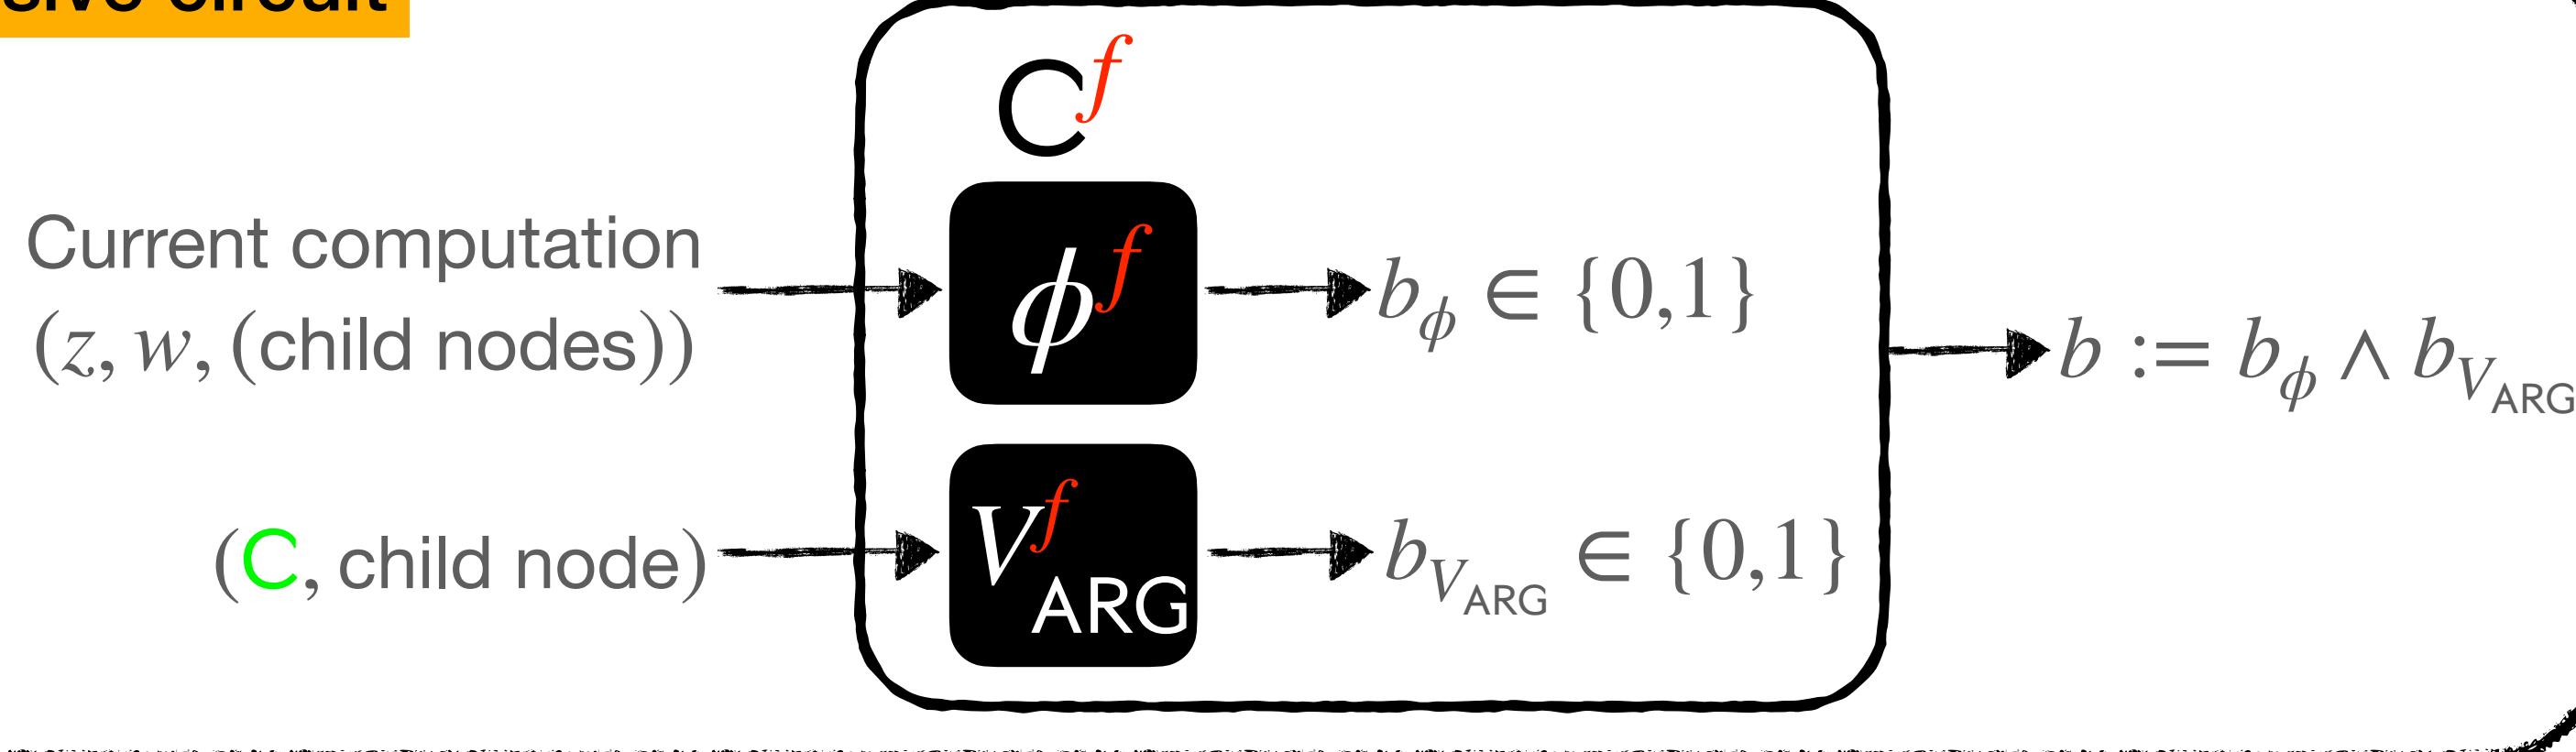
Can't we use the previous recursive composition?

Recursive circuit



Can't we use the previous recursive composition?

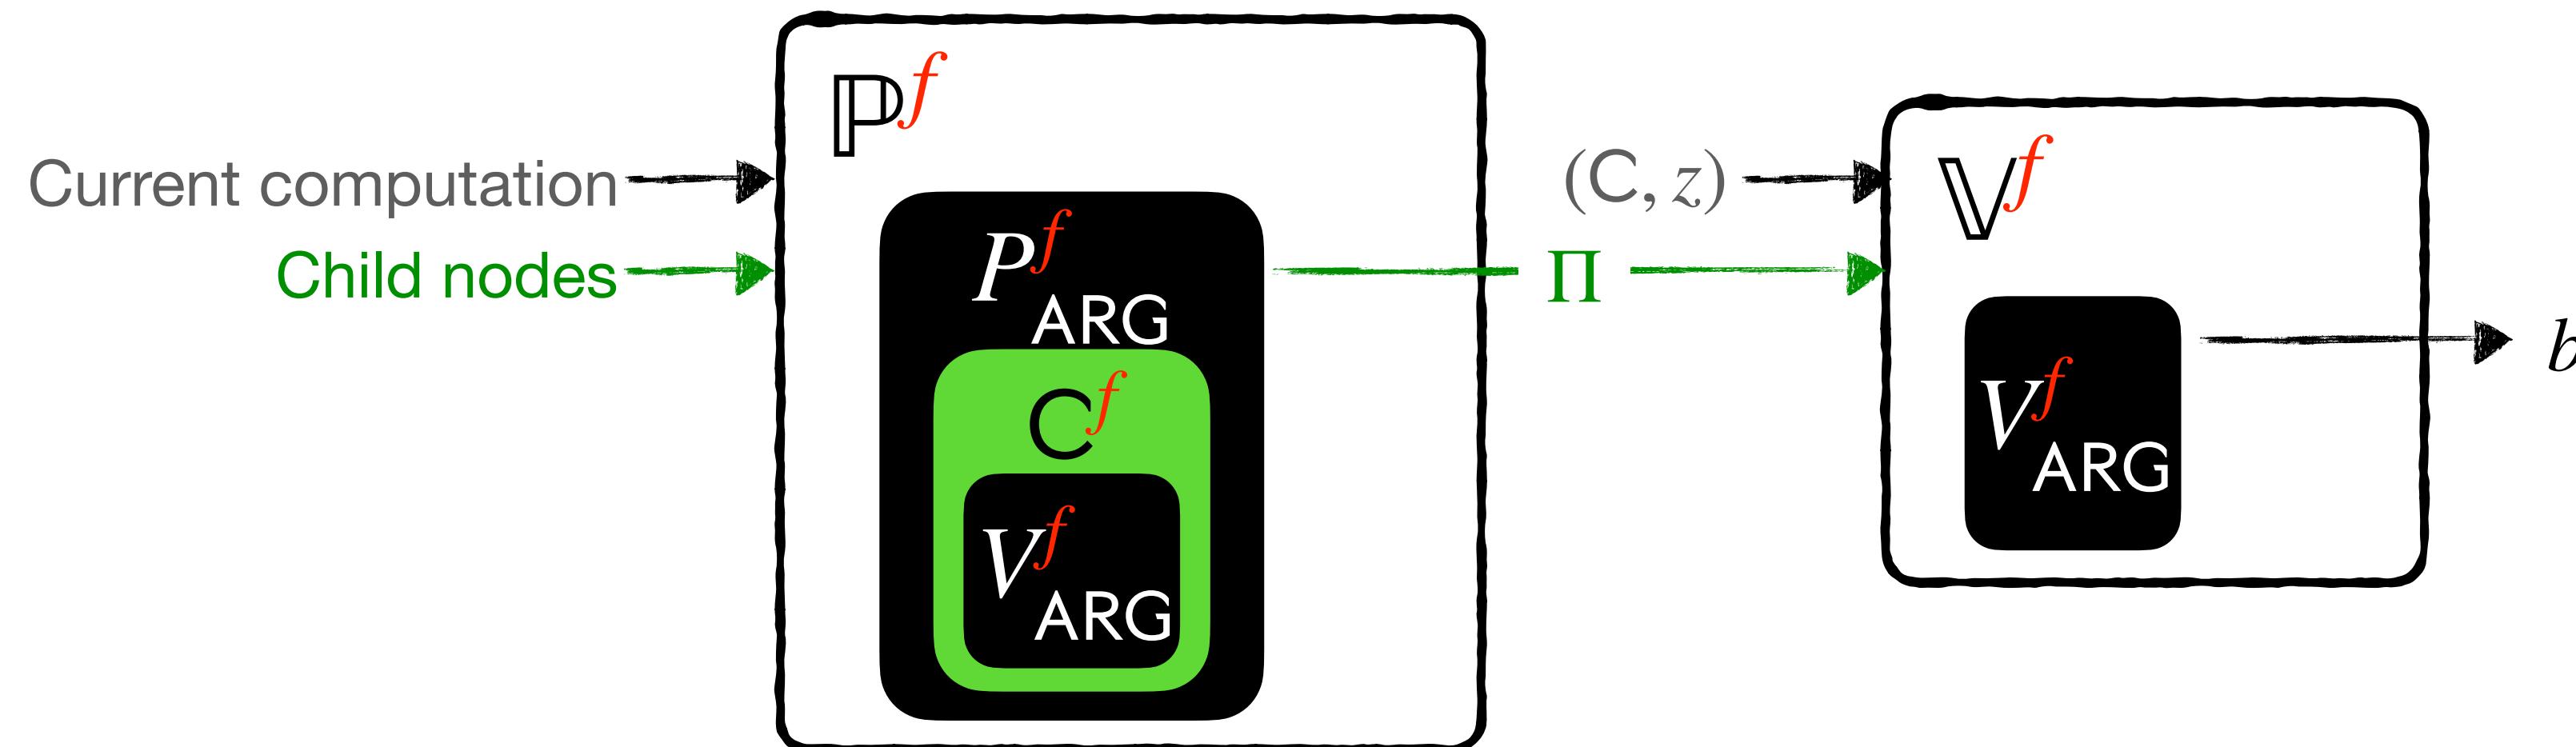
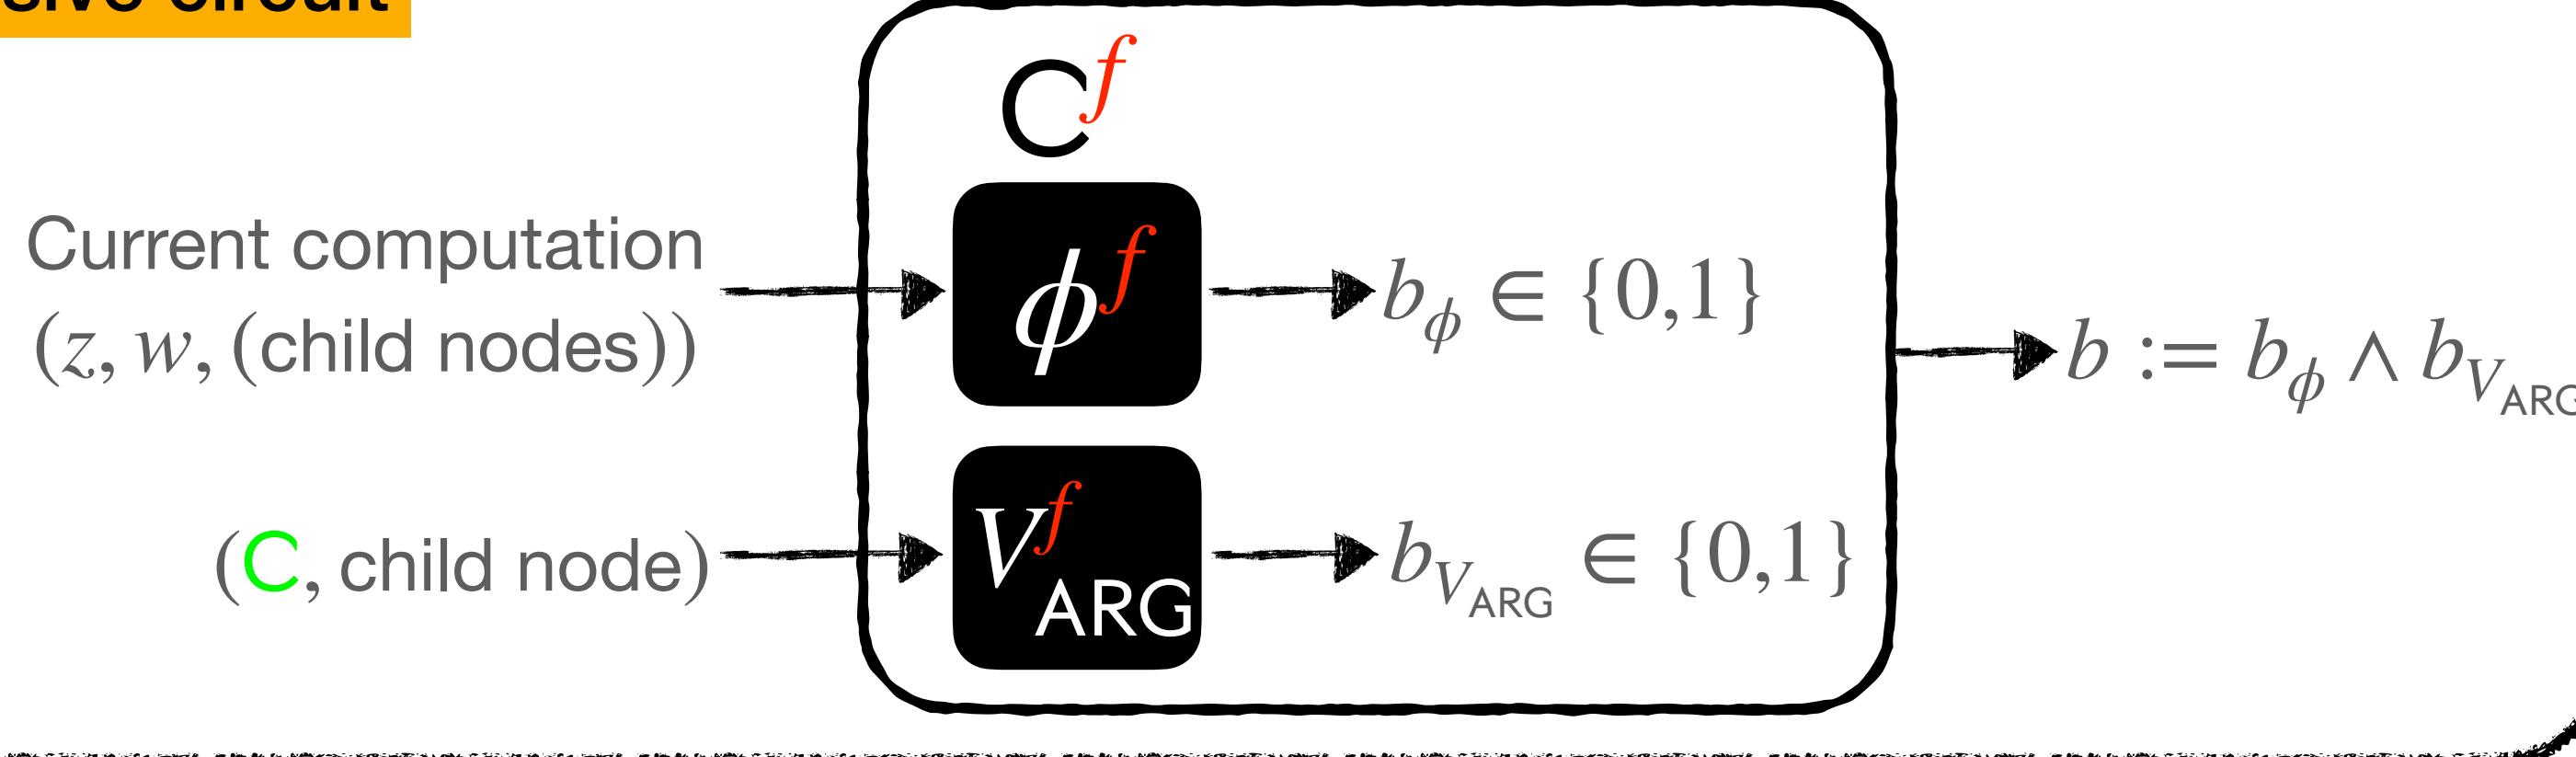
Recursive circuit



ISSUE! C has oracle access to f .

Can't we use the previous recursive composition?

Recursive circuit

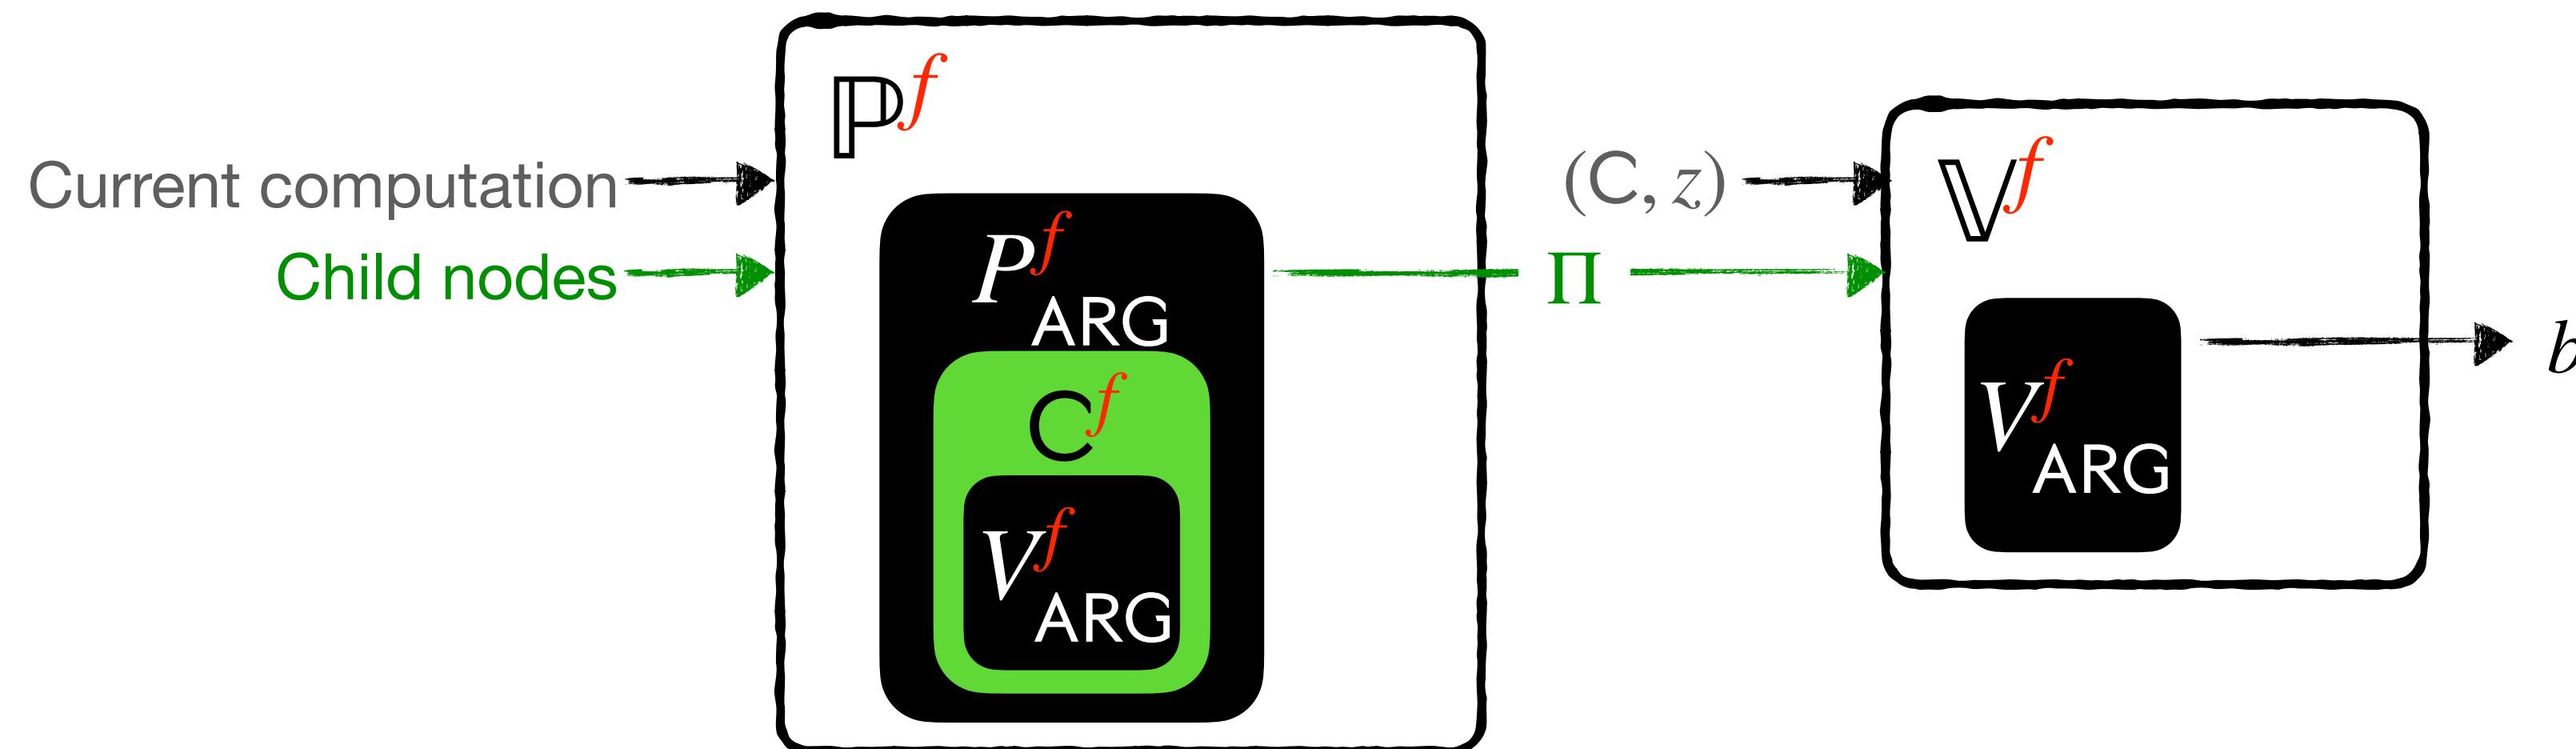
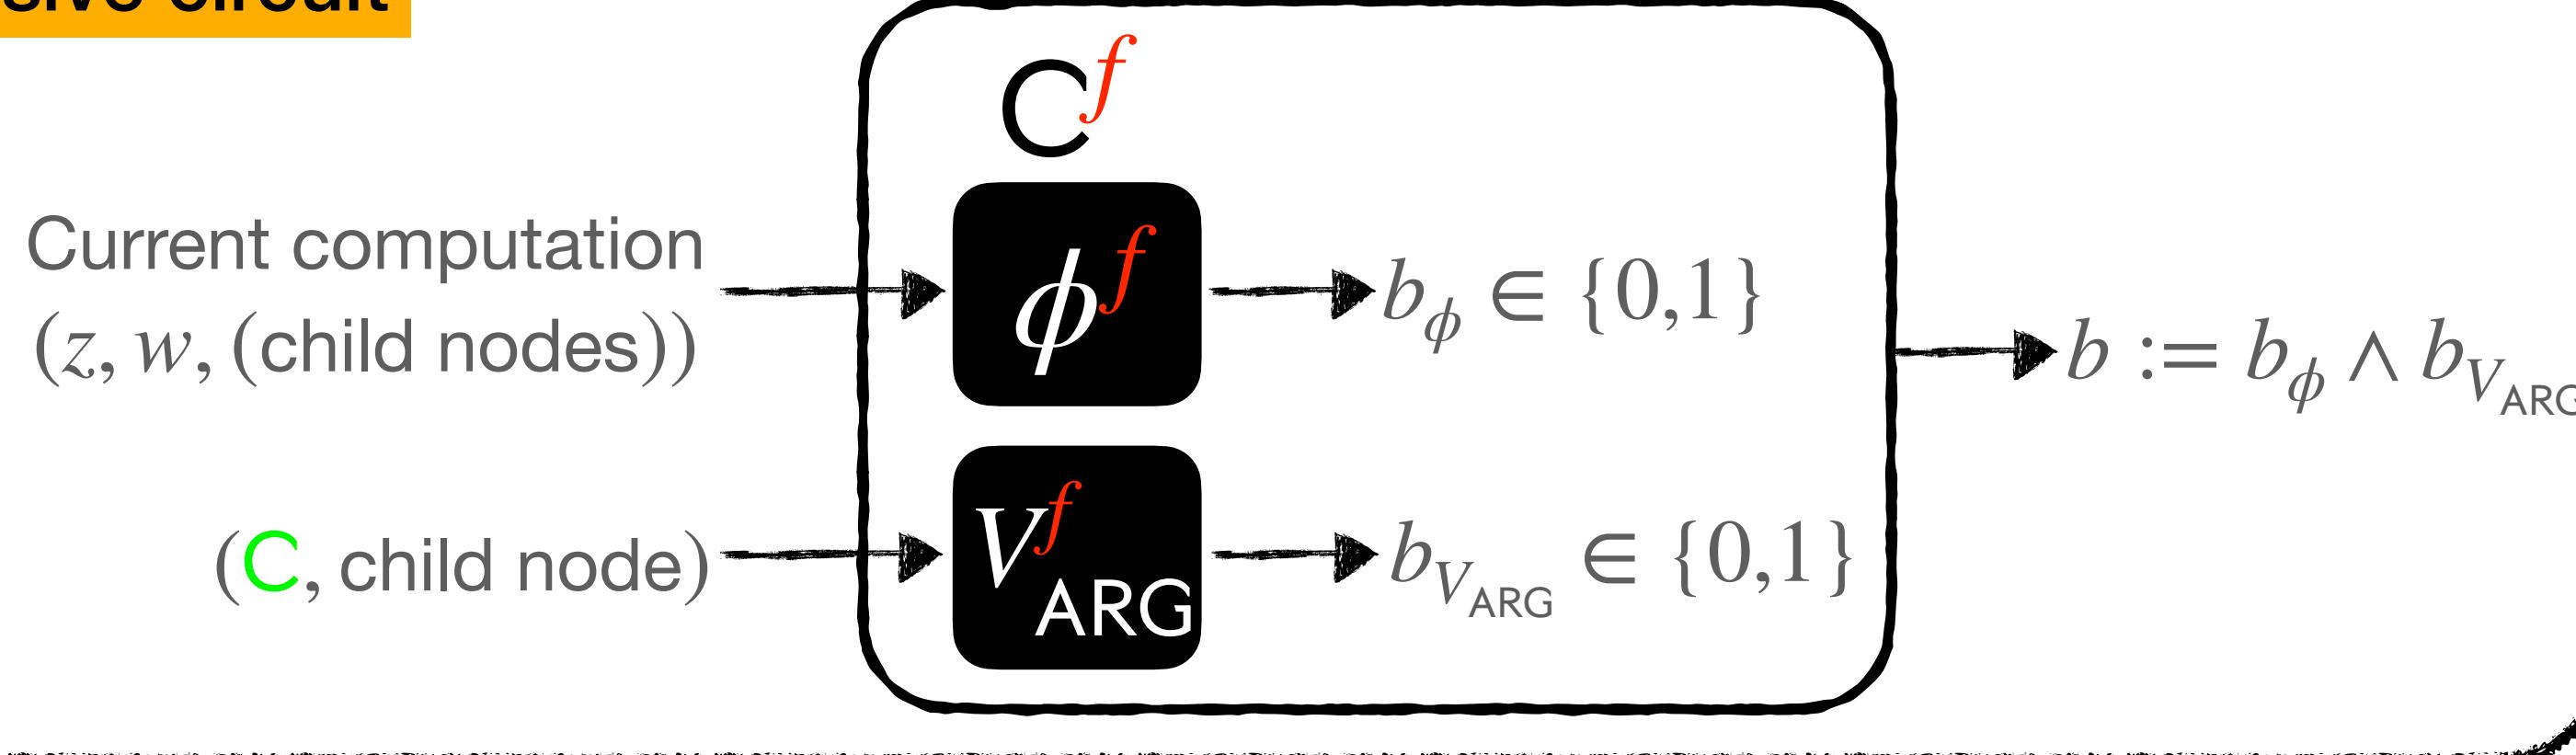


ISSUE! C has oracle access to f .

$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$

Can't we use the previous recursive composition?

Recursive circuit



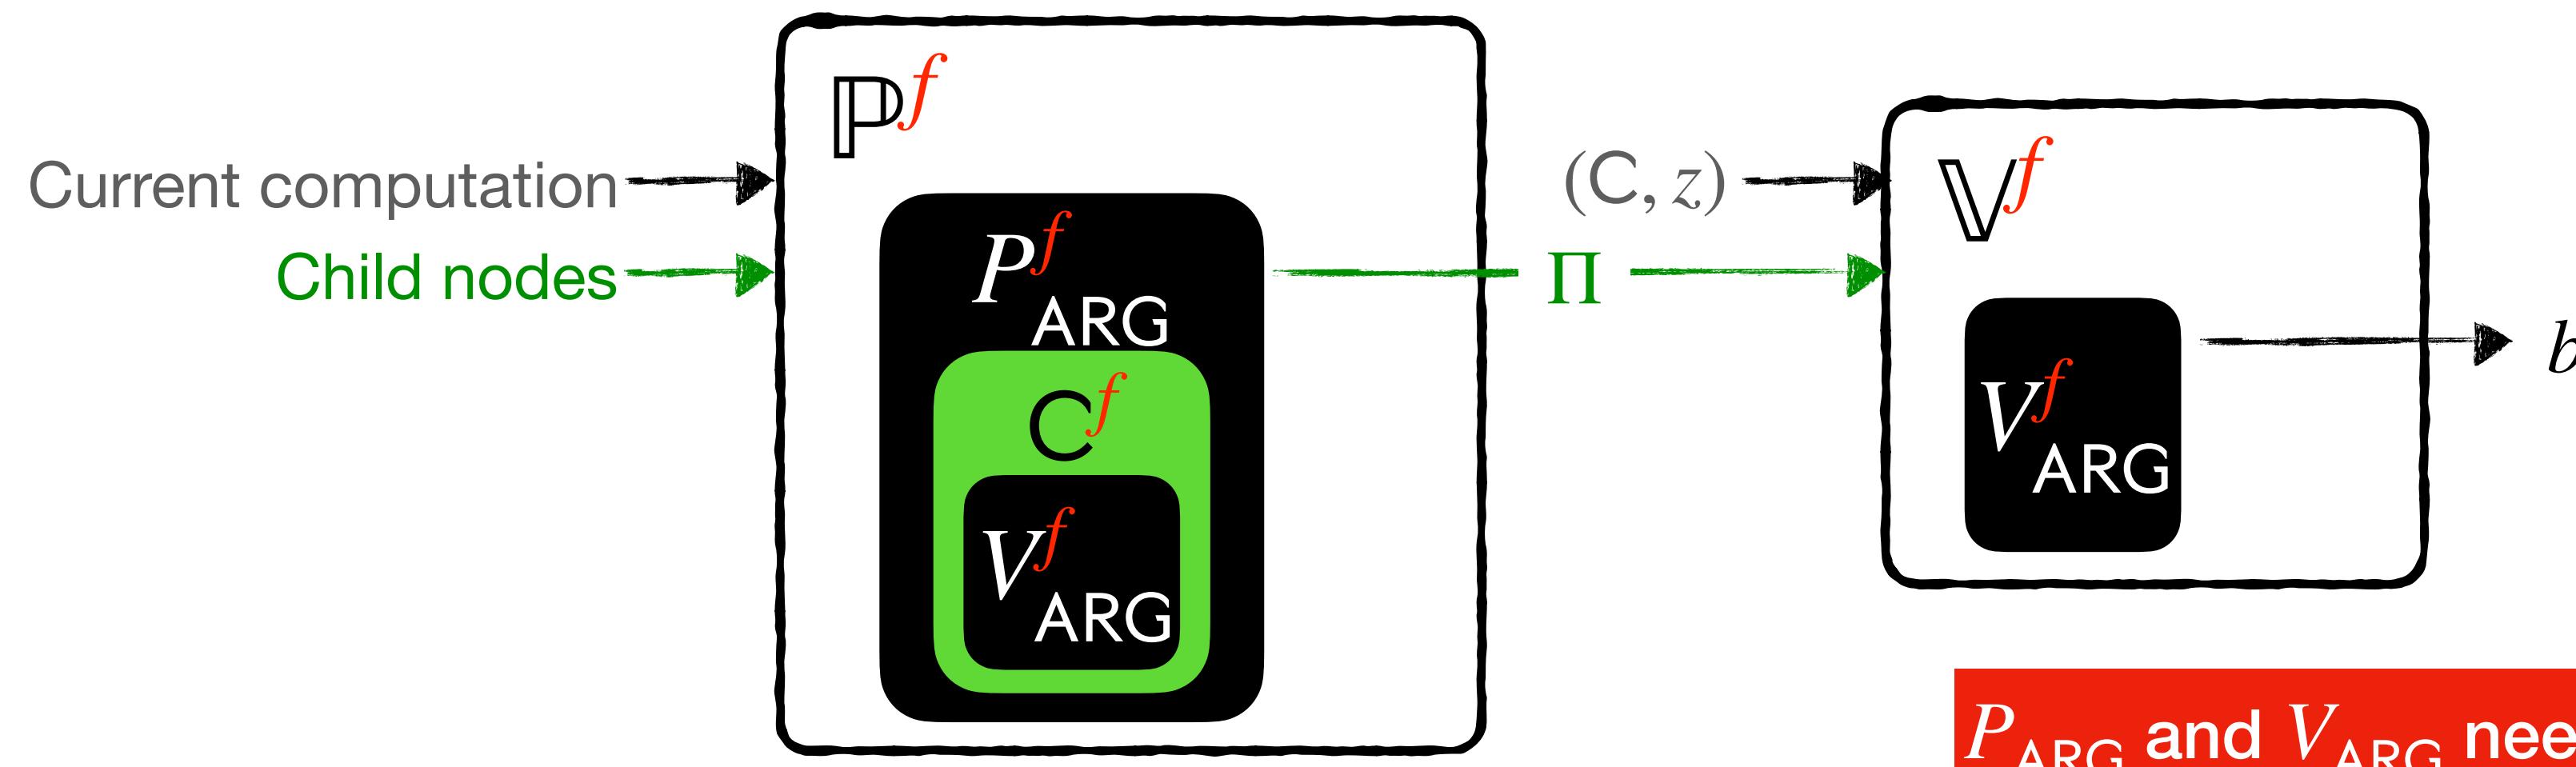
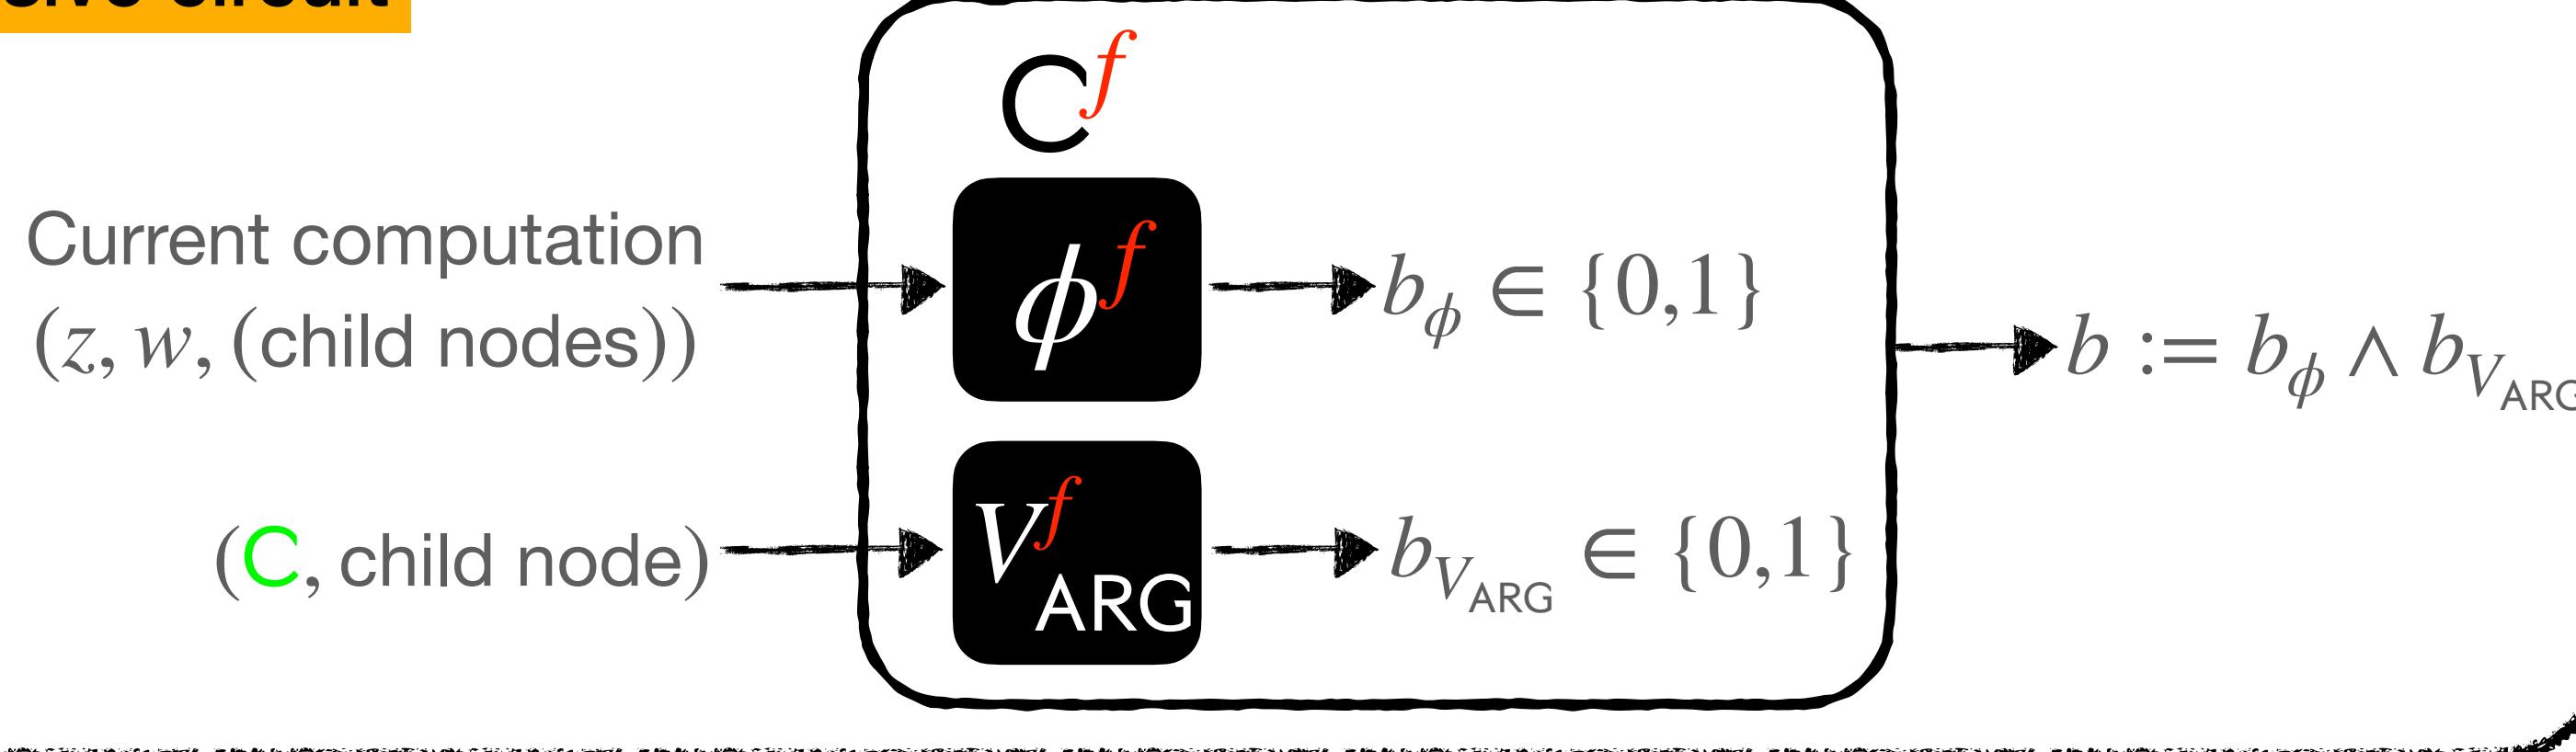
ISSUE! C has oracle access to f .

$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$

$$\text{CSAT}^f := \{(C, a) : C^f(a) = 1\}$$

Can't we use the previous recursive composition?

Recursive circuit



ISSUE! C has oracle access to f .

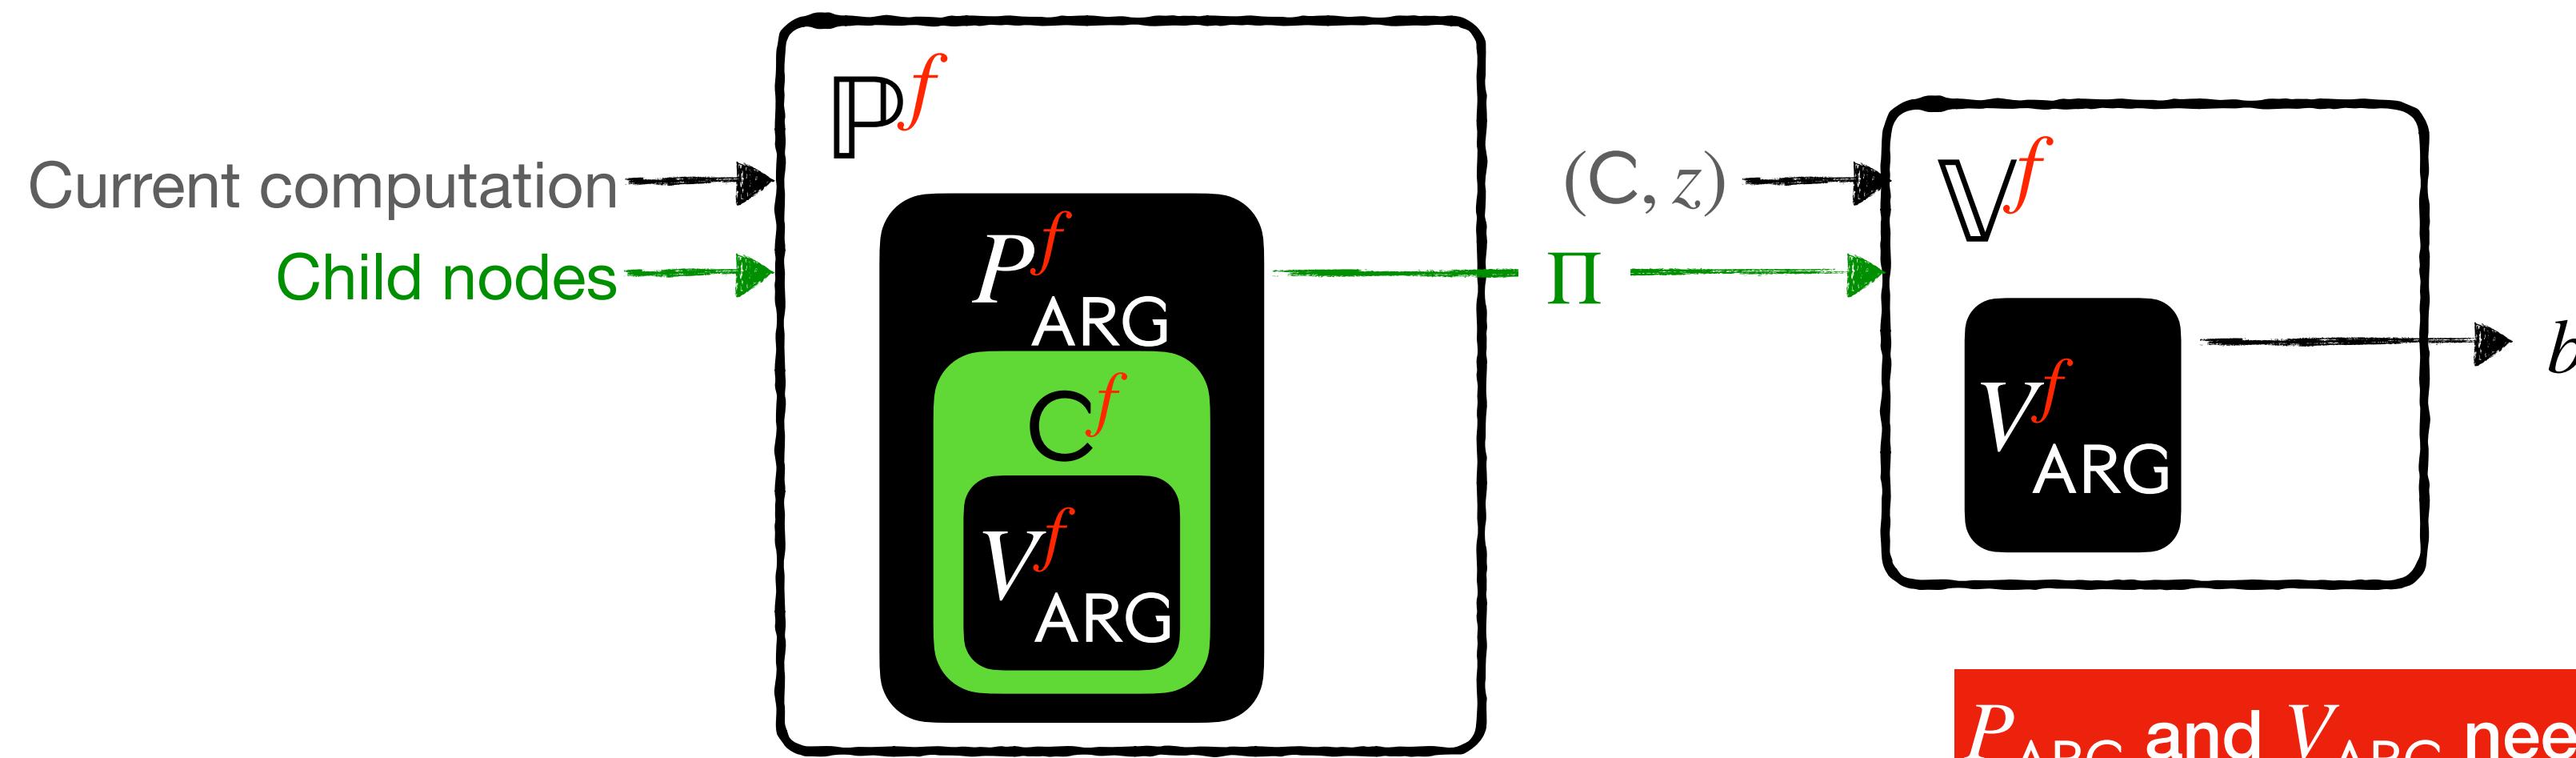
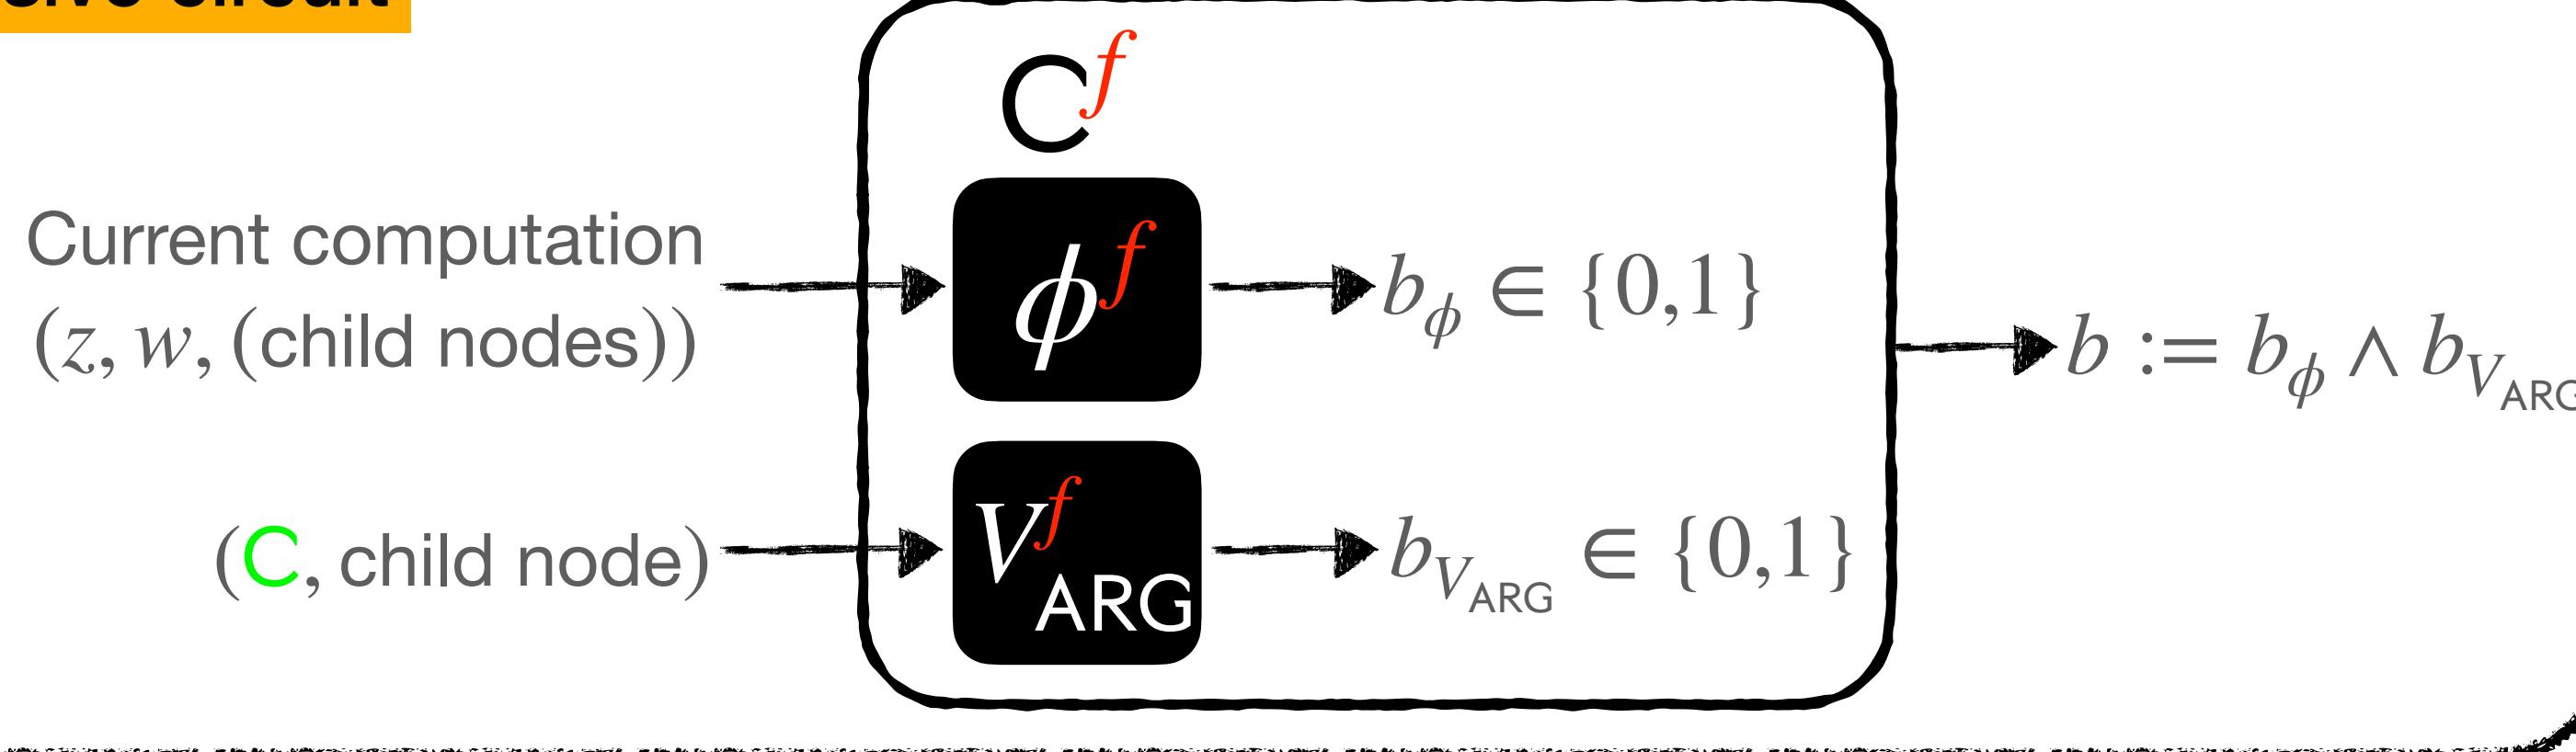
$\text{CSAT} := \{(C, a) : C(a) = 1\}$

$\text{CSAT}^f := \{(C, a) : C^f(a) = 1\}$

P_{ARG} and V_{ARG} need to prove computations involving oracle f .

Can't we use the previous recursive composition?

Recursive circuit



ISSUE! C has oracle access to f .

$$\text{CSAT} := \{(C, a) : C(a) = 1\}$$

$$\text{CSAT}^f := \{(C, a) : C^f(a) = 1\}$$

P_{ARG} and V_{ARG} need to prove computations involving oracle f .

Solution: Relativized SNARKs!

Relativized SNARKs in an oracle model

Relativized SNARKs in an oracle model

$\text{CSAT}^f := \{(C, a) : C^f(a) = 1\}$



Relativized SNARKs in an oracle model

$\text{CSAT}^f := \{(C, a) : C^f(a) = 1\}$



PCD straightline knowledge soundness: \exists deterministic extractor \mathbb{E} , \forall bounded adversary \tilde{P} , the following happens with probability at most $\kappa(\lambda, q, N)$.

λ : security parameter

N : maximum transcript size

q : adversary query bound

Relativized SNARKs in an oracle model

$\text{CSAT}^f := \{(C, a) : C^f(a) = 1\}$

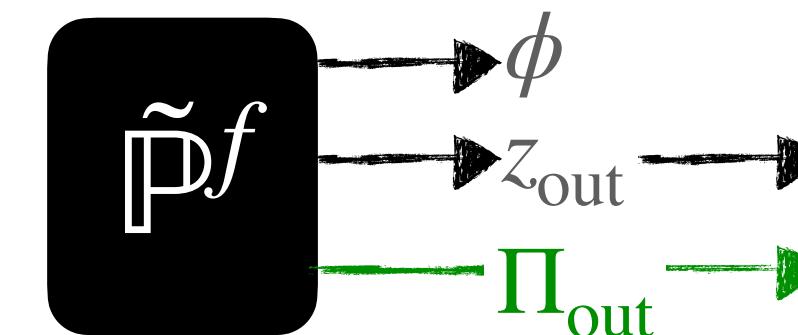


PCD straightline knowledge soundness: \exists deterministic extractor \mathbb{E} , \forall bounded adversary \tilde{P} , the following happens with probability at most $\kappa(\lambda, q, N)$.

λ : security parameter

N : maximum transcript size

q : adversary query bound



Relativized SNARKs in an oracle model

$\text{CSAT}^f := \{(C, a) : C^f(a) = 1\}$

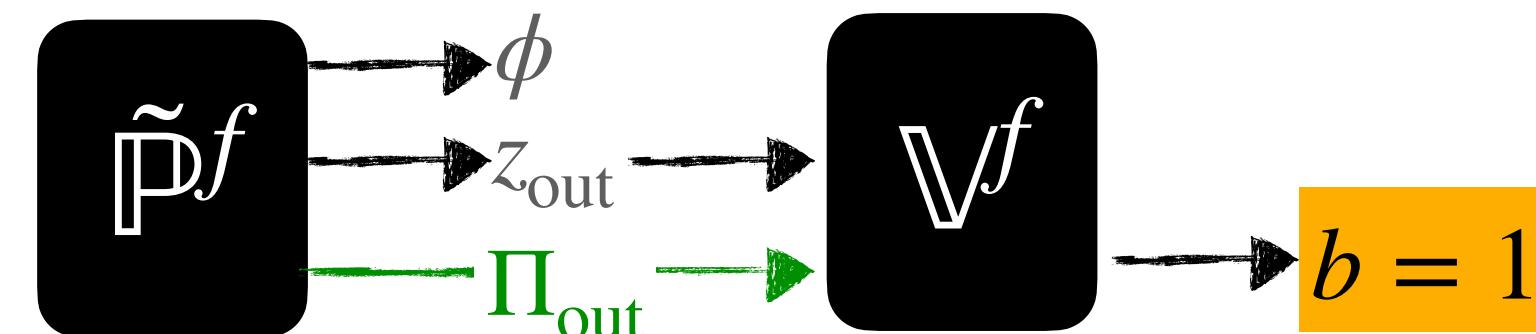


PCD straightline knowledge soundness: \exists deterministic extractor \mathbb{E} , \forall bounded adversary \tilde{P} , the following happens with probability at most $\kappa(\lambda, q, N)$.

λ : security parameter

N : maximum transcript size

q : adversary query bound



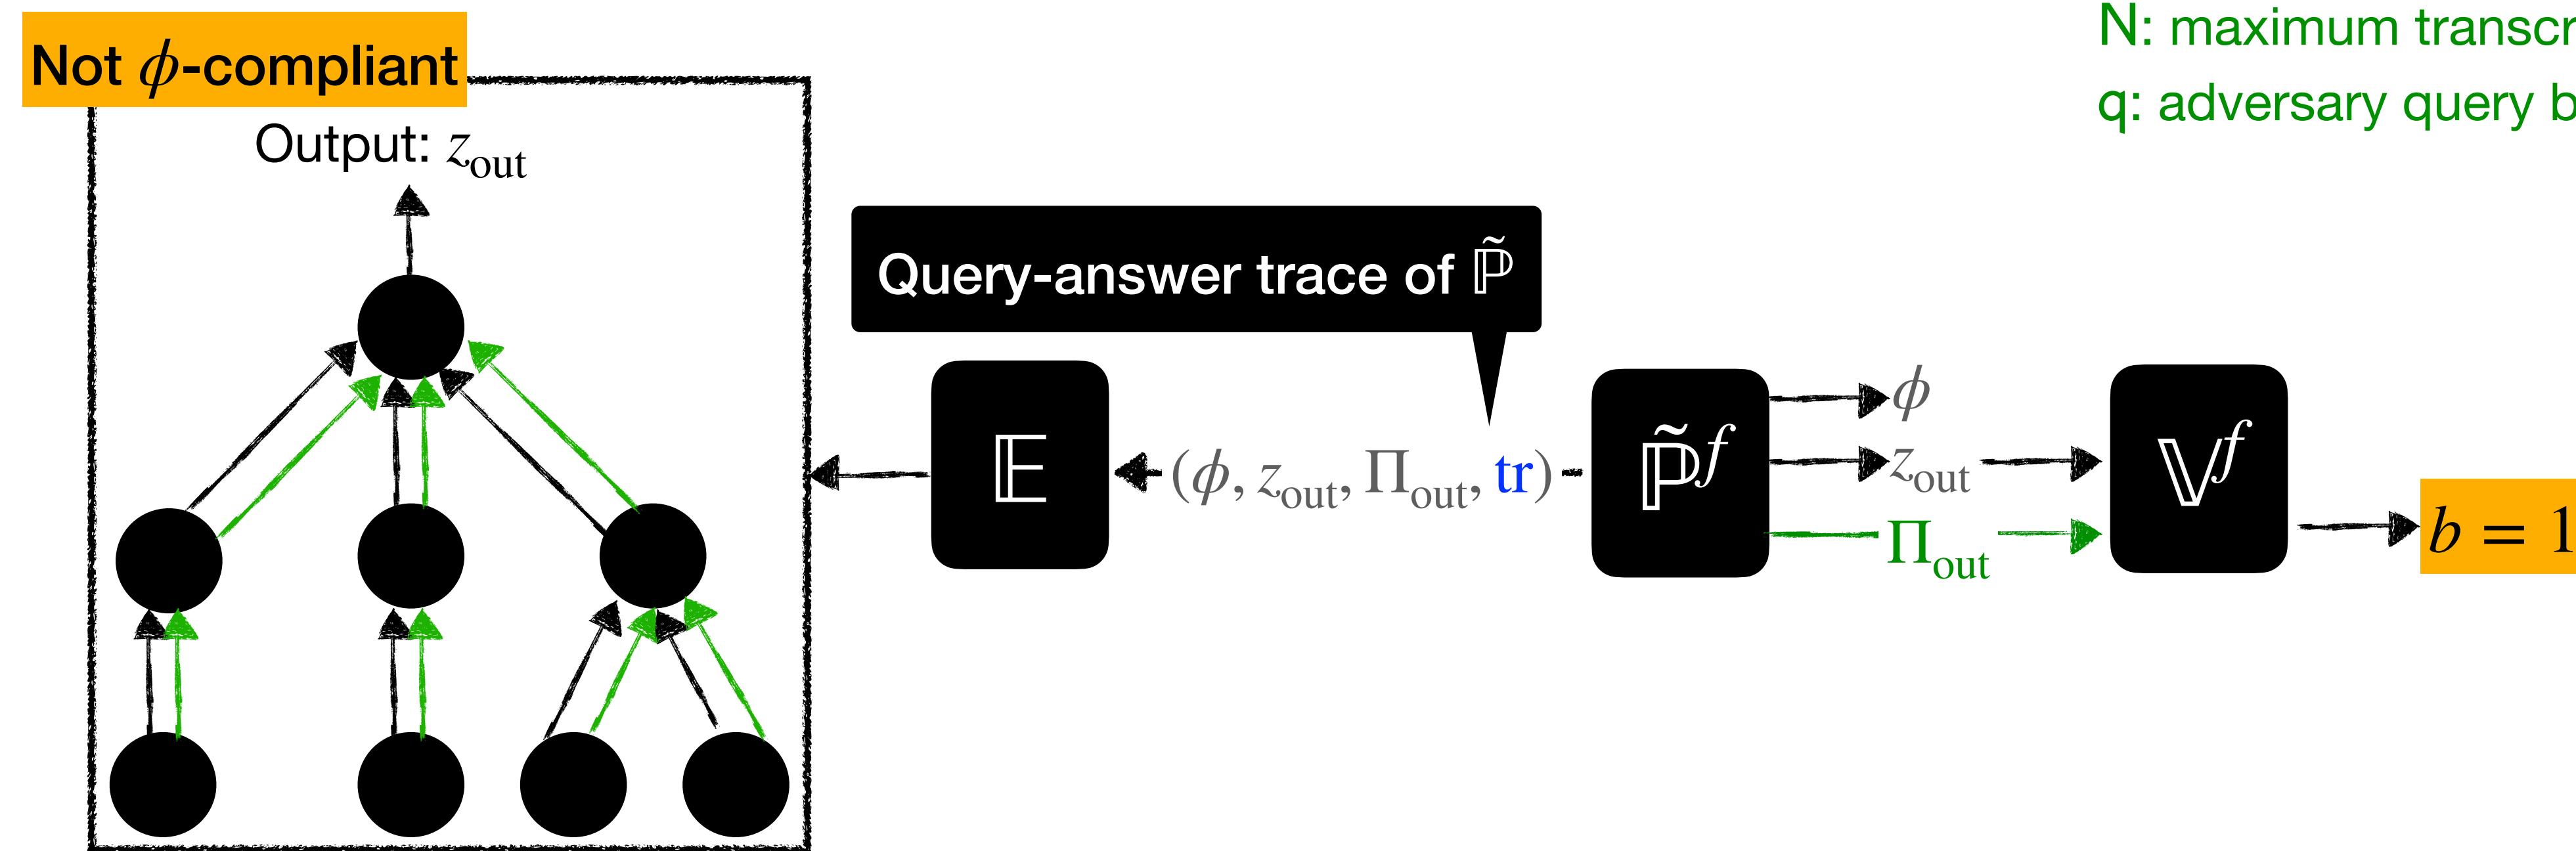
Relativized SNARKs in an oracle model

$\text{CSAT}^f := \{(C, a) : C^f(a) = 1\}$



PCD straightline knowledge soundness: \exists deterministic extractor \mathbb{E} , \forall bounded adversary \tilde{P} , the following happens with probability at most $\kappa(\lambda, q, N)$.

λ : security parameter
 N : maximum transcript size
 q : adversary query bound



Application: Set security for hash-based PCD

Warm-up: analyzing hash-based SNARKs

Warm-up: analyzing hash-based SNARKs

Three-step recipe:

Warm-up: analyzing hash-based SNARKs

Three-step recipe:

Step 1. Model the hash function as "ideal": a random function.

Warm-up: analyzing hash-based SNARKs

Three-step recipe:

Step 1. Model the hash function as "ideal": a random function.

- the hash-based SNARK is idealized as **a SNARK in the random oracle model (ROM-SNARK)**.

Warm-up: analyzing hash-based SNARKs

Three-step recipe:

Step 1. Model the hash function as "ideal": a random function.

- the hash-based SNARK is idealized as **a SNARK in the random oracle model (ROM-SNARK)**.



Warm-up: analyzing hash-based SNARKs

Three-step recipe:

Step 1. Model the hash function as "ideal": a random function.

- the hash-based SNARK is idealized as **a SNARK in the random oracle model (ROM-SNARK)**.

Step 2. Establish **concrete** security bounds for the ROM-SNARK.



Warm-up: analyzing hash-based SNARKs

Three-step recipe:

Step 1. Model the hash function as "ideal": a random function.

- the hash-based SNARK is idealized as **a SNARK in the random oracle model (ROM-SNARK)**.

Step 2. Establish **concrete** security bounds for the ROM-SNARK.

Step 3. Set security parameters of the hash-based SNARK accordingly.



Warm-up: analyzing hash-based SNARKs

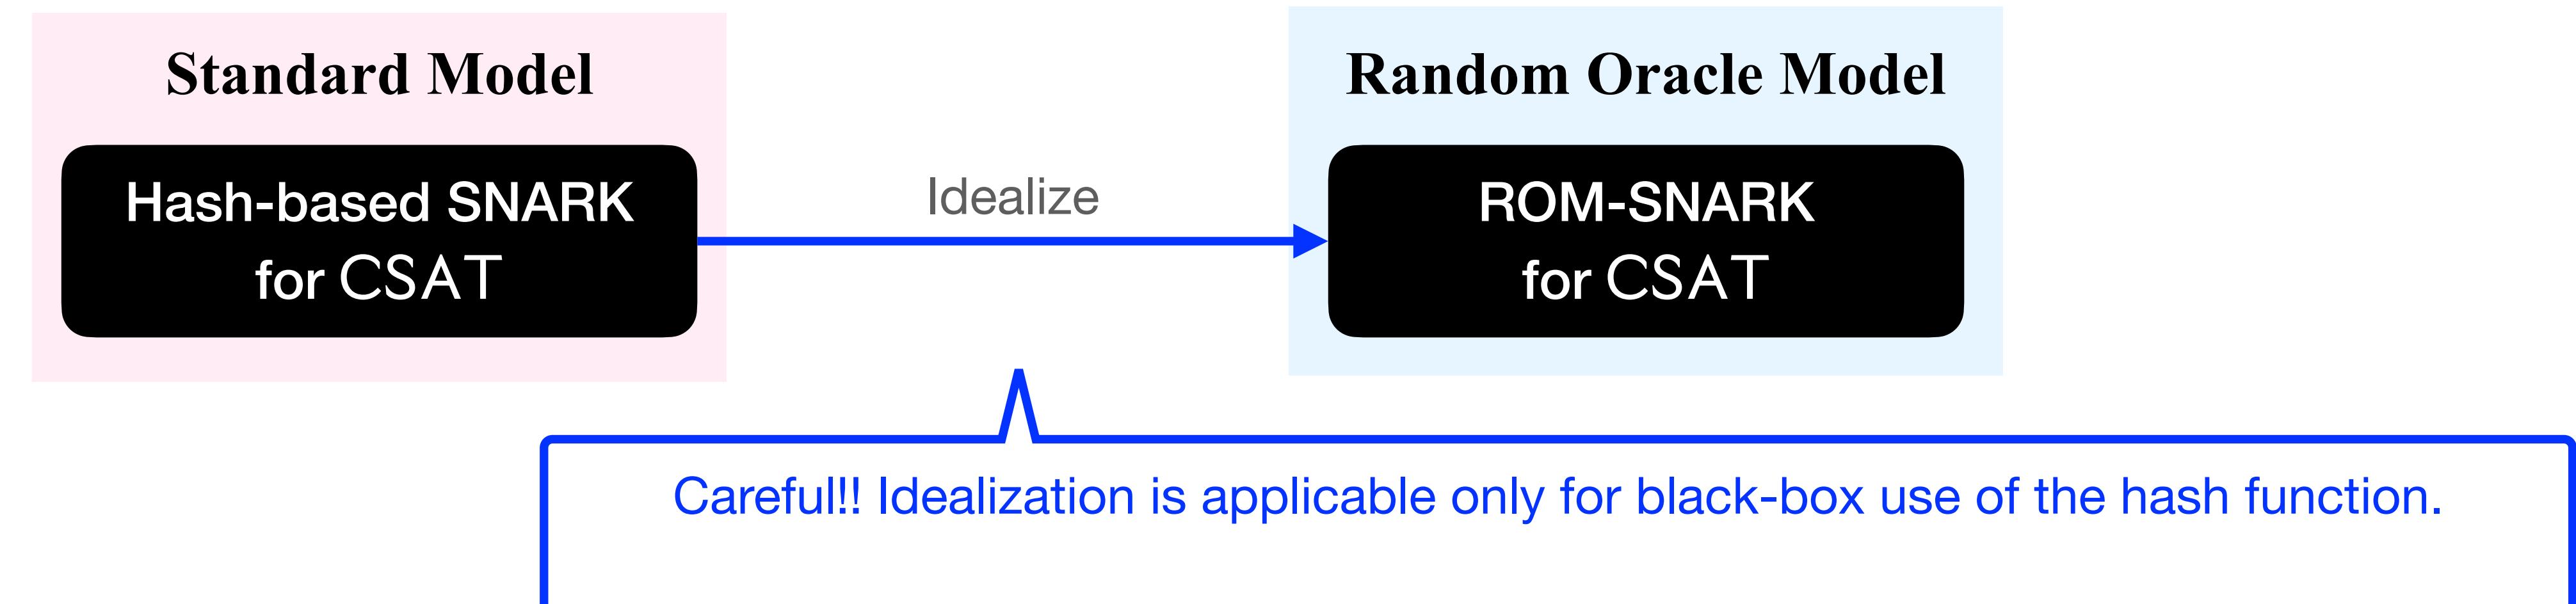
Three-step recipe:

Step 1. Model the hash function as "ideal": a random function.

- the hash-based SNARK is idealized as **a SNARK in the random oracle model (ROM-SNARK)**.

Step 2. Establish **concrete** security bounds for the ROM-SNARK.

Step 3. Set security parameters of the hash-based SNARK accordingly.



Warm-up: analyzing hash-based SNARKs

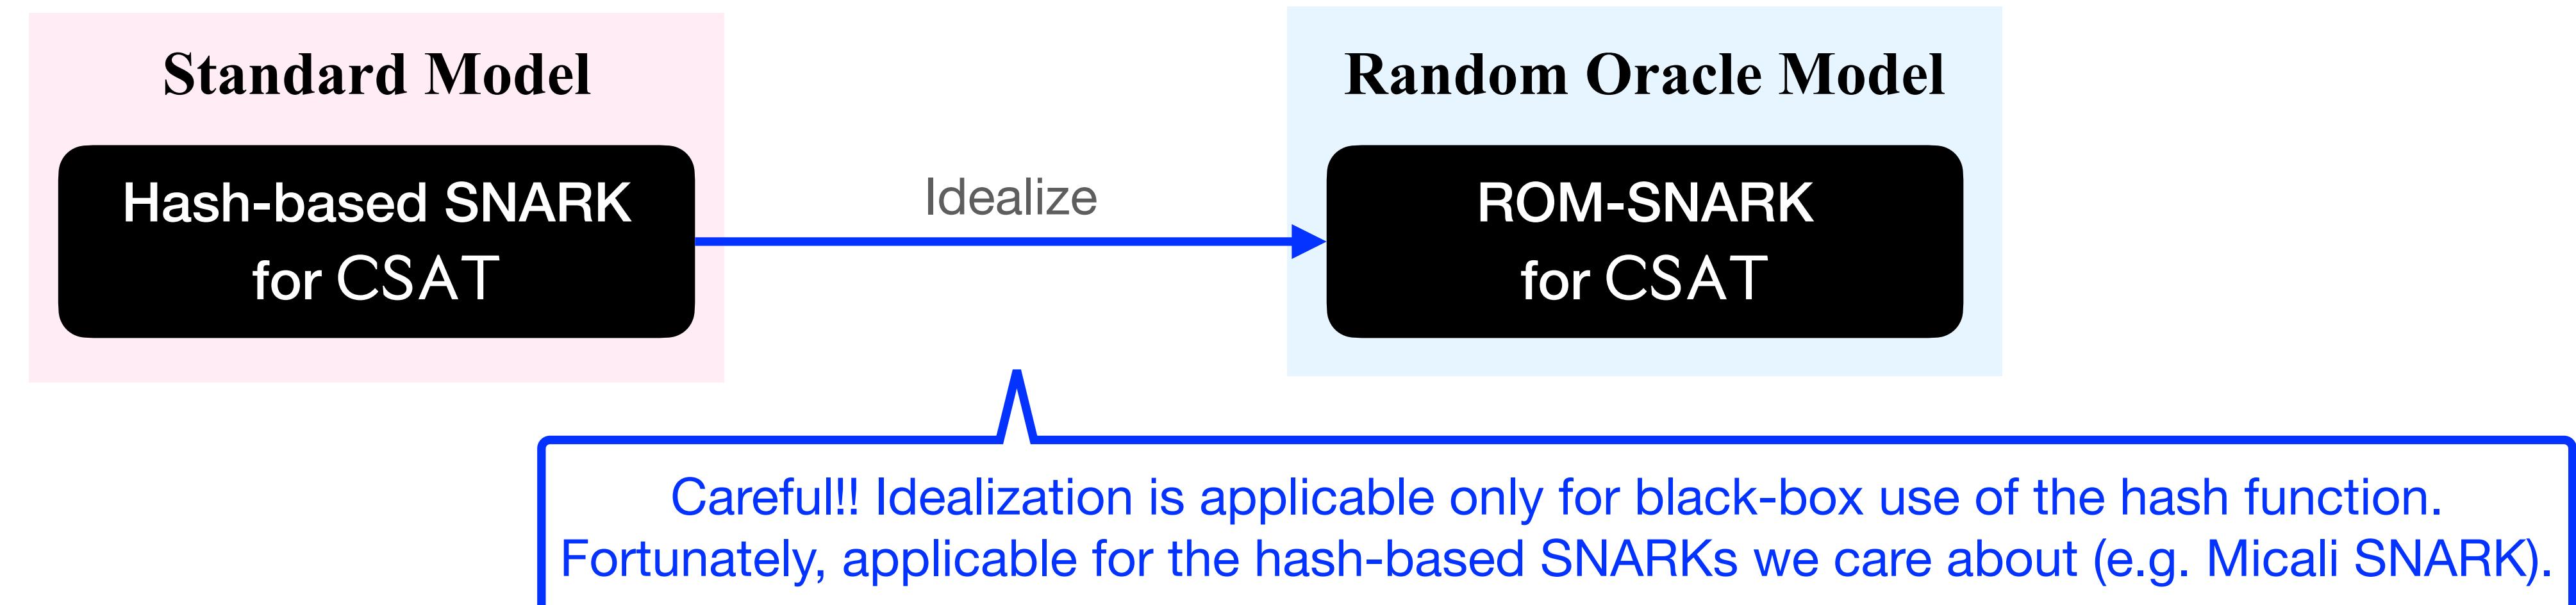
Three-step recipe:

Step 1. Model the hash function as "ideal": a random function.

- the hash-based SNARK is idealized as **a SNARK in the random oracle model (ROM-SNARK)**.

Step 2. Establish **concrete** security bounds for the ROM-SNARK.

Step 3. Set security parameters of the hash-based SNARK accordingly.



First attempt for idealization of hash-based PCD

First attempt for idealization of hash-based PCD

PCDs are deployed based on various approaches. A popular approach is **hash-based PCD**.

First attempt for idealization of hash-based PCD

PCDs are deployed based on various approaches. A popular approach is **hash-based PCD**.

Standard Model

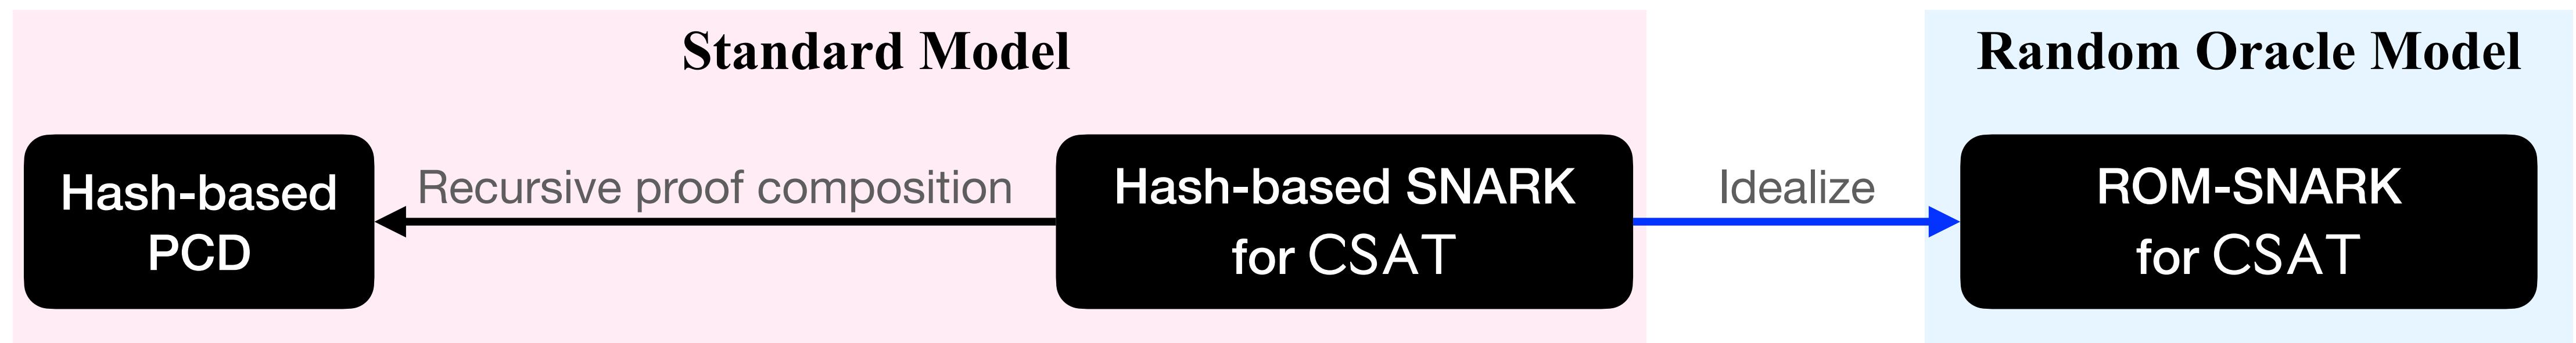
First attempt for idealization of hash-based PCD

PCDs are deployed based on various approaches. A popular approach is **hash-based PCD**.



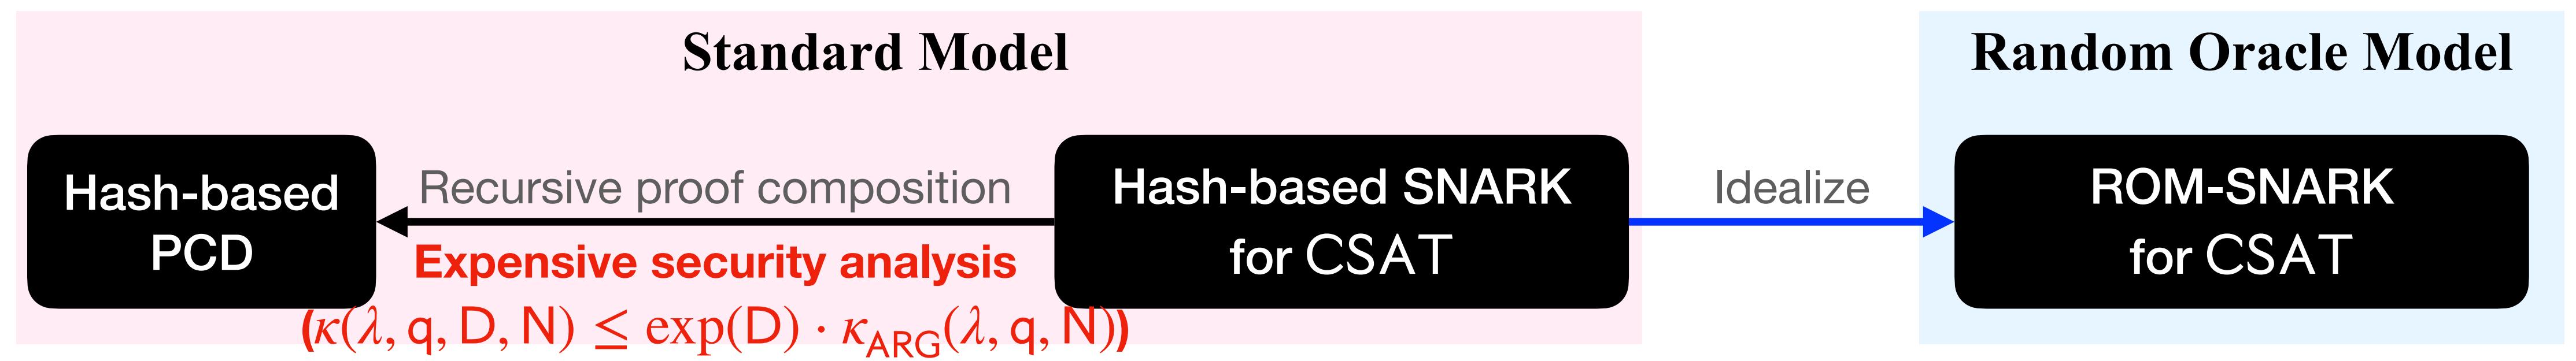
First attempt for idealization of hash-based PCD

PCDs are deployed based on various approaches. A popular approach is **hash-based PCD**.



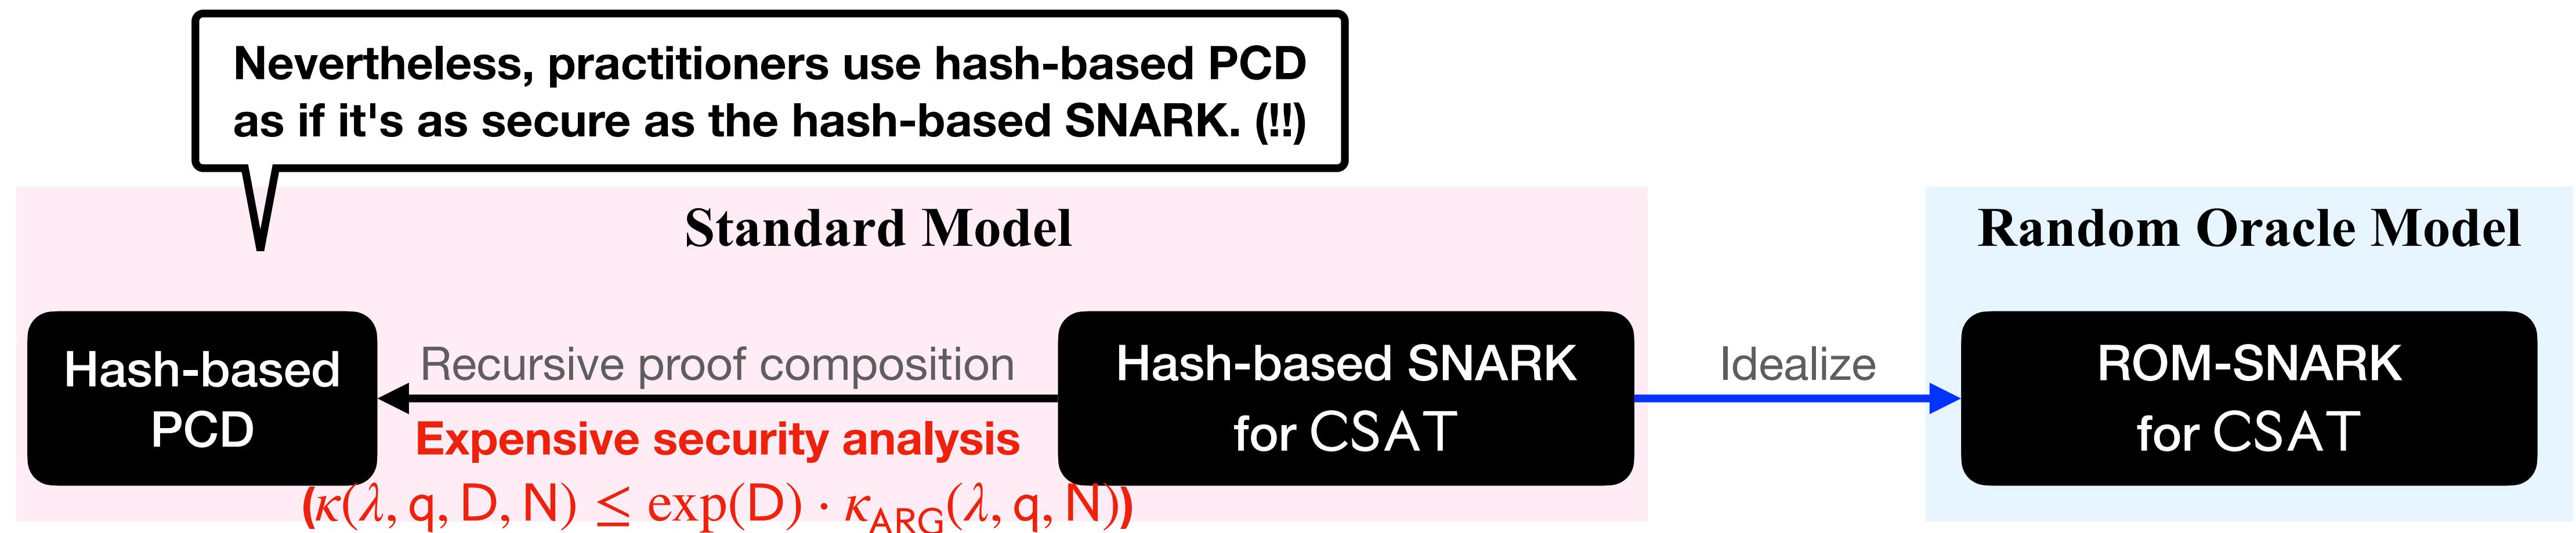
First attempt for idealization of hash-based PCD

PCDs are deployed based on various approaches. A popular approach is **hash-based PCD**.



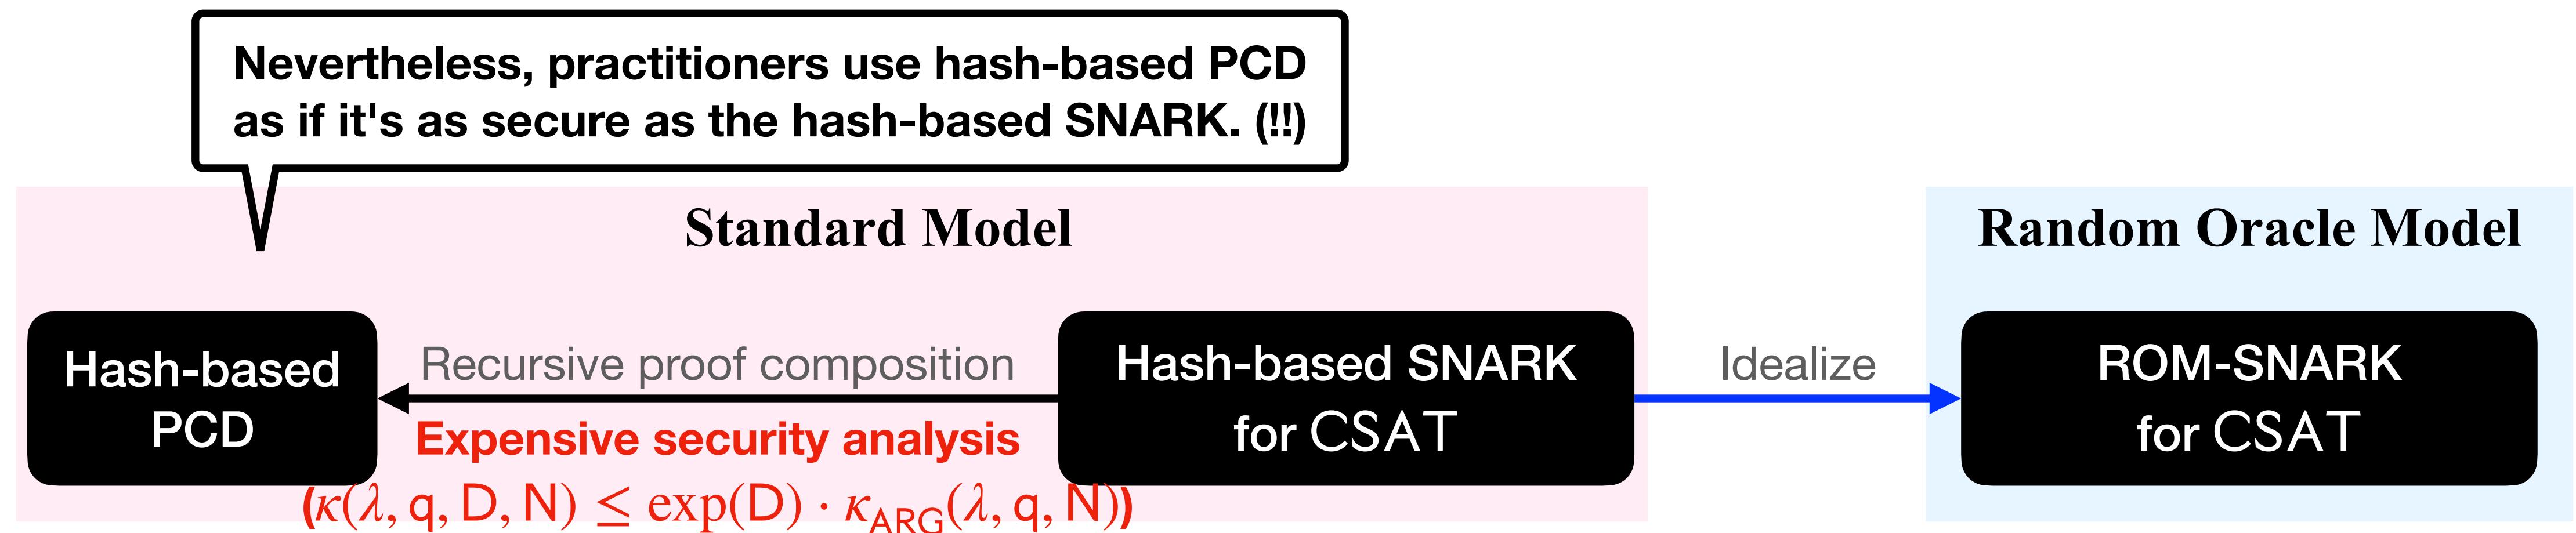
First attempt for idealization of hash-based PCD

PCDs are deployed based on various approaches. A popular approach is **hash-based PCD**.



First attempt for idealization of hash-based PCD

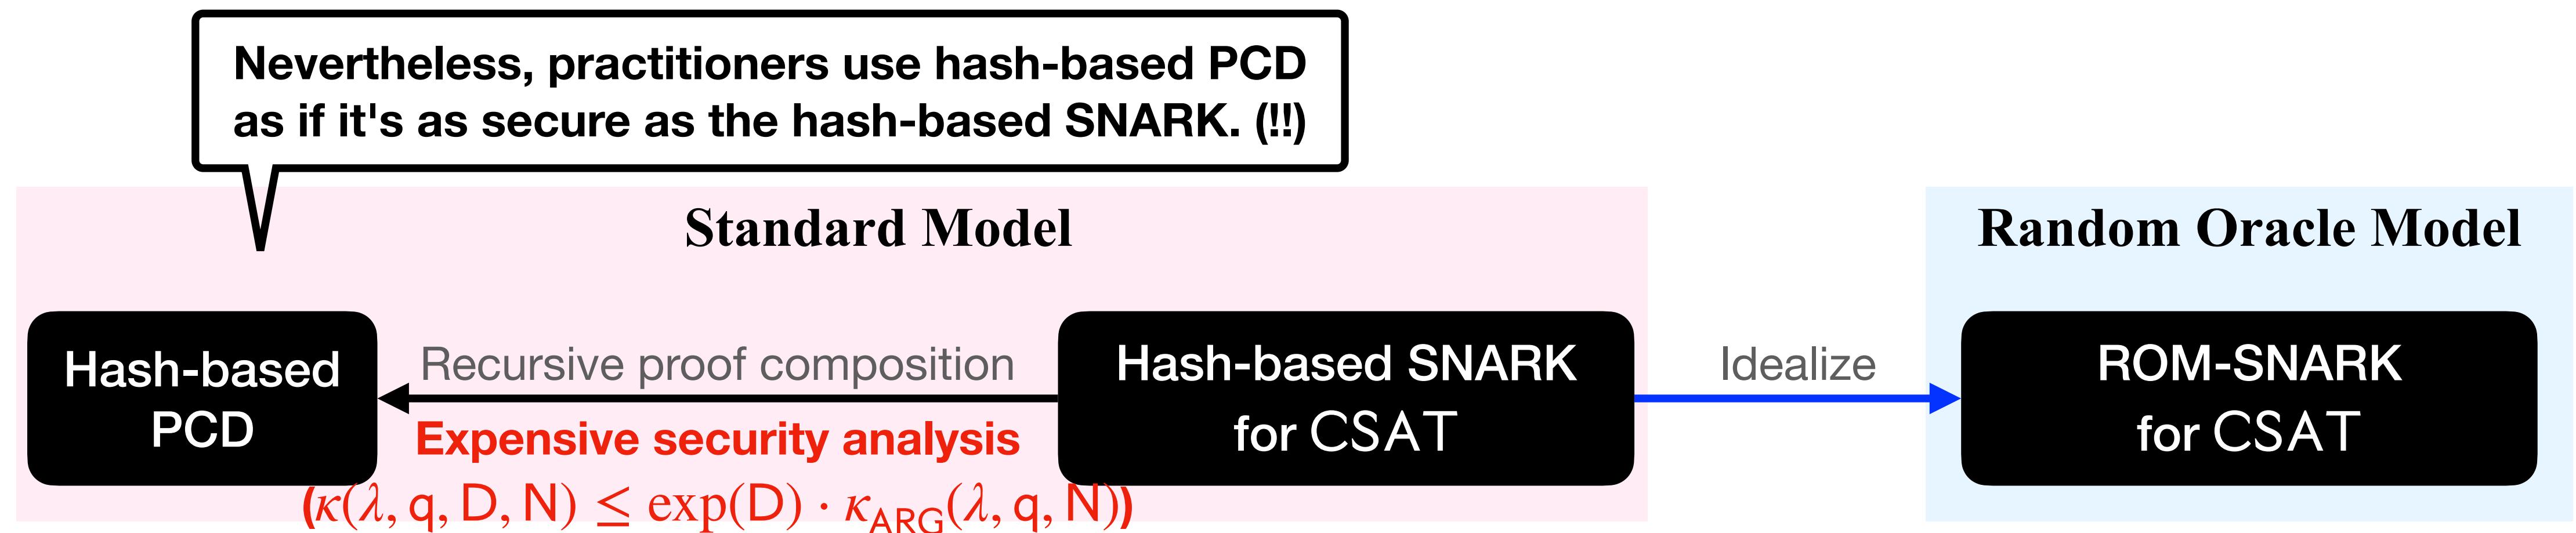
PCDs are deployed based on various approaches. A popular approach is **hash-based PCD**.



Can our new analysis justify the above practice?

First attempt for idealization of hash-based PCD

PCDs are deployed based on various approaches. A popular approach is **hash-based PCD**.



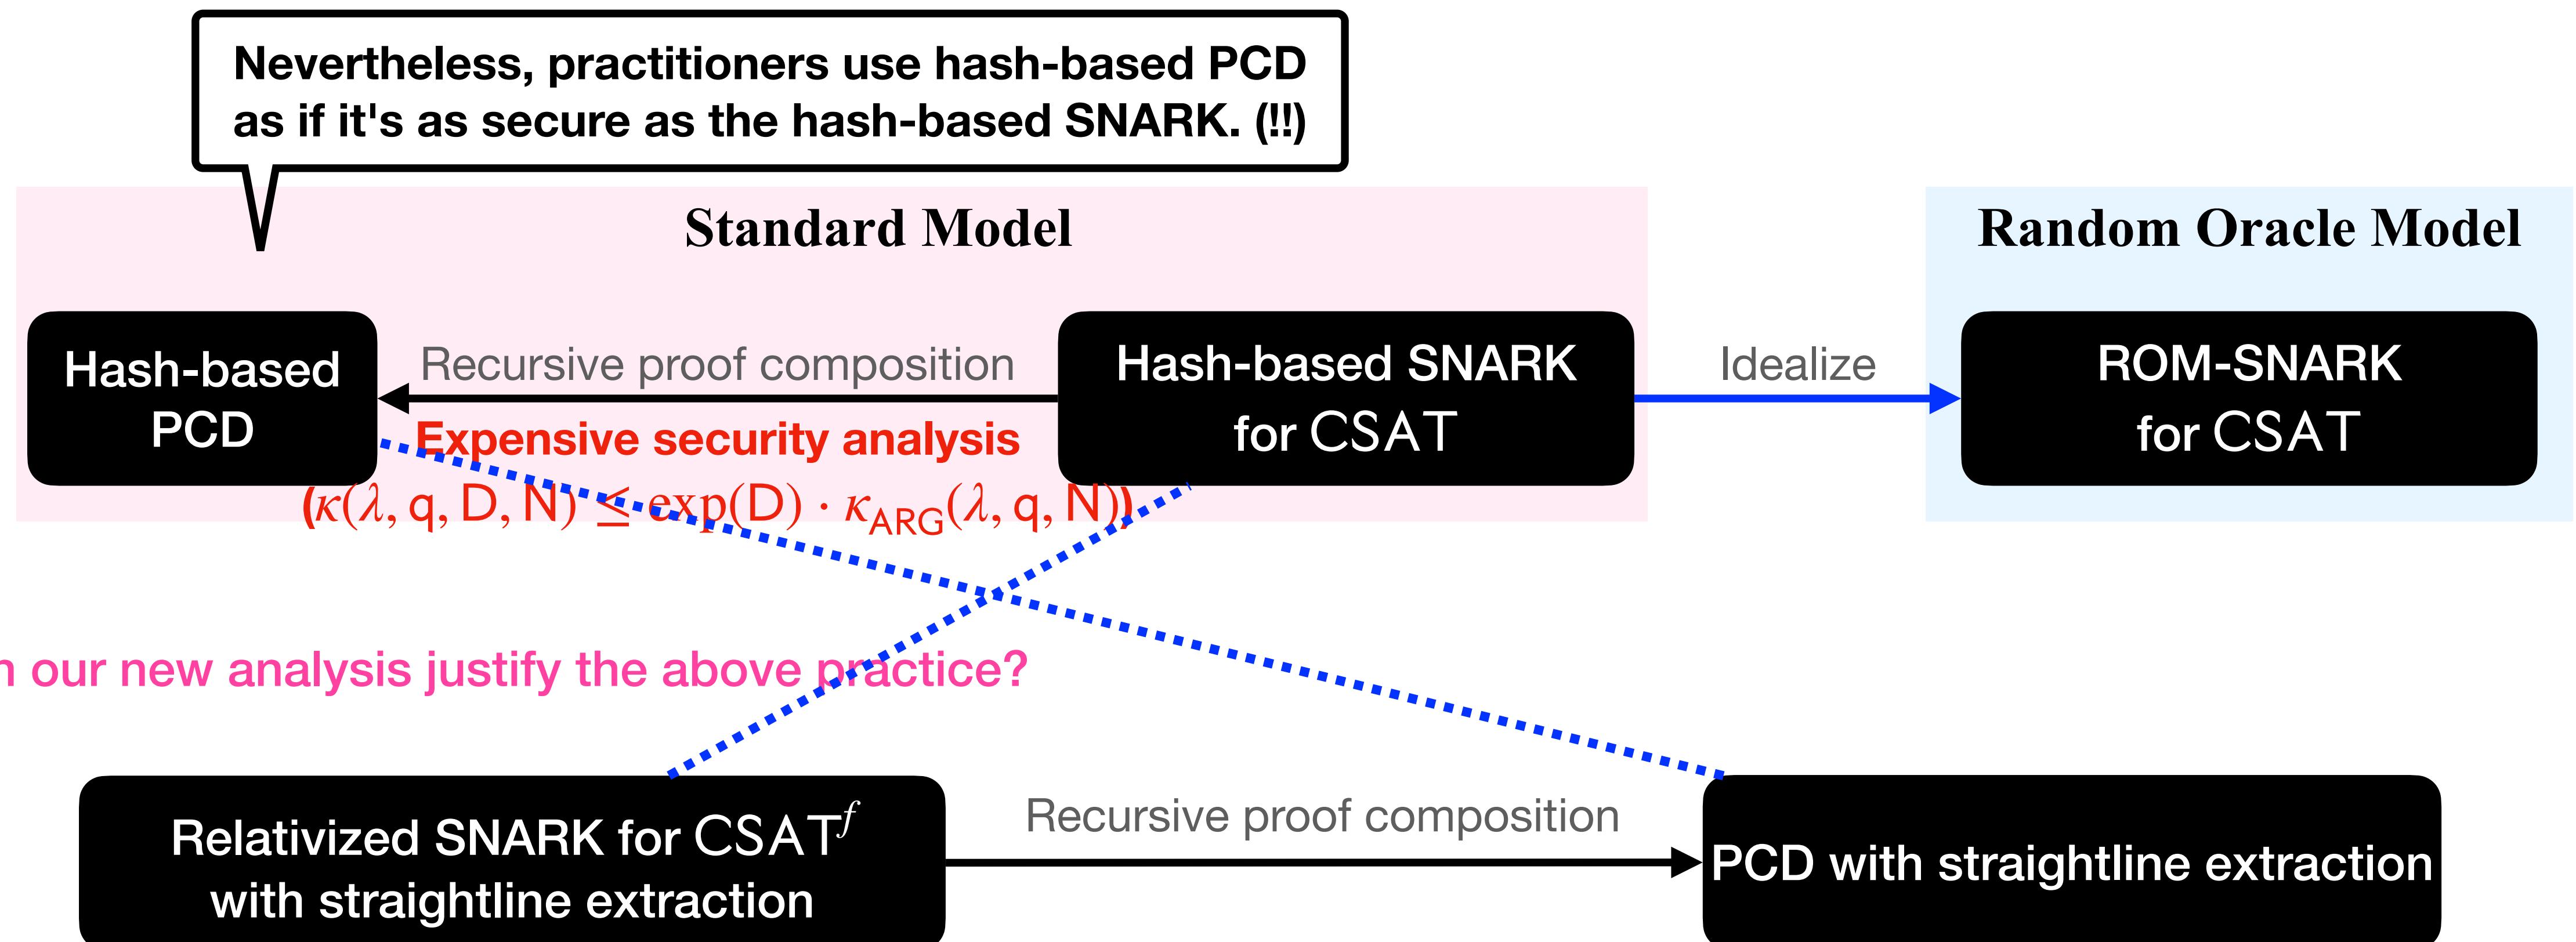
Can our new analysis justify the above practice?



Our theorem: $\kappa(\lambda, q, D, N) \leq \kappa_{\text{ARG}}(\lambda, q, N)$

First attempt for idealization of hash-based PCD

PCDs are deployed based on various approaches. A popular approach is **hash-based PCD**.



Second attempt for idealization of hash-based PCD

Second attempt for idealization of hash-based PCD

What we hope to do

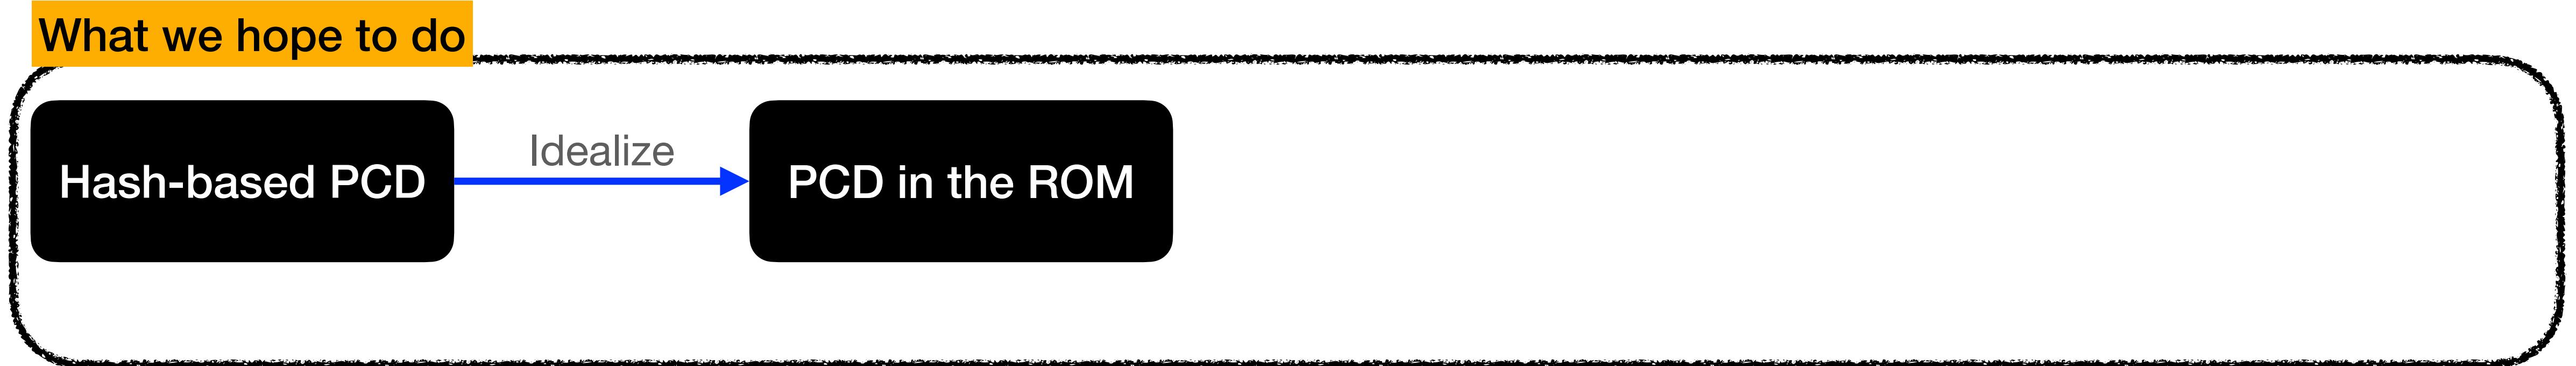


Second attempt for idealization of hash-based PCD

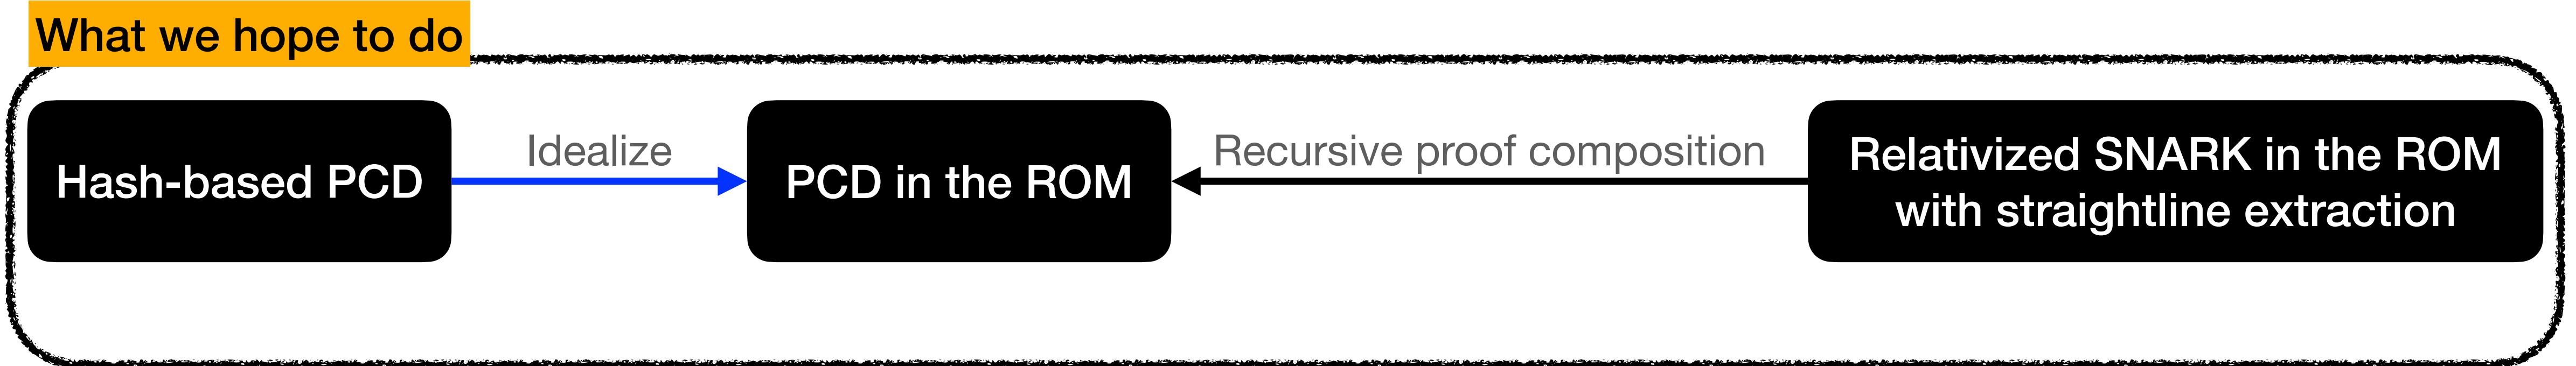
What we hope to do

Hash-based PCD

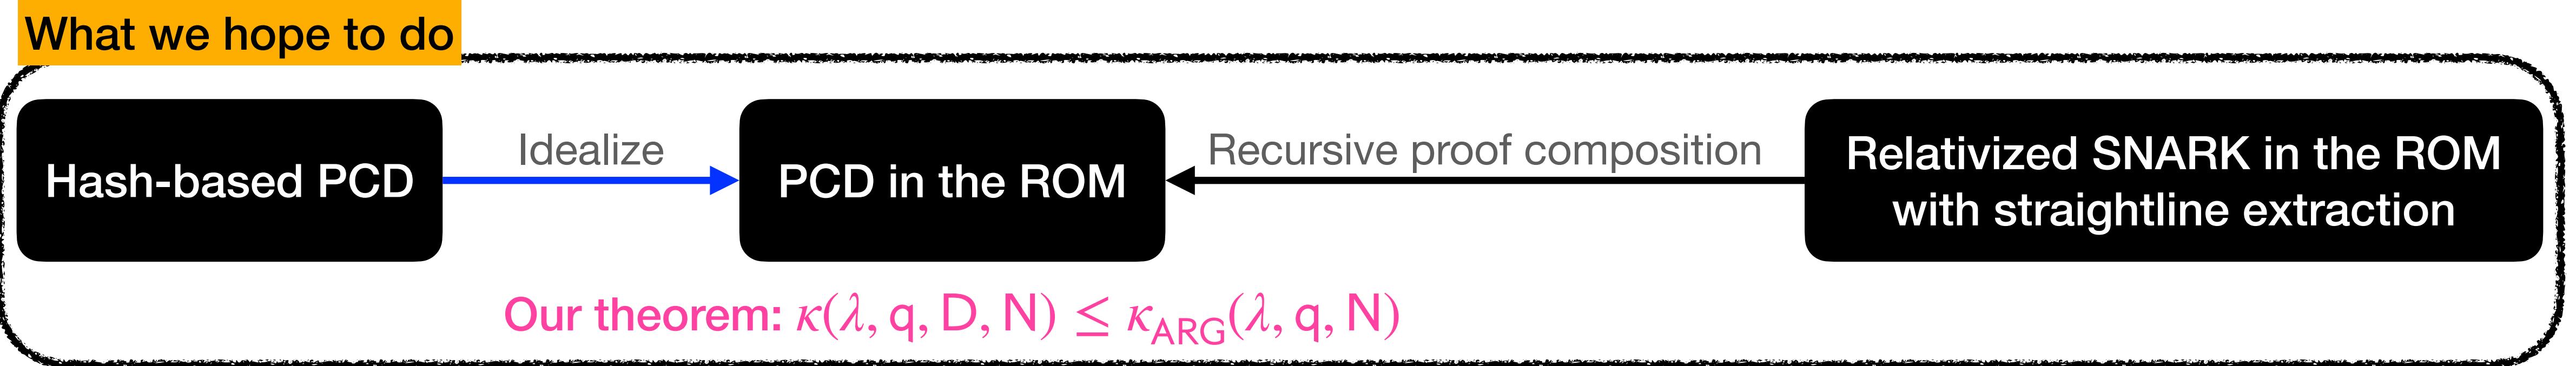
Second attempt for idealization of hash-based PCD



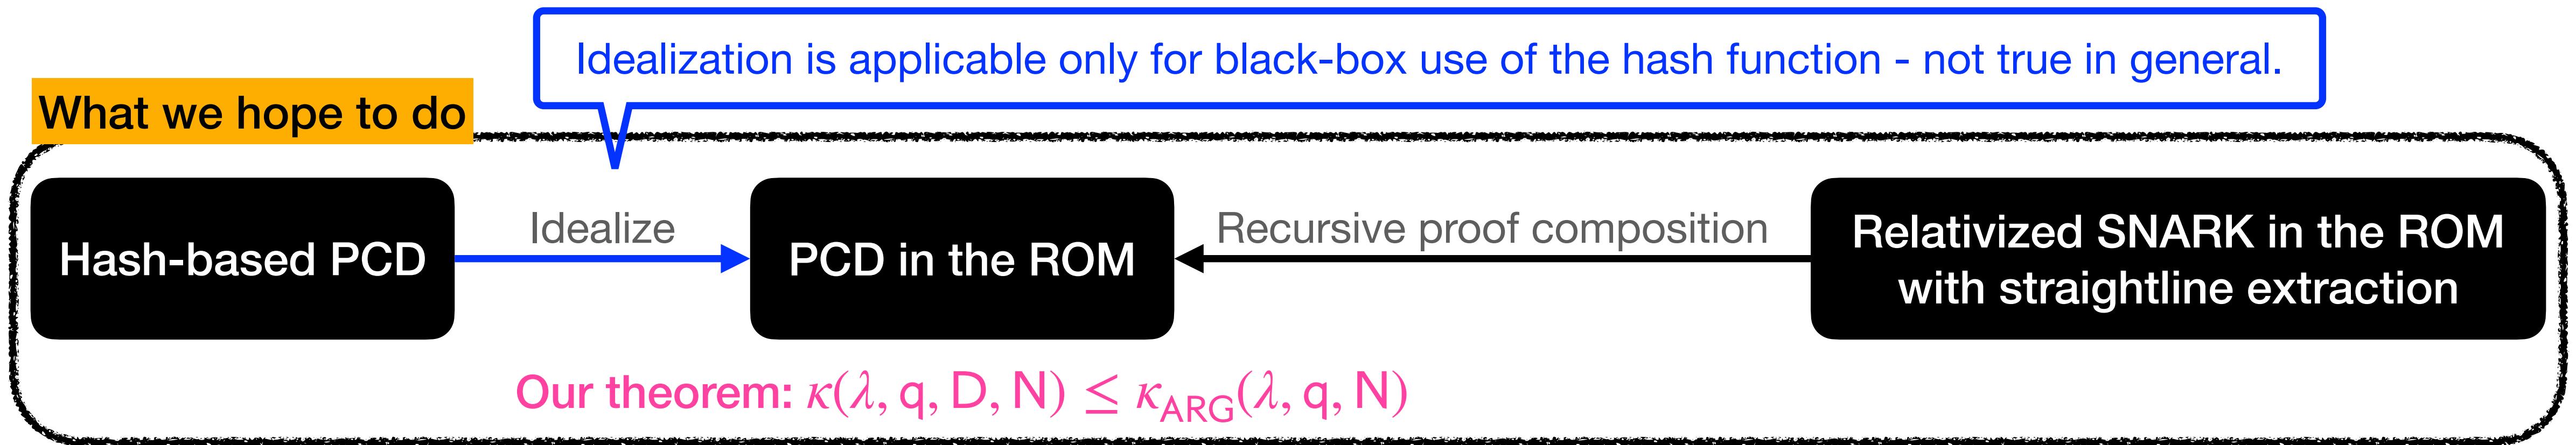
Second attempt for idealization of hash-based PCD



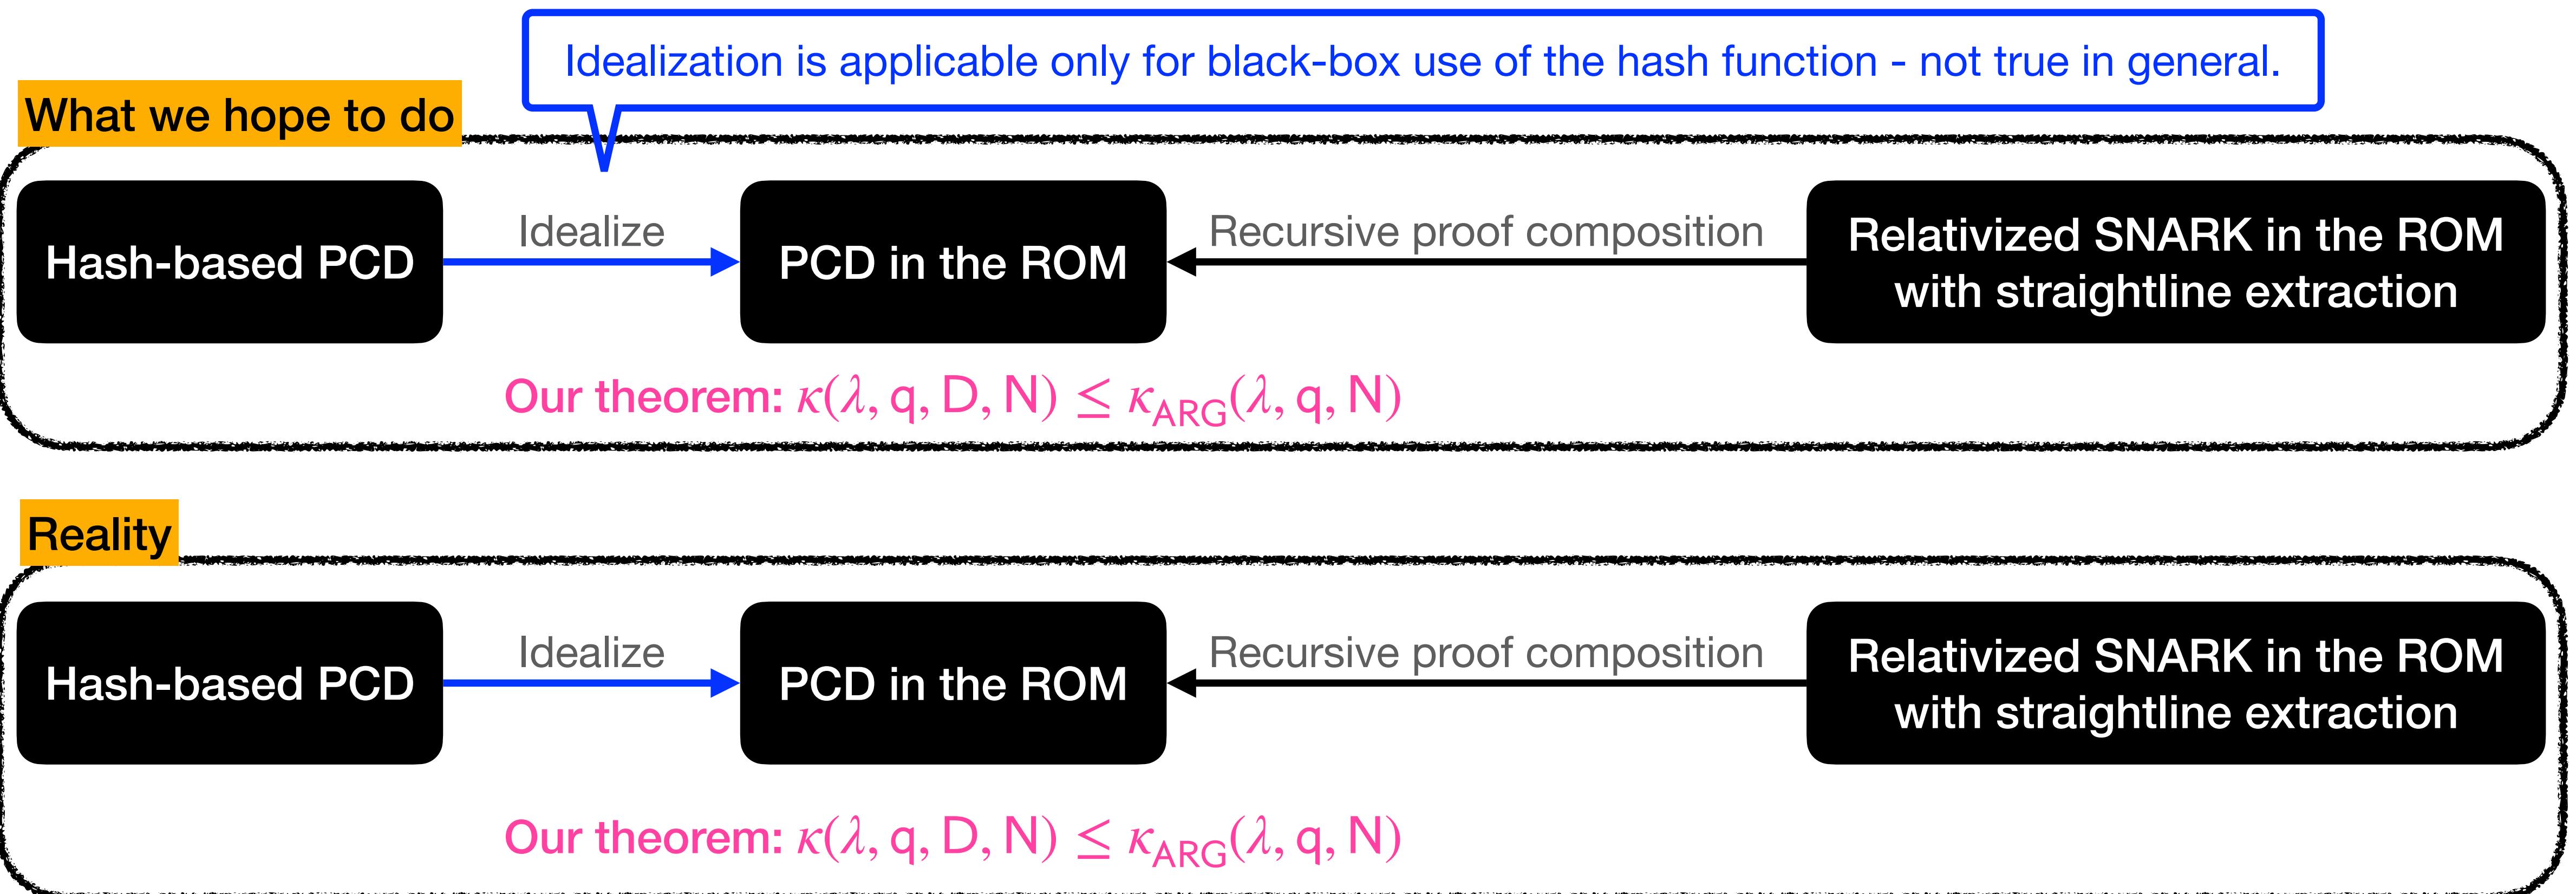
Second attempt for idealization of hash-based PCD



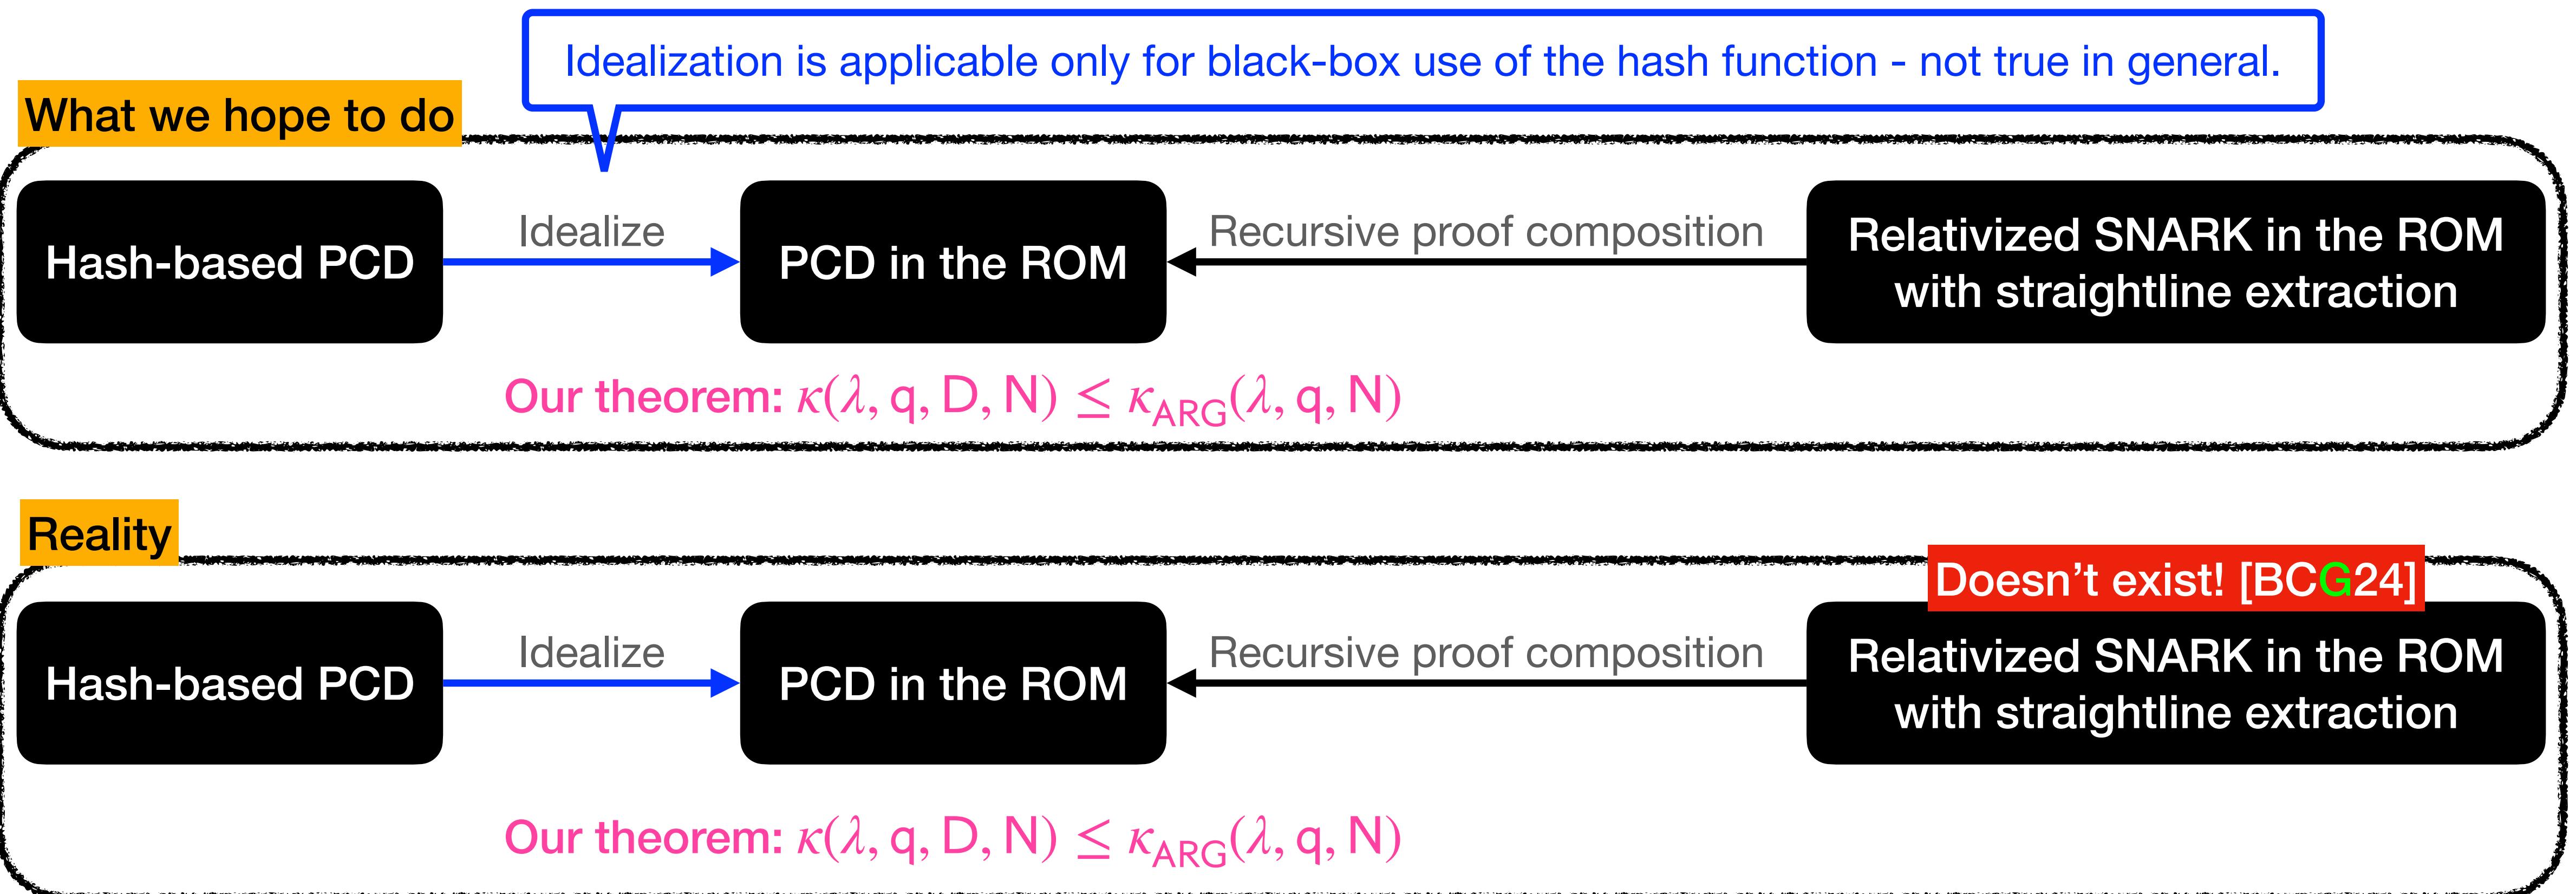
Second attempt for idealization of hash-based PCD



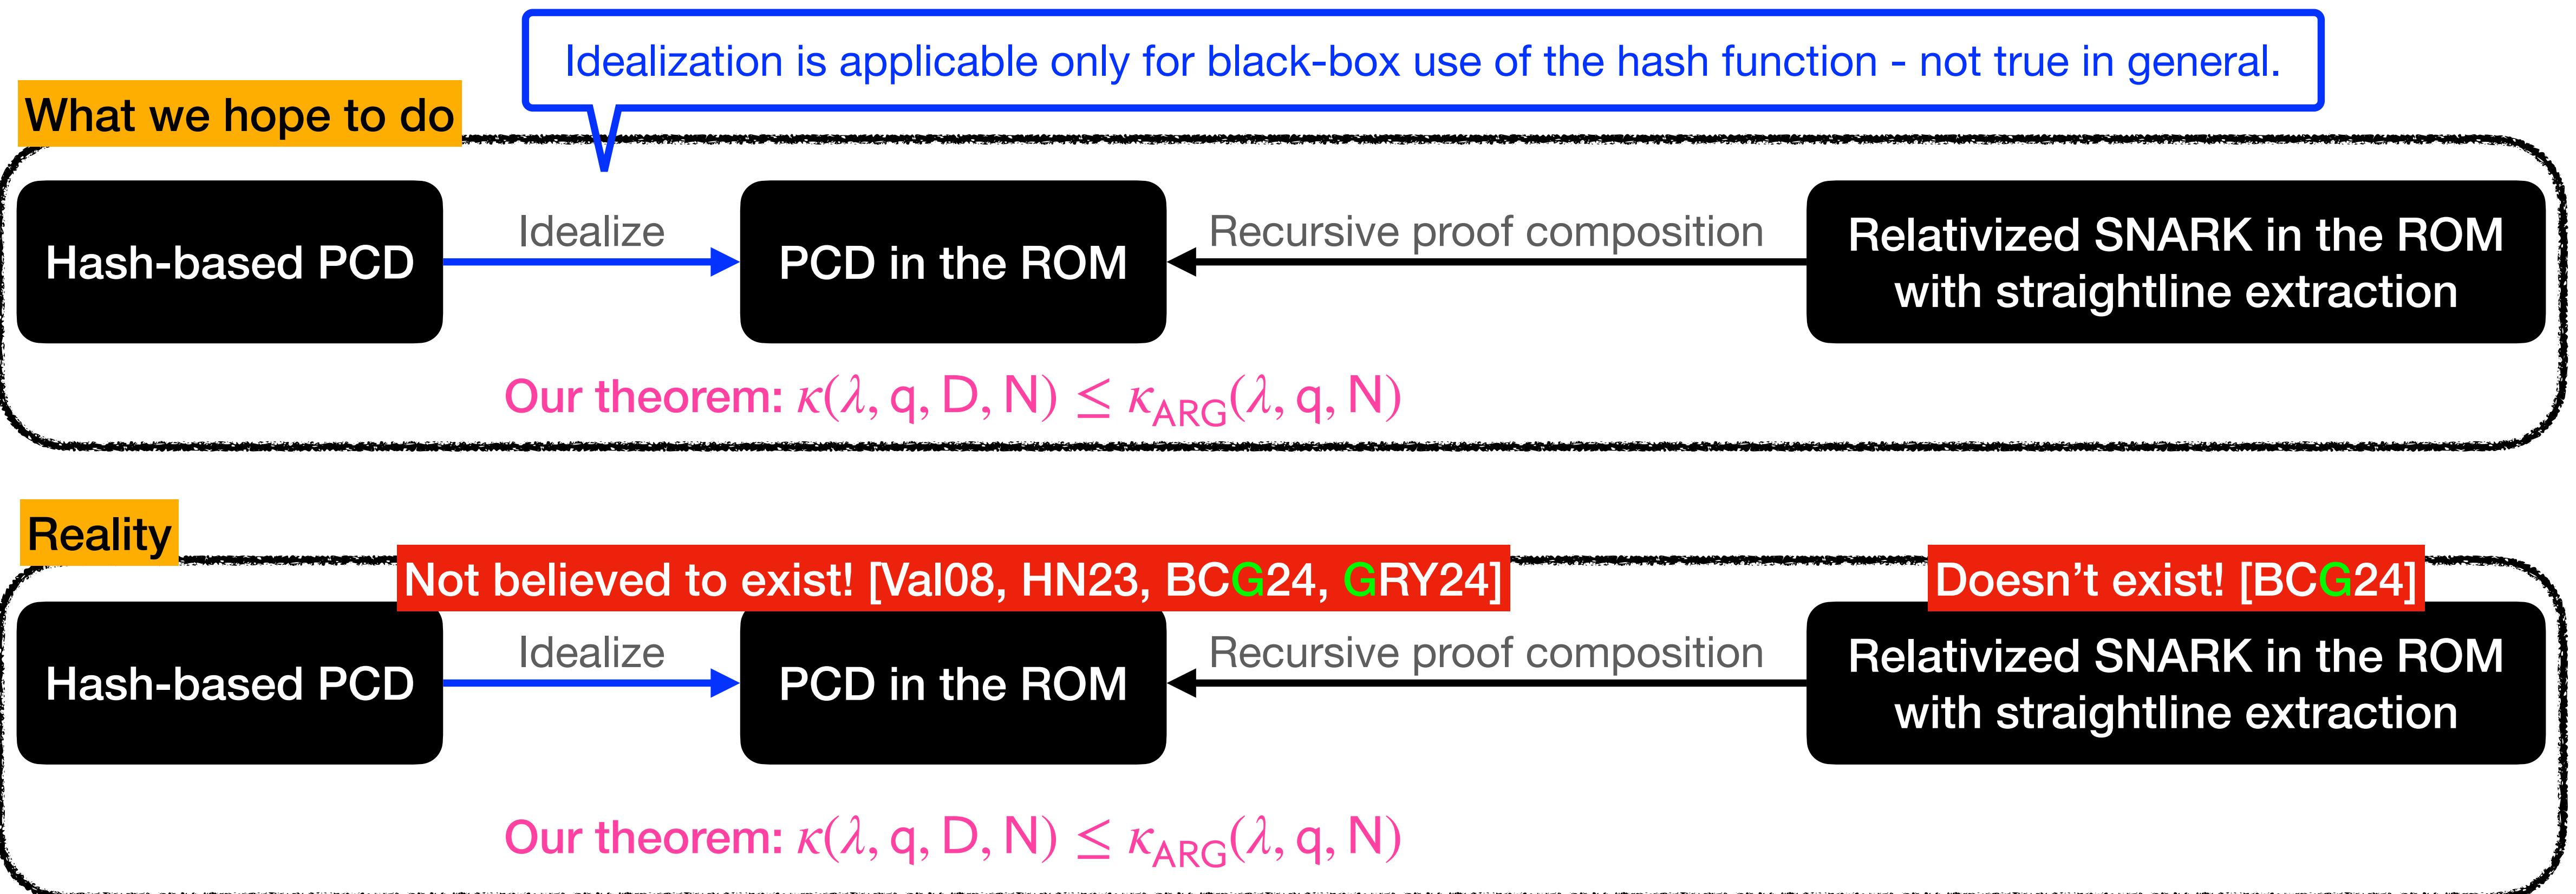
Second attempt for idealization of hash-based PCD



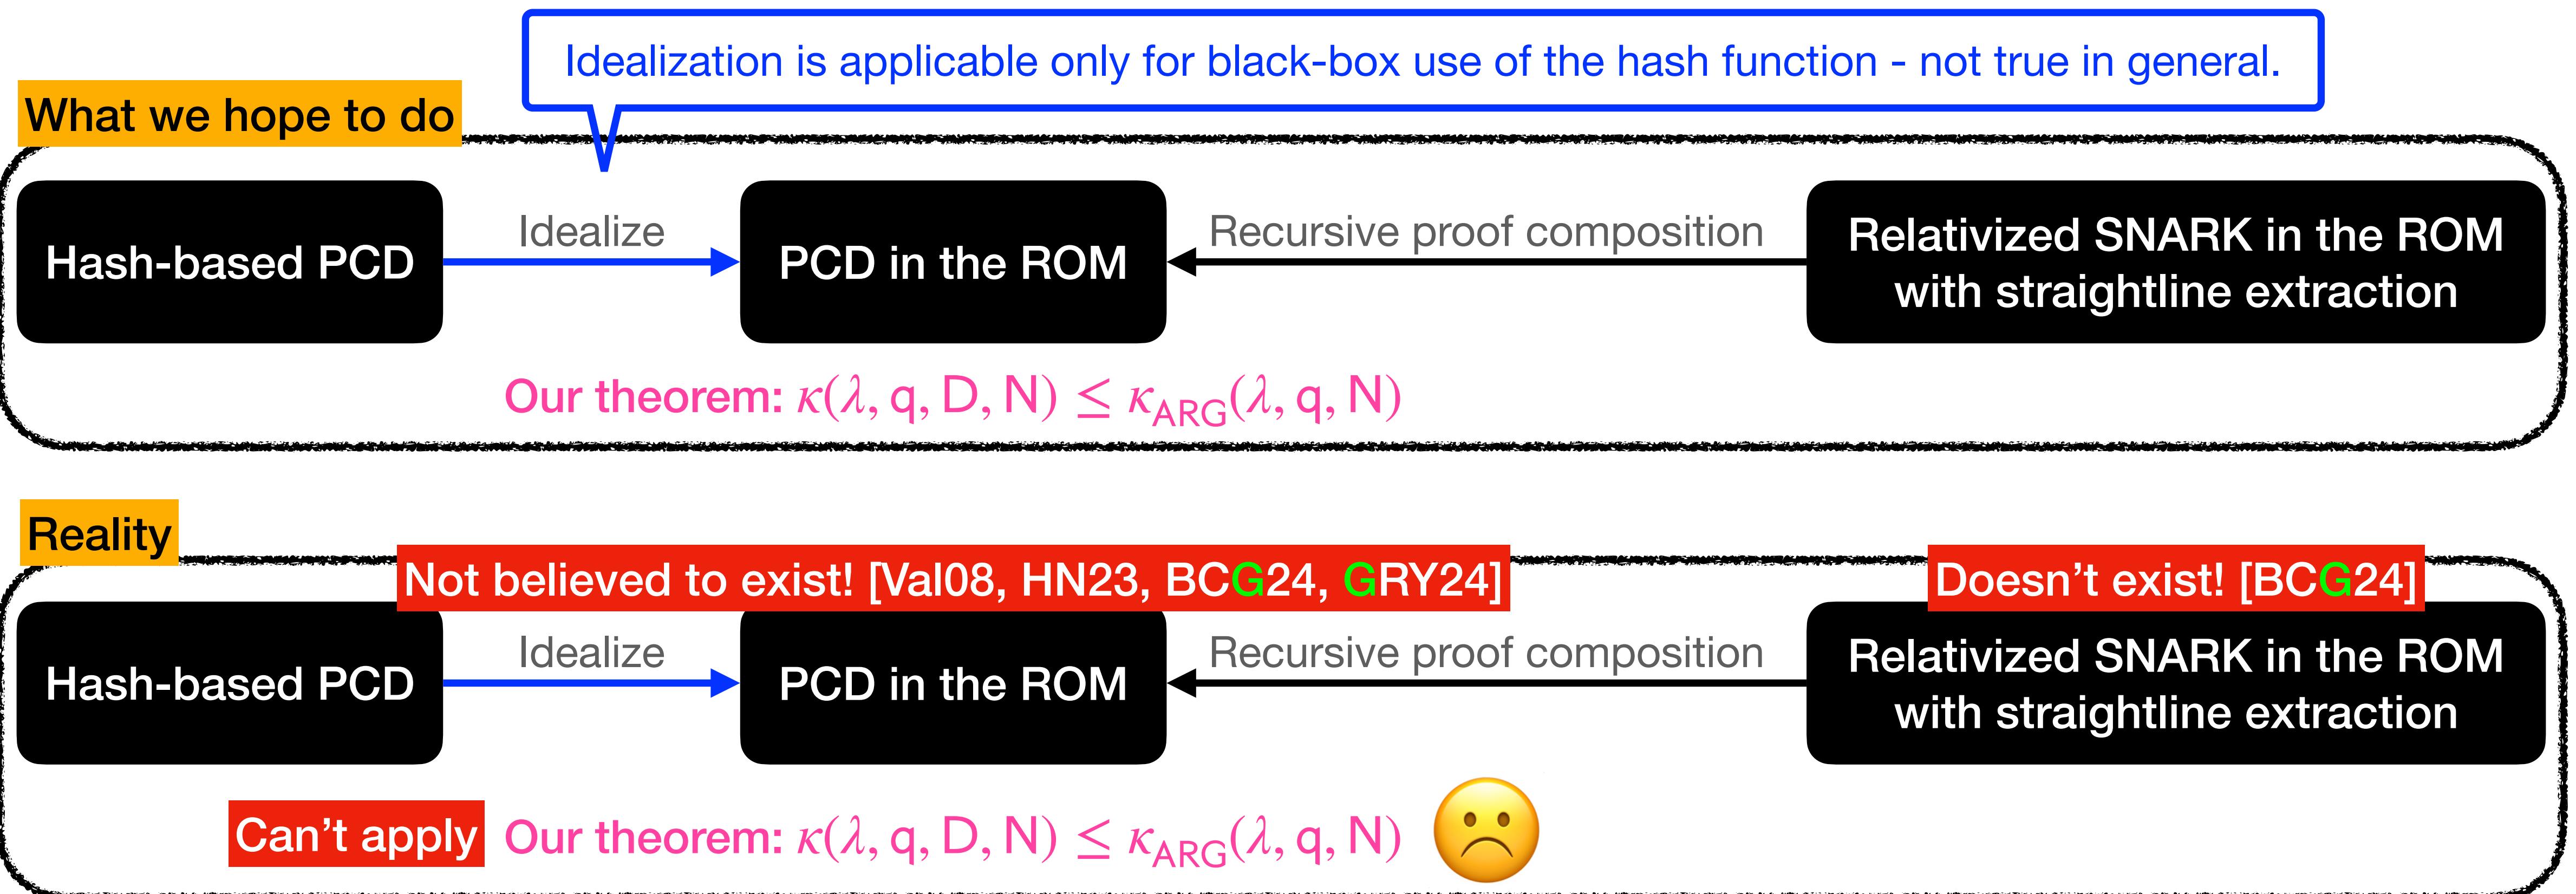
Second attempt for idealization of hash-based PCD



Second attempt for idealization of hash-based PCD



Second attempt for idealization of hash-based PCD



Our idealization for hash-based PCD

Our idealization for hash-based PCD

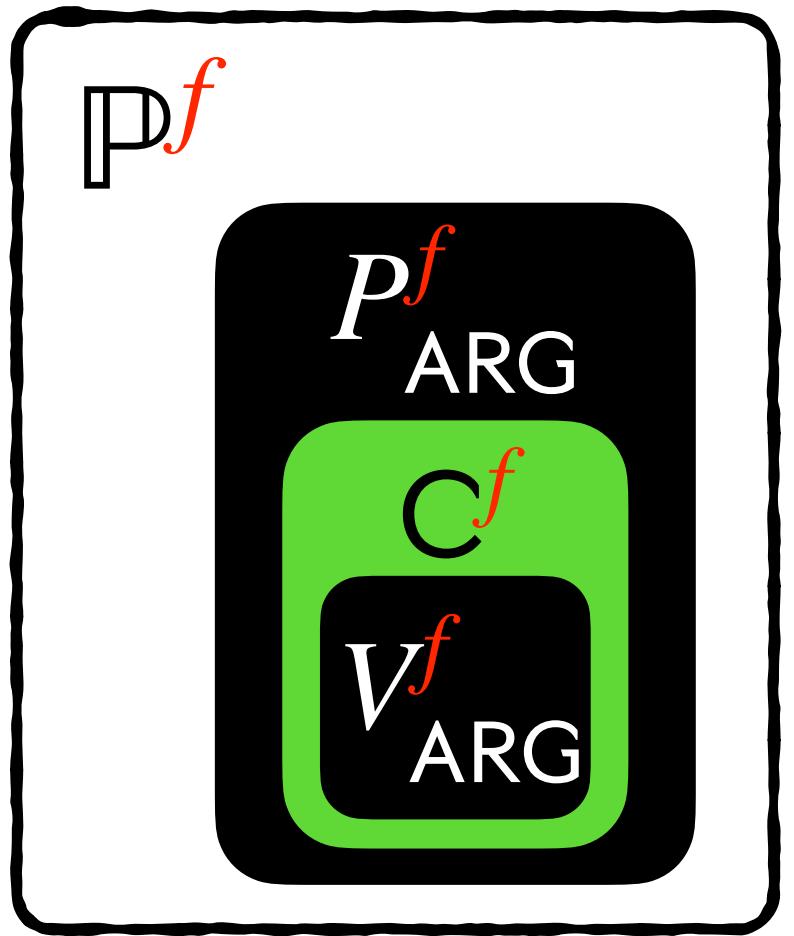
Issue: Hash-based PCD uses hash function in a non-black-box way.

Our idealization for hash-based PCD

Issue: Hash-based PCD uses hash function in a non-black-box way.

Observation 1: PCD looks at hash function to check the correctness, it doesn't "destroy" the hash function.

Our idealization for hash-based PCD

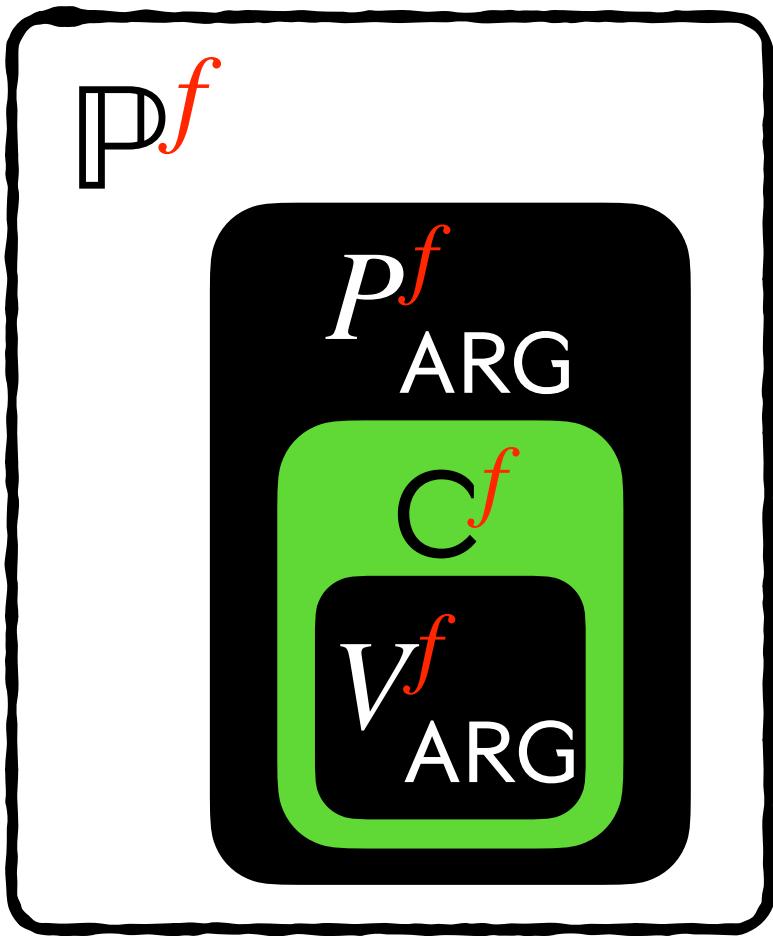


Issue: Hash-based PCD uses hash function in a non-black-box way.

Observation 1: PCD looks at hash function to check the correctness, it doesn't “destroy” the hash function.

Observation 2: C is an oracle circuit because V_{ARG} make oracle queries.

Our idealization for hash-based PCD



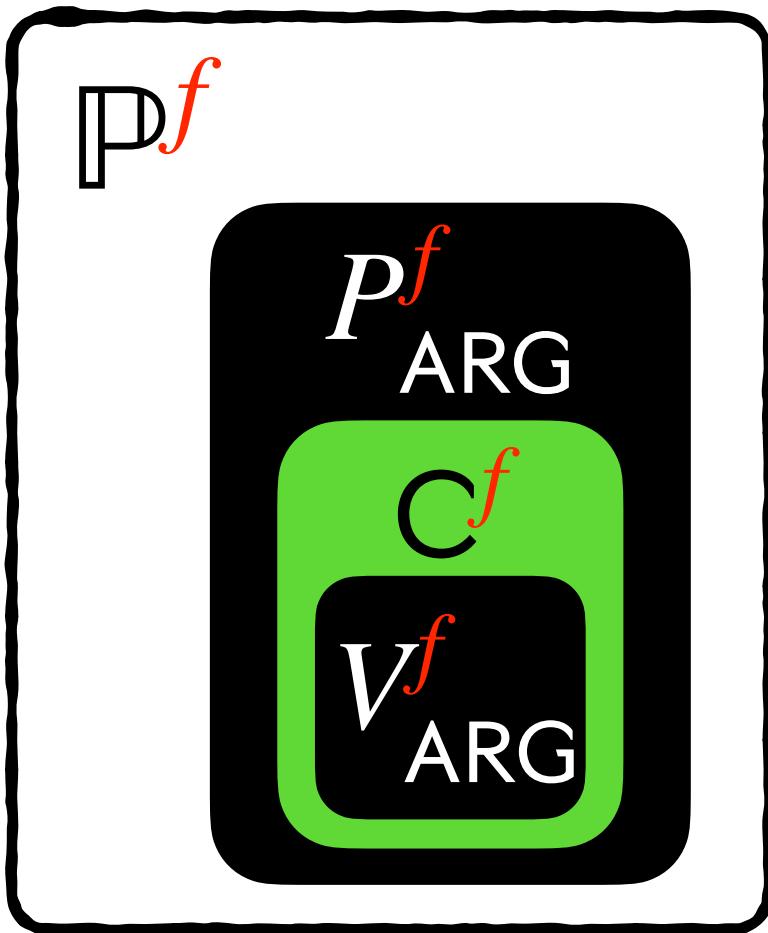
Issue: Hash-based PCD uses hash function in a non-black-box way.

Observation 1: PCD looks at hash function to check the correctness, it doesn't "destroy" the hash function.

Observation 2: C is an oracle circuit because V_{ARG} make oracle queries.

Solution: Forward all the queries of C by asking P_{ARG} to attach C 's "query-answer trace" in the proof.

Our idealization for hash-based PCD



Issue: Hash-based PCD uses hash function in a non-black-box way.

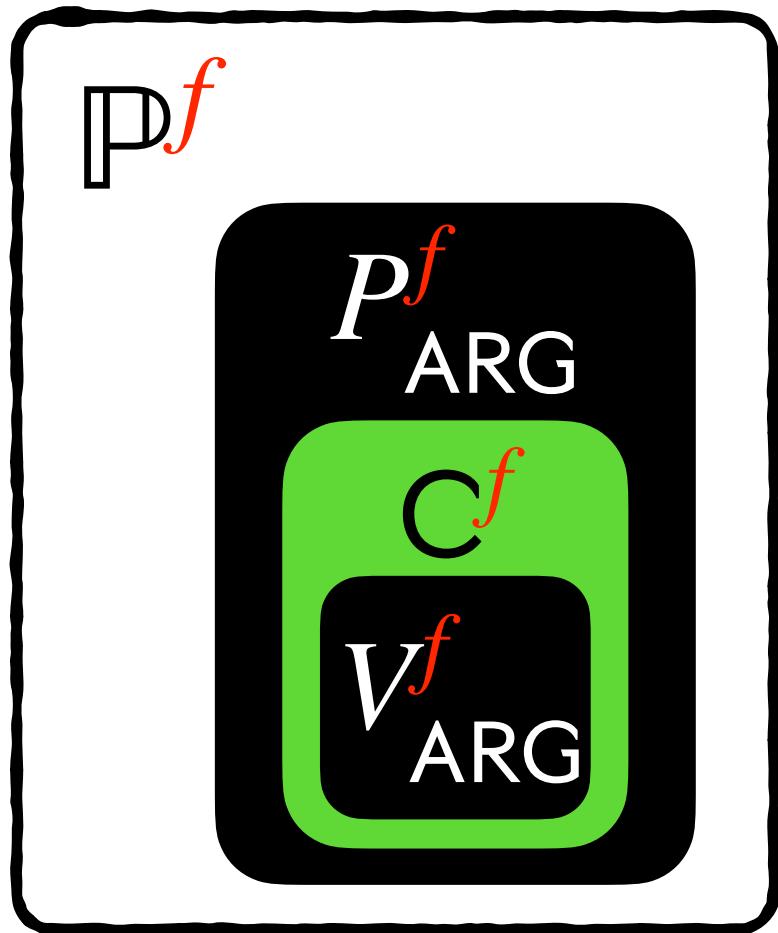
Observation 1: PCD looks at hash function to check the correctness, it doesn't "destroy" the hash function.

Observation 2: C is an oracle circuit because V_{ARG} make oracle queries.

Solution: Forward all the queries of C by asking P_{ARG} to attach C 's "query-answer trace" in the proof.

Forwarding the queries makes
the proof non-succinct

Our idealization for hash-based PCD



Issue: Hash-based PCD uses hash function in a non-black-box way.

Observation 1: PCD looks at hash function to check the correctness, it doesn't "destroy" the hash function.

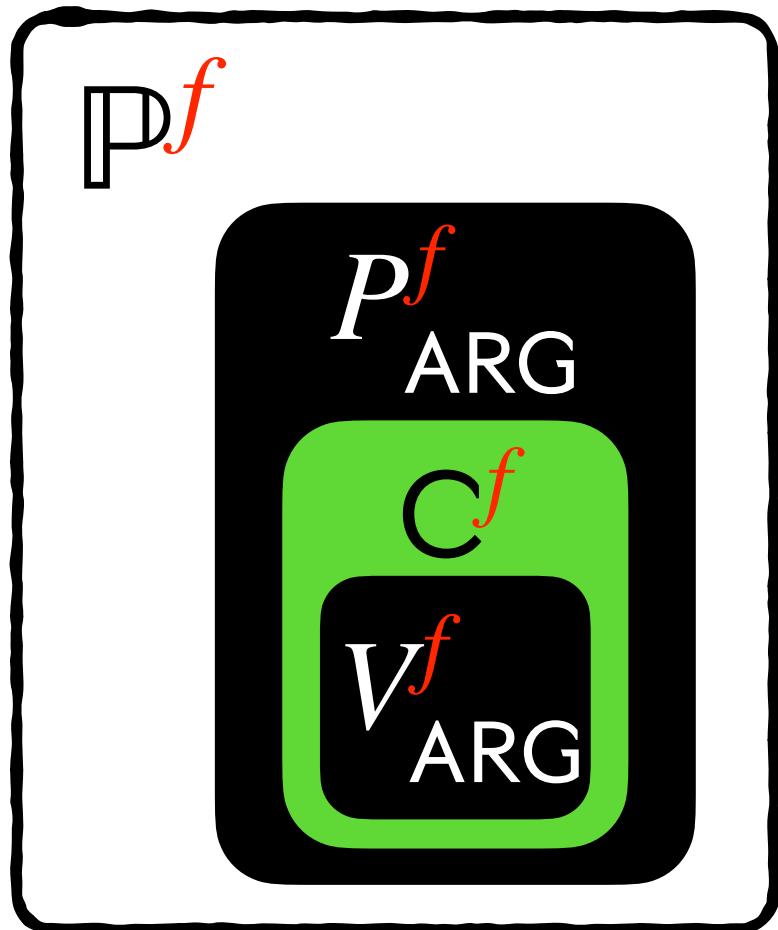
Observation 2: C is an oracle circuit because V_{ARG} make oracle queries.

Solution: Forward all the queries of C by asking P_{ARG} to attach C 's "query-answer trace" in the proof.

Forwarding the queries makes
the proof non-succinct

Hash-based PCD

Our idealization for hash-based PCD



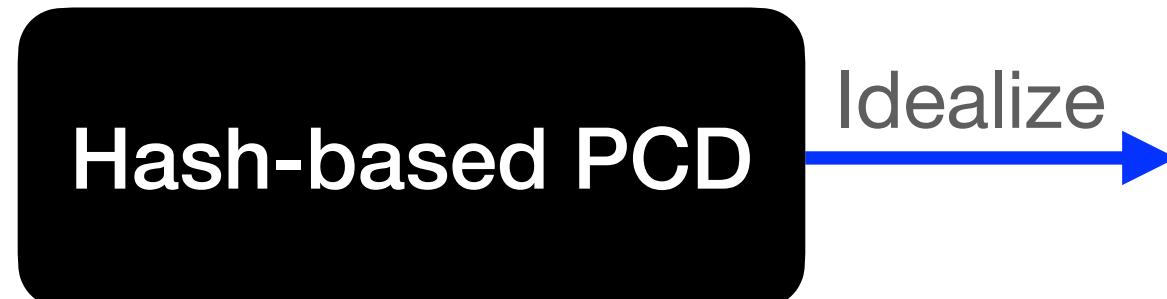
Issue: Hash-based PCD uses hash function in a non-black-box way.

Observation 1: PCD looks at hash function to check the correctness, it doesn't "destroy" the hash function.

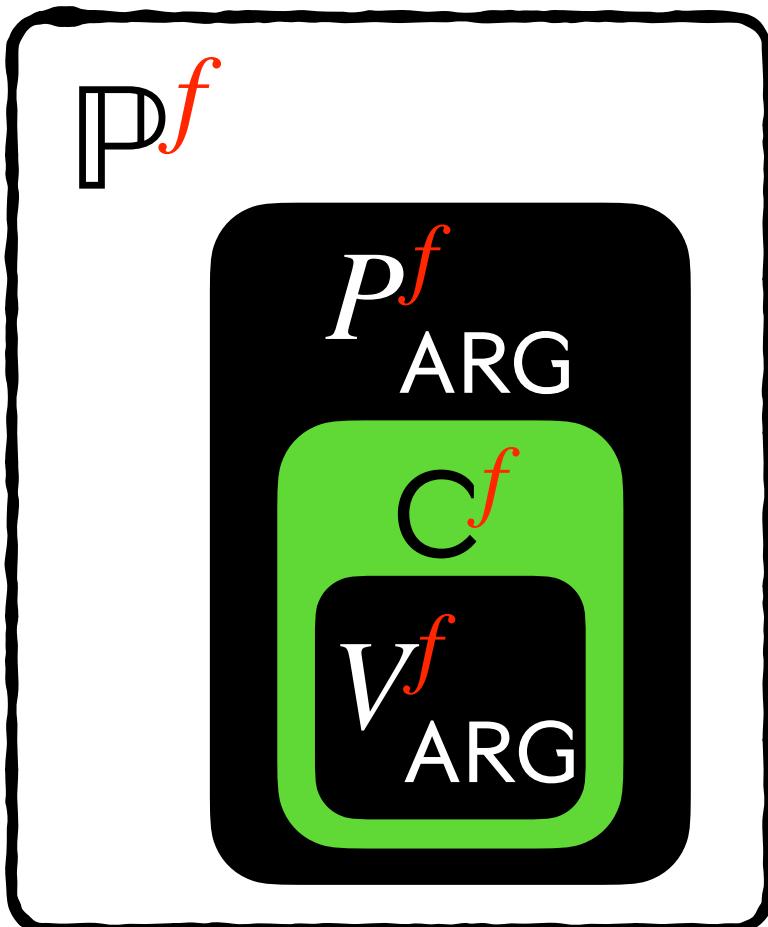
Observation 2: C is an oracle circuit because V_{ARG} make oracle queries.

Solution: Forward all the queries of C by asking P_{ARG} to attach C 's "query-answer trace" in the proof.

Forwarding the queries makes
the proof non-succinct



Our idealization for hash-based PCD



Issue: Hash-based PCD uses hash function in a non-black-box way.

Observation 1: PCD looks at hash function to check the correctness, it doesn't "destroy" the hash function.

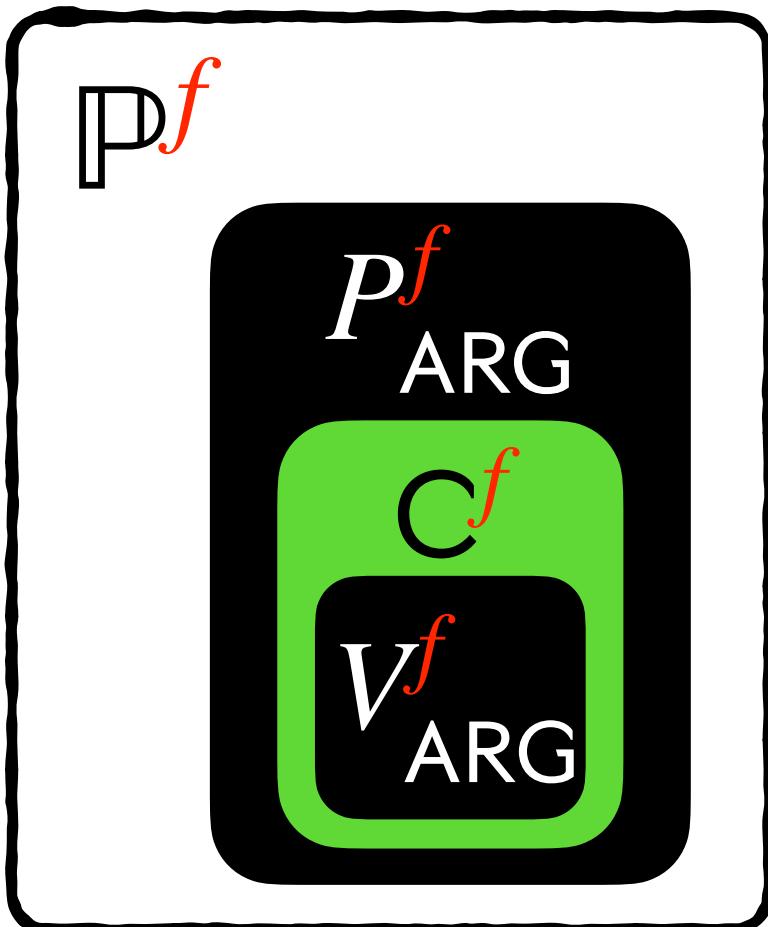
Observation 2: C is an oracle circuit because V_{ARG} make oracle queries.

Solution: Forward all the queries of C by asking P_{ARG} to attach C 's "query-answer trace" in the proof.

Forwarding the queries makes
the proof non-succinct



Our idealization for hash-based PCD



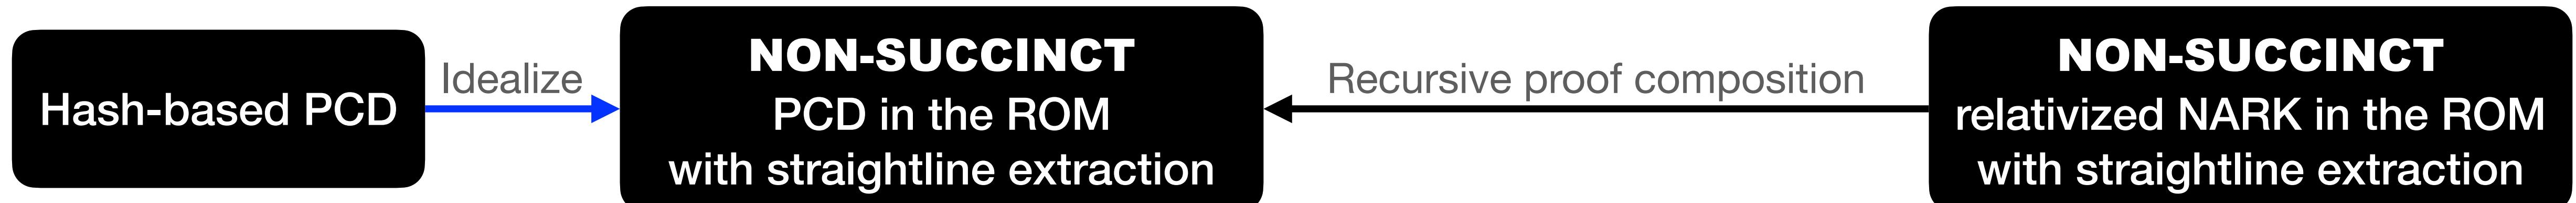
Issue: Hash-based PCD uses hash function in a non-black-box way.

Observation 1: PCD looks at hash function to check the correctness, it doesn't "destroy" the hash function.

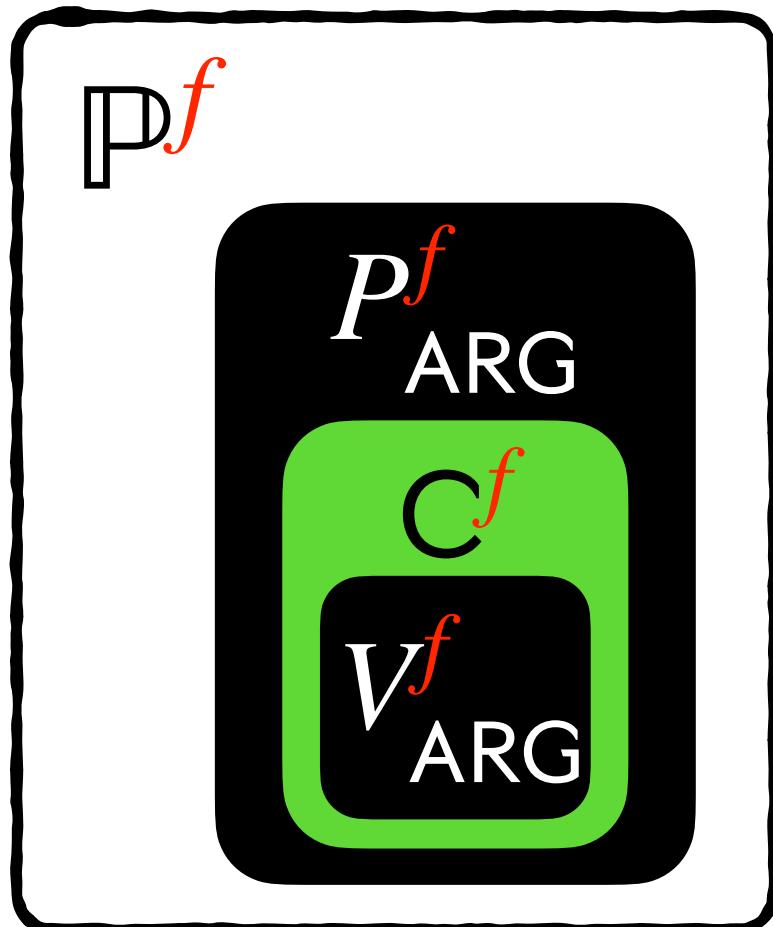
Observation 2: C is an oracle circuit because V_{ARG} make oracle queries.

Solution: Forward all the queries of C by asking P_{ARG} to attach C 's "query-answer trace" in the proof.

Forwarding the queries makes
the proof non-succinct



Our idealization for hash-based PCD



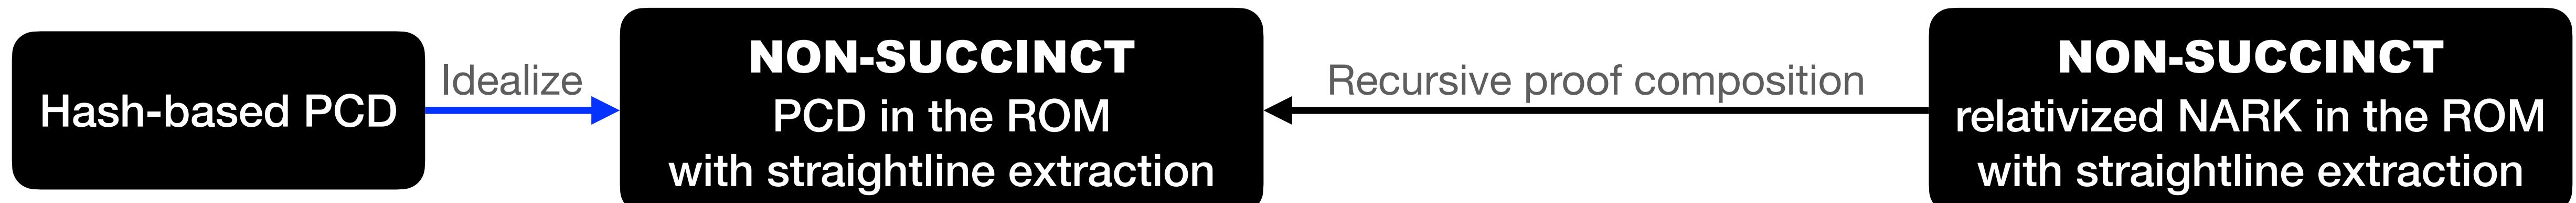
Issue: Hash-based PCD uses hash function in a non-black-box way.

Observation 1: PCD looks at hash function to check the correctness, it doesn't "destroy" the hash function.

Observation 2: C is an oracle circuit because V_{ARG} make oracle queries.

Solution: Forward all the queries of C by asking P_{ARG} to attach C 's "query-answer trace" in the proof.

Forwarding the queries makes
the proof non-succinct



Our theorem: $\kappa(\lambda, q, D, N) \leq \kappa_{\text{ARG}}(\lambda, q, N) = \kappa_{\text{ARG}}(\lambda, q)$

TL;DR

What is *proof-carrying data* (PCD)?

- Recursive compositions of SNARKs.
- It's useful for efficiently verifying distributed computations.

Problem:

- PCD is deployed under the assumption "security of PCD" = "security of underlying SNARK".
- BUT existing security analyses show a huge gap in security ("PCD is far less secure than underlying SNARK").

This work:

- We propose **an idealized PCD** that models hash-based PCD in practice.
- We prove that this idealized PCD is **as secure as its underlying SNARK**.

Thank you!

